

*Updated for
Version 9*

A Handbook of
**Statistical
Analyses**
Using **Stata**

Fourth Edition

Sophia Rabe-Hesketh
Brian S. Everitt



Chapman & Hall/CRC
Taylor & Francis Group

A Handbook of
**Statistical
Analyses**
Using **Stata**
Fourth Edition

Sophia Rabe-Hesketh
Brian S. Everitt



Chapman & Hall/CRC
Taylor & Francis Group

Boca Raton London New York

Chapman & Hall/CRC is an imprint of the
Taylor & Francis Group, an Informa business

Chapman & Hall/CRC
Taylor & Francis Group
6000 Broken Sound Parkway NW, Suite 300
Boca Raton, FL 33487-2942

© 2007 by Taylor & Francis Group, LLC
Chapman & Hall/CRC is an imprint of Taylor & Francis Group, an Informa business

No claim to original U.S. Government works.
Printed in the United States of America on acid-free paper
30987654321

International Standard Book Number-10: 1-58488-756-7 (Softcover)
International Standard Book Number-13: 978-1-58488-756-0 (Softcover)

This book contains information obtained from authentic and highly regarded sources. Reprinted material is quoted with permission, and sources are indicated. A wide variety of references are listed. Reasonable efforts have been made to publish reliable data and information, but the author and the publisher cannot assume responsibility for the validity of all materials or for the consequences of their use.

No part of this book may be reprinted, reproduced, transmitted, or utilized in any form by any electronic, mechanical, or other means, now known or hereafter invented, including photocopying, microfilm, and recording, or in any information storage or retrieval system, without written permission from the publishers.

For permission to photocopy or use material electronically from this work, please access www.copyright.com (<http://www.copyright.com>) or contact the Copyright Clearance Center, Inc. (CCC) 222 Rosewood Drive, Danvers, MA 01923, 978-750-8400. CCC is a not-for-profit organization that provides licenses and registration for a variety of users. For organizations that have been granted a photocopy license by the CCC, a separate system of payment has been arranged.

Trademark Notice: Product or corporate names may be trademarks or registered trademarks, and are used only for identification and explanation without intent to infringe.

Library of Congress Cataloging-in-Publication Data

Rabe-Hesketh, S.
A Handbook of statistical analyses using Stata / Sophia Rabe-Hesketh, Brian S. Everitt. — 4th ed.
p. cm.
Includes bibliographical references and index.
ISBN 1-58488-756-7 (acid-free paper)
1. Stata. 2. Mathematical statistics—Data processing. 1. Everitt, Brian. II. Title.

QA276.4.R33 2006
519.5'0285'536—dc22

2006049170

Visit the Taylor & Francis Web site at
<http://www.taylorandfrancis.com>

and the CRC Press Web site at
<http://www.crcpress.com>

Dedication

To my parents, Birgit and Georg Rabe
Sophia Rabe-Hesketh

To my wife, Mary Elizabeth
Brian S. Everitt

Preface

Stata is an exciting statistical package that offers all standard and many non-standard methods of data analysis. In addition to general methods such as linear, logistic and Poisson regression, and generalized linear models, Stata provides many more specialized analyses, such as generalized estimating equations from biostatistics and the Heckman selection model from econometrics. Stata has extensive capabilities for the analysis of survival data, time series, panel (or longitudinal) data, and complex survey data. For all estimation problems, inferences can be made more robust to model misspecification using bootstrapping or robust standard errors based on the sandwich estimator. In each new release of Stata, its capabilities are significantly enhanced by a team of excellent statisticians and developers at StataCorp.

Although extremely powerful, Stata is easy to use, either by point-and-click or through its intuitive command syntax. Applied researchers, students, and methodologists therefore all find Stata a rewarding environment for manipulating data, carrying out statistical analyses, and producing publication quality graphics.

Stata also provides a powerful programming language making it easy to implement a 'tailor-made' analysis for a particular application or to write more general commands for use by the wider Stata community. In fact we consider Stata an ideal environment for developing and disseminating new methodology. First, the elegance and consistency of the programming language is appealing for methodologists. Second, it is simple to make new commands behave in every way like Stata's own commands, making them accessible to applied researchers and students. Third, Stata's email listserver *Statalist*, *The Stata Journal*, the Stata Users' Group Meetings, and the Statistical Software Components (SSC) archive on the internet all make exchange and discussion of new commands extremely easy. For these reasons Stata is constantly kept

up-to-date with recent developments, not just by its own developers, but also by a very active Stata community.

This handbook follows the format of its two predecessors, *A Handbook of Statistical Analysis Using S-PLUS* and *A Handbook of Statistical Analysis Using SAS*. Each chapter deals with the analysis appropriate for a particular application. A brief account of the statistical background is included in each chapter including references to the literature, but the primary focus is on how to use Stata, and how to interpret results. Our hope is that this approach will provide a useful complement to the excellent but very extensive Stata manuals. The majority of the examples are drawn from areas in which the authors have most experience, but we hope that current and potential Stata users from outside these areas will have little trouble in identifying the relevance of the analyses described for their own data.

In the fourth edition, we have added many new exercises based on new datasets. For exercises marked with the symbol •, answers are provided in the appendix. For the remaining exercises, a solutions manual is available from Chapman & Hall/CRC for course instructors.

Particular thanks are due to Nick Cox who provided us with extensive general comments for the second, third, and fourth editions of our book, and also gave us clear guidance as to how best to use a number of Stata commands. We are also grateful to Anders Skrondal for commenting on several drafts of the third edition. Various people at StataCorp have been very helpful in preparing the second, third, and fourth editions of this book. We would also like to acknowledge the usefulness of the Stata NetCourses in the preparation of the first edition of this book.

All the datasets can be downloaded from:

- <http://www.stata.com/texts/stas4>

Individual datasets can also be read directly into Stata from the above site by specifying the full path. For example, to read the data `wagepan.dta` for Exercise 1.2, use the following command:

```
use http://www.stata.com/texts/stas4/wagepan
```

S. Rabe-Hesketh
B. S. Everitt
Berkeley and London

Contents

1 A Brief Introduction to Stata.....	1
1.1 Getting help and information 1	
1.2 Running Stata 2	
1.3 Conventions used in this book 9	
1.4 Datasets in Stata 9	
1.5 Stata commands 13	
1.6 Data management 19	
1.7 Estimation 22	
1.8 Graphics 24	
1.9 Stata as a calculator 30	
1.10 Matrix calculations using Mata 32	
1.11 Brief introduction to programming 34	
1.12 Keeping Stata up to date 39	
1.13 Exercises 40	
2 Data Description and Simple Inference: Female Psychiatric Patients.....	43
2.1 Description of data 43	
2.2 Group comparison and correlations 46	
2.3 Analysis using Stata 47	
2.4 Exercises 57	
3 Multiple Regression: Determinants of Pollution in U.S. Cities	61
3.1 Description of data 61	
3.2 The multiple regression model 63	
3.3 Analysis using Stata 64	
3.4 Exercises 82	

4	Analysis of Variance I: Treating Hypertension	85
4.1	Description of data	85
4.2	Analysis of variance model	85
4.3	Analysis using Stata	87
4.4	Exercises	96
5	Analysis of Variance II: Effectiveness of Slimming Clinics	101
5.1	Description of data	101
5.2	Analysis of variance model	102
5.3	Analysis using Stata	104
5.4	Exercises	108
6	Logistic Regression: Treatment of Lung Cancer and Diagnosis of Heart Attacks	111
6.1	Description of data	111
6.2	The logistic regression model	112
6.3	Analysis using Stata	116
6.4	Exercises	129
7	Generalized Linear Models: Australian School Children	133
7.1	Description of data	133
7.2	Generalized linear models	134
7.3	Analysis using Stata	139
7.4	Exercises	153
8	Summary Measure Analysis of Longitudinal Data: Treatment of Post-Natal Depression.....	157
8.1	Description of data	157
8.2	The analysis of longitudinal data	159
8.3	Analysis using Stata	159
8.4	Exercises	170
9	Random Effects Models: Thought Disorder and Schizophrenia	173
9.1	Description of data	173
9.2	Random effects models	173
9.3	Analysis using Stata	178
9.4	Thought disorder data	190
9.5	Exercises	199
10	Generalized Estimating Equations: Epileptic Seizures and Chemotherapy	201
10.1	Description of data	201

10.2 Generalized estimating equations	203
10.3 Analysis using Stata	205
10.4 Exercises	218
11 Some Epidemiology	221
11.1 Description of data	221
11.2 Introduction to epidemiology	222
11.3 Analysis using Stata	228
11.4 Exercises	236
12 Survival Analysis: Retention of Heroin Addicts in Methadone Maintenance Treatment	239
12.1 Description of data	239
12.2 Survival analysis	242
12.3 Analysis using Stata	245
12.4 Exercises	258
13 Maximum Likelihood Estimation: Age of Onset of Schizophrenia	263
13.1 Description of data	263
13.2 Finite mixture distributions	263
13.3 Analysis using Stata	264
13.4 Exercises	277
14 Principal Components Analysis: Hearing Measurement Using an Audiometer	281
14.1 Descriptions of data	281
14.2 Principal component analysis	283
14.3 Analysis using Stata	284
14.4 Exercises	291
15 Cluster Analysis: Tibetan Skulls and Determinants of Pollution in U.S. Cities	295
15.1 Description of data	295
15.2 Cluster analysis	297
15.3 Analysis using Stata	298
15.4 Exercises	311
Appendix: Answers to Selected Exercises	315
References	327
Index	335

Chapter 1

A Brief Introduction to Stata

1.1 Getting help and information

Stata is a general purpose statistics package developed and maintained by StataCorp. There are several forms or "flavors" of Stata, the standard Intercooled Stata, the more limited Small Stata, Stata/SE (Special Edition) which can handle extremely large datasets, and Stata/MP (Multiple Processors) which runs in parallel on up to 32 processors. Each flavor exists for Windows (2000, XP, and later versions), Unix platforms, and the Macintosh. Almost all Stata features discussed in this book are common across platforms.

The base documentation set for Stata consists of eight manuals (StataCorp 2005a-h): *Getting Started with Stata*, *Stata User's Guide*, *Base Reference Manuals* (three volumes), *Data Management Reference Manual*, *Graphics Reference Manual*, and *Quick Reference and Index*. In addition there are more specialized reference manuals such as the *Stata Programming Reference Manual* and the *Stata Longitudinal/Panel Data Reference Manual*. The reference manuals provide extremely detailed information on each command while the *User's Guide* describes Stata more generally. Features that are specific to the operating system are described in the appropriate *Getting Started* manual, e.g., *Getting Started with Stata for Windows*.

Each Stata command has associated with it a help file that may be viewed within a Stata session using the help facility. Both the help-files and the manuals refer to the *Base Reference Manuals* by [R] name of entry, to the *User's Guide* by [U] chapter or section number and

name, the *Graphics Manual* by [G] name of entry, etc. (see *Stata Getting Started* manual, immediately after the table of contents, for a complete list).

There are an increasing number of general introductory books on Stata, including the book you are reading now, Kohler and Kreuter (2005), and Acock (2006). In addition, there are books on Stata for particular types of analysis such as categorical data analysis (Long and Freese, 2006), survival analysis (Cleves, Gould and Gutierrez, 2004), generalized linear models (Hardin and Hilbe, 2006), and multilevel and longitudinal models (Rabe-Hesketh and Skrondal, 2005). The web site <http://www.stata.com/bookstore/statabooks.html> provides up-to-date information on these and other books.

The Stata web page at <http://www.stata.com> offers much useful information for learning Stata including an extensive series of "frequently asked questions" (FAQs). Stata also offers Internet courses, called *NetCourses*. These courses take place via a temporary mailing list for course organizers and "attenders". Each week, the course organizers send out lecture notes and exercises which the attenders can discuss with each other until the organizers send out the answers to the exercises and to the questions raised by attenders.

The UCLA Academic Technology Services offer useful textbook and paper examples at <http://www.ats.ucla.edu/stat/stata/>, showing how analyses can be carried out using Stata. Also very helpful for learning Stata are the regular columns *Speaking Stata* and *Stata Tips* in *The Stata Journal*; see <http://www.stata-journal.com>. It is possible to purchase individual issues, or a compilation of Stata tips by Newton and Cox (2006).

One of the exciting aspects of being a Stata user is being part of a very active Stata community as reflected in the busy *Statalist* mailing list, Stata Users' Group meetings taking place every year in the UK, USA and various other countries, and the large number of user-contributed Stata programs; see also Section 1.12. *Statalist* also functions as a technical support service with Stata staff and expert users such as Nick Cox offering very helpful responses to questions.

1.2 Running Stata

This section gives an overview of what happens in a typical Stata session, referring to subsequent sections for more details. We are using the Windows version here and some features may be different in Stata for other platforms. We therefore recommend consulting the *Getting Started With Stata* manual for your platform.

1.2.1 Stata windows

When Stata is started, a screen opens as shown in Figure 1.1 containing four windows labeled:

- Command: here commands are issued interactively
- Results: here results are displayed
- Review: here all commands issued within the current Stata session are shown
- Variables: here the variables of the current dataset are listed

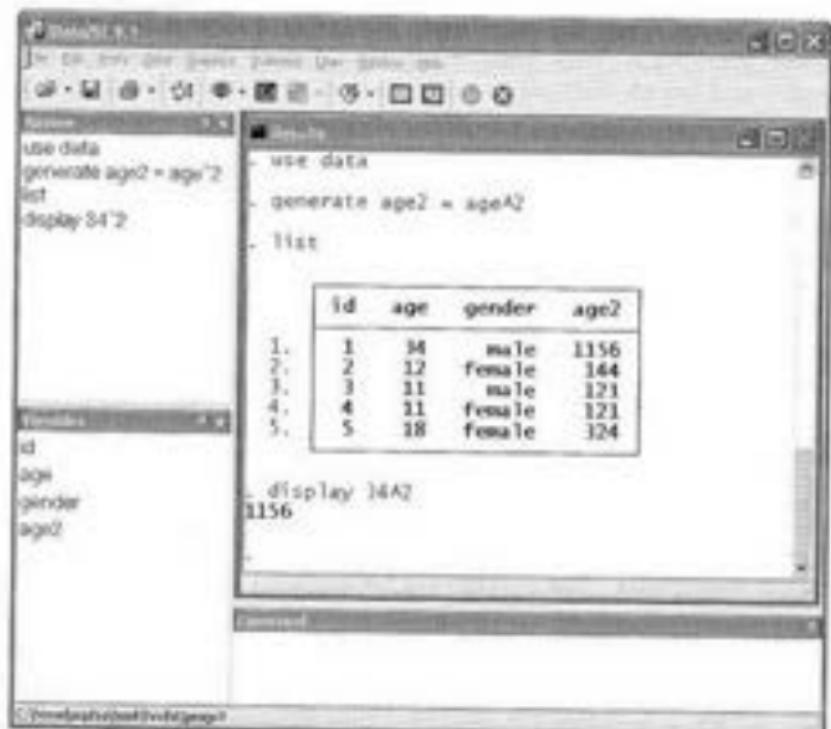


Figure 1.1: Stata windows.

Each of the Stata windows can be resized and moved around in the usual way; the Command, Review, and Variables windows can also be moved outside the main window (undocked) in which case they will not move along with the main Stata window. To bring an undocked

window forward that may be obscured by other windows, make the appropriate selection in the **Window** menu. To dock a window, drag it back into the main window. A transparent blue box appears in place of the window being dragged and docking guides appear at the center and edges of the main window. Release the mouse button when the transparent blue box is on the appropriate docking guide, for instance on the arrow pointing down, to dock the window at the bottom of the main Stata window.

The fonts in a window can be changed by clicking the right mouse button over the window. All these settings are automatically saved when Stata is closed. Use the **Manage Preferences** selection from the **Prefs** menu to save and load specific settings, for instance a large font setting for teaching, or to reload the factory (or default) settings.

Three other types of windows can be created within a Stata session: Viewer windows to view help or log files, Graph windows to display graphs, and Do-file Editors to build and run scripts (called do-files).

1.2.2 Datasets

Stata datasets have the .dta extension and can be loaded into Stata in the usual way through the **File** menu (for reading other data formats; see Section 1.4.1). As in other statistical packages, a dataset is a matrix where the columns represent variables (with names and labels) and the rows represent observations. When a dataset is open, the variable names and variable labels appear in the Variables window. The dataset may be viewed as a spreadsheet by opening the Data Browser with the  button and edited by clicking  to open the Data Editor. Both the Data Browser and the Data Editor can also be opened through the **Window** menu. Note, however, that nothing else can be done in Stata while the Data Browser or Data Editor is open (e.g., the Command window disappears). See Section 1.4 for more information on datasets.

1.2.3 Commands and output

Until release 8.0, Stata was entirely command-driven and many users still prefer using commands as follows: a command is typed in the Command window and executed by pressing the *Return* (or *Enter*) key. The command then appears next to a full stop (period) in the Stata Results window, followed by the output.

If the output produced is longer than the Results window, --more-- appears at the bottom of the screen. Pressing any key scrolls the output forward one screen. The scroll-bar may be used to move up and down previously displayed output. However, only a certain amount of

past output is retained in the Results window. For this reason and to save output for later, it is useful to open a log file; see Section 1.2.6. It is possible to copy and print selected output from the Results window. **Edit → Copy Table** can be used to copy and paste tables so that columns are separated by tabs making it easy to produce a nicely formatted table for instance in Word.

Stata is ready to accept a new command when the prompt (a period) appears at the bottom of the screen. If Stata is not ready to receive new commands because it is still running or has not yet displayed all the current output, it may be interrupted by holding down *Ctrl* and pressing the *Pause/Break* key or by pressing the red Break button .

A previous command can be accessed using the *PgUp* and *PgDn* keys or by selecting it from the Review window where all commands from the current Stata session are listed (see Figure 1.1). The command may then be edited if required before pressing *Return* to execute the command.

Most Stata commands refer to a list of variables, the basic syntax being **command varlist**. For example, if the dataset contains variables *x*, *y*, and *z*, then

```
list x y
```

lists the values of *x* and *y*. Other components may be added to the command; for example, adding *if exp* after *varlist* causes the command to process only those observations satisfying the logical expression *exp*. Options are separated from the main command by a comma. The complete command structure and its components are described in Section 1.5.

1.2.4 GUI versus commands

Since release 8.0, Stata has a Graphical User Interface (GUT) that allows all non-programming commands to be accessed via point-and-click. Simply start by clicking into the **Data**, **Graphics**, or **Statistics** menus, make the relevant selections, fill in a dialog box, and click **OK**. Stata then behaves exactly as if the corresponding command had been typed with the command appearing in the Results and Review windows and being accessible via *PgUp* and *PgDn*.

A great advantage of the menu system is that it is intuitive so that a complete novice to Stata could learn to run a linear regression in a few minutes. A disadvantage is that pointing and clicking can be time-consuming and cannot be automated. Commands, on the other hand, can be saved in a file (called a do-file in Stata) and run again at a later time. In our opinion, the menu system is a great device for

finding out which command is needed and learning how it works, but serious statistical analysis is best undertaken using commands. In this book we therefore say very little about the menus and dialogs (they are largely self-explanatory after all), but see Section 1.8 for an example of creating a graph through the dialogs.

It is easy to move back and forth between a dialog box and help for the corresponding command; to move to help, click into at the bottom-left of the dialog box; to move to the dialog box, click on the link at the top-right of the help viewer.

1.2.5 Do-files

It is useful to build up a file containing the commands necessary to carry out a particular data analysis. This may be done using Stata's Do-file Editor or any other editor. The Do-file Editor may be opened by clicking . Commands can then be typed in and run as a batch either by clicking into in the Do-file Editor or by using the command

```
do dofile
```

Alternatively, a subset of commands can be highlighted and executed by clicking into . The do-file can be saved for use in a future Stata session. To open a do-file, select **Do...** from the **File** menu or open the do-file editor and use its **File** menu. See Section 1.11 for more information on do-files.

1.2.6 Log files

It is useful to open a log file at the beginning of a Stata session. Press the button , type a filename into the dialog box, and choose **Save**. By default, this produces a SMCL (Stata Markup and Control Language, pronounced "smcl") file with extension **.smcl**, but a plain text (ASCII) file can be produced by selecting the **.log** extension. If the file already exists, another dialog opens to allow you to decide whether to overwrite the file with new output or to append new output to the existing file.

The log file can be viewed in the Stata Viewer during the Stata session (again through). For long log files, it can be useful to click into in the Stata Viewer to search for a piece of text. The log file is automatically saved when it is closed. Log files can also be opened, viewed, and closed by selecting **Log** from the **File** menu, followed by **Begin...**, **View...**, or **Close**, respectively. The following commands

can be used to open and close a log file `mylog`, replacing the old one if it already exists:

```
log using mylog, replace
log close
```

To view a log file produced in a previous Stata session, select **File** → **Log** → **View...** and specify the full path of the log file. The log may then be printed by selecting **Print** → **Viewer** from the **File** menu.

It is also possible to translate SMCL files to plain text files and vice versa using the **File** menu or the **translate** command. To save the output in the Results window as a plain text log file, type

```
translate @Results mylog.txt
```

1.2.7 Getting help

Help may be obtained by clicking on **Help** which brings up the menu shown in Figure 1.2. To get help on a Stata command, assuming the



Figure 1.2: Menu for help.

command name is known, select **Stata Command...**. To find the appropriate Stata command first, select **Search...**, which opens up the dialog in Figure 1.3. For example, to find out how to fit a Cox regression model, type "survival" or "cox" under **Keywords:** and press **OK**. This opens a Stata Viewer containing a list of relevant command names with a brief description. In this case `stcox` is the command we need. Also listed are topics for which Frequently Asked Questions (FAQs) or examples are available on the web, and user contributed commands published in the Stata Journal (abbreviated SJ) or its predecessor, the Stata Technical Bulletin (abbreviated STB). Each entry in this list

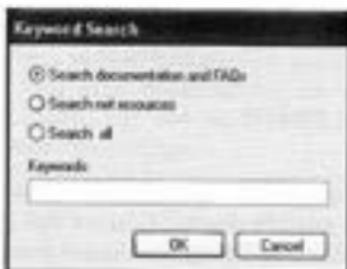


Figure 1.3: Dialog for search.

includes a blue keyword (a *hyperlink*) that may be selected to view the appropriate help file or web site. Each help file contains hyperlinks to other relevant help files.

The search can also be performed via the commands

`search survival`

and help on the command `stcox` can be found using

`help stcox`

Help will then appear in the Stata Results window (instead of the Stata Viewer) where words displayed in blue also represent hyperlinks to other files.

If the computer running Stata is connected to the internet, you can also search through "unofficial" materials on the Internet, to find, for instance, user-contributed programs not published in the Stata Journal or Stata Technical Bulletin (see 1.12 for more information). This is accomplished by selecting "Search net resources" or "Search all" in the search dialog box. The latter is equivalent to using the `findit` keyword command. More refined searches can be carried out using the `search` command (see `help search`).

The other selections in the help dialog, News, Official Updates, SJ and User-written Programs, and Stata Web Site, all provide access to the relevant web sites.

1.2.8 Closing Stata

Stata can be closed in three ways:

- click on the Close button  at the top right-hand corner of the Stata screen

- select **Exit** from the **File** menu
- type **exit**, **clear** in the Stata Commands window, and press **Return**.

1.3 Conventions used in this book

In this book we will use typewriter font like this for anything that could be typed into the Stata Command window or a do-file, that is, command names, options, variable names, etc. In contrast, italicized words are not supposed to be typed; they should be substituted by another word. For example, **summarize varname** means that *varname* should be substituted by a specific variable name, such as *age*, giving **summarize age**. We will usually display sequences of commands as follows:

```
summarize age
drop age
```

If a command continues over two lines, we use **///** at the end of the first line to make Stata ignore the line break. An alternative is to use **/*** at the end of the first line and ***/** at the beginning of the second line to "comment out" the linebreak. Note that these methods are for use in a do-file and do not work in the Command window where they would result in an error. In the Command window, commands can wrap over several lines.

Output taking little space is displayed immediately following the commands but without indentation and in a smaller font:

```
display 1
```

1

Output taking up more space is shown in a numbered display floating in the text. Some commands produce little notes, for example, the **generate** command prints out how many missing values are generated. We will usually not show such notes.

1.4 Datasets in Stata

1.4.1 Data input and output

Stata has its own data format with default extension **.dta**. Reading and saving a Stata file are straightforward. If the filename is **wagepan.dta**, the commands are

```
use wagepan
save wagepan
```

If the data are not stored in the current directory, then the complete path must be specified, as in the command

```
use c:\user\data\wagepan
```

(If the path contains spaces, it must be enclosed in quotes, e.g., "c:\my own data\wagepan".) However, the least error-prone way of keeping all the files for a particular project in one directory is to change to that directory and refer to all files without specifying the path:

```
cd c:\user\data
use wagepan
save wagepan
```

Note that the datasets of this book can also be read from a web site by specifying the path <http://www.stata.com/texts/stas4>, e.g.,

```
use http://www.stata.com/texts/stas4/wagepan
```

Data supplied with Stata can be read in using the `sysuse` command. For instance, the famous `auto.dta` data, which are often used in the Stata manuals, can be read using

```
sysuse auto
```

Before reading a file into Stata, all data already in memory need to be cleared, either by running `clear` before the `use` command or by using the option `clear` as follows:

```
use wagepan, clear
```

If we wish to save data under an existing filename, this results in an error message unless we use the option `replace` as follows:

```
save wagepan, replace
```

For large datasets it is sometimes necessary to increase the amount of memory Stata allocates to its data areas. For example, when no dataset is loaded (e.g., after issuing the command `clear`), set the memory to 2 megabytes using

```
set memory 2m
```

The `memory` command without arguments gives information on how much memory is being used and how much is available.

If the data are not available in Stata format, they may be converted to Stata format using another package (e.g., Stat/Transfer) or saved as an ASCII file (although the latter option means losing all the labels). When saving data as ASCII, missing values should be replaced by some numerical code.

There are three commands available for reading different types of ASCII data: `insheet` is for files containing one observation (on all variables) per line with variables separated by tabs or commas, where the first line may contain the variable names; `infile` with `evarlist` (free format) allows line breaks to occur anywhere and variables to be separated by spaces, commas, or tabs; `infix` is for files with fixed columns format but a single observation can go over several lines; `infile` with a dictionary (fixed format) is the most flexible command since the dictionary can specify exactly what lines and columns contain what information.

Data can be saved as ASCII using `outfile` or `outsheet`. Finally, the `odbc` command can be used to load, write, or view data from Open Data Base Connectivity (ODBC) sources. See `help infiling` or [U]

21 Inputting data for an overview of commands for reading data.

Only one dataset may be loaded at any given time but a dataset may be combined with the currently loaded dataset using the command `merge` or `append` to add observations or variables, respectively; see also Section 1.6.2.

1.4.2 Variables

There are essentially two kinds of variables in Stata: `string` and `numeric`. Each variable can be one of a number of storage types that require different numbers of bytes. The storage types are `byte`, `int`, `long`, `float`, and `double` for numeric variables and `str1` to `str244` for string variables of different lengths. Besides the storage type, variables have associated with them a name, a label, and a format. The name of a variable `y` can be changed to `x` using

```
rename y x
```

The variable label can be defined using

```
label variable x "cost in pounds"
```

and the format of a numeric variable can be set to "general numeric" with two decimal places using

```
format x %7.2g
```

Numeric variables

A missing value in a numeric variable is represented by a period ".," (system missing values), or by a period followed by a letter, such as .a, .b, etc., codes that can be used for distinguishing between different kinds of missing values. Missing values are interpreted as very large positive numbers with $. < \text{.a} < \text{.b}$, etc. Note that this can lead to mistakes in logical expressions; see also Section 1.5.2. Numerical missing value codes (such as "-99") may be converted to missing values (and vice versa) using the command `mvdecode`. For example,

```
mvdecode x, mv(-99)
```

replaces all values of variable `x` equal to -99 by periods and

```
mvencode x, mv(-99)
```

changes the missing values back to -99.

Numeric variables can be used to represent categorical or continuous variables including dates. For categorical variables it is not always easy to remember which numerical code represents which category. Value labels can therefore be defined as follows:

```
label define s 1 married 2 divorced 3 widowed 4 single
label values marital s
```

The categories can also be recoded. For example, the command

```
recode marital 2/3=2 4=3
```

merges categories 2 and 3 into category 2 and changes category 4 to 3.

Dates are defined as the number of days since 1/1/1960 and can be displayed using a date format such as `%d`. For example, listing the variable `time` in `%7.0g` format gives

```
list time
```

time	
1.	16978
2.	200

which is not as easy to interpret as

```
format time %d
list time
```

time	
1.	01jan2001
2.	19jan1960

See `help dfmt` for other date formats.

String variables

String variables are typically used for categorical variables or identifiers and in some cases for dates (e.g., if the file was saved as an ASCII file from SPSS or Excel). In Stata, it is generally advisable to represent these variables by numeric variables, and conversion from string to numeric is straightforward. A categorical string variable (or identifier) can be converted to a numeric variable using the command `encode` which replaces each unique string by a different integer and uses that string as the label for the corresponding integer value. The command `decode` converts the labeled numeric variable back to a string variable. The command `destring` can be used to convert several (or all) variables from string to numeric by interpreting the strings as numbers. This is useful if numeric variables saved from another program are interpreted by Stata as string variables, for instance due to missing values being represented by blanks.

A string variable `string1` representing dates can be converted to numeric using the function `date(string1, string2)` where `string2` is a permutation of "day" to specify the order of the day, month, and year in `string1`. For example, the commands

```
display date("30/1/1930", "dmy")
```

and

```
display date("january 30, 1930", "mdy")
```

both return the negative value -10928 because the date is 10928 days before 1/1/1960.

1.5 Stata commands

Typing `help language` gives the following generic command structure for most Stata commands:

```
[prefix:] command [varlist] [* exp] [if] [in]
[weight] [using filename] [, options]
```

The help file contains links to information on each of the components, and we will briefly describe them here.

[prefix:] could be a number of different things; see `help prefix`. One example is by `varlist`: which instructs Stata to repeat the command for each combination of values in the list of variables `varlist`.

`command` is the name of the command and can often be abbreviated; for example, the command `display` can be abbreviated as `dis`.

{`varlist`} is the list of variables to which the command applies.

[=exp] is an expression.

[if] if `exp` restricts the command to that subset of the observations that satisfies the logical expression `exp`.

[in] in `range` restricts the command to those observations whose indices lie in a particular range `range`.

[weight] allows weights to be associated with observations (see Section 1.7).

[using `filename`] specifies the filename to be used.

[, options] a comma is only needed if `options` are used; options are specific to the command and can often be abbreviated.

For any given command, some of these components may not be available; for example, `list` does not allow [`using filename`]. The help file for a specific command specifies which components are available, using the same notation as above, with square brackets enclosing components that are optional. For example, `help list` gives

```
list [varlist] [if] [in] [, options]
```

implying that [prefix:] and various other components are not allowed and that all permissible components are optional.

The syntax for `varlist`, `exp`, and `range` is described in the next three subsections, followed by information on how to loop through sets of variables or observations.

1.5.1 Varlist

The simplest form of `varlist` is a list of variable names separated by spaces. Variable names may also be abbreviated as long as this is unambiguous, e.g., `x1` may be referred to by `x` only if there is no other variable name starting with `x` such as `x` itself or `x2`. A set of adjacent variables such as `m1, m2`, and `x` may be referred to as `m1-x`. All variables

starting with the same set of letters can be represented by that set of letters followed by a wild card *, so that `m*` may stand for `m1 m6 mother`. The wild card can appear anywhere within the variable name; for instance `my*var` could refer to `mylongvar` and `myshortvar`. The set of all variables is referred to by `.all` or *. Examples of a varlist are

```
x y
xi-x16
al-a3 my* sex age
```

Note that Stata is case sensitive.

Long variable names are abbreviated in Stata's output by replacing a middle section of the name by a ~. This method of abbreviating variables can also be used in a varlist as long as the abbreviation is unambiguous. For instance, `my~var` would not work if there are variables `mylongvar` and `myshortvar` in the dataset, whereas `my~var` would work.

A useful command for finding variables in a large dataset is `lookfor` which searches for a keyword among the variable names and labels.

1.5.2 Expressions

There are logical, algebraic, and string expressions in Stata. Logical expressions evaluate to 1 (true) or 0 (false) and use the operators < and <= for "less than" and "less than or equal to", respectively. Similarly, > and >= are used for "greater than" and "greater than or equal to". The symbols == and != (or ~=) stand for "equal to" and "not equal to", and the characters ! (or ~), &, and | represent "not", "and", and "or", respectively, so that

```
if (y!=2 & z>x) | x==1
```

means "if `y` is not equal to 2 and `z` is greater than `x` or if `x` equals 1". In fact, expressions involving variables are evaluated for each observation so that the expression really means

$$(y_i \neq 2 \& z_i > x_i) \mid x_i == 1$$

where i is the observation index.

Great care must be taken in using the > or >= operators when there are missing data. For example, if we wish to delete all subjects older than 16, the command

```
drop if age>16
```

will also delete all subjects for whom age is missing since a missing value (represented by ".", ".a", ".b", etc.) is interpreted as a very large number. It is always safer to allow for missing values explicitly using for instance

```
drop if age>16 & age<.
```

Note that this is safer than specifying `age!=.`, which would not exclude missing values coded as ".a", ".b", etc.

Algebraic expressions use the usual operators +, -, *, /, and ^ for addition, subtraction, multiplication, division, and powering, respectively. Stata also has many mathematical functions such as `sqrt()`, `exp()`, `log()`, etc. and statistical functions such as `chi2tail()` and `normal()` for cumulative distribution functions and `invnormal()`, etc., for inverse cumulative distribution functions. Pseudo-random numbers with a uniform distribution on the [0,1] interval may be generated using `uniform()`. Examples of algebraic expressions are

```
y + x  
(y + x)^3 + a/b  
invnormal(uniform())+2
```

where `invnormal(uniform())` returns a (different) draw from the standard normal distribution for each observation.

Finally, string expressions mainly use special string functions such as `substr(str,n1,n2)` to extract a substring from `str` starting at `n1` for a length of `n2`. The logical operators == and != are also allowed with string variables and the operator + concatenates two strings. For example, the combined logical and string expression

```
("moon") + (substr("sunlight",4,5)) == "moonalight"
```

returns the value 1 for "true".

For a list and explanation of all functions, use `help functions`.

1.5.3 Observation indices and ranges

Each observation has an index associated with it. For example, the value of the third observation on a particular variable `x` may be referred to as `x[3]`. The macro `_n` takes on the value of the running index and `_N` is equal to the number of observations. We can therefore refer to the previous observation of a variable as `x[_n-1]`.

An indexed variable is only allowed on the right-hand side of an assignment. If we wish to replace `x[3]` by 2, we can do this using the

syntax

```
replace x = 2 if _n==3
```

We can refer to a range of observations either using `if` with a logical expression involving `_n` or, more easily, by using `in range`. The command above can then be replaced by

```
replace x = 2 in 3
```

More generally, `range` can be a range of indices specified using the syntax `f/l` (for "first to last") where `f` and/or `l` may be replaced by numerical values if required, so that `5/12` means "fifth to twelfth" and `f/10` means "first to tenth", etc. Negative numbers are used to count from the end, for example

```
list x in -10/1
```

lists the last 10 observations.

1.5.4 Looping through variables or observations

Explicitly looping through observations is often not necessary because expressions involving variables are automatically evaluated for each observation. It may however be required to repeat a command for subsets of observations and this is what the prefix `by varlist:` is for. Before using `by varlist:`, however, the data must be sorted using

```
sort varlist
```

where `varlist` includes the variables to be used in `by varlist:`. Note that if `varlist` contains more than one variable, ties in the earlier variables are sorted according to the next variable(s). For example,

```
sort school class  
by school class: summarize test
```

give the summary statistics of `test` for each class. If `class` is labeled from 1 to n_i for the i th school, then not using `school` in the above commands would result in the observations for all classes with the same label being grouped together. To avoid having to sort the data first, the `sort` option can be used as follows:

```
by school class, sort: summarize test
```

A very useful feature of `by varlist:` is that it causes the observation index `_n` to count from 1 within each of the groups defined by the

distinct combinations of the values of *varlist*. The macro *N* represents the number of observations in each group. For example,

```
sort group age
by group: list age if _n==_N
```

lists *age* for the last observation in each group where the last observation in this case is the observation with the highest *age* within its group. The same can be achieved in a single command using the *sort* option:

```
by group (age), sort: list age if _n==_N
```

where the variable in parentheses is used to sort the data but does not contribute to the definition of the subgroups of observations to which the *list* command applies.

We can loop through a list of variables or other objects using *foreach*. The simplest syntax is

```
foreach variable in v1 v2 v3 {
    list `variable'
}
```

This syntax uses a *local* macro (see also Section 1.9) *variable* which takes on the (string) values *v1*, then *v2*, and finally *v3* inside the loop. (Local macros can also be defined explicitly using *local variable v1*.) Enclosing the local macro name in `` is equivalent to typing its contents, i.e., ``*variable*'' evaluates to *v1*, then *v2*, and finally *v3* so that each of these variables is listed in turn.

In the first line above we listed each variable explicitly. We can instead use the more general *varlist* syntax by specifying that the list is of type *varlist* as follows:

```
foreach variable of varlist v* {
    list `variable'
}
```

Numeric lists can also be specified using *foreach*. For instance, the command

```
foreach number of numlist 1 2 3 {
    display `number'
}
```

produces the output

```
1
2
3
```

Numeric lists may be abbreviated by "first/last", here 1/3 or "first(increment)last", for instance 1(2)7 for the list 1 3 5 7. See help foreach for other list types.

For numeric lists, a simpler syntax is **forvalues**. To produce the output above, use:

```
forvalues i=1/3 {
    display `i'
}
```

The same output can also be produced using **while** as follows:

```
local i = 1
while i<=3 {
    disp `i'
    local i = `i' + 1
}
```

Here the local macro *i* was defined using **local i = 1** and then incremented by 1 using **local i = `i' + 1**. See also Section 1.11 on programming. Cox (2002b) gives a useful tutorial on by varlist: and Cox (2002a; 2003) discusses **foreach** and **forvalues** in detail.

1.6 Data management

1.6.1 Generating and changing variables

New variables may be generated using the commands **generate** or **egen**. The command **generate** creates a new variable, evaluates an expression for each observation, and places the result into the new variable. For example,

```
generate x = 1
```

creates a new variable called *x* and sets it equal to one. When **generate** is used together with *if exp* or *in range*, the remaining observations are set to missing. For example,

```
generate percent = 100*(old - new)/old if old>0
```

generates the variable *percent* and sets it equal to the percentage decrease from *old* to *new* where *old* is positive and equal to missing otherwise. The command **replace** works in the same way as **generate** except that it allows an existing variable to be changed. For example,

```
replace percent = 0 if old<=0
```

changes the missing values in the variable `percent` to zeros. The two commands above could be replaced by the single command

```
generate percent = cond(old>0, 100*(old-new)/old, 0)
```

where `cond()` evaluates to the second argument if the first argument is true and to the third argument otherwise.

The command `egen` provides extensions to `generate`. One advantage of `egen` is that some of its functions accept a variable list as an argument, whereas the functions for `generate` can only take simple expressions as arguments. For example, we can form the average of 100 variables `m1` to `m100` using

```
egen average = rowmean(m1-m100)
```

where missing values are ignored. Other functions for `egen` operate on groups of observations. For example, if we have the income (variable `income`) for members within families (variable `family`), we may want to compute the total income of each member's family using

```
egen faminc = total(income), by(family)
```

An existing variable can be replaced using `egen` functions only by first deleting it using

```
drop x
```

Another way of dropping variables is using `keep varlist` where `varlist` is the list of all variables not to be dropped.

A very useful command for changing categorical numeric variables is `recode`. For instance, to merge the first three categories and recode the fourth to "2", type

```
recode categ 1/3 = 1 4 = 2
```

If there are any other values, such as missing values, these will remain unchanged. See `help recode` for more information.

1.6.2 Changing the shape of the data

It is frequently necessary to change the shape of data, the most common application being grouped data, in particular repeated measures or panel data. If we have measurement occasions j for subjects i , this may be viewed as a multivariate dataset in which each occasion j is

represented by a variable x_{ij} , and the subject identifier is in the variable `subj`. However, for some statistical analyses we may need one single, long, response vector containing the responses for all occasions for all subjects, as well as two variables `subj` and `occ` to represent the indices i and j , respectively. The two "data shapes" are called wide and long, respectively. We can convert from the wide shape with variables x_{ij} and `subj` given by

```
list
```

	<code>x1</code>	<code>x2</code>	<code>subj</code>
1.	2	3	1
2.	4	5	2

to the long shape with variables `x`, `occ`, and `subj` using the syntax

```
reshape long x, i(subj) j(occ)
```

(note: $j = 1/2$)

Data	wide	->	long
Number of obs.	2	->	4
Number of variables	3	->	3
j variable (2 values)		->	<code>occ</code>
n_{ij} variables:		<code>x1 x2</code>	-> <code>x</code>

The data now look like this:

```
list
```

	<code>subj</code>	<code>occ</code>	<code>x</code>
1.	1	1	3
2.	1	2	3
3.	2	1	4
4.	2	2	5

We can change the data back again using

```
reshape wide x, i(subj) j(occ)
```

For data in the long shape, it may be required to collapse the data so that each group is represented by a single summary measure. For example, for the data above, we may want to create new variables `meanx`, `sdx`, and `num` containing the mean, standard deviation, and the number of nonmissing responses, respectively. This can be achieved using

```
collapse (mean) means=x (sd) std=x (count) num=x, by(subj)
list
```

	subj	means	std	num
1.	1	2.5	.707107	2
2.	2	4.5	.707107	2

Since it is not possible to convert back to the original format in this case, the data may be preserved before running `collapse` and restored again later using the commands `preserve` and `restore`.

Other ways of changing the shape of data include dropping observations using

```
drop in 1/10
```

to drop the first 10 observations or

```
by group (weight), sort: keep if _n==1
```

to drop all but the lightest member of each group. Sometimes it may be necessary to transpose the data, converting variables to observations and vice versa. This may be done and undone using `xpose`.

If each observation represents a number of units (as after `collapse`), it may sometimes be required to replicate each observation by the number of units, `num`, that it represents. This may be done using

```
expand num
```

If there are two datasets, `subj.dta`, containing subject-specific variables, and `occ.dta`, containing occasion-specific variables for the same subjects, then if both files contain the same sorted subject identifier `subj_id` and `subj.dta` is currently loaded, the files may be merged as follows:

```
merge subj_id using occ
```

resulting in the variables from `subj.dta` being expanded as in the `expand` command above and the variables from `occ.dta` being added.

1.7 Estimation

All estimation commands in Stata, for example `regress`, `logistic`, `poisson`, and `glm`, follow the same syntax and share many of the same

options.

The estimation commands also produce the same kind of output and save the same kind of information. The stored information may be processed using the same set of post-estimation commands.

The basic command structure is

```
[prefix:] command depvar [indepvars] [if] [in] [weights], options
```

The response variable is specified by *depvar* and the explanatory variables by *indepvars*. If dummy variables and interactions are required for categorical explanatory variables, using the *xi:* ("interaction expansion") prefix enables special notation to be used by *indepvars*. For example,

```
xi: regress resp i.x
```

creates dummy variables for each value of *x* except the lowest value and includes these dummy variables as predictors in the model.

```
xi: regress resp i.x*y z
```

fits a regression model with the main effects of *x*, *y*, and *z* and their interaction *x*x* where *x* is treated as categorical and *y* and *z* as continuous (see *help xi* for further details).

The syntax for the [*weights*] option is

```
weighttype = name
```

where *weighttype* depends on the reason for weighting the data. If the data are in the form of a table where each observation represents a group containing a total of *freq* observations, using [*fweight=freq*] is equivalent to running the same estimation command on the expanded dataset where each observation has been replicated *freq* times. If the observations have different standard deviations, for example, because they represent averages of different numbers of observations, then *aweights* is used with weights proportional to the reciprocals of the standard deviations. Finally, *pweights* is used for probability weighting where the weights are equal to the inverse probability that each observation was sampled. (Another type of weights, *iweights*, is available for some estimation commands, mainly for use by programmers.)

All the results of an estimation command are stored and can be processed using post-estimation commands. For example, *predict* may be used to compute predicted values or different types of residuals for the observations in the present dataset and the commands *test*, *testparm*,

`lrtest`, and `lincom` for inferences based on previously estimated models. It is easy to find out what post-estimation commands are available for a given estimation command: simply click into "post-estimation" at the top-right of the help file.

The saved results can also be accessed directly using the appropriate names. For example, the regression coefficients are stored in global macros called `_b[tablename]`. To display the regression coefficient of `x`, simply type

```
display _b[x]
```

To access the entire parameter vector, use `e(b)`. Many other results may be accessed using the `e(name)` syntax. See the "Saved Results" section of the entry for the estimation command in the *Stata Reference Manuals* to find out under what names particular results are stored. After estimation, the command

```
ereturn list
```

lists the names and contents of all estimation results accessible via `e(name)`.

Note that "r-class" results produced by commands that are not estimation commands can be accessed using `r(name)`. For example, after `summarize`, the mean can be accessed using `r(mean)`. The command

```
return list
```

list the names and contents of all "r-class" results currently available.

1.8 Graphics

The graphical user interface (GUI) makes it extremely easy to produce a very attractive graph with different line-styles, legends, etc. To demonstrate this, we first simulate some data as follows:

```
clear
set obs 100
set seed 13211
gen x = invnormal(uniform())
gen y = 2 + 3*x + invnormal(uniform())
```

To produce a scatterplot of `y` versus `x` via the GUI, select **Two-way graph (scatterplot, line, etc.)** from the **Graphics** menu and click into **Scatter** under **Plot type** to obtain the dialog box shown in Figure 1.4. Specify `x` and `y` in the boxes labeled **X variable:** and **Y**

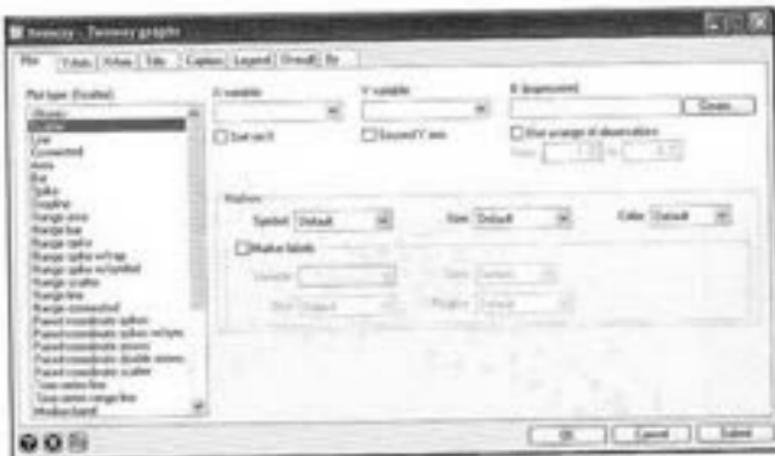


Figure 1.4: Dialog box for twoway graph.

variable: This can be done either by typing or by pressing the little down arrow \downarrow to select among the variables in the dataset. To add a label to the x -axis, click into the tab labeled **X-Axis** and type "Simulated x" in the **Title** box. Similarly, type "Simulated y" in the **Title** box in the **Y-Axis** tab. Finally, click **OK** to produce the graph shown in Figure 1.5. To change for instance the symbol we would have to plot the graph again, this time selecting a different option in the box labeled **Symbol** under the heading **Markers** in the dialog box (it is not possible to edit a graph). The following command appears in the output:

```
twoway (scatter y x), ytitle(Simulated y) xtitle(Simulated x)
```

The command **twoway**, short for **graph twoway**, can be used to plot scatterplots, lines or curves, and many other plots requiring an x and y -axis. Here the **plottype** is **scatter** which requires a y and x variable to be specified. Details such as axis labels are given after the comma. Help on scatterplots can be found (either in the manual or using **help**) under "graph twoway scatter". Help on options for **graph twoway** can be found under "twoway_options".

To produce several graphs, each displayed in a separate window, the graphs must be given different names. In the GUI this can be achieved by clicking into the **Overall** tab of the graph dialog box and typing a name into the box labeled **Name of graph**: If a graph of the same name already exists, use a different name to open a new window or click into the **Replace** box to replace the current graph. In the

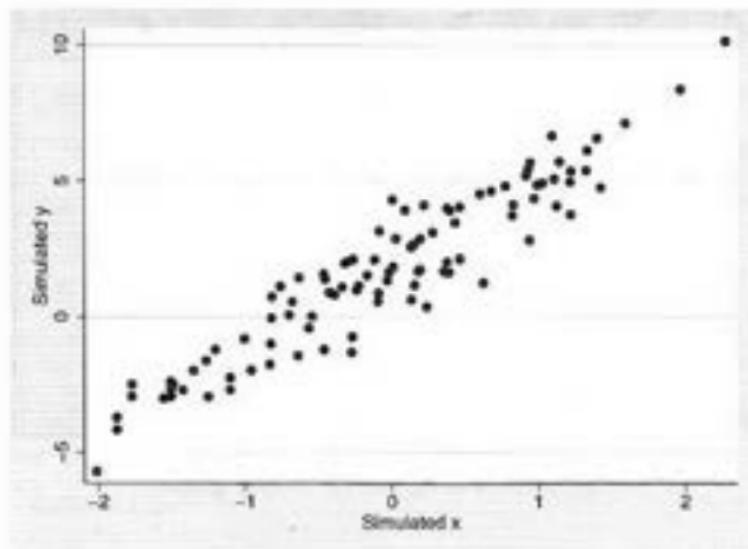


Figure 1.5: Scatterplot of simulated data.

graph command, names can be assigned using the option `name(name of graph)` together with the `replace` option if required. To view a particular graph that may be hidden behind other windows, use either the **Window** menu or click into the arrow in .

We can use a single `graph twoway` command to produce a scatterplot with a regression line superimposed:

```
twoway (scatter y x) (lfit y x), ///
    ytitle(Simulated y) xtitle(Simulated x) ///
    legend(order(1 "Observed" 2 "Fitted"))
```

giving the graph in Figure 1.6. Inside each pair of parentheses is a command specifying a plot to be added to the same graph. The options applying to the graph as a whole appear after these individual plots preceded by a comma as usual. Here the `legend()` option was used to specify labels for the legend; see the manual or help for "legend.option".

Each plot can have its own `if exp` or `in range` restrictions as well as various options. To demonstrate this, we first create a new variable, `group`, taking the values 1 and 2 and add 2 to the `y`-values of group 2:

```
gen group = cond(_n < 50, 1, 2)
replace y = y + 2 if group==2
```

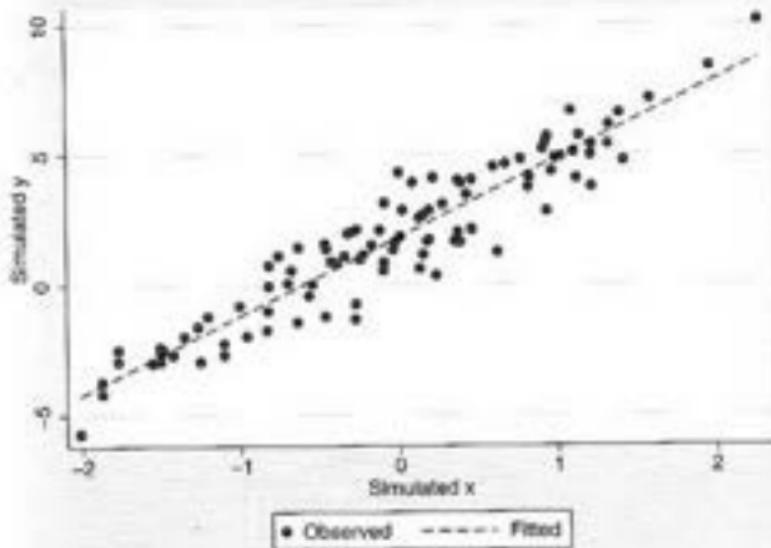


Figure 1.6: Scatterplot and fitted regression line.

Now produce a scatterplot with different symbols for the two groups and separate regression lines using

```
twoway (scatter y x if group==1, msymbol(0))      ///
    (lfit y x if group==1, lpatt(solid))           ///
    (scatter y x if group==2, msymbol(Oh))          ///
    (lfit y x if group==2, lpatt(dash)),            ///
    ytitle(Simulated y) xtitle(Simulated x)        ///
    legend(order(1 2 "Group 1" 3 4 "Group 2"))
```

giving the graph shown in Figure 1.7. Here, the options `msymbol(0)` and `msymbol(Oh)` produce solid and hollow circles, respectively, whereas `lpatt(solid)` and `lpatt(dash)` produce solid and dashed lines, respectively. These options are inside the parentheses for the corresponding plots. The options referring to the graph as a whole, `xtitle()`, `ytitle()`, and `legend()`, appear after the individual plots have been specified. Just before the final comma, we could also specify `if exp` or `in range` restrictions for the graph as a whole.

Some people find it more convenient to separate plots using `||` instead of enclosing them in parentheses, for instance replacing the first two lines of the command above by

```
twoway scatter y x if group==1, ms(0) ||     ///
```

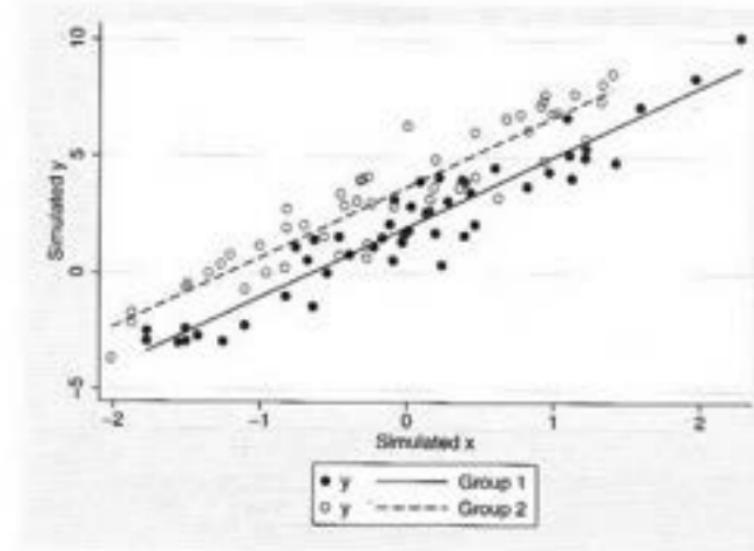


Figure 1.7: Scatterplot and fitted regression line.

```
lfit y x if group==1, clpatt(solid)
```

The `by()` option can be used to produce separate plots (each with their own sets of axes) in the same graph. For instance

```
label define gr 1 "Group 1" 2 "Group 2"
label values group gr
twoway scatter y x, by(group)
```

produces the graph in Figure 1.8. Here the value labels of `group` are used to label the individual panels.

Other useful graphics commands include `graph twoway function` for plotting a function without having to define any new variables, `graph matrix` for scatterplot matrices, `graph box` for boxplots, `graph bar` for bar charts, `histogram` for histograms, `kdensity` for kernel density plots, and `qnorm` for Q-Q plots.

For `graph box` and `graph bar`, we may wish to plot different variables, referred to as `years` in Stata, for different subgroups or categories of individuals, specified using the `over()` option. For example,

```
replace x = x + 1
graph bar y x, over(group)
```

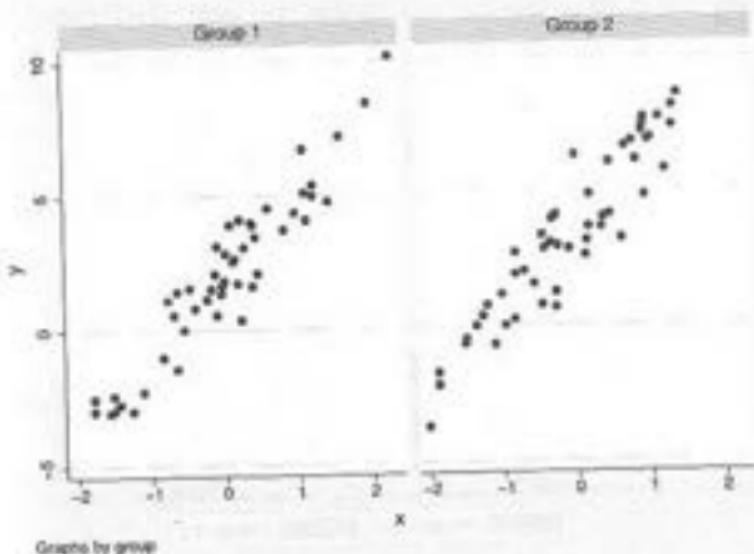


Figure 1.8: Separate scatterplot produced using `by()`.

results in the bar chart in Figure 1.9. For more information on changing the labeling and presentation of the bars, see `yvar_options` and `group_options` in [G] `graph bar`.

The general appearance of graphs is defined in `schemes`. In this book we use scheme `sj` (Stata Journal) by issuing the command

```
set scheme sj
```

at the beginning of each Stata session. See [G] `schemes intro` or `help schemes` for a complete list and description of schemes available.

Graphs can be saved in Stata's .gph format and read back in using

```
graph save mygraph  
graph using mygraph
```

Graphs can also be exported as encapsulated postscript or PNG files using, respectively,

```
graph export mygraph.eps  
graph export mygraph.png
```

See `help graph export` for a list of all available storage types.

We find the GUI interface particularly useful for learning about these and other graphics commands and their options. Mitchell (2004)

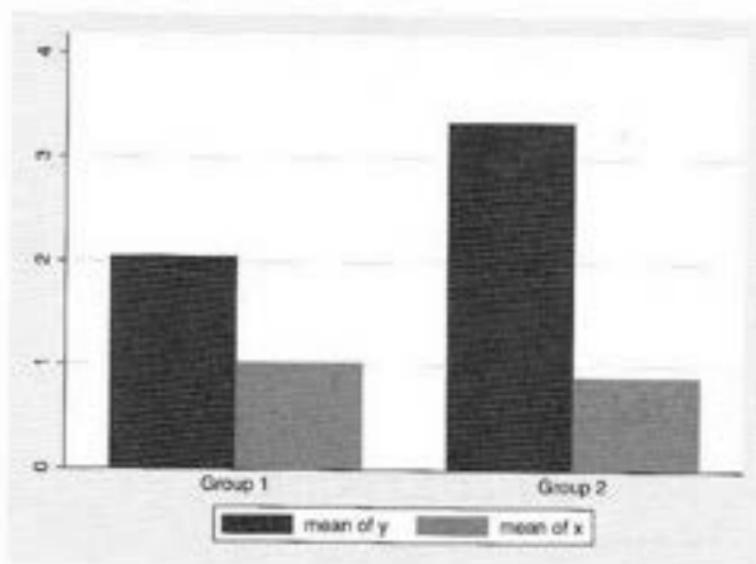


Figure 1.9: Bar chart.

is a useful reference book which contains a large collection of graphs and the commands used to create them.

1.9 Stata as a calculator

Stata can be used as a simple calculator using the command `display` followed by an expression, e.g.,

```
display sqrt(5*(11-3^2))
```

```
3.1622777
```

There are also a number of statistical commands that can be used without reference to any variables. These commands end in `i`, where `i` stands for `immediate` command. For example, we can calculate the sample size required for an independent samples *t*-test to achieve 80% power to detect a difference at the 1% level (2-sided) if the population means are 1 and 2 and the within population standard deviation is 1 using `sampsiz` as follows:

```
sampsiz 1 2, sd(1) power(.8) alpha(0.01)
```

(see Display 1.1). Similarly, `ttesti` can be used to carry out a *t*-test

```
Estimated sample size for two-sample comparison of means
Test Ho: m1 = m2, where m1 is the mean in population 1
and m2 is the mean in population 2
Assumptions:
alpha = 0.0100 (two-sided)
power = 0.8000
m1 = 1
m2 = 2
sd1 = 1
sd2 = 1
n2/n1 = 1.00
Estimated required sample sizes:
n1 = 24
n2 = 24
```

Display 1.1

if the means, standard deviations, and sample sizes are given.

As briefly shown in Section 1.5.4, results can be saved in local macros using the syntax

```
local a = exp
```

The result may then be used again by enclosing the local macro name in single quotes `' (using two different keys on the keyboard). For example,

```
local a = 5
display sqrt(`a')
```

```
2.236068
```

Matrices can also be defined and matrix algebra carried out interactively. The following matrix commands define a matrix `a`, display it, and give its trace and its eigenvalues:

```
matrix a = (1,2\ 2,4)
matrix list a

symmetric a[2,2]
 c1  c2
r1  1
r2  2  4

display trace(a)
```

5

```

matrix symeigen x v = a
matrix list v

v[1,2]
    e1   e2
r1   6   0

```

For more powerful matrix calculations, use the Mata programming language as described in the next section.

1.10 Matrix calculations using Mata

Mata is a fast matrix language resembling C. Here we show how Mata can be used for performing calculations interactively; Mata can also be used to write programs that are automatically compiled by Stata and hence run faster than programs written in Stata's usual programming language. See [M] *Mata Reference Manual* and **help mata** for full details.

To enter the Mata environment, type **mata** in the Command window and press *Return*. To exit Mata, type **end** followed by *Return*. You can tell if you are within Mata because the prompt changes from a period “.” to a colon “:”. An example of a Mata session, as it would appear in the output, is given in Display 1.2. We see that typing an expression displays the result. Variables can be defined using the = operator and subsequently used in expressions. Here we define a 2×2 matrix **A** using commas to separate columns and \ to separate rows as in Stata's matrix commands.

The variables defined in this session will not be cleared when we exit Mata; they will still be there and can be listed using **mata describe** if we re-enter Mata:

mata		
mata describe		
# bytes	type	name and extent
32	real matrix	4[2,2]
8	real scalar	x

We see that associated with each object is a *type* defined by two aspects which Stata calls *stotype*, here **real** and *ortotype*, here **matrix** and **scalar**. To clear Mata without disturbing Stata, use the command

```
: mata
: 2 + 3
5
: x = 2 + 3
: x
5
: A = (1, 2 \ 3, 4)
: A
1 2
1 1 2
2 3 4
: A + A
1 2 3 4
1 1 2 1 2
2 3 4 3 4
: A \ A
1 2
1 1 2
2 3 4
3 1 2
4 3 4
: end
```

Display 1.2

```
mata clear
```

In Mata, we can evaluate complex matrix expressions using very intuitive notation. For instance, the least squares estimator $(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y}$ is evaluated using

```
b = invsym(X'X)*X'y
```

where `invsym()` stands for "inverse of a symmetric matrix". Here `*` performs matrix multiplication. To perform element-by-element multiplication, use the colon operator:

```
. mata
: A
: 1 2
: 1 1 2
: 2 3 4
: A*A
: 1 2
: 1 7 10
: 2 15 22
: A:A
: 1 2
: 1 4
: 2 9 16
: end
```

1.11 Brief introduction to programming

So far we have described commands as if they would be run interactively. However, in practice, it is always useful to be able to repeat the entire analysis using a single command. This is important, for example, when a data entry error is detected after most of the analysis has already been carried out. It is also important to keep a record of all data manipulation so that it can be checked and corrected later. In Stata, a set of commands stored as a do-file, called for example, `analysis.do`, can be executed using the command

```
do analysis
```

We strongly recommend that readers create do-files for any work in Stata, e.g., for the exercises of this book.

One way of generating a do-file is to carry out the analysis interactively and save the commands, for example, by right-mouse clicking into the Review window and clicking into **Save Review Contents...**. Stata's **Do-file Editor**, which is opened by clicking into , can also be used to create or edit a do-file. One way of trying out commands interactively and building up a do-file is to run commands in the Commands window and copy them into the **Do-file Editor** after checking that they work. Another possibility is to type commands into the **Do-file Editor** and try them out individually by highlighting the commands and clicking into  or selecting **Tools → Do Selection**. Alternatively, any text editor may be used to create a do-file. The following is a useful template for a do-file:

```
/* comment describing what the file does */
version 9.2
capture log close
log using filename, replace
set more off

command 1
command 2
etc.

log close
exit
```

We will explain each line in turn.

1. The "brackets" /* and */ cause Stata to ignore everything between them. Another way of commenting out lines of text is to start the lines either with an asterisk * or with a double forward slash //.
2. The command **version 9.2** causes Stata to interpret all commands as if we were running Stata version 9.2 even if, in the future, we have actually installed a later version in which some of these commands do not work any more.
3. The **capture** prefix causes the do-file to continue running even if the command results in an error. The **capture log close** command therefore closes the current log file if one is open or returns an error message. (Another useful prefix is **quietly** which suppresses all output, except error messages.)
4. The command **log using filename, replace** opens a log file, replacing any file of the same name if it already exists.

5. The command `set more off` causes all the output to scroll past automatically instead of waiting for the user to scroll through it manually. This is useful if the user intends to look at the log file for the output.
6. After the analysis is complete, the log file is closed using `log close`.
7. The last statement, `exit`, is not necessary at the end of a do-file but may be used to make Stata stop running the do-file wherever it is placed.

Variables, global macros, local macros, and matrices can be used for storing and referring to data and these are used extensively in programs. For example, we may wish to subtract the mean of `x` from `x`. Interactively, we would use

```
summarize x
```

to find out what the mean is and then subtract that value from `x`. However, we should not type the value of the mean into a do-file because the result would no longer be valid if the data change. Instead, we can access the mean computed by `summarize` using `r(mean)`:

```
quietly summarize x, meanonly
gen xnew = x-r(mean)
```

(If all that is required from `summarize` is the mean or the sum, it is more efficient to use the `meanonly` option.) Most Stata commands are `r`-class, meaning that they store results that may be accessed using `r()` with the appropriate name inside the brackets. Estimation commands store the results in `e()`. To find out under what names results are stored, see the "Stored Results" section for the command of interest in the *Stata Reference Manuals*. Alternatively, execute the command and then issue the command `return list` for a list of all results stored in `r()` or `ereturn list` for a list of all results stored in `e()`.

If a local macro is defined without using the `=` sign, anything can appear on the right-hand side and typing the local macro name in single quotes has the same effect as typing whatever appeared on the right-hand side in the definition of the macro. For example, if we have a variable `y`, we can use the commands

```
local a y
display "'a'[1] = " `a'[1]
y[1] = 4.616998
```

Local macros are only "visible" within the do-file or program in which they are defined. Global macros may be defined using

```
global a = 1
```

and accessed by prefixing them with a dollar sign, for example,

```
gen b = $a
```

Sometimes it is useful to have a general set of commands (or a program) that may be applied in different situations. It is then essential that variable names and parameters specific to the application can be passed to the program. If the commands are stored in a do-file, the "arguments" with which the do-file will be used are referred to as '1', '2' etc. inside the do-file. For example, a do-file *filename.do* containing the command

```
list `1' `2'
```

may be run using

```
do filename x1 x2
```

to cause *x1* and *x2* to be listed. Alternatively, we can define a program which can be called without using the *do* command in much the same way as Stata's own commands. This is done by enclosing the set of commands by

```
program programme  
end
```

After running the program definition, we can run the program by typing the program name and arguments.

Most programs require things to be done repeatedly by looping through some list of objects. This can be achieved using *foreach* and *forvalues*. For example, we define a program called *mylist* that lists the first three observations of each variable in a variable list:

```
program mylist  
version 9.2  
syntax varlist  
foreach var in `varlist' { /* outer loop: variables */  
    display "`var'"  
    forvalues i=1/3 { /* inner loop: observations */  
        display "`var'`[`i']"  
    }  
    display ""  
}  
end
```

We can run the program using the command

```
mylist x y z
```

Here the `syntax` command defines the syntax to be

```
mylist varlist
```

(no options allowed), issues an error message if `varlist` is not valid, for example if any of the variables do not exist, and places the variable names into the local macro `varlist` (see `help syntax` and [P] `syntax`). The outer `foreach` loop repeats everything within the outer braces for each variable in `varlist`. Within this loop, the "current" variable is placed in the local macro `var`. For each variable, the inner `forvalues` loop repeats the `display` command for `i` equal to 1, 2, and 3.

A program may be defined by typing it into the Commands window. However, this is almost never done in practice. A more useful method is to define the program within a do-file where it can easily be edited. Note that once the program has been loaded into memory (by running the `program` command), it has to be cleared from memory using `program drop` before it can be redefined. It is therefore useful to have the command

```
capture program drop mylist
```

in the do-file before the `program` command, where `capture` ensures that the do-file continues running even if `mylist` does not yet exist.

A program may also be saved in a separate file (containing only the program definition) of the same name as the program itself and having the extension `.ado`. If the ado-file (automatic do-file) is in a directory in which Stata looks for ado-files, for example the current directory, it can be executed simply by typing the name of the file. There is no need to load the program first (by running the program definition). To find out where Stata looks for ado-files, type

```
adopath
```

This lists various directories including `\ado\personal\`, the directory where personal ado-files may be stored. Many of Stata's own commands are actually ado-files stored in the `ado` subdirectory of the directory where the Stata executable (e.g., `wstata.exe`) is located.

The [P] *Programming Reference Manual* gives detailed information on the programming commands mentioned here and many more. Type `help dialog programming` for information on programming your own dialogs. The Stata Plugin Interface (SPI) which allows compiled C programs to be called from a Stata program is described in detail at

<http://www.stata.com/plugins>. While this is useful if the C-program already exists, it will often be easier to write functions in Mata than in C.

Chapter 13 of this book gives some examples of maximizing your own likelihood using the `m1` command, and this is discussed in detail in Gould et al. (2006).

1.12 Keeping Stata up to date

StataCorp continually updates the current version of Stata. If the computer is connected to the Internet, Stata can be updated by issuing the command

```
update all
```

Ado-files are then downloaded and stored in the correct directory. If the executable has changed since the last update, a new executable (e.g., `wstata.bin`) is also downloaded. This file should be used to overwrite the old executable (e.g., `wstata.exe`) after saving the latter under a new name (e.g., `wstata.old`). A quick and easy way of achieving this is to issue the command `update swap` within Stata. The command `help whatnew` lists all the changes since the release of the present version of Stata.

In addition to Stata's official updates to the package, users are continuously creating and updating their own commands and making them available to the Stata community. Articles on user-written programs are published in a peer-reviewed journal called *The Stata Journal* (SJ) which replaced the *Stata Technical Bulletin* (STB) at the end of 2001 and is indexed in the Science Citation Index. These and other user-written programs can be downloaded by clicking into **Help → SJ & User-written Programs**, and selecting one of a number of sites including sites for the SJ and STB. A large repository for user-written Stata programs is the *Statistical Software Components* (SSC) archive at <http://ideas.repec.org/s/boc/bocode.html> maintained by Kit Baum (the archive is part of IDEAS which uses the RePEc database). These programs can be downloaded using the `ssc` command. To find out about commands for a particular problem (user-written or part of Stata), use the `findit` command. For example, running

```
findit meta
```

brings up the Stata Viewer with a long list of entries including one on STB-42:

```
STB-42 sbe16.1 . . . . . New syntax and output for the meta-analysis command
```

`help meta if installed) S. Sharp and J. Sterne
3/08 pp.6--8; STB Reprints Vol 7, pp.106--108`

which reveals that STB-42 has a directory in it called `sbe16.1` containing files for "New syntax and output for the meta-analysis command" and that help on the new command may be found using `help meta`, but only after the program has been installed. The authors are S. Sharp and J. Sterne. The command can be installed by clicking into the corresponding hyperlink in the **Stata Viewer** (or going through **Help → SJ & User-written Programs**, clicking on **STB**, then **stb42**, then **sbe16.1**) and selecting (**click here to install**). The program can also be installed using the commands

```
net stb 42 sbe16_1
```

(`ssc help net`). Note that `findit` first lists programs that have been published in the SJ and STB, followed by programs from other sites such as the SSC. This order often does not reflect the (reverse) chronological order of versions of a given program since the SSC usually has the most up-to-date version (look for `rePEC` in the URL). The most reliable way of installing a program from the SSC is using the command (here illustrated for the program `gllamm`)

```
ssc install gllamm
```

See `help ssc`.

1.13 Exercises

1.1 • Some data manipulation

1. Use a text editor (e.g., Notepad, PFE, or the Stata Do-file Editor) to generate the dataset `test.dat` given below, where the columns are separated by tabs (make sure to save it as a text only or ASCII file).

v1	v2	v3
1	8	5
2	16	3
5	12	2

2. Read the data into Stata using `insheet` (`ssc help insheet`).
3. Click into the Data Editor and type in the variable `sex` with values 1, 2, and 1.
4. Define value labels for `sex` (1=male, 2=female).

5. Use `generate` to generate `id`, a subject index (from 1 to 3).
6. Use `rename` to rename the variables `v1` to `v3` to `time1` to `time3`. Also try doing this in a loop using `forvalues`.
7. Use `reshape` to convert the dataset to long shape.
8. Generate a variable `d` that is equal to the squared difference between the variable `time` at each occasion and the average of `time` for each subject.
9. Drop the observation corresponding to the third occasion for `id=2`.

1.2 Wage increases

Here we use panel data for 545 American young males taken from the National Longitudinal Survey (Youth Sample) for the period 1980-1987. The data have been analyzed by Vella and Verbeek (1998) and Wooldridge (2002). The subset of variables in `wagepan.dta` considered here are:

- `year`: calendar year 1980 to 1987
 - `lwage`: natural log of hourly wage in US \$
 - `black`: dummy variable for being black
 - `hispanic`: dummy variable for being Hispanic
1. Create a new variable equal to the exponential of `lwage`.
 2. Collapse the data to obtain the mean hourly wages by year and ethnic/racial group (black, Hispanic, other).
 3. Produce a line graph (using `twoway line`) showing the mean wages over time, separately for the ethnic/racial groups.
 4. Improve the graph by defining labels, line patterns, legends, etc., using the GUI if preferred.

1.3 Finding information

Without consulting the manuals, use Stata's help facilities and/or GUI to find out the following:

1. The name and syntax of the Stata function that calculates the inverse cumulative F distribution
2. The name of the Stata command to produce a LOWESS curve
3. The option for the `regress` command that will cause the intercept to be omitted
4. How to get adjusted means for a given regression model (Hint: this is a post-estimation problem)
5. How to plot a histogram with a dashed normal density curve superimposed

Chapter 2

Data Description and Simple Inference: Female Psychiatric Patients

2.1 Description of data

The data to be used in this chapter consist of observations on 8 variables for 118 female psychiatric patients and are available in Hand *et al.* (1994). The variables are as follows:

- age: age in years
- iq: intelligence score
- anxiety: anxiety (1=nose, 2=mild, 3=moderate, 4=severe)
- depress: depression (1=nose, 2=mild, 3=moderate, 4=severe)
- sleep: can you sleep normally? (1=yes, 2=no)
- sex: have you lost interest in sex? (1=no, 2=yes)
- life: have you thought recently about ending your life? (1=no, 2=yes)
- weight: increase in weight over last six months (in lbs)

The data are given in Table 2.1; missing values are coded as -99. There are a variety of questions that might be addressed by these data; for example, do women who have recently contemplated suicide differ in any respects from those who have not? Also of interest are the correlations between anxiety and depression and between weight change, age,

and IQ. It should be noted, however, that any associations found from cross-sectional observational data like these are at best suggestive of causal relationships.

Table 2.1 Data in fem.dat

id	age	IQ	anx	depress	sleep	sex	life	weight
1	39	94	2	2	2	2	2	4.9
2	41	89	2	2	2	2	2	2.2
3	42	83	3	3	3	2	2	4.0
4	30	99	2	2	2	2	1	-0.3
5	35	94	2	1	1	2	1	0.9
6	44	90	-99	1	2	1	2	-1.5
7	31	94	2	2	-99	2	1	3.5
8	39	87	3	2	2	2	2	-1.2
9	35	-99	3	2	2	2	2	0.8
10	33	92	2	2	2	1	1	-1.9
11	38	92	2	1	1	-99	1	5.5
12	31	94	2	2	2	2	1	2.7
13	40	91	3	2	2	2	2	4.4
14	44	86	2	2	2	2	2	3.2
15	43	90	3	2	1	2	1	-1.5
16	32	-99	1	2	2	-99	1	-1.9
17	32	91	1	3	2	2	2	8.3
18	43	82	4	2	2	2	2	3.6
19	46	86	3	2	2	2	1	1.4
20	30	88	2	2	-99	2	2	-99.0
21	34	97	3	2	2	2	1	-99.0
22	37	96	3	1	2	2	1	-1.0
23	35	95	2	2	2	2	2	6.5
24	45	87	2	2	2	2	1	-2.1
25	36	103	2	2	2	2	1	-0.4
26	31	-99	2	2	2	2	1	-1.9
27	32	91	2	2	2	2	2	3.7
28	44	87	2	3	2	2	2	4.5
29	40	94	3	3	-99	2	2	4.2
30	42	89	3	3	2	2	2	-99.0
31	36	92	3	-99	2	2	2	1.7
32	42	84	3	3	2	2	2	4.8
33	46	94	2	-99	2	2	1	1.7
34	41	92	2	1	2	2	2	-3.0
35	30	96	-99	2	2	2	1	0.8
36	39	96	2	2	2	2	2	1.5
37	40	86	2	3	2	2	1	1.3
38	42	92	3	2	2	2	2	3.0
39	35	102	2	2	2	2	1	1.0
40	31	82	2	2	2	2	2	1.5
41	33	92	3	3	2	2	2	3.4
42	43	90	-99	-99	2	2	1	-99.0
43	37	92	2	1	1	1	1	-99.0

Table 2.1 Data in fem.dat (continued)

44	32	88	4	2	2	2	1	-99.0
45	34	98	2	2	2	2	2	0.6
46	34	93	3	2	2	2	1	3.3
47	42	90	2	1	1	1	1	4.8
48	41	91	2	1	2	2	1	-2.2
49	31	-99	3	1	2	2	2	1.0
50	32	92	3	2	2	1	2	-1.2
51	29	92	2	2	2	2	2	4.0
52	41	91	2	2	2	2	2	5.9
53	39	91	2	2	2	2	1	0.2
54	41	86	2	1	1	2	1	3.5
55	34	95	2	1	1	1	1	2.9
56	39	91	1	1	2	1	1	-0.6
57	35	96	3	2	2	1	1	-0.6
58	31	100	2	2	2	2	2	-2.5
59	32	99	4	3	2	2	2	3.2
60	41	89	2	1	2	1	1	3.2
61	41	89	3	2	2	2	2	2.1
62	44	98	3	2	2	2	2	3.8
63	35	98	2	2	2	2	1	-2.4
64	41	103	2	2	2	2	2	-0.8
65	41	91	3	1	2	2	1	5.8
66	42	91	4	3	-99	-99	2	2.5
67	33	94	2	2	2	2	1	-1.8
68	41	91	2	1	2	2	1	4.3
69	43	85	2	2	2	1	1	-99.0
70	37	92	1	1	2	2	1	1.0
71	36	96	3	3	2	2	2	3.5
72	44	90	2	-99	2	2	2	3.3
73	42	87	2	2	2	1	2	-0.7
74	31	95	2	3	2	2	2	-1.6
75	29	95	3	3	2	2	2	-0.2
76	32	87	1	1	2	2	1	-3.7
77	35	95	2	2	2	2	2	3.8
78	43	88	1	1	1	2	1	-1.0
79	32	94	2	2	2	2	1	4.7
80	39	-99	3	-99	2	2	1	-99.0
81	34	-99	3	3	2	2	1	2.2
82	34	87	3	3	2	1	1	5.0
83	42	92	1	1	2	2	2	0.4
84	43	86	2	3	2	2	2	-1.2
85	31	93	-99	2	2	2	1	-1.1
86	31	92	2	2	2	1	2	-1.0
87	36	100	2	2	2	2	2	4.2
88	37	93	2	2	2	2	1	2.4
89	43	95	2	2	2	2	2	4.9
90	32	95	3	2	2	2	2	3.0
91	32	92	-99	-99	-99	2	2	3.0

Table 2.1 Data in fem.dat (continued)

92	32	98	2	2	2	2	2	-0.3
93	43	92	2	2	2	2	2	1.2
94	41	88	2	2	2	2	1	2.6
95	43	85	1	1	2	2	1	1.9
96	39	92	2	2	2	2	1	3.5
97	41	84	2	2	2	2	2	-0.6
98	41	92	2	1	2	2	1	1.4
99	32	91	2	2	2	2	2	5.7
100	44	86	3	2	2	2	2	4.6
101	42	92	3	2	2	2	1	-99.0
102	39	89	2	2	2	2	1	2.0
103	45	-99	2	2	2	2	2	0.6
104	39	96	3	-99	2	2	2	-99.0
105	31	97	2	-99	-99	-99	2	2.8
106	34	92	3	2	2	2	2	-2.1
107	41	92	2	2	2	2	2	-2.5
108	33	98	3	2	2	2	2	5.7
109	34	91	2	1	1	2	1	2.4
110	42	91	3	3	2	2	1	1.5
111	40	89	3	1	1	1	1	1.7
112	35	94	3	3	2	2	2	2.5
113	41	90	3	2	2	2	1	-99.0
114	32	96	2	1	1	2	1	-99.0
115	39	87	2	2	2	1	2	-1.0
116	41	86	3	2	1	1	1	6.5
117	33	89	1	1	1	1	2	4.9
118	42	-99	3	2	2	2	2	

2.2 Group comparison and correlations

The data in Table 2.1 contain a number of interval scale or continuous variables (weight change, age, and IQ), ordinal variables (anxiety and depression), and dichotomous variables (sex and sleep) that we wish to compare between two groups of women: those who have thought about ending their lives and those who have not.

For interval scale variables, the most common statistical test is the *t*-test which assumes that the observations in the two groups are independent and are sampled from two populations each having a normal distribution and equal variances. A nonparametric alternative (which does not rely on the latter two assumptions) is the *Mann-Whitney U*-test.

For ordinal variables, either the *Mann-Whitney U*-test or a chi-squared test may be appropriate depending on the number of levels of the ordinal variable. The latter test can also be used to compare

dichotomous variables between the groups.

Continuous variables can be correlated using the *Pearson correlation*. If we are interested in the question whether the correlations differ significantly from zero, then a hypothesis test is available that assumes bivariate normality. A significance test not making this distributional assumption is also available; it is based on the correlation of the ranked variables, the *Spearman rank correlation*. Finally, if variables have only few categories, *Kendall's tau-b* provides a useful measure of correlation (see, e.g., Sprent and Smeeton, 2001). More details of these tests and correlation coefficients can be found in Altman (1999) and Agresti (2002).

2.3 Analysis using Stata

Assuming the data have been saved from a spreadsheet or statistical package (for example SAS or SPSS) as a tab-delimited ASCII file, *fem.dat*, they can be read using the instruction

```
inset using fem.dat, clear
```

There are missing values which have been coded as -99. We replace these with Stata's missing value code ".," using

```
mvdecode _all, mv(-99)
```

The variable *sleep* has been entered incorrectly as "3" for subject 3. Such data entry errors can be detected using the command

```
codebook
```

which displays information on all variables; the output for *sleep* is shown below:

<i>sleep</i>	SLEEP	
	type: numeric (byte)	
	range: (1,3)	units: 1
unique values:	3	missing ..: 5/118
tabulation:	Freq. Value	
	14 1	
	98 2	
	1 3	
	5 .	

Alternatively, we can detect errors using the *assert* command. For *sleep*, we would type

```
assert sleep==1|sleep==2|sleep==".
```

```
i contradiction in 118 observations
assertion is false
```

Since we do not know what the correct code for `sleep` should have been, we can replace the incorrect value of 3 by "missing"

```
replace sleep=. if sleep==3
```

In order to have consistent coding for "yes" and "no", we recode the variable `sleep`

```
recode sleep 1=2 2=1
```

and to avoid confusion in the future, we label the values as follows:

```
label define yn 1 no 2 yes
label values sex yn
label values life yn
label values sleep yn
```

The last three commands could also have been carried out in a `foreach` loop:

```
foreach x in sex life sleep {
    label values `x' yn
}
```

First, we could compare the suicidal women with the non-suicidal by simply tabulating suitable summary statistics in each group. To obtain means and standard deviations for the continuous variables, we can use the `tabstat` command:

```
tabstat weight age iq, by(life) statistics(mean sd)
Summary statistics: mean, sd
by categories of: life (LIFE)
```

life	weight	age	iq
no	1.408889 2.609234	36.94231 4.295065	91.27083 3.757203
yes	1.731148 2.826629	37.92308 5.078471	92.09836 5.02223
Total	1.59434 2.727805	37.48718 4.751797	91.73394 4.508815

A more formal approach to comparing the two groups on say weight gain over the last six months might involve an independent samples *t*-test. First, however, we need to check whether the assumptions needed for the *t*-test appear to be satisfied for weight gain. One way this can be done is by plotting the variable `weight` as a boxplot for each group:

```
graph box weight, by(life) box(1, bfcolor(mono))    ///
  box(2, bfcolor(mono)) yline(0) medtyps(lime)    ///
  ytitle(weight change in last six months)
```

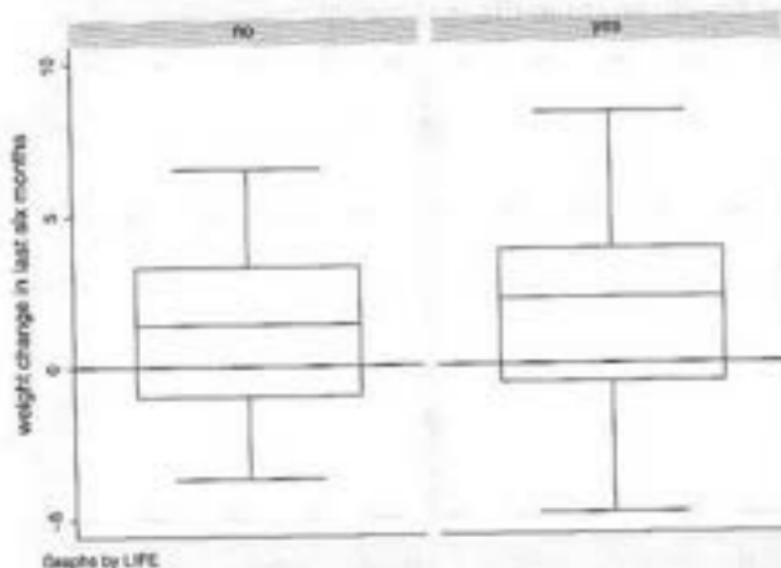


Figure 2.1: Boxplot of weight by group.

giving the graph shown in Figure 2.1. The `yline(0)` option has placed a horizontal line at 0. (Note that in the instructions above, the forward slashes `///` were used to make Stata ignore the line breaks in the middle of the `graph box` command in a do-file, but this should not be used in the Stata Command window, where commands can wrap over multiple lines.) The groups do not seem to differ much in their median weight change and the assumptions for the *t*-test seem reasonable because the distributions are symmetric with similar spread (box heights represent interquartile ranges).

We can also check the assumption of normality more formally by plotting a normal quantile plot of suitably defined residuals. Here the difference between the observed weight changes and the group-specific mean weight changes can be used. If the normality assumption is satisfied, the quantiles of the residuals should be linearly related to the quantiles of the normal distribution with the same mean and standard

deviation. The residuals can be computed and plotted using

```
egen res=mean(weight), by(life)
replace res=weight-res
qnorm res, title("Normal Q-Q plot") saving(qnorm,replace) ///
ytitle(residuals for weight change)
```

The points in the Q-Q plot in Figure 2.2 appear to be sufficiently close to the straight line to justify the normality assumption.

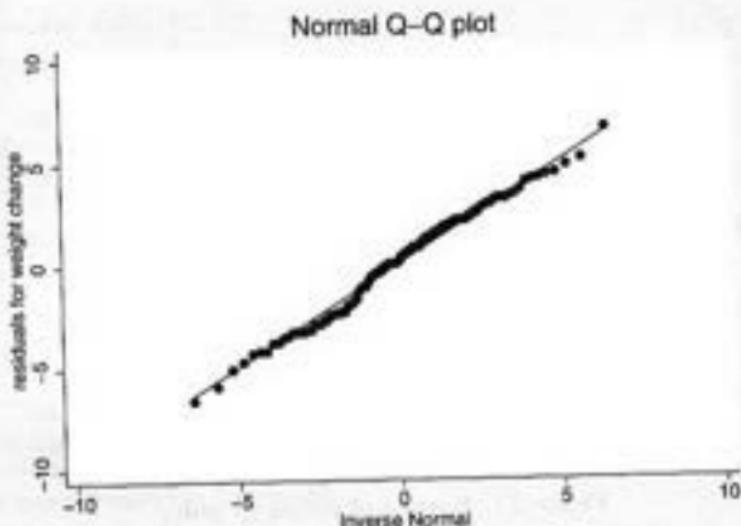


Figure 2.2: Normal Q-Q plot of residuals of weight change.

We could also test whether the variances differ significantly using

```
robvar weight, by(life)
```

giving the output shown in Display 2.1. Here the first test is Levene's test, and this test indicates that there is no evidence that the variances differ ($F(2, 104) = 1.37, p = 0.26$). The W50 test statistic replaces the means in Levene's test by medians and the W10 statistics replaces the means by 10% trimmed means. All these tests are more robust to non-normality than the conventional homogeneity of variance test produced by the `adtest` command.

Having found no strong evidence that the assumptions of the t-test are not valid for weight gain, we can proceed to apply a t-test:

Variance ratio test

LIFE	Summary of WEIGHT		
	Mean	Std. Dev.	Freq.
no	1.4088889	2.6092338	45
yes	1.7311478	2.8256293	61
Total	1.5943396	2.727805	106
$\chi^2 = 1.371606$	df(2, 104)	Pr > F = .25826967	
$\chi^2 = 1.2167221$	df(2, 104)	Pr > F = .30038041	
$\chi^2 = 1.308366$	df(2, 104)	Pr > F = .27447211	

Display 2.1

ttest weight, by(life)

Two-sample t test with equal variances

Group	Obs	Mean	Std. Err.	Std. Dev.	[95% Conf. Interval]
no	45	1.4088889	.3889616	2.609234	.6249983 2.19279
yes	61	1.7311478	.3617847	2.825629	1.00747 2.454825
combined	106	1.5943396	.2649478	2.727805	1.068997 2.119682
diff		-0.3222587	.5376605		-1.388499 .743982
diff = mean(no) - mean(yes)					t = -0.5993
H0: diff = 0					degrees of freedom = 104
Ha: diff < 0					Ha: diff > 0
Pr(T < t) = 0.3751					Pr(T > t) = 0.7249

Display 2.2

Display 2.2 shows that the difference in means is estimated as -0.32 with a 95% confidence interval from -1.39 to 0.74. The two-tailed *p*-value is 0.55, so there is no evidence that the populations differ in their mean weight change. (The unequal option could be used to relax the assumption of equal population variances.)

Now suppose we wish to compare the prevalence of depression between suicidal and non-suicidal women. The two categorical variables can be cross-tabulated and the appropriate chi-squared statistic calculated using a single command:

`tabulate life depress, row chi2`
 The output is shown in Display 2.3. Here the `row` option was used

Key		DEPRESS			Total
	frequency	1	2	3	
life	no	26	24	1	51
	yes	50.98	47.06	1.96	100.00
yes	0	42	16	58	
	0.00	72.41	27.59	100.00	
Total	26	66	57	109	
	33.85	60.55	15.60	100.00	

Pearson chi2(2) = 43.8758 Pr = 0.000

Display 2.3

to display row-percentage, making it easier to compare the groups of women. For example, 50.98% of non-suicidal women are not depressed at all, compared with 0% of suicidal women. The value of the chi-squared statistic implies that there is a highly significant association between depression and suicidal thoughts ($\chi^2 = 43.9$, d.f.=2, $p < 0.001$). Note that this test does not take account of the ordinal nature of depression and is therefore likely to be less sensitive than, for example, ordinal regression (see Chapter 6). Since some cells in the cross classification have only small counts, we might want to use Fisher's exact test (see Everitt, 1992) rather than the chi-squared test. (We could first use the `expected` option in the `tabulate` command to check if the expected counts are small.) The necessary command (without reproducing the table) is as follows:

```
tabulate life depress, exact noref  

Fisher's exact = 0.000
```

Again we find strong evidence for a relationship between depression and suicidal thoughts (Fisher's exact test, $p < 0.001$).

A useful display for two-way tables is a bar chart which can be produced as follows:

```

graph bar (count) id,
    over(depress, relabel(1 "none" 2 "mild" 3 "moderate")) ///
    over(life, relabel(1 "non-suicidal" 2 "suicidal")) ///
    ytitle(Percentages by group (suicidal versus not)) ///
    asyvars percent showyvars legend(off)

```

Here we used `(count)` `id` to plot the number of non-missing values of `id` and two `over()` options to specify the grouping variables `depress` and `life`. Within these `over()` options, we defined the labels to be printed for the categories of `depress` and `life`. To display the bars for the first grouping variable in different colors, we used the `asyvars` option. The `asyvars` option also allows us to use the `percent` option to convert the counts for the different depression categories to percentages within the two groups (suicidal versus not). Finally, we used `showyvars` and `legend(off)` to place the labels for the depression categories underneath the corresponding bars instead of within a legend. The graph is shown in Figure 2.3.

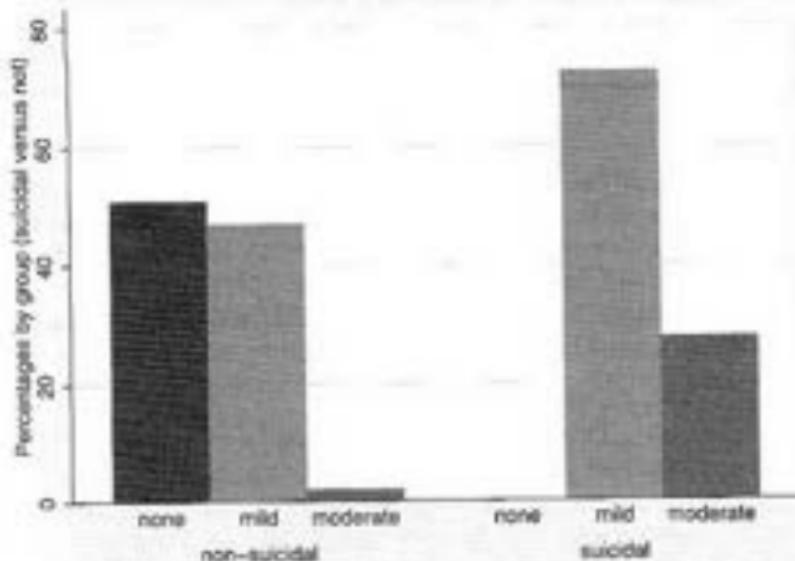


Figure 2.3: Bar chart of percentages of women in different depression categories by group (suicidal versus non-suicidal).

We now test the null hypothesis that the proportion of women who have lost interest in sex does not differ between the populations of suicidal and non-suicidal women. We can obtain the relevant table and both the chi-squared test and Fisher's exact tests using

```
tabulate life sex, row chi2 exact
```

with results shown in Display 2.4. Therefore, those who have thought

Key				
		SEX		
		no	yes	Total
no	12	38		50
	24.00	76.00		100.00
yes	5	58		63
	7.94	92.06		100.00
Total	17	96		113
	15.04	84.96		100.00
Pearson's chi2(1) =		6.6279	Pr = 0.018	
Fisher's exact =			0.032	
1-sided Fisher's exact =			0.017	

Display 2.4

about ending their lives are more likely to have lost interest in sex than those who have not (92% compared with 76%) and the association is significant (Fisher's exact test, $p = 0.032$).

The correlations between the three variables weight, iq, and age can be found using the correlate command.

```
corr weight iq age
```

(see Display 2.5). This correlation matrix has been evaluated for those 100 women who had complete data on all three variables. An alternative approach is to use the pwcorr command to include, for each correlation, all observations that have complete data for the corresponding pair of variables, resulting in different sample sizes for different correlations. These pairwise correlations can be obtained together with the sample sizes and p-values using

```
pwcorr weight iq age, obs sig
```

(obs=100)

	weight	iq	age
weight	1.0000		
iq	-0.2920	1.0000	
age	0.4131	-0.4363	1.0000

Display 2.5

(see Display 2.6).

	weight	iq	age
weight	1.0000		
iq	-0.2920	1.0000	
age	0.4131	-0.4363	1.0000
	0.0032	100	110
	107		
	0.0000	110	118
	118		

Display 2.6

The corresponding scatterplot matrix is obtained using `graph matrix` as follows:

```
graph matrix weight iq age, half jitter(1) msymbol(Oh) ///
  msize(small) diagonal("Weight change" "IQ" "Age")
```

where `jitter(1)` randomly moves the points by a very small amount to stop them overlapping completely due to the discrete nature of age and IQ. The resulting graph is shown in Figure 2.4. We see that older and less intelligent women tend to put on more weight than younger and more intelligent ones. However, older women in this sample also tended to be less intelligent so that age and intelligence are confounded.

It is of some interest to assess whether age and weight change have the same relationship in suicidal as in non-suicidal women. We shall do this informally by constructing a single scatterplot of weight change against age in which the women in the two groups are represented by different symbols. This is easily done by simply specifying two overlaid

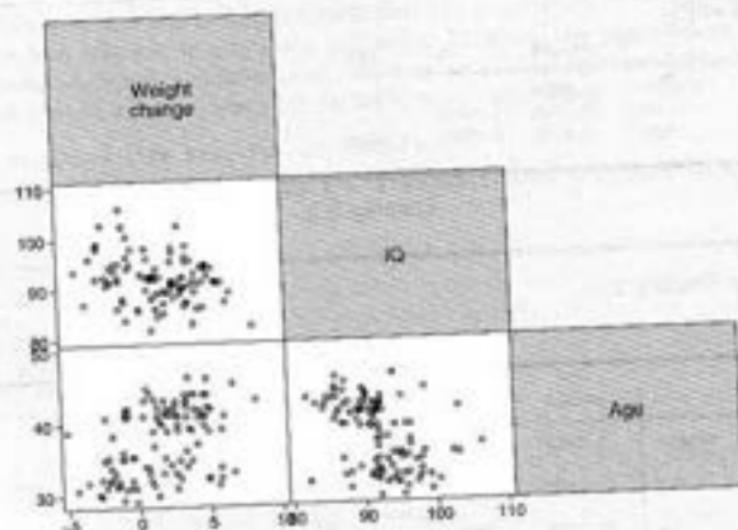


Figure 2.4: Scatterplot matrix for weight, IQ, and age.

scatterplots as follows:

```
twoway (scatter weight age if life==1,
        msymbol(O) mcolor(black) jitter(2))    ///
        (scatter weight age if life==2,
        msymbol(O) mcolor(black)                ///
        jitter(2)), legend(order(1 "no" 2 "yes"))
```

The resulting graph in Figure 2.5 shows that within both groups, higher age is associated with larger weight increases, and the groups do not form distinct "clusters".

Finally, an appropriate correlation between the ordinal variables depression and anxiety is Kendall's tau-b which can be obtained using

```
ktau depress anxiety
Number of obs =      107
Kendall's tau-a =    0.2827
Kendall's tau-b =    0.4951
Kendall's score =   1603
SE of score =     289.279  (corrected for ties)
Test of H0: depress and anxiety are independent
Prob > |z| =       0.0000  (continuity corrected)
```

giving a value of 0.50 with an approximate p-value of $p < 0.001$. Depression and anxiety are clearly related in these psychiatrically ill women.

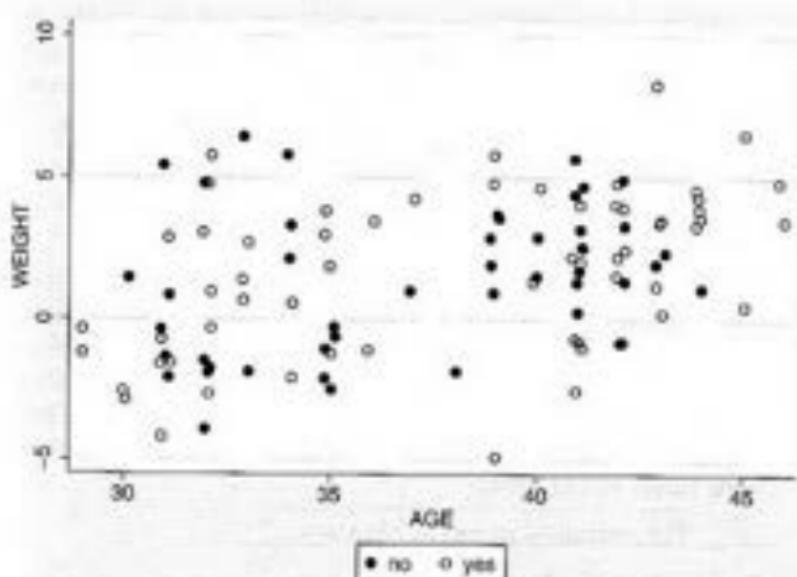


Figure 2.5: Scatterplot of weight against age.

2.4 Exercises

2.1 • Female psychiatric patients

1. Tabulate the mean weight change by level of depression.
2. By looping through the variables `age`, `iq`, and `weight` using `foreach`, tabulate the means and standard deviations for each of these variables by `life`.
3. Produce a bar chart analogous to the one in Figure 2.3 but for `sex` and `life`.
4. Use `search nonparametric` or `search mann` or `search whitney` to find help on how to run the Mann-Whitney U -test.
5. Compare the weight changes between the two groups using the Mann-Whitney U -test.
6. Form a scatterplot for `iq` and `age` using different symbols for the two groups (`life=1` and `life=2`). Explore the use of the option `jitter(#)` for different integers `#` to stop symbols overlapping.
7. Find the command for the Spearman correlation coefficient and use it to find the Spearman correlation between `age` and `iq`.
8. Having tried out all these commands interactively, create a

do-file containing these commands and run the do-file. In the graph commands, use the option `saving(filename, replace)` to save the graphs in the current directory and view the graphs later using the command `graph use filename`.

See also Exercises in Chapter 6.

2.2 Australians going metric

Shortly after metric units of length were officially introduced in Australia, each of a group of 44 students was asked to guess, to the nearest meter, the width of the lecture hall in which they were sitting. Another group of 69 students in the same room was asked to guess the width in feet, to the nearest foot. The true width of the hall was 13.1 meters (43.0 feet). The data come from Hand et al.(1994).

The variables in `meter.dta` are:

- `id`: student identifier
 - `meters`: dummy variable for guess being in meters
 - `guess`: guesses
1. Investigate, by the use of suitable graphics, significance tests and estimation procedures whether there is any evidence of a systematic difference in the guesses made in meters and those made in feet.

2.3 Mortality from skin cancer

Here we consider a dataset from van Belle et al. (2004) which consist of mortality rates due to malignant melanoma of the skin for white males during the period 1950-1969, for each state on the U.S. mainland.

The variables in `mortality.dta` are:

- `state`: name of the state
- `mortality`: mortality rate (in deaths per 10 million per year)
- `latitude`: latitude of the center of each state
- `longitude`: longitude of the center of the state
- `population`: population (in millions)
- `ocean`: dummy variable for state being contiguous with an ocean

1. Construct some suitable graphics for investigating how mortality is related to latitude and longitude and how any relationship between these variables is affected by being an ocean

state.

2.4 Invasion of acacia trees by ants

The data in the 2×2 contingency table below (from Sokal and Rohlf, 1981) record the results on an experiment with acacia ants. All but 28 trees of two species of acacia (A and B) were cleared from an area in Central America, and these 28 trees were cleared of ants using insecticide. Sixteen colonies of a particular species of ant were obtained from other trees of species A. The colonies were placed roughly equidistant from the 28 trees and allowed to invade them.

Acacia Species	Not invaded	Invaded	Total
A	2	13	15
B	10	3	13
Total	12	16	28

1. Produce a table containing the percentage of trees invaded by tree type and the expected frequencies under the null hypothesis that there is no association between type of tree and invasion by ants. (Hint: use the `tabi` command to specify the frequencies within the Stata command instead of entering the data.)
2. Investigate whether there is any evidence that the invasion probability differs between the two species of acacia tree.
3. Obtain an approximate 95% confidence interval for the relevant difference in proportions (Hint: Use the `cst` command).

2.5 Sexual satisfaction

Hout, Duncan and Sobel (1987) investigated the relative sexual satisfaction of married couples, by asking each member of 91 married couples to rate the degree to which they agreed with the statement "Sex is fun for me and my partner" on a four-point scale ranging from "never or occasionally", to "almost always". The data for 30 couples are given in `satisfaction.dta`. The variables are:

- `couple`: couple identifier
 - `husband`: satisfaction score of husband
 - `wife`: satisfaction score of wife
1. Carry out an appropriate significance test to investigate whether there is any evidence that men and women differ in their mean sexual satisfaction.

2. Construct a crosstabulation of husband's and wife's sexual satisfaction and calculate a suitable measure of correlation between the two ratings.
3. Construct a 95% confidence interval for the true mean difference between the sexual satisfaction scores of husbands and wives.

2.6 Crowd reactions to threatened suicide

Mann (1981) conducted a study to investigate the causes of jeering or baiting behavior by a crowd when a person is threatening to commit suicide by jumping from a high building. The data given below result from the classification of threatened suicides by two factors, the time of year and whether or not baiting occurred.

	Baiting	non-baiting
June-September	8	4
October-May	2	7

1. A hypothesis is that baiting is more likely to occur in warm weather. (The data come from the northern hemisphere, so June-September are the warm months). Produce a table of counts and percentages for assessing this hypothesis. (You can use the `tabi` command which allows you to list the cell frequency in the command, rather than entering them as a dataset.)
2. Produce a table of expected frequencies under the null hypothesis that there is no association between season and baiting behavior.
3. Test the null hypothesis that there is no association between season and baiting behavior.

Chapter 3

Multiple Regression: Determinants of Pollution in U.S. Cities

3.1 Description of data

Data on air pollution in 41 U.S. cities were collected by Sokal and Rohlf (1981) from several U.S. government publications and are reproduced here in Table 3.1. (The data are also given in Hand *et al.*, 1994.) There is a single dependent variable, so_2 , the annual mean concentration of sulphur dioxide, in micrograms per cubic meter. These data are means for the three years 1969 to 1971 for each city. The values of six explanatory variables, two of which concern human ecology and four climate, are also recorded; details are as follows:

- temp: average annual temperature in °F
- manuf: number of manufacturing enterprises employing 20 or more workers
- pop: population size (1970 census) in thousands
- wind: average annual wind speed in miles per hour
- precip: average annual precipitation in inches
- days: average number of days with precipitation per year

The main question of interest about these data is how the pollution level as measured by sulphur dioxide concentration is determined by the six explanatory variables. The central method of analysis will be multiple regression.

Table 3.1 Data in usair.dat

Town	SO2	temp	manuf	pop	wind	precip	days
Phoenix	10	70.3	213	582	6.0	7.05	36
Little Rock	13	61.0	91	132	8.2	48.52	100
San Francisco	12	56.7	453	716	8.7	20.66	67
Denver	17	51.9	454	515	9.0	12.95	86
Hartford	56	49.1	412	158	9.0	43.37	127
Wilmington	36	54.0	80	80	9.0	40.25	114
Washington	29	57.3	434	757	9.3	38.89	113
Jackson	14	68.4	136	529	8.8	54.47	116
Miami	10	75.5	207	335	9.0	59.80	128
Atlanta	24	61.5	368	497	9.1	48.34	115
Chicago	110	50.6	3344	3369	10.4	34.44	122
Indianapolis	28	52.3	361	746	9.7	38.74	121
Des Moines	17	49.0	104	201	11.2	30.85	103
Wichita	8	56.6	125	277	12.7	30.58	82
Louisville	30	55.6	291	593	8.3	43.11	123
New Orleans	9	68.3	204	361	8.4	56.77	113
Baltimore	47	55.0	625	905	9.6	41.31	111
Detroit	35	49.9	1064	1513	10.1	30.96	129
Minneapolis	29	43.5	699	744	10.6	25.94	137
Kansas	14	54.5	381	507	10.0	37.09	99
St. Louis	56	55.9	775	622	9.5	35.89	105
Omaha	14	51.5	181	347	10.9	30.18	98
Albuquerque	11	56.8	46	244	8.9	7.77	58
Albany	46	47.6	44	116	8.8	33.36	135
Buffalo	11	47.1	391	463	12.4	36.11	166
Cincinnati	23	54.0	462	453	7.1	39.04	132
Cleveland	65	49.7	1007	751	10.9	34.99	155
Columbia	26	51.5	266	540	8.6	37.01	134
Philadelphia	69	54.6	1692	1960	9.6	39.93	115
Pittsburgh	61	50.4	347	520	9.4	36.22	147
Providence	94	50.0	343	179	10.6	42.75	125
Memphis	10	61.6	337	624	9.2	49.10	105
Nashville	18	59.4	275	448	7.9	46.00	119
Dallas	9	60.2	641	844	10.9	35.94	78
Houston	10	68.9	721	1233	10.8	48.19	103
Salt Lake City	28	51.0	137	176	8.7	35.17	89
Norfolk	31	59.3	96	308	10.6	44.68	116
Richmond	26	57.8	197	209	7.6	42.59	115
Seattle	29	51.1	379	531	9.4	38.79	164
Charleston	31	55.2	35	71	6.5	40.75	148
Milwaukee	16	45.7	569	717	11.8	29.07	123

3.2 The multiple regression model

The multiple regression model has the general form

$$y_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \cdots + \beta_p x_{pi} + \epsilon_i \quad (3.1)$$

where y_i is a continuous response (or dependent) variable for the i th member of the sample, $x_{1i}, x_{2i}, \dots, x_{pi}$ are a set of explanatory (or independent) variables or covariates, $\beta_0, \beta_1, \beta_2, \dots, \beta_p$ are regression coefficients, and ϵ_i is a residual or error term with zero mean that is uncorrelated with the explanatory variables. It follows that the expected value of the response for given values of the covariates is

$$E(y_i | \mathbf{x}_i) = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \cdots + \beta_p x_{pi},$$

where $\mathbf{x}'_i = (x_{1i}, \dots, x_{pi})$. This is also the value we would predict for a new individual with covariate values \mathbf{x}_i if we knew the regression coefficients.

Each regression coefficient represents the mean change in the response variable when the corresponding explanatory variable increases by one unit and all other explanatory variables remain constant. The coefficients therefore represent the effects of each explanatory variable, controlling for all other explanatory variables in the model, giving rise to the term "partial" regression coefficients. The residual is the difference between the observed value of the response and the expected value based on the explanatory variables.

The regression coefficients β_0, \dots, β_p are generally estimated by least squares; in other words the estimates $\hat{\beta}_0, \dots, \hat{\beta}_p$ minimize the sum of the squared differences between observed and predicted responses, or the sum of squared estimated residuals,

$$\sum_{i=1}^n [y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_{1i} + \hat{\beta}_2 x_{2i} + \cdots + \hat{\beta}_p x_{pi})]^2. \quad (3.2)$$

Significance tests for the regression coefficients can be derived by assuming that the error terms are independently normally distributed with zero mean and constant variance σ^2 .

For n observations of the response and explanatory variables, the regression model may be written concisely as

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon} \quad (3.3)$$

where \mathbf{y} is the $n \times 1$ vector of responses, \mathbf{X} is an $n \times (p+1)$ matrix of known constants, the first column containing a series of ones corre-

sponding to the term β_0 in (3.1), and the remaining columns values of the explanatory variables. The elements of the vector β are the regression coefficients β_0, \dots, β_p , and those of the vector ϵ , the error terms $\epsilon_1, \dots, \epsilon_n$. The least squares estimates of the regression coefficients can then be written as

$$\hat{\beta} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y}$$

and the variances and covariances of these estimates can be found from

$$\mathbf{S}_{\hat{\beta}} = s^2(\mathbf{X}'\mathbf{X})^{-1},$$

where s^2 is an estimate of the residual variance σ^2 given by the sum of squared estimated residuals in equation (3.2) divided by $n-p-1$.

The coefficient of determination, R^2 , represents the portion of the total variance of the response variable that is explained by the explanatory variables. Alternatively, it can be interpreted as the proportional reduction in prediction error variance of the model compared with the constant-only model (without covariates). R , also known as the *multiple correlation coefficient*, is just the correlation between the observed responses y_i and the predicted responses \hat{y}_i . For full details of multiple regression see, for example, Rawlings *et al.* (1998).

3.3 Analysis using Stata

Assuming the data are available as an ASCII file `usair.dat` in the current directory and that the file contains city names (abbreviated versions of those in Table 3.1), they may be read in for analysis using the following instruction:

```
infile str10 town so2 temp manuf pop      ///
    wind precip days using usair.dat, clear
```

Here we had to declare the "type" for the string variable `town` as `str10` which stands for "string variable with 10 characters".

Before undertaking a formal regression analysis of these data, it will be helpful to examine them graphically using a scatterplot matrix. Such a display is useful for assessing the general relationships between the variables, for identifying possible outliers, and for highlighting potential collinearity problems amongst the explanatory variables. The basic plot can be obtained using

```
graph matrix so2 temp manuf pop wind precip days
```

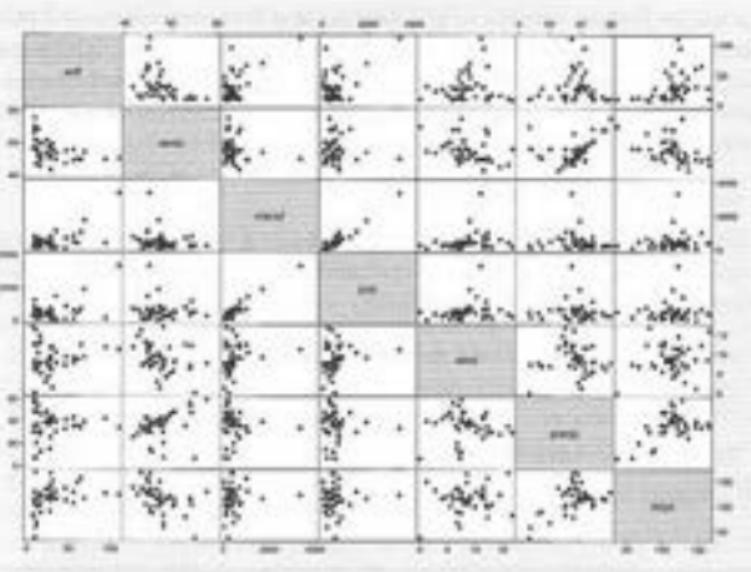


Figure 3.1: Scatterplot matrix.

The resulting diagram is shown in Figure 3.1. Several of the scatterplots show evidence of outliers, and the relationship between `manuf` and `pop` is very strong suggesting that using both as explanatory variables in a regression analysis may lead to problems (see later). The relationships of particular interest, namely those between `so2` and the explanatory variables (the relevant scatterplots are those in the first row of Figure 3.1), indicate some possible nonlinearity. A more informative, although slightly more "messy" diagram can be obtained if the plotted points are labeled with the associated town names. We first create a variable containing the first three characters of the strings in `town` using the function `substr()`

```
generate twn = substr(town,1,3)
```

We then create a scatterplot matrix with these three-character town labels using

```
graph matrix so2-days, msymbol(none) mlabel(twn) ///
    mlabposition(0)
```

The `mlabel()` option labels the points with the names in the `twn` variable. By default, a "marker symbol" would also be plotted and this can be suppressed using `msymbol(none)`; `mlabposition(0)` centers the labels where the symbol would normally go. The resulting diagram

appears in Figure 3.3. Clearly, Chicago and to a lesser extent Philadelphia might be considered outliers. Chicago has such a high degree of pollution compared with the other cities that it should perhaps be considered as a special case and excluded from further analysis. We can remove Chicago using

```
drop if town=="Chicago"
```

The command `regress` may be used to fit a basic multiple regression model. The necessary Stata command for regressing sulphur dioxide concentration on the six explanatory variables is

```
regress so2 temp manuf pop wind precip days  
or, alternatively,
```

```
regress so2 temp-days  
(see Display 3.1).
```

Source	SS	df	MS	Number of obs = 40			
Model	8203.60623	6	1367.26754	F(6, 33) = 6.20			
Residual	7282.29477	33	220.575599	Prob > F = 0.0002			
Total	15485.9	39	397.074358	R-squared = 0.5297			
				Adj R-squared = 0.4442			
				Root MSE = 14.855			
so2	Coeff.	Std. Err.	t	P> t	[95% Conf. Interval]		
temp	-1.268452	.6305250	-2.01	0.052	-2.561266	.0143631	
manuf	.0664927	.0181777	3.60	0.001	.0285008	.1024756	
pop	-.030431	.0185342	-2.54	0.016	-.0710357	-.0078284	
wind	-3.198267	1.888713	-1.72	0.095	-6.981881	.5853469	
precip	.5136844	.3687273	1.39	0.173	-.2364906	1.263866	
days	-.05332051	.1653678	-0.32	0.760	-.3896277	.3832175	
_cons	111.8709	48.07439	2.33	0.026	14.06278	209.679	

Display 3.1

The main features of interest in the output in Display 3.1 are the analysis of variance table and the parameter estimates. In the former, the ratio of the model mean square to the residual mean square gives an *F*-test for the hypothesis that all the regression coefficients in the fitted model are zero (except the constant β_0). The resulting *F*-statistic with 6 and 33 degrees of freedom takes the value 6.20 and is shown on the right-hand side; the associated *p*-value is very small. Consequently, the hypothesis is rejected. The square of the multiple

correlation coefficient (R^2) is 0.53 showing that 53% of the variance of sulphur dioxide concentration is accounted for by the six explanatory variables of interest.

The adjusted R^2 statistic is an estimate of the population R^2 taking account of the fact that the parameters were estimated from the same data for which R^2 is evaluated. The statistic is calculated as

$$\text{adj } R^2 = 1 - \frac{(n-1)(1-R^2)}{n-p} \quad (3.4)$$

where n is the number of observations used in fitting the model. The root MSE is simply the square root of the residual mean square in the analysis of variance table, which itself is an estimate of the parameter σ^2 . The estimated regression coefficients give the estimated change in the mean of the response variable produced by a unit change in the corresponding explanatory variable with the remaining explanatory variables held constant.

One concern generated by the initial graphical material on this data was the strong relationship between the two explanatory variables manuf and pop. The correlation of these two variables is obtained by using,

```
correlate manuf pop
(n=40)

```

	manuf	pop
manuf	1.0000	
pop	0.8906	1.0000

The strong linear dependence might be a source of collinearity problems and can be investigated further by calculating what are known as *covariance inflation factors* for each of the explanatory variables. These are given by

$$\text{VIF}(x_k) = \frac{1}{1 - R_k^2} \quad (3.5)$$

where $\text{VIF}(x_k)$ is the variance inflation factor for explanatory variable x_k , and R_k^2 is the square of the multiple correlation coefficient obtained from regressing x_k on the remaining explanatory variables. The variance inflation factor represents the squared standard error (or sampling variance) of $\hat{\beta}_k$ in the estimated model divided by the squared standard error that would be obtained if x_k were uncorrelated with the remaining variables.

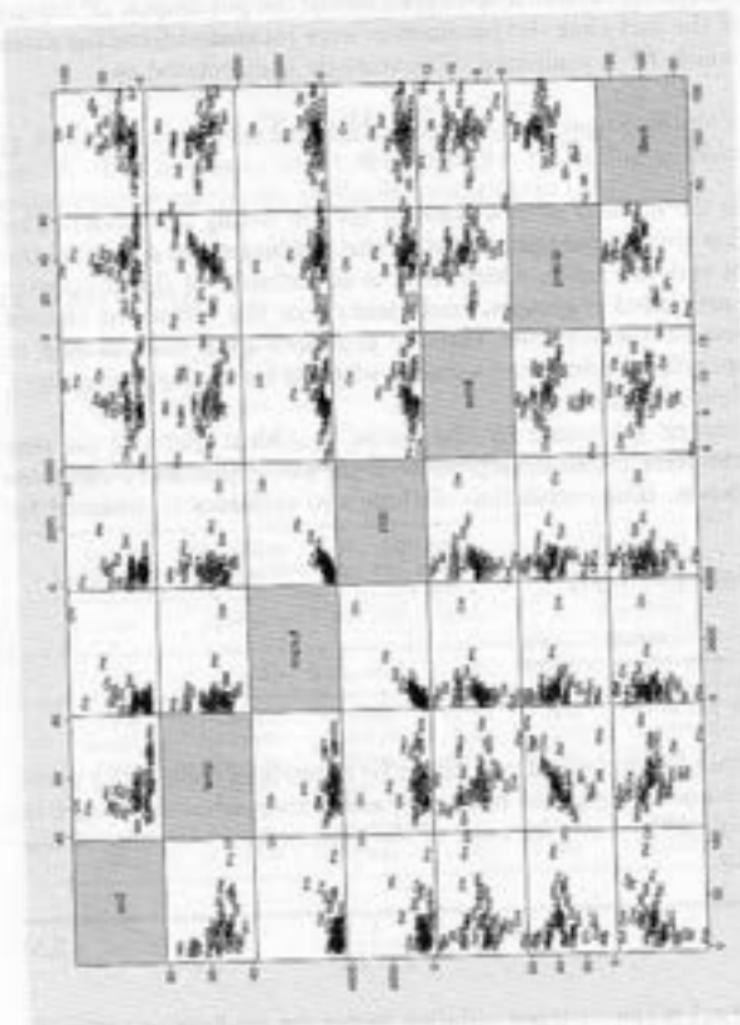


Figure 3.2: Scatterplot matrix with city labels.

The variance inflation factors can be obtained using the `estat vif` command after `regress`:

```
estat vif
(see Display 3.2).
```

Variable	VIF	1/VIF
manuf	6.28	0.159275
pop	6.13	0.163165
temp	3.72	0.269156
days	3.47	0.287962
precip	3.41	0.293125
wind	1.26	0.790619
Mean VIF	4.05	

Display 3.2

Chatterjee *et al.* (2006) give the following “rules-of-thumb” for evaluating these factors:

- Values larger than 10 give evidence of collinearity.
- A mean of the VIF factors considerably larger than one suggests collinearity.

Here there are no values greater than 10 (as an exercise we suggest readers also calculate the VIFs when the observations for Chicago are included), but the mean value of 4.05 gives some cause for concern. A simple (although not necessarily the best) way to proceed is to drop one of `manuf` or `pop`. Another possibility is to replace `manuf` by a new variable equal to `manuf` divided by `pop`, representing the number of large manufacturing enterprises per thousand inhabitants (see Exercise 3.1). However, we shall simply exclude `manuf` and repeat the regression analysis using the five remaining explanatory variables:

```
regress so2 temp pop wind precip days
```

The output is shown in Display 3.3.

Now recompute the variance inflation factors:

```
estat vif
```

The variance inflation factors in Display 3.4 are now satisfactory.

The very general hypothesis concerning all regression coefficients mentioned previously is not usually of great interest in most applications of multiple regression because it is most unlikely that all the

Source	SS	df	MS	Number of obs =	40
Model	5339.03466	5	1067.80693	F(5, 34) =	3.58
Residual	10146.9654	34	298.437216	Prob > F =	0.0106
Total	15485.9	39	397.074269	R-squared =	0.3448
				Adj R-squared =	0.2486
				Root MSE =	17.275
so2	Coeff.	Std. Err.	t	P> t	[95% Conf. Interval]
temp	-1.867665	.7072827	-2.64	0.012	-.306037 -.430294
pop	.0113089	.0075627	1.51	0.141	-.0039723 .0267661
wind	-3.126429	2.16257	-1.45	0.157	-.7.8213 1.268443
precip	-.6021106	.4270489	1.41	0.168	-.2673827 1.471604
days	-.020142	.1920012	-0.10	0.917	-.4103424 .3700445
_cons	138.8568	55.36787	2.45	0.019	33.33529 248.3778

Display 3.3

Variable	VIF	1/VIF
days	3.46	0.288750
temp	3.46	0.289282
precip	3.40	0.295429
wind	1.26	0.790710
pop	1.07	0.931016
Mean VIF	3.63	

Display 3.4

chosen explanatory variables will be unrelated to the response variable. The more interesting question is whether a subset of the regression coefficients is zero, implying that not all the explanatory variables are of use in predicting the response variable. A preliminary assessment of the likely importance of each explanatory variable can be made using the table of estimated regression coefficients and associated statistics. Using a conventional 5% criterion, the only "significant" coefficient is that for the variable `temp`. Unfortunately, this very simple approach is not in general suitable, since in most cases the explanatory variables are correlated, and the *t*-tests will not be independent of each other. Consequently, removing a particular variable from the regression will alter both the estimated regression coefficients of the remaining variables and their standard errors. A more involved approach to identifying important subsets of explanatory variables is therefore required. A number of procedures are available.

1. *Confirmatory approach:* A small set of explanatory variables are included as suggested by substantive theory, or to allow testing of particular *a priori* hypotheses. The model is typically modified somewhat by removing some variables, considering interactions, etc. to achieve a better fit to the data.
2. *Exploratory approach:* Automatic selection methods, which are of the following types:
 - a. *Forward selection:* This method starts with a model containing none of the explanatory variables and then considers variables one by one for inclusion. At each step, the variable added is the one that results in the biggest increase in the model or regression sum of squares. An *F*-type statistic is used to judge when further additions would not represent a significant improvement in the model.
 - b. *Backward elimination:* Here variables are considered for removal from an initial model containing all the explanatory variables. At each stage, the variable chosen for exclusion is the one leading to the smallest reduction in the regression sum of squares. Again, an *F*-type statistic is used to judge when further exclusions would represent a significant deterioration in the model.
 - c. *Stepwise regression:* This method is essentially a combination of the previous two. The forward selection procedure is used to add variables to an existing model and, after each addition, a backward elimination step is introduced to assess whether variables entered earlier might now be removed because they no longer contribute significantly to the model.

It is clear that the automatic selection methods are based on a large number of significance tests, one for each variable considered for inclu-

sion or exclusion in each step. It is well known that the probability of a false positive result or Type I error increases with the number of tests. The chosen model should therefore be interpreted with extreme caution, particularly if there were a large number of candidate variables. Another problem with the three automatic procedures is that they often do not lead to the same model; see also Harrell (2001) for a discussion of model selection strategies. Although we would generally not recommend automatic procedures, we will use them here for illustration.

First we will take a confirmatory approach to investigate if climate (temp, wind, precip, days) or human ecology (pop) or both are important predictors of air pollution. We treat these groups of variables as single terms, allowing either all variables in a group to be included or none. This can be done by enclosing the variables in parentheses in the following command:

```
stepwise, pe(0.05): regress so2 (temp wind precip days) (pop)
```

(see Display 3.5). Here the prefix command `stepwise` is used with the

begin with empty model						
p = 0.0119 < 0.0500	adding temp wind precip days				Number of obs =	40
Source	SS	df	MS		F(4, 35) =	3.77
Model	4661.27545	4	1165.31886		P>F =	0.0119
Residual	50024.4346	35	1429.274987		R-squared =	0.3010
Total	15485.9	39	397.074359		Adj R-squared =	0.2211
					Root MSE =	17.986
so2	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
temp	-1.609548	.7069204	-2.38	0.023	-3.131063	-.2486329
wind	-2.309449	2.13119	-1.08	0.285	-8.636996	2.017097
precip	.5241586	.4323535	1.21	0.234	-.3535647	1.401884
days	.0119373	.1942809	0.06	0.951	-.2834123	.4052876
_cons	133.5942	55.75296	2.32	0.033	10.41091	236.7778

Display 3.5

`pe()` ("probability to enter") option to indicate that forward selection should be used with a significance level of 0.05; terms with a p-value less than 0.05 will be included. Here, only the climate variables are shown since they are jointly significant ($p = 0.0119$) using an F -test.

As a further illustration of automatic selection procedures, the following Stata instruction applies the backward elimination method, with

explanatory variables whose *F*-values for removal have associated *p*-values greater than 0.2 being removed:

`stepwise, pr(0.2): regress so2 temp pop wind precip days`
 (see Display 3.6). Here, the `pr()` option indicates that backward selection

begin with full model removing days						
Source	SS	df	MS	Number of obs = 40 F(4, 35) = 4.60 Prob > F = 0.0043 R-squared = 0.3448 Adj R-squared = 0.2896 Root MSE = 17.03		
Model	5335.74801	4	1333.893			
Residual	10150.152	35	286.004343			
Total	15485.9	39	397.074259			
so2	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
temp	-1.810123	.4404001	-4.11	0.000	-2.704183	-.9160636
pop	.0113080	.0074091	1.53	0.136	-.0037323	.0263501
wind	-3.086264	2.096471	-1.47	0.190	-.7341347	1.170778
precip	.5660172	.2908601	2.26	0.030	.0587441	1.07529
_cons	131.3308	34.32034	3.83	0.001	61.66458	301.0126

Display 3.6

tion should be used with a "probability to remove" of 0.2. With this significance level, only the variable *days* is excluded.

The next stage in the analysis should be an examination of the residuals from the chosen model; that is, the differences between the observed and fitted values of sulphur dioxide concentration. Such a procedure is vital for assessing model assumptions, identifying any unusual features in the data indicating outliers, and suggesting possibly simplifying transformations. The most useful ways of examining the residuals are graphical, and the most commonly used plots are as follows:

- A plot of the residuals against each explanatory variable in the model. The presence of a curvilinear relationship, for example, would suggest that a higher-order term, perhaps a quadratic in the explanatory variable, should be added to the model.
- A plot of the residuals against predicted values of the response variable. If the variance of the residuals appears to increase or decrease with the predicted value, a transformation of the response may be in order.

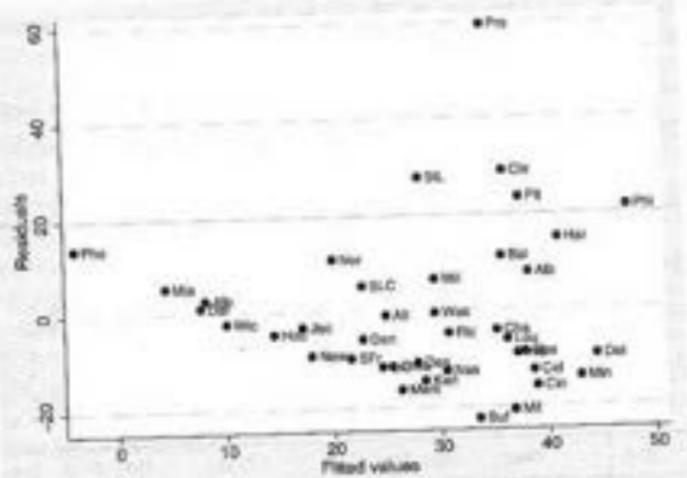


Figure 3.3: Residuals against predicted response.

- A normal probability plot of the residuals—after all systematic variation has been removed from the data, the residuals should look like a sample from the normal distribution. A plot of the ordered residuals against the expected order statistics from a normal distribution (with mean and variance equal to the sample estimates) provides a graphical check on this assumption.

The first two plots can be obtained after estimation with the `regress` command using the `rvfplot` ("residual versus predictor") and `rvfplot` ("residual versus fitted") instructions. For example, for the model chosen by the backward selection procedure, a plot of residuals against predicted values with the first three letters of the town name used to label the points is obtained using the command

```
rvfplot, mlabel(twn)
```

The resulting plot is shown in Figure 3.3, and indicates a possible problem, namely the apparently increasing variance of the residuals as the fitted values increase (see also Chapter 7). Perhaps some thought needs to be given to the possible transformations of the response variable (see exercise 3.1).

Next, graphs of the residuals plotted against each of the four explanatory variables can be obtained using the following `foreach` loop:

```
foreach x in pop temp wind precip {
    rvfplot `x', mlabel(twn)
```

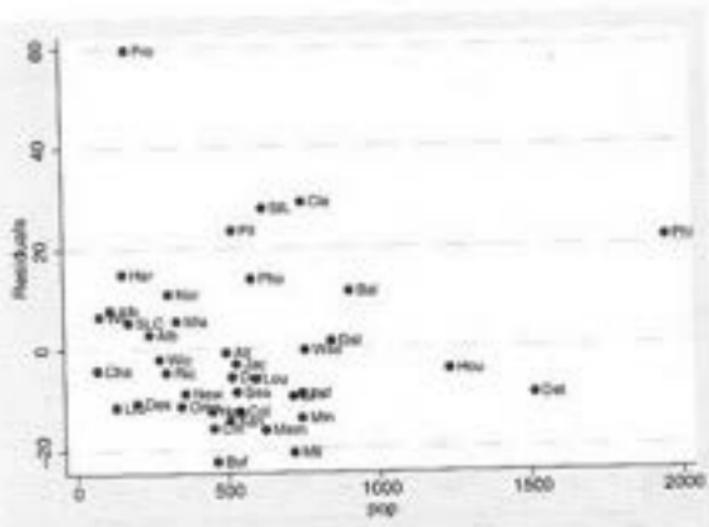


Figure 3.4: Residuals against population size.

NOTE

Here more causes Stata to pause after each graph has been plotted until the user presses any key. The resulting graphs are shown in Figures 3.4 to 3.7. In each graph the point corresponding to the town *Providence* is somewhat distant from the bulk of the points, and the graph for wind has perhaps a "hint" of a curvilinear structure. Note that the appearance of these graphs could be improved using the `slabvpos(varname)` option to specify the "clock positions" (e.g., 12 is straight above) of the labels relative to the points.

The simple residuals plotted by `rvfplot` and `rvpplot` have a distribution that is scale dependent because the variance of each is a function of both σ^2 and the diagonal values of the so-called "hat" matrix, \mathbf{H} , given by

$$\mathbf{H} = \mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}' \quad (3.6)$$

(see Cook and Weisberg (1982) for a full explanation of the hat matrix). Consequently, it is often more useful to work with a standardized version

$$r_i = \frac{y_i - \hat{y}_i}{s\sqrt{1 - h_{ii}}} \quad (3.7)$$

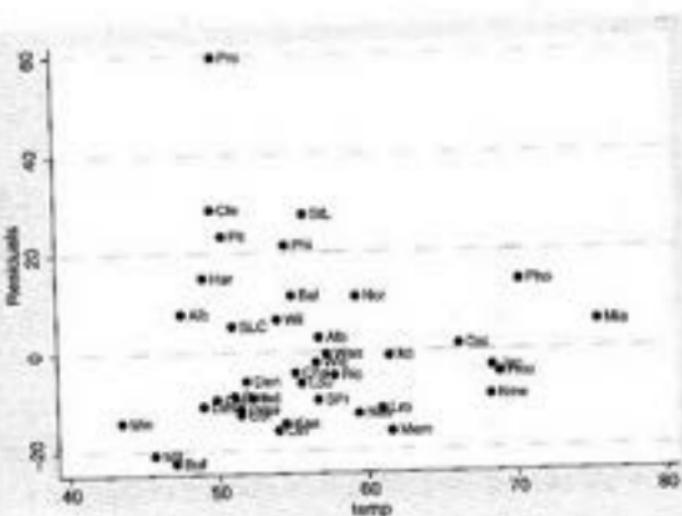


Figure 3.5: Residuals against temperature.

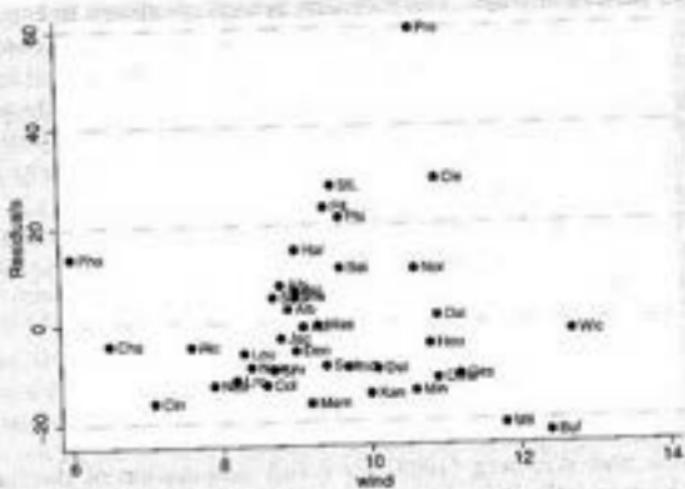


Figure 3.6: Residuals against wind speed.

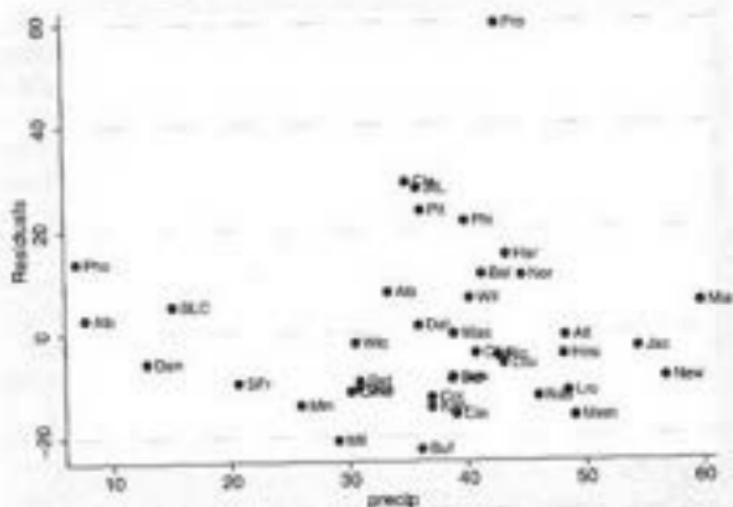


Figure 3.7: Residuals against precipitation.

where s^2 is the estimate of σ^2 , \hat{y}_i is the predicted value of the response, and h_{ii} is the i th diagonal element of \mathbf{H} .

These standardized residuals can be obtained using the `predict` command. For example, to obtain a normal probability plot of the standardized residuals and to plot them against the fitted values requires the following instructions:

```
predict fit
predict sdres, rstandard
ppplot sdres
twoway scatter sdres fit, xlabel(twn)
```

The first instruction stores the fitted values in the variable `fit`, the second stores the standardized residuals in the variable `sdres`, the third produces a normal probability plot (Figure 3.8), and the last instruction produces the graph of standardized residuals against fitted values, which is shown in Figure 3.9.

The normal probability plot indicates that the distribution of the residuals departs somewhat from normality. The pattern in the plot shown in Figure 3.9 is very similar to that in Figure 3.3 but here values outside $(-2, 2)$ indicate possible outliers, in this case the point corresponding to the town Providence. Analogous plots to those in Figures 3.4 to 3.7 could be obtained in the same way.

A rich variety of other diagnostics for investigating fitted regression models has been developed and many of these are available after estima-

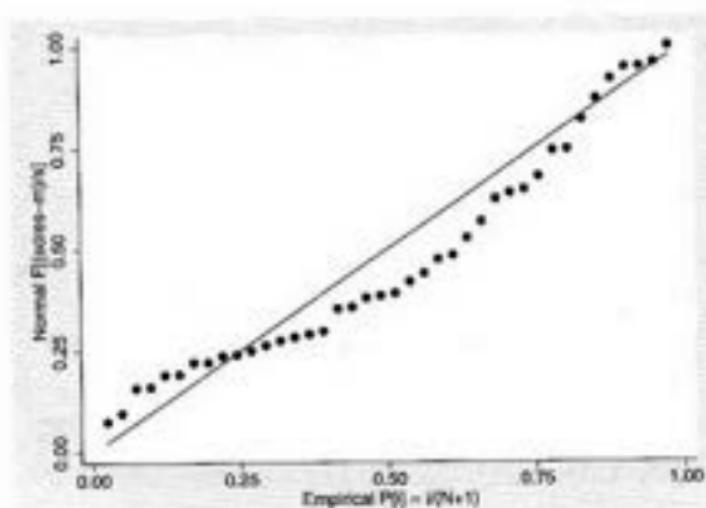


Figure 3.8: Normal probability plot of standardized residuals.

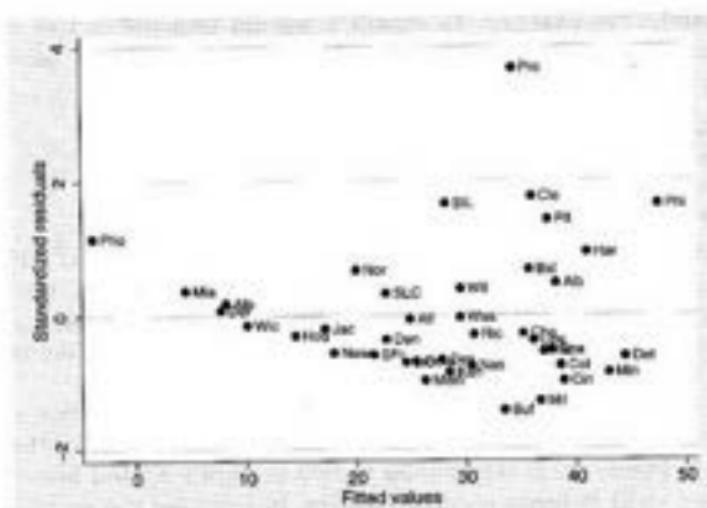


Figure 3.9: Standardized residuals against predicted values.

tion with the `regress` procedure (see `help regress postestimation`). Illustrated here is the use of two of these, namely the partial residual plot (Mallows, 1973) and Cook's distance (Cook, 1977, 1979). The former are useful in identifying whether, for example, quadratic or higher order terms are needed for any of the explanatory variables; the latter measures the change to the estimates of the regression coefficients that results from deleting each observation and can be used to indicate those observations that may be having an undue influence on the estimates.

The partial residual plots are obtained using the `cprplot` ("component plus residual") command. For the four explanatory variables in the selected model for the pollution data, the required plots are obtained as follows:

```
foreach x in pop temp wind precip {
    cprplot `x', lowess
    more
}
```

The `lowess` option produces a locally weighted regression curve or *lowess*. The resulting graphs are shown in Figures 3.10 to 3.13. The graphs have to be examined for nonlinearities and for assessing whether the regression line, which has slope equal to the estimated effect of the corresponding explanatory variable in the chosen model, fits the data adequately. The added *lowess* curve is generally helpful for both. None of the four graphs gives any obvious indication of nonlinearity.

The Cook's distances are found using the `predict` command with the `cooksd` option; the following calculates these statistics for the chosen model for the pollution data and lists the observations where the statistic is greater than $4/40$ ($4/n$), which is usually the value regarded as indicating possible problems.

```
predict cook, cooksd
list town so2 cook if cook>4/40
```

	town	so2	cook
1.	Phoenix	10	.2643286
28.	Philed	69	.3686437
30.	Provld	94	.2839324

The first instruction stores the Cook's distance statistics in the variable `cook`, and the second lists details of those observations for which the statistic is above the suggested cut-off point.

There are three influential observations. Several of the diagnostic procedures used previously also suggest these observations as possibly giving rise to problems, and some consideration should be given to re-

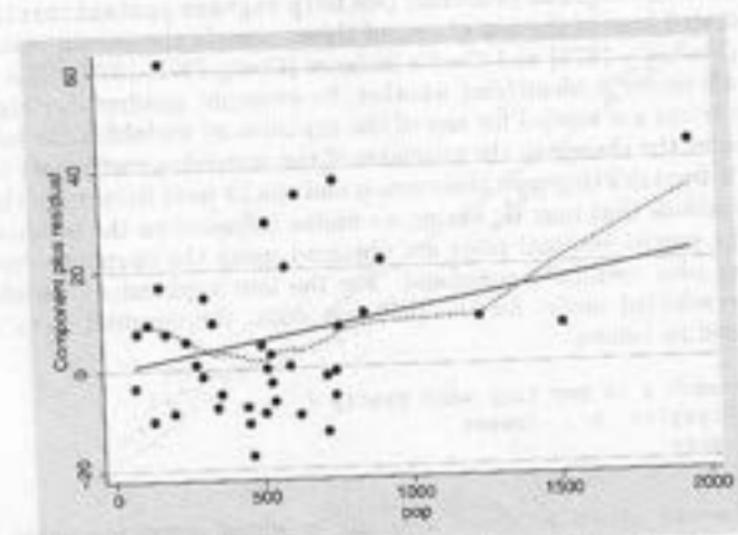


Figure 3.10: Partial residual plot for population size.

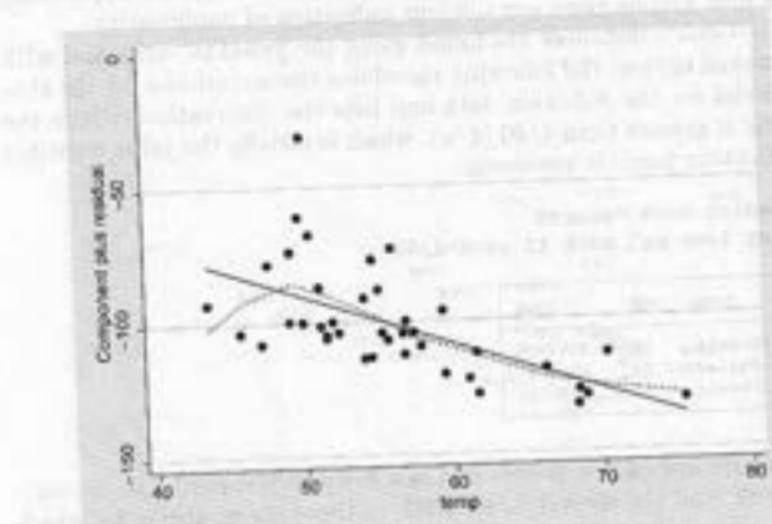


Figure 3.11: Partial residual plot for temperature.

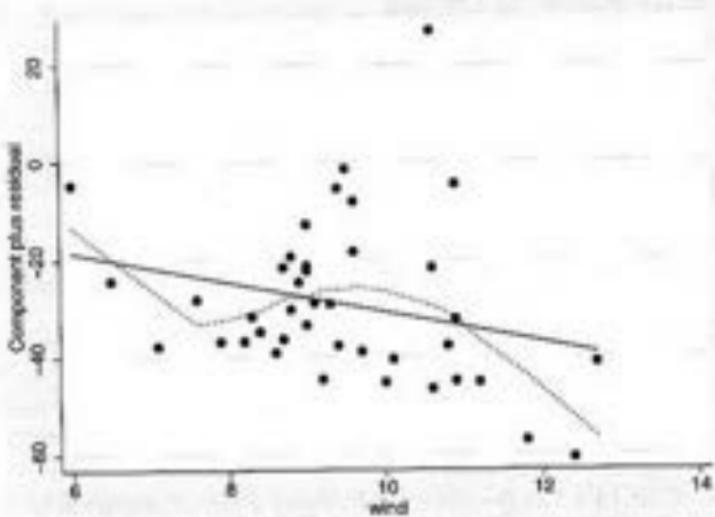


Figure 3.12: Partial residual plot for wind speed.

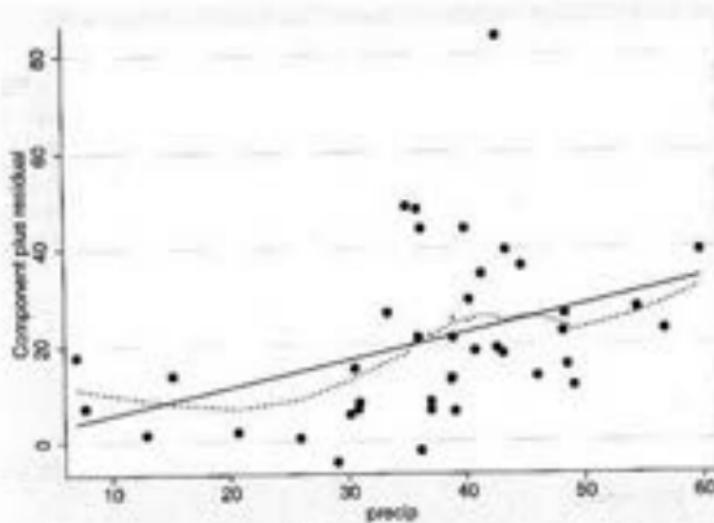


Figure 3.13: Partial residual plot for precipitation.

peating the analyses with these three observations removed in addition to the initial removal of Chicago.

3.4 Exercises

3.1 Determinants of pollution in U.S. cities

1. Repeat the analyses described in this chapter after removing the three possible outlying observations identified by Cook's distances.
2. The solution to the high correlation between the variables manuf and pop adopted in the chapter was simply to remove the former. Investigate other possibilities such as defining a new variable manuf/pop in addition to pop to be used in the regression analysis.
3. Consider the possibility of taking a transformation of sulphur dioxide pollution before undertaking any regression analyses. For example, try a log transformation.
4. Explore the use of the many other diagnostic procedures available with the regress procedure.

See also Exercises in Chapter 14.

3.2 Extroversion and car care

Miles and Shevlin (2001) describe a dataset collected in an investigation of how people project their self-image through objects they own, in this case their cars. The main question is how a person's extroversion affects the amount of time spent looking after his or her car. But since it is known that extroversion is related to both gender and age, the latter two variables need to be controlled for.

The variables in `extroversion.dta` are:

- sex: sex of respondent (0=female, 1=male)
 - age: age (in years)
 - ex: extroversion score
 - car: time respondent spends looking after car (in minutes per week)
1. Fit a suitable regression model to address the main research question stated above.
 2. Interpret the estimated regression coefficients.
 3. Perform some residual diagnostics.

3.3 Mortality from skin cancer

1. For the malignant melanoma data given in Exercise 2.3 of the previous chapter, fit the multiple regression model of mortality on latitude, longitude, population size, and ocean state.
2. Try to find a more parsimonious model (one with fewer explanatory variables) that fits the data adequately.
3. Investigate the assumptions of the model by constructing suitable residual plots or other diagnostic plots.

3.4 Water hardness

Data were collected on 61 large towns in England and Wales to investigate the environmental causes of disease (see Hand et al., 1994). Here we consider the annual mortality per 100,000 for males, averaged over the years 1958–1964, and the calcium concentration in parts per million in the drinking water supply. (The higher the calcium concentration, the harder the water.) Towns at least as far north as Derby are considered northern towns.

The variables in *water.dta* are:

- town: string variable (N=northern town, S=southern town)
 - mortality: mortality per 100,000 for males per year
 - calcium: calcium concentration in parts per million
1. How are mortality and water hardness related, and is there a geographical factor in the relationship?
 2. For your chosen regression model, plot predicted mortality versus calcium concentration with separate regression lines for northern and southern towns.
 3. Superimpose LOWESS curves onto the predicted regression lines. Which assumptions does this graph allow you to assess?

Chapter 4

Analysis of Variance I: Treating Hypertension

4.1 Description of data

Maxwell and Delaney (1990) describe a study in which the effects of three possible treatments for hypertension were investigated. The details of the treatments are as follows:

Treatment	Description	Levels
drug	medication	drug X, drug Y, drug Z
biofeed	biofeedback	present, absent
diet	special diet	present, absent

All 12 combinations of the three treatments were included in a $3 \times 2 \times 2$ design. Seventy-two subjects suffering from hypertension were recruited, and six were allocated randomly to each combination of treatments. Blood pressure measurements were made on each subject leading to the data shown in Table 4.1. Questions of interest concern differences in mean blood pressure for the different levels of the three treatments and the effects of interactions between the treatments on blood pressure.

4.2 Analysis of variance model

A suitable model for these data is

Table 4.1 Data in bp.raw

Biofeedback			No Biofeedback		
drug X	drug Y	drug Z	drug X	drug Y	drug Z
Diet absent					
170	186	180	173	189	202
175	194	187	194	194	228
165	201	199	197	217	190
180	215	170	190	206	206
160	219	204	176	199	224
158	209	194	198	195	204
Diet present					
161	164	162	164	171	205
173	166	184	190	173	199
157	159	183	169	196	170
152	182	156	164	199	160
181	187	180	176	180	179
190	174	173	175	203	179

$$y_{ijk} = \mu + \alpha_i + \beta_j + \gamma_k + (\alpha\beta)_{ij} + (\alpha\gamma)_{ik} + (\beta\gamma)_{jk} + (\alpha\beta\gamma)_{ijk} + \epsilon_{ijk} \quad (4.1)$$

where y_{ijk} represents the blood pressure of the i th subject for the j th drug, the j th level of biofeedback, and the k th level of diet, μ is the overall mean, α_i , β_j , and γ_k are the main effects for drugs, biofeedback, and diets, $(\alpha\beta)_{ij}$, $(\alpha\gamma)_{ik}$, and $(\beta\gamma)_{jk}$ are the first-order interaction terms, $(\alpha\beta\gamma)_{ijk}$ is a second-order interaction term, and ϵ_{ijk} are the residual or error terms assumed to be normally distributed with zero mean and variance σ^2 .

To identify the model, some constraints have to be imposed on the parameters. The standard constraints are:

$$\sum_i \alpha_i = \sum_j \beta_j = \sum_k \gamma_k = 0,$$

$$\begin{aligned} \sum_i (\alpha\beta)_{ij} &= \sum_j (\alpha\beta)_{ij} = \sum_i (\alpha\gamma)_{ik} = \sum_k (\alpha\gamma)_{ik} \\ &= \sum_j (\beta\gamma)_{jk} = \sum_k (\beta\gamma)_{jk} = 0, \end{aligned}$$

and

$$\sum_i (\alpha\beta\gamma)_{ijk} = \sum_j (\alpha\beta\gamma)_{ijk} = \sum_k (\alpha\beta\gamma)_{ijk} = 0.$$

We can test the following null hypotheses:

$$H_0^{(1)} : \text{No drug effect} : \alpha_1 = \alpha_2 = \alpha_3 = 0$$

$$H_0^{(2)} : \text{No biofeedback effect} : \beta_1 = \beta_2 = 0$$

$$H_0^{(3)} : \text{No diet effect} : \gamma_1 = \gamma_2 = 0$$

$$H_0^{(4)} : \text{No drug by biofeedback interaction} : \\ (\alpha\beta)_{ij} = 0, \quad i = 1, 2, 3; \quad j = 1, 2$$

$$H_0^{(5)} : \text{No drug by diet interaction} : (\alpha\gamma)_{ik} = 0, \quad i = 1, 2, 3; \quad k = 1, 2$$

$$H_0^{(6)} : \text{No biofeedback by diet interaction} : \\ (\beta\gamma)_{jk} = 0, \quad j = 1, 2; \quad k = 1, 2$$

$$H_0^{(7)} : \text{No drug by biofeedback by diet interaction} : \\ (\alpha\beta\gamma)_{ijk} = 0, \quad i = 1, 2, 3; \quad j = 1, 2; \quad k = 1, 2$$

Since there are an equal number of observations in each cell of Table 4.1, the total variation in the responses can be partitioned into non-overlapping parts (an orthogonal partition) representing main effects and interactions and residual variation as shown in Figure 4.1. *F*-tests can then be constructed for each hypothesis described above. More details can be found in Everitt (2001).

4.3 Analysis using Stata

Assuming the data are in an ASCII file `bp.raw`, exactly as shown in Table 4.1, i.e., 12 rows, the first containing the observations 170 186 180 173 189 202, they can be read into Stata by producing a dictionary file `bp.dct` containing the following statements:

```
dictionary using bp.raw {
    _column(6) int bp11
    _column(14) int bp12
    _column(22) int bp13
    _column(30) int bp01
```

Source	SS	df	$MS = \frac{SS}{df}$
Model	MSS	$abc - 1$	MMS (between cells)
A	SSA	$a - 1$	MSA
B	SSB	$b - 1$	MSB
C	SSC	$c - 1$	MSC
AB	SSAB	$(a - 1)(b - 1)$	MSAB
AC	SSAC	$(a - 1)(c - 1)$	MSAC
BC	SSBC	$(b - 1)(c - 1)$	MSBC
ABC	SSABC	$(a - 1)(b - 1)(c - 1)$	MSABC
Error (Residual)	SSE	$abc(n - 1)$	MSE (within cells) (RMS)
Total	TSS	$abcn - 1$	

Figure 4.1: ANOVA table for three-way factorial design with factors A, B, and C having a , b , and c levels, respectively and with n observations per cell

```
_column(38) int bp02
_column(46) int bp03
```

```
}
```

and using the following command

```
infile using bp, clear
```

Note that it was not necessary to define a dictionary here since the same result could have been achieved using a simple `infile` varlist command (see exercises). Here the variable names end on two digits, the first standing for the levels of biofeedback (1: present, 0: absent), and the second for the levels of drug (1,2,3 for X,Y,Z). The final dataset should have a single variable, `bp`, that contains all the blood pressures, and three additional variables, `drug`, `biofeed`, and `diet`, representing the corresponding levels of drug, biofeedback, and diet.

First, create `diet` which should take on one value for the first six rows and another for the following rows. This is achieved using the commands

```
generate diet = 0 if _n <= 6
replace diet = 1 if _n > 6
```

or, more concisely, using

```
generate diet = (_n > 6)
```

Now use the `reshape long` command to stack the columns on top of each other. If we specify `bp0` and `bp1` as the variable names in the `reshape` command, then `bp01`, `bp02`, and `bp03` are stacked into one column with variable name `bp0` (and similarly for `bp1`) and another

variable is created that contains the suffixes 1, 2, and 3. We ask for this latter variable to be called drug using the option j(drug) as follows:

```
generate id = _n
reshape long bp0 bp1, i(id) j(drug)
list in 1/9
```

(see Display 4.1). Here, id was generated because we needed to specify the row indicator in the i() option.

	id	drug	bp1	bp0	diet
1.	1	1	170	173	0
2.	1	2	186	189	0
3.	1	3	180	202	0
4.	2	1	176	194	0
5.	2	2	194	194	0
6.	2	3	187	226	0
7.	3	1	165	197	0
8.	3	2	201	217	0
9.	3	3	199	190	0

Display 4.1

We now run the reshape long command again to stack up the columns bp0 and bp1 and generate the variable biofeed. The instructions to achieve this and to label all the variables are given below.

```
replace id = _n
reshape long bp, i(id) j(biofeed)
replace id = _n

label define d 0 "absent" 1 "present"
label values diet d
label values biofeed d
label define dr 1 "Drug X" 2 "Drug Y" 3 "Drug Z"
label values drug dr
```

To begin, it will be helpful to look at some summary statistics for each of the cells of the design. A simple way of obtaining the required summary measures is to use the table command.

```
table drug, contents(freq mean bp median bp sd bp) ///
by(diet biofeed)
```

diet, biofeed and drug	Freq.	mean(bg)	med(bg)	sd(bg)
absent:				
absent:				
Drug X	6	188	193	10.66278
Drug Y	6	200	197	10.67968
Drug Z	6	209	205	14.3527
present:				
present:				
Drug X	6	168	167.5	8.602325
Drug Y	6	204	205	12.68069
Drug Z	6	189	190.5	12.61745
present:				
absent:				
Drug X	6	173	172	9.797959
Drug Y	6	187	188	14.01428
Drug Z	6	182	179	17.1114
present:				
present:				
Drug X	6	169	167	14.81891
Drug Y	6	172	170	10.93618
Drug Z	6	173	176.5	11.6619

Display 4.2

The standard deviations in Display 4.2 indicate that there are considerable differences in the within cell variability. This may have implications for the analysis of variance of these data: one of the assumptions made is that the observations within each cell have the same population variance. To begin, however, we will fit the model specified in Section 3.2 to the raw data using the *anova* command.

```
anova bp drug diet biofeed diet*drug diet*biofeed ///
    drug*biofeed drug*diet*biofeed
```

The resulting ANOVA table is shown in Display 4.3.

Source	Number of obs = 72		R-squared = 0.5840		Prob > F
	Partial SS	df	MS	F	
Model	13594	11	1199.45455	7.68	0.0000
drug	3675	2	1837.5	11.73	0.0001
diet	5202	1	5202	33.20	0.0000
biofeed	2048	1	2048	13.07	0.0005
diet*drug	903	2	451.5	2.88	0.0538
diet*biofeed	32	1	32	0.20	0.6529
drug*biofeed	259	2	129.5	0.83	0.4425
drug*diet*biofeed	1075	3	637.5	3.43	0.0388
Residual	9400	60	156.66667		
Total	23594	71	318.225352		

Display 4.3

The Root MSE is simply the square root of the residual mean square, with R-squared and Adj R-squared being as described in Chapter 3. The F-statistic of each effect represents the mean sum of squares for that effect, divided by the residual mean square, given under the heading MS. There are highly significant main effects of drug ($F_{2,60} = 11.73$, $p < 0.001$), diet ($F_{1,60} = 33.20$, $p < 0.001$), and biofeed ($F_{1,60} = 13.07$, $p < 0.001$). The two-way interactions are not significant at the 5% level but the three-way interaction drug by diet by biofeed is ($F_{2,60} = 3.43$, $p = 0.04$). The existence of a three-way interaction complicates the interpretation of the other terms in the model; it implies that the interaction between any two of the factors is different at the different levels of the third factor. Perhaps the best way of trying to understand the meaning of the three-way interaction is to plot a number of interaction diagrams; that is, plots of mean values for a factor

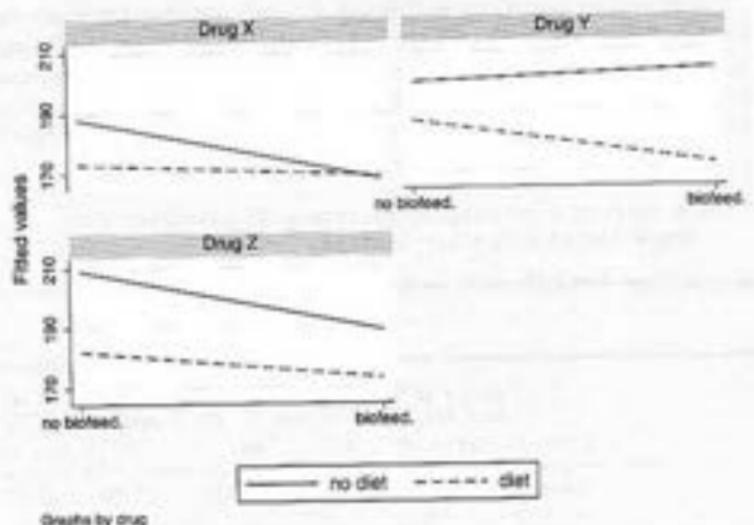


Figure 4.2: Interaction diagrams showing the interaction between diet and biofeedback for each level of drug.

at the different levels of the other factors.

This can be done by first creating a variable `predbp` containing the predicted means (which in this case coincide with the observed cell means because the model fitted is saturated, i.e., the number of parameters is equal to the number of cell means) using the command

```
predict predbp
```

Plots of `predbp` against `biofeed` for each level of drug with separate lines for `diet` can be obtained using the command

```
twoway (line predbp biofeed if diet==0, sort)
       (line predbp biofeed if diet==1, lpat(dash) sort), ///
       by(drug) xlabel(0 "no biofeed." 1 "biofeed.") ///
       ylabel(170 190 210) xtitle(" ")
       legend(order(1 "no diet" 2 "diet"))
```

The resulting interaction diagrams are shown in Figure 4.2. For drug Y, the presence of biofeedback increases the effect of diet (the vertical distance between the solid and dashed lines), whereas for drug Z the effect of diet is hardly altered by the presence of biofeedback and for drug X the effect is decreased.

Tables of the cell means plotted in the interaction diagrams, as well as the corresponding standard deviations, are produced for each drug using the following command:

```
table diet biofeed, contents(means bp sd bp) by(drug)
giving the output shown in Display 4.4.
```

drug and diet	biofeed	
	absent	present
Drug X	188 10.86278	168 8.602226
	173 9.797959	169 14.51891
Drug Y	200 10.07968	204 12.68069
	187 14.01428	172 10.93618
Drug Z	209 14.2527	189 12.65745
	182 17.11114	173 11.6619

Display 4.4

As mentioned previously, the observations in the 12 cells of the $3 \times 2 \times 2$ design have variances that differ considerably. Consequently, an analysis of variance of the data transformed in some way might be worth considering. For example, to analyze the log transformed observations, we can use the following commands:

```
generate lbp = log(bp)
anova lbp drug diet biofeed diet*drug diet*biofeed ///
drug*biofeed drug*diet*biofeed
```

The resulting analysis of variance table is shown in Display 4.5.

The results are similar to those for the untransformed blood pressures. The three-way interaction is only marginally significant. If no substantive explanation of this interaction is available, it might be bet-

Source	Number of obs =		72	R-squared	=	0.5776
	Root MSE	=	.068013	Adj R-squared	=	0.5002
	Partial SS	df	MS	F		Prob > F
Model	.379634762	11	.034503316	7.46		0.0000
diet	.149561559	1	.149561559	32.33		0.0000
drug	.107061236	2	.053530618	11.57		0.0001
biofeed	.081475807	1	.081475807	13.29		0.0006
diet*drug	.026011594	2	.013005797	2.60		0.0830
diet*biofeed	.000657678	1	.000657678	0.14		0.7975
drug*biofeed	.006467873	2	.0003233936	0.70		0.5050
diet*drug*biofeed	.030299315	2	.015149557	3.28		0.0447
Residual	.277545987	60	.004625766			
Total	.657060749	71	.009254558			

Display 4.5

ter to interpret the results in terms of the very significant main effects. The relevant summary statistics for the log transformed blood pressures can be obtained using the following instructions:

```
table drug, contents(mean lbp sd lbp)
```

```
table diet, contents(mean lbp sd lbp)
```

```
table biofeed, contents(mean lbp sd lbp)
```

giving the tables in Displays 4.6 to 4.8.

drug	mean(lbp)	sd(lbp)
Drug X	5.189352	.075965
Drug Y	5.247087	.0903679
Drug Z	5.232964	.0998621

Display 4.6

diet	mean(lbp)	sd(lbp)
absent	8.250651	.0915982
present	8.167498	.0781688

Display 4.7

biofeed	mean(lbp)	sd(lbp)
absent	8.242295	.0890138
present	8.183854	.0953618

Display 4.8

Drug X appears to produce lower blood pressures as does the special diet and the presence of biofeedback. Readers are encouraged to try other transformations.

Note that it is easy to estimate the model with main effects only using regression with dummy variables. Since drug has three levels and therefore requires two dummy variables, we save some time by using the *xi* prefix as follows:

```
xi: regress lbp i.drug i.diet i.biofeed
```

leading to the results shown in Display 4.9. The coefficients represent the mean differences between each level compared with the reference level (the omitted categories: drug X, diet absent, and biofeedback absent) when the other variables are held constant. The p-values are equal to those of ANOVA with main effects only, except that no overall p-value for drug is given. This can be obtained using

```
testparm _Idrug*
```

```
:1) _Idrug_2 = 0  
:2) _Idrug_3 = 0  
F( 2,    67) =   10.58  
Prob > F =    0.0001
```

The *F*-statistic is different from that in the last *anova* command because no interactions were included in the model; hence the residual degrees of freedom and the residual sum of squares have both increased.

i.drug	_Idrug_1-3	(naturally coded; _Idrug_1 omitted)			
i.diet	_Idiet_0-1	(naturally coded; _Idiet_0 omitted)			
i.biofeed	_Ibiofeed_0-1	(naturally coded; _Ibiofeed_0 omitted)			
Source	SS	df			
Model	.318098302	4			
Residual	.338982447	67			
Total	.657080749	71			
	MS				
		Number of obs = 72			
		F(4, 67) = 15.72			
		P>F = 0.0000			
		R-squared = 0.4841			
		Adj R-squared = 0.4633			
		Root MSE = .07113			
lbp	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
_Idrug_2	.0679354	.0205334	4.28	0.000	.0469506 -.1289303
_Idrug_3	.0738315	.0205334	3.60	0.001	.0328467 .1148163
_Idiet_1	-.0915536	.0167654	-5.44	0.000	-.1246176 -.0575896
_Ibiofeed_1	-.0584406	.0167654	-3.49	0.001	-.0919046 -.0249767
_cons	5.2330459	.0167643	279.23	0.000	5.196535 5.271363

Display 4.9

4.4 Exercises

4.1 • Treating hypertension

1. Reproduce the result of the command `infile` using `bp` without using the dictionary, and follow the `reshape` instructions to generate the required dataset.
2. Produce three diagrams with boxplots of blood pressure: (1) for each level of drug, (2) for each level of diet, and (3) for each level of biofeedback.
3. Investigate other possible transformations of the response variable and compare the resulting analyses of variance with those given in the text.
4. Suppose that in addition to the blood pressure of each of the individuals in the study, the investigator had also recorded their ages in the file `age.dat` as shown in Table 4.2 (but with data on one person per row). Reanalyze the data using `age` as a covariate (see `help merge` and `help anova`).

4.2 Auto pollution filter noise

The data used here are from Lewin and Shukun (1976) and the Data and Story Library (lib.stat.cmu.edu/DASL). They were originally used as part of a statement by Texaco to the Air and Water Pollution Subcommittee of the Senate Public Works Committee on June 26, 1973. Mr. John McKinley, President of Tex-

**Table 4.2 Data in
age.dat**

id	age	id	age
1	39	37	45
2	39	38	58
3	61	39	61
4	50	40	47
5	51	41	67
6	43	42	49
7	59	43	54
8	50	44	48
9	47	45	46
10	60	46	67
11	77	47	56
12	57	48	54
13	62	49	66
14	44	50	43
15	63	51	47
16	77	52	35
17	56	53	50
18	62	54	69
19	44	55	73
20	61	56	46
21	66	57	59
22	52	58	65
23	53	59	49
24	54	60	52
25	40	61	40
26	62	62	80
27	68	63	46
28	63	64	63
29	47	65	56
30	70	66	58
31	57	67	53
32	51	68	56
33	70	69	64
34	57	70	57
35	64	71	60
36	66	72	48

aco, cited an automobile filter developed by Associated Octel Company as effective in reducing pollution. However, questions had been raised about the effects of filters on vehicle performance, fuel consumption, exhaust gas back pressure, and silencing. On the last question, he referred to the data included here as evidence that the silencing properties of the Octel filter were at least equal to those of standard silencers.

The variables in *filters.dta* are:

- **noise**: noise level (in decibels)
 - **size**: vehicle size (1=small, 2=medium, 3=large)
 - **type**: type of filter (1=standard silencer, 2=Octel filter)
 - **side**: side of car (1=right side, 2=left side)
1. Produce tables of means and standard deviations of **noise** by **type** and **size**.
 2. Fit a two-way ANOVA model with **noise** as response variable and **type** and **size** as factors.
 3. Produce appropriate interaction diagrams and use them to interpret the results of the two-way ANOVA model.
 4. Now fit a three-way ANOVA model with **size**, **type**, and **side** as factors.
 5. Produce appropriate interaction diagrams and use them to interpret the results of the three-way ANOVA model.

4.3 Efficiency of cycling

Kapor (1981) investigated the effect of knee-joint angle on the efficiency of cycling. Efficiency was measured in terms of distance pedalled on an ergocycle until exhaustion. The experimenter selected three knee-joint angles of particular interest: 50, 70, and 90 degrees. Thirty subjects were available for the experiment and 10 subjects were randomly allocated to each angle. The drag of the ergocycle was kept constant at 14.7N, and subjects were instructed to pedal at a constant speed of 20km/h.

The variables in *cycling.dta* are:

- **id**: subject identifier
 - **group**: knee angle group (1=50 degrees, 2=70 degrees, 3=90 degrees)
 - **km**: distance pedalled (in km)
1. Carry out an initial graphical inspection of the data to assess whether there are any aspects of the observations that might be a cause for concern in later analyses.

2. Derive the appropriate analysis of variance table for the data.
3. Investigate differences in means between the three age populations in more detail using a suitable multiple comparison test (Hint: see help oneway).

4.4 Maternal behavior in rats

Here we consider data collected to study the maternal behavior of laboratory rats (Everitt, 2001). The response variable was the time (in seconds) required for the mother to retrieve the pup to the nest, after being moved a fixed distance away. In the study, three independent groups of pups of different ages (5 days, 20 days, and 35 days) were used.

The variables in `maternal.dta` are:

- `mother`: rat-mother identifier
 - `age`: age of pup (1=5 days, 2=20 days, 3=35 days)
 - `time`: time to retrieve the pup (in seconds)
1. Produce a table of means and standard deviations of `time` by `age`.
 2. Carry out a one way analysis of variance of the data.
 3. Use an orthogonal polynomial approach to investigate whether there is any evidence of a linear or quadratic trend in the group means. If the model is

$$y_{ij} = \mu + \alpha_i + \epsilon_{ij},$$

the linear contrast is $\alpha_3 - \alpha_1$ and the quadratic contrast is $\alpha_2 - (\alpha_1 + \alpha_3)/2$. Use the `anova` command, followed by `test`, `showorder` to find out the order of the columns in the design matrix. Then define a one-row matrix for each contrast with elements equal to the required contrast coefficients. Use the command `test` with the `mat` option to test the null hypotheses that the contrasts are zero (i.e., no linear trend and no quadratic trend, respectively).

Chapter 5

Analysis of Variance II: Effectiveness of Slimming Clinics

5.1 Description of data

Slimming clinics aim to help people lose weight by offering encouragement and support about dieting through regular meetings. A study was carried out to assess their effectiveness. Half of the clients participating in the study were randomly selected to receive a technical manual containing slimming advice based on psychological behaviorist theory to investigate if this would help them to control their diet. Some of the clients had previously tried to slim whereas others were novices. The data collected are shown in Table 5.1. (They are also given in Hand et al., 1994.) The response variable `resp` was defined as follows:

$$\frac{\text{weight after three months of treatment} - \text{ideal weight}}{\text{initial weight} - \text{ideal weight}} \quad (5.1)$$

The design can be thought of as a 2×2 factorial design where `manual` (1: received a manual, 2: did not) is crossed with `exper` (1: previous slimming experience, 2: novice). The number of observations in each cell of the design is not the same, so this is an example of an unbalanced 2×2 design.

Table 5.1 Data in slim.dat

exper	manual	resp	exper	manual	resp
1	1	-14.67	1	1	-1.83
1	1	-8.55	1	1	-23.03
1	1	11.61	1	2	0.81
1	2	2.38	1	2	2.74
1	2	3.36	1	2	2.10
1	2	-0.83	1	2	-3.05
1	2	-5.98	1	2	-3.64
1	2	-7.38	1	2	-3.60
1	2	-0.94	2	1	-3.39
2	1	-4.00	2	1	-2.31
2	1	-3.60	2	1	-7.69
2	1	-13.92	2	1	-7.64
2	1	-7.59	2	1	-1.62
2	1	-12.21	2	1	-8.85
2	2	5.84	2	2	1.71
2	2	-4.10	2	2	-5.19
2	2	0.00	2	2	-2.80

5.2 Analysis of variance model

A suitable analysis of variance model for the data is

$$y_{ijk} = \mu + \alpha_i + \beta_j + \gamma_{ij} + \epsilon_{ijk} \quad (5.2)$$

where y_{ijk} represents the weight change of the k th individual having experience status j and manual condition i , μ is the overall mean, α_i represents the effect of manual condition i , β_j the effect of experience status j , γ_{ij} the experience \times manual interaction, and ϵ_{ijk} the errors. The errors are assumed to have a normal distribution with mean zero and variance σ^2 .

The unbalanced nature of the slimming data presents some difficulties for analysis not encountered in factorial designs having the same number of observations in each cell (see the previous chapter). If the data were balanced, the among cells sum of squares would partition orthogonally into three component sums of squares representing the two main effects and their interaction. However, with unbalanced data there is no unique way of finding a "sum of squares" corresponding to each main effect and their interactions, because these effects are no longer independent of one another. Several methods have been proposed for dealing with this problem and each leads to a different partition of the overall sum of squares. The different methods for arriving at the sums of squares for unbalanced designs can be explained in terms of the comparisons of different sets of specific models. For a

design with two factors A and B, Stata can calculate sequential sums of squares or unique sums of squares as described in the next subsections.

5.2.1 Sequential sums of squares

Sequential sums of squares (also known as hierarchical sums of squares) represent the effect of adding a term to an existing model. So, for example, a set of sequential sums of squares such as

Source	SS
A	SS(A)
B	SS(B A)
AB	SS(AB A,B)

represent a comparison of the following models:

- SS(AB|A,B)—model including an interaction and main effects compared with one including only main effects.
- SS(B|A)—model including both main effects, but with no interaction, compared with one including only the main effects of factor A.
- SS(A)—model containing only the main effect of A compared with one containing only the overall mean.

The use of these sums of squares in a series of tables in which the effects are considered in different orders (see later) will often provide the most satisfactory way of deciding which model is most appropriate for the observations. (These are SAS Type I sums of squares—see Der and Everitt, 2002.)

5.2.2 Unique sums of squares

By default, Stata produces unique sums of squares that represent the contribution of each term to a model including all the other terms. So, for a two-factor design, the sums of squares represent the following.

Source	SS
A	SS(A B,AB)
B	SS(B A,AB)
AB	SS(AB A,B)

(These are SAS Type III sums of squares.) Note that these sums of squares generally do not add up to the model sums of squares.

5.2.3 Regression

As we have shown in Chapter 4, ANOVA models may also be estimated using regression by defining suitable dummy variables. Assume that A is represented by a single dummy variable. The regression coefficient for A represents the partial contribution of that variable, adjusted for all other variables in the model, say B. This is equivalent to the contribution of A to a model already including B. A complication with regression models is that, in the presence of an interaction, the p-values of the terms depend on the exact coding of the dummy variables (see Aitkin, 1978). The unique sums of squares correspond to regression where dummy variables are coded in a particular way, for example a two-level factor can be coded as -1, 1.

There have been numerous discussions over which sums of squares are most appropriate for the analysis of unbalanced designs. The Stata manual appears to recommend its default for general use. Nelder (1977) and Aitkin (1978), however, are strongly critical of "correcting" main effects for an interaction term involving the same factor; their criticisms are based on both theoretical and pragmatic arguments and seem compelling. A frequently used approach is therefore to test the highest order interaction adjusting for all lower order interactions and not vice versa. Both Nelder and Aitkin prefer the use of Type I sums of squares in association with different orders of effects as the procedure most likely to identify an appropriate model for a data set. For a detailed explanation of the various types of sums of squares, see Boniface (1995).

5.3 Analysis using Stata

The data can be read in from an ASCII file `slim.dat` in the usual way using

```
infile manual exper resp using slim.dat
```

A table showing the unbalanced nature of the 2×2 design can be obtained from

manual	exper		Total
	1	2	
1	5	12	17
2	11	6	17
Total	16	18	34

We now use the `anova` command with no options specified to obtain the unique (Type III) sums of squares:

`anova resp manual exper manual*exper`
 (see Display 5.1).

Source	Number of obs =		34	R-squared	= 0.2103
	Partial SS	df	MS	F	Prob > F
Model	287.331861	3	95.7439637	2.66	0.0699
manual	2.19850409	1	2.19850409	0.06	0.8064
exper	265.871053	1	265.871053	7.39	0.0108
manual*exper	.130318264	1	.130318264	0.00	0.9524
Residual	1078.54812	30	35.961604		
Total	1366.07999	33	41.2863631		

Display 5.1

Our recommendation is that the sums of squares shown in this table are not used to draw inferences because the main effects have been adjusted for the interaction.

Instead we prefer an analysis that consists of obtaining two sets of sequential sums of squares, the first using the order `manual exper manual*exper` and the second the order `exper manual manual*exper`. The necessary instructions are

`anova resp manual exper manual*exper, sequential`
 (see Display 5.2).

`anova resp exper manual manual*exper, sequential`

(see Display 5.3). The sums of squares corresponding to model and residuals are, of course, the same in both tables, as is the sum of squares for the interaction term. What differ are the sums of squares in the `manual` and `exper` rows in the two tables. The terms of most interest are the sum of squares of `exper|manual` which is obtained from the table as 265.91, and the sum of squares of `manual|exper` which is 2.13. These sums of squares are less than the sums of squares for `exper` and `manual` alone (284.97 and 21.19, respectively), by an amount of 19.06, a portion of the model sums of squares which cannot be uniquely attributed to either of the variables. The associated *F*-tests in the two tables make it clear that there is no interaction effect and that `exper|manual` is significant but `manual|exper` is not. The conclusion is

Source			R-squared	Adj R-squared	F	Prob > F
	Seq. SS	df				
Model	387.235861	3	98.7439537	0.66	0.0859	
manual	21.1878098	1	21.1878098	0.59	0.4487	
exper	265.913733	1	265.913733	7.39	0.0106	
manual*exper	.130318264	1	.130318264	0.00	0.9624	
Residual	1078.84812	30	35.961604			
Total	1366.07998	33	41.3963631			

Display 5.2

Source			R-squared	Adj R-squared	F	Prob > F
	Seq. SS	df				
Model	387.235861	3	98.7439537	0.66	0.0859	
exper	264.971071	1	264.971071	7.92	0.0065	
manual	2.13047169	1	2.13047169	0.06	0.8004	
manual*exper	.130318264	1	.130318264	0.00	0.9624	
Residual	1078.84812	30	35.961604			
Total	1366.07998	33	41.3963631			

Display 5.3

that only *exper*, i.e., whether the woman had been slimming for over one year, is important in determining weight change. Provision of the manual appears to have no discernible effect. Figure 5.1 illustrates the two ways of partitioning the model sums of squares into components due to *exper* (large circle) and *manual* (small circle), depending on the order in which the main effects are entered.

Results equivalent to the unique (Type III) sums of squares can be obtained using regression:

```
generate manual1 = manual
recode manual1 1=-1 2=1
generate exper1 = exper
recode exper1 1=-1 2=1
generate exp_man = manual1*exper1
regress res manual1 exper1 exp_man
```

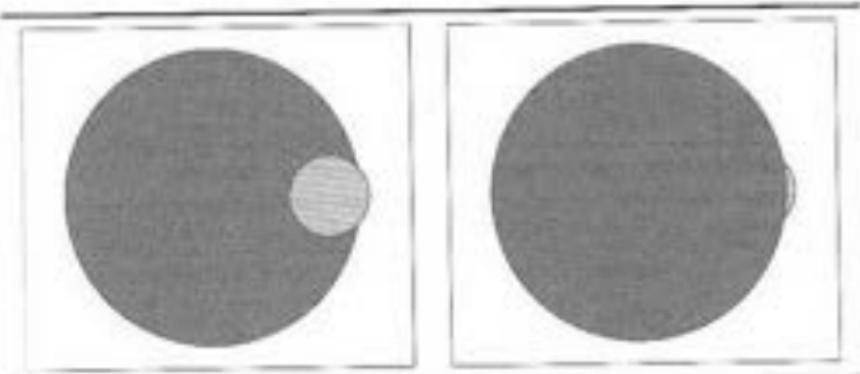


Figure 5.1: Venn diagrams showing sequential sums of squares for weight loss data. The large circle represents `exper` and the small circle `manual`. In the left panel `manual` enters the model first, and in the right panel `exper` enters the model first.

(see Display 5.4). The p-values agree with those based on unique sums

Source	SS	df	MS	Number of obs =	34
Model	287.231861	3	95.7439537	F(3, 30) =	2.66
Residual	1078.84812	30	35.943604	Prob > F =	0.0659
Total	1366.07998	33	41.3963631	R-squared =	0.2103
				Adj R-squared =	0.1313
				Root MSE =	5.5968
resp	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
manual	-2726251	1.102609	-0.26	0.806	-1.979204 2.524454
exper1	2.888043	1.102609	2.72	0.011	.746213 5.24987
exp_mn	-0.000375	1.102609	-0.05	0.952	-2.318204 2.188454
cons	-3.962958	1.102609	-3.59	0.001	-6.212787 -1.70913

Display 5.4

of squares. However, these results differ from the regression used by Stata's `anova` with the option `regress`:

```
anova resp manual exper manual*exper, regress
```

SOURCE	SS	df	MS				Number of obs = 34
Model	287.231863	3	95.7439537	P(3, 30) =	2.66		
Residual	5078.84812	30	35.961604	Prob > F =	0.0659	R-squared =	0.2100
Total	1366.07993	33	41.3963631	Adj R-squared =	0.1313	Root MSE =	5.9968
resp	Coeff.	Std. Err.	t	P> t		[95% Conf. Interval]	
cons	-75655666	2.448183	-3.1	0.759	-5.756534	4.24319	
manual							
1	-4125001	2.9884	-0.14	0.891	-6.536049	5.711049	
2	(dropped)						
exper							
1	-5.863333	3.043493	-1.83	0.064	-12.07897	.3523044	
2	(dropped)						
manual*exper							
1 1	-28555002	4.410437	-0.66	0.952	-9.372815	8.741815	
1 2	(dropped)						
2 1	(dropped)						
2 2	(dropped)						

Display 5.5

(see Display 5.5) because this uses different dummy variables, coded as 1 for the levels shown on the left of the reported coefficient and 0 otherwise, i.e., the dummy variable for manual*exper is 1 when exper and manual are both 1.

A table of mean values helpful in interpreting these results can be found using

```
table manual exper, content(mean resp) row col f(18.2f)
```

manual	exper		
	1	2	Total
1	-7.30	-1.17	-2.97
2	-6.62	-0.76	-4.55
Total	-6.93	-1.03	-3.76

The means demonstrate that experienced slimmers achieve the greatest weight reduction.

5.4 Exercises

5.1 • Effectiveness of slimming clinics

1. Investigate what happens to the sequential sums of squares if the `manual*exper` interaction term is given before the main effects `manual` and `exper` in the `anova` command with the `sequential` option.
2. Use `regress` to reproduce the analysis of variance by coding both `manual` and `exper` as (0,1) dummy variables and creating an interaction variable as the product of these dummy variables.
3. Use `regress` in conjunction with `xi:` to fit the same model without the need to generate any dummy variables.
4. Reproduce the results of `anova resp manual exper manual *exper`, `regress` using `regress` by making `xi:` omit the last category instead of the first (see `help xi`, under "Summary of controlling the omitted dummy").

See also the Exercises in Chapters 7 and 13.

5.2 Systolic blood pressure

Boniface (1995) provides data from Maxwell and Delaney (1990) on systolic blood pressure of individuals, classified according to their smoking status and family history of circulation and heart problems.

The variables in the dataset `systolic.dta` are:

- `history`: family history (0=no, 1=yes)
- `smoking`: smoking status
(1=non-smoker, 2=ex-smoker, 3=current smoker)
- `systolic`: systolic blood pressure

1. Carry out an analysis of variance of the data, retaining the interaction only if it is significant at the 5% level.
2. Produce an appropriate graph for interpreting the analysis of variance results and state your conclusions.
3. Examine the residuals from fitting what you consider the most suitable model for the data, and use various plots to assess the assumptions of the analyses you have performed (see `help anova postestimation`).

5.3 Role-taking in young children

Klemchuk et al. (1990) studied role-taking in children. In their study, children between the ages of 2 and 5 years were administered a battery of role-taking tasks. Participants were classified into a group who had had no previous daycare experience and a

group who had extensive daycare experience. Children were also sorted into two age groups (2 to 3 years and 4 to 5 years). The investigators' hypothesis was that children with daycare experience would perform better on role-taking tasks than would children without daycare experience because of the former group's greater opportunity for social development. There was also interest in the effect of age and the possible interaction of age with daycare group. The dependent variable was a role-taking score with higher values representing better performance.

The variables in `role.dta` are:

- `id`: child identifier
 - `daycare`: dummy variable for having had previous experience of daycare
 - `age`: age group (0=2 to 3 years, 1=4 to 5 years)
 - `role`: score for performance on role-playing tasks
1. Generate a variable equal to the mean of `role` for each child's age and daycare group and produce line graphs of these means versus age with separate lines for the daycare and no daycare groups.
 2. Fit a two-way ANOVA model to the data. Simplify the model as much as possible using a 5% significance level.
 3. For the chosen model, plot a line graph of the model-implied means (analogous to question 1). Also add the corresponding sample means represented by dots.

Chapter 6

Logistic Regression: Treatment of Lung Cancer and Diagnosis of Heart Attacks

6.1 Description of data

Two datasets will be analyzed in this chapter. The first dataset shown in Table 6.1 originates from a clinical trial in which lung cancer patients were randomized to receive two different kinds of chemotherapy (sequential therapy and alternating therapy). The outcome was classified into one of four categories: progressive disease, no change, partial remission, or complete remission. The data were published in Holt-brugge and Schumacher (1991) and also appear in Hand et al. (1994). The central question is whether there is any evidence of a difference in the outcomes achieved by the two types of therapy.

Table 6.1 Lung cancer data in tumor.dat

Therapy	Sex	Progressive disease	No change	Partial remission	Complete remission
Sequential	Male	28	45	20	26
	Female	4	12	5	2
Alternating	Male	41	44	20	29
	Female	12	7	3	1

Table 6.2 Data in sck.dat

Maximum CK level	Infarct present	Infarct absent
0 - 30	3	88
40 - 79	13	26
80 - 119	30	8
120 - 159	30	5
160 - 199	21	0
200 - 239	19	1
240 - 279	18	1
280 - 319	13	1
320 - 359	19	0
360 - 399	15	0
400 - 439	7	0
440 - 479	8	0
480 -	35	0

The second dataset to be used in this chapter arises from a study investigating the use of serum creatine kinase (CK) levels for the diagnosis of myocardial infarction (heart attack). Patients admitted to a coronary care unit because they were suspected of having had a myocardial infarction within the last 48 hours had their CK levels measured on admission and the next two mornings. A clinician who was "blind" to the CK results came to an independent "gold standard" diagnosis using electrocardiograms, clinical records, and autopsy reports. The maximum CK levels for 360 patients are given in Table 6.2 together with the clinician's diagnosis. The table was taken from Sackett et al. (1991) (with permission of the publisher, Little Brown & Company), where only the ranges of CK levels were given, not their precise values.

The main questions of interest for this second dataset are how well CK discriminates between those with and without myocardial infarction, and how diagnostic tests perform that are based on applying different thresholds to CK.

6.2 The logistic regression model

6.2.1 Binary responses

Dichotomous or binary responses arise when the outcome is presence or absence of a characteristic or event, for example myocardial infarction in the second dataset. What we would like to do is to investigate the effects of a number of explanatory variables on this binary response variable. This appears to be the same aim as for multiple regression.

discussed in Chapter 3, where the model for a response y_i and explanatory variables x_{i1} to x_{ip} can be written as

$$\mu_i \equiv E(y_i | \mathbf{x}_i) = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} \quad (6.1)$$

$$y_i \sim N(\mu_i, \sigma^2)$$

Binary responses are typically coded 1 for the event of interest, such as infarct present, and 0 for the opposite event. In this case the expected value is simply the probability π_i that the event of interest occurs. This raises the first problem for applying the model above to a binary response variable, namely that

1. the predicted probability must satisfy $0 \leq \pi_i \leq 1$ whereas the linear model above can yield any value from minus infinity to plus infinity.
- A second problem with using linear regression is that
2. the observed values of y_i do not follow a normal distribution with mean π_i , but rather a Bernoulli (or Binomial(1, π_i)) distribution.

Consequently a new approach is needed, and this is provided by logistic regression.

In logistic regression, the first problem is addressed by replacing the probability $\pi_i = E(y_i | \mathbf{x}_i)$ on the left-hand side of equation (6.1) by the logit of the probability, giving

$$\text{logit}(\pi_i) = \log(\pi_i / (1 - \pi_i)) = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_p x_{ip}. \quad (6.2)$$

The logit of the probability is simply the log of the odds of the event of interest. Writing $\boldsymbol{\beta}$ and \mathbf{x}_i for the column vectors $(\beta_0, \beta_1, \dots, \beta_p)'$ and $(1, x_{i1}, \dots, x_{ip})'$, respectively, the probability as a function of the covariates is

$$\pi_i = \frac{\exp(\mathbf{x}_i' \boldsymbol{\beta})}{1 + \exp(\mathbf{x}_i' \boldsymbol{\beta})} = \frac{1}{1 + \exp(-\mathbf{x}_i' \boldsymbol{\beta})}. \quad (6.3)$$

When the logit takes on any real value, this probability always satisfies $0 \leq \pi_i \leq 1$. This is illustrated for a single covariate in Figure 6.1 where the logistic function is shown along with a linear function. For π between 0.2 and 0.8, both functions are similar, but as π approaches 0 and 1, the logistic function curve flattens, producing an "S" shape.

The second problem relates to the estimation procedure. Whereas maximum likelihood estimation in conventional linear regression leads to least squares, this is not the case in logistic regression. In logistic regression the log likelihood is maximized numerically using an iterative algorithm. For full details of logistic regression, see for example

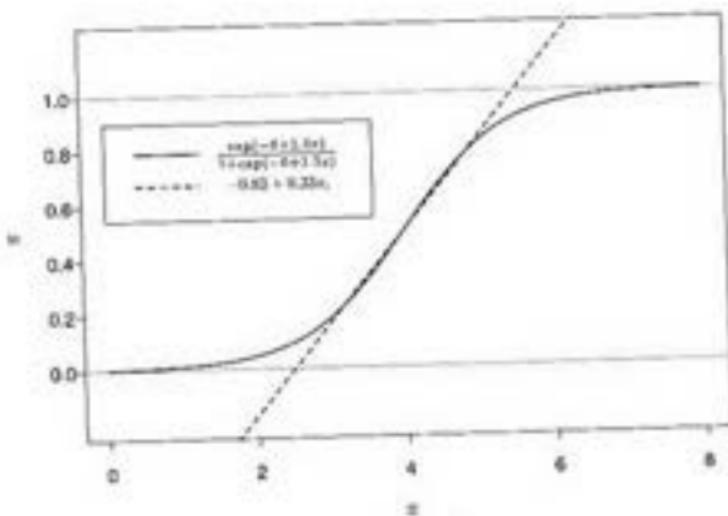


Figure 6.1: Linear and logistic functions of x .

Collett (2002), Agresti (1996), and Long and Freese (2006). (The last reference provides a comprehensive discussion of regression models for categorical variables using Stata.)

6.2.2 Ordinal responses

Logistic regression can be generalized to the situation where the response variable has more than two ordered response categories. In the latent response formulation, we think of the ordered categories as representing successive intervals of an underlying latent (unobserved) continuous response. If there are S response categories labeled a_1, \dots, a_S , the relationship between the observed and latent response can be formulated as a threshold model:

$$y_i = \begin{cases} a_1 & \text{if } y_i^* \leq \kappa_1 \\ a_2 & \text{if } \kappa_1 < y_i^* \leq \kappa_2 \\ \vdots & \vdots \\ a_S & \text{if } \kappa_{S-1} < y_i^*. \end{cases}$$

where κ_s , $s = 1, \dots, S - 1$ are threshold or cut-point parameters. The latent response is modeled as a linear regression

$$y_i^* = \mathbf{x}_i' \boldsymbol{\beta} + \epsilon_i,$$

where ϵ_i has a logistic distribution,

$$\Pr(\epsilon_i \leq t) = \frac{\exp(t)}{1 + \exp(t)}, \quad \Pr(\epsilon_i > t) = \frac{\exp(-t)}{1 + \exp(-t)}.$$

The latent response and threshold model imply a logistic model for the cumulative probabilities. The cumulative probability γ_{is} that the response y_i takes on a value greater than a_s becomes

$$\begin{aligned} \gamma_{is} &\equiv \Pr(y_i > a_s) = \Pr(y_i^* > \kappa_s) = \Pr(\underbrace{y_i^* - \mathbf{x}_i' \boldsymbol{\beta}}_{\epsilon_i} > \kappa_s - \mathbf{x}_i' \boldsymbol{\beta}) \\ &= \frac{\exp(\mathbf{x}_i' \boldsymbol{\beta} - \kappa_s)}{1 + \exp(\mathbf{x}_i' \boldsymbol{\beta} - \kappa_s)}, \quad s = 1, \dots, S - 1. \end{aligned} \quad (6.4)$$

It is clear from this expression that we could not simultaneously estimate the constant β_0 in the model for y_i^* and all threshold parameters since we could increase the constant and all thresholds by the same amount without changing the model. In Stata the constant is therefore set to zero for identification.

The model is also called the proportional odds model because the log odds that $y_i > a_s$ are

$$\log \left(\frac{\gamma_{is}}{1 - \gamma_{is}} \right) = \mathbf{x}_i' \boldsymbol{\beta} - \kappa_s \quad (6.5)$$

so that the log odds ratio for two units i and j is $(\mathbf{x}_i - \mathbf{x}_j)' \boldsymbol{\beta}$ which is independent of s . Therefore, $\exp(\beta_k)$ represents the odds ratio that $y > a_s$ for any s when x_k increases by one unit if all other covariates remain the same.

In binary logistic regression for dichotomous responses, $a_1 = 0$, $a_2 = 1$, $\kappa_1 = 0$, and $\exp(\beta_k)$ is the odds ratio that $y = 1$ when x_k increases by one unit and all other covariates remain the same. Note that a different identifying restriction is used than for ordinal responses: the threshold κ_1 is set to zero instead of the constant β_0 in the model for y_i^* .

The probit and ordinal probit models correspond to logistic and ordinal logistic regression models with the cumulative distribution function in (6.4) replaced by the standard normal cumulative distribution function. More information on models for ordinal data can be found in Agresti (1996) and Long and Freese (2006).

6.3 Analysis using Stata

6.3.1 Chemotherapy treatment of lung cancer

The ASCII file `tumor.dat` contains the four by four matrix of frequencies shown in Table 6.1. First we read the data and generate variables for therapy and sex using the `ege`n function `seq()`:

```
infile fr1 fr2 fr3 fr4 using tumor.dat
ege therapy = seq(), from(0) to(1) block(2)
ege sex = seq(), from(1) to(2) by(therapy)
label define t 0 seq 1 alt
label values therapy t
label define s 1 male 2 female
label values sex s
```

`block(2)` causes the number in the sequence (from 0 to 1) to be repeated in blocks of two, whereas `by(therapy)` causes the sequence to start from the lower limit every time the value of `therapy` changes.

We next reshape the data to long, placing the four levels of the outcome, represented by `f1` to `f4` into a variable `outc`,

```
reshape long fr, i(therapy sex) j(outc)
```

and expand the dataset by replicating each observation `freq` times so that we have one observation per subject

```
expand fr
```

We can check that the data conversion is correct by tabulating these data as in Table 6.1:

```
table sex outc, contents(freq) by(therapy)
```

giving the table in Display 6.1.

therapy and sex	outc			
	1	2	3	4
seq				
male	28	45	29	26
female	4	12	5	2
alt				
male	41	44	20	20
female	12	7	3	1

Display 6.1

To use ordinary (binary) logistic regression, we must dichotomize the outcome, for example, by considering partial and complete remission to be an improvement and the other categories to be no improvement. The new outcome variable can be generated as follows:

```
recode outc 1/2=0 3/4=1, generate(improve)
```

or using

```
generate improve = outc>2
```

The command **logit** for logistic regression behaves the same way as **regress** and all other estimation commands. For example, automatic selection procedures can be carried out using the **stepwise** prefix and post-estimation commands such as **testparm** and **predict** are available. First, include **therapy** as the only explanatory variable:

```
logit improve therapy
```

(see Display 6.2). The algorithm takes two iterations to converge. The

Iteration 0:	log likelihood = -194.40888				
Iteration 1:	log likelihood = -192.30753				
Iteration 2:	log likelihood = -192.30471				
<hr/>					
Logistic regression					
Number of obs = 299					
LR chi2(1) = 4.21					
Prob > chi2 = 0.0403					
Pseudo R2 = 0.0108					
<hr/>					
improve	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
therapy	-0.4988993	.2443506	-2.04	0.041	-.977618 -.0197805
_cons	-.361502	.1654236	-2.19	0.029	-.6857263 -.0372777

Display 6.2

coefficient of **therapy** represents the difference in the log odds (of an improvement) between the alternating and sequential therapies. The negative value indicates that sequential therapy is superior to alternating therapy. The *p*-value of the coefficient is 0.041 in the table. This was derived from the *z* statistic, which is given by the coefficient divided by its asymptotic standard error (Std. Err.). Under the null hypothesis that the true coefficient is zero, the statistic has a standard normal distribution, and its square, the Wald statistic, has a χ^2 -distribution with one degree of freedom. This *p*-value from this Wald test is less reliable than the *p*-value based on the likelihood ratio

between the model including only the constant and the current model. Under the null hypothesis that the constant-only model is correct, minus twice the likelihood ratio has an approximate χ^2 -distribution with one degree of freedom (because there is one additional parameter in the competing model). Here the likelihood ratio statistic is equal to 4.21 giving a p -value of 0.0402, very similar to that based on the Wald test.

The coefficient of therapy represents the difference in log odds between the therapies and is not easy to interpret apart from the sign. Exponentiating the coefficient gives the odds ratio and exponentiating the 95% confidence limits gives the confidence interval for the odds ratio. Fortunately, the or option can be used to obtain the required odds ratio and its confidence interval directly (alternatively, we could use the logistic command):

```
logit improve therapy, or
```

(see Display 6.3). The standard error now represents the approximate

Logistic regression				Number of obs	=	299
				LR chi2(1)	=	4.21
				Prob > chi2	=	0.0402
				Pseudo R2	=	0.0108
Log likelihood = -192.30471						
improve	Odds Ratio	Std. Err.	Z	P> z	[95% Conf. Interval]	
therapy	.6073201	.1483991	-2.04	0.041	.3762061	.9604138

Display 6.3

standard error of the odds ratio (calculated using the delta method, see, e.g., Agresti, 2002). Since the sampling distribution of the odds ratio is not well approximated by a normal distribution, the Wald statistic and confidence interval are derived using the log odds and its standard error. Alternating therapy is associated with a $100(1 - 0.6073201)\% = 39\%$ reduction in the odds of an improvement compared with sequential therapy (95% confidence interval from 2% to 62%).

To test whether the inclusion of sex in the model significantly increases the likelihood, the current likelihood (and all the estimates) can be saved using

```
estimates store model1
Including sex
logit improve therapy sex, or
```

gives the output shown in Display 6.4. The *p*-value of *sex* based on

Logistic regression					Number of obs	=	299
					LR chi2(2)	=	7.55
					Prob > chi2	=	0.0229
					Pseudo R2	=	0.0194
Log likelihood = -190.43171							
improve	Odds Ratio	Std. Err.	Z	P> z	[95% Conf. Intervall]		
therapy	.0051969	.1486907	-2.04	0.041	.3739054	.9796537	
sex	.5197993	.1930918	-1.76	0.078	.2509785	1.076551	

Display 6.4

the Wald-statistic is 0.078, and a *p*-value for the likelihood-ratio test is obtained using lrtest

```
lrtest model1 .
Likelihood-ratio test
Assumption: model1 nested in .
LR chi2(1) =      3.35
Prob > chi2 =    0.0674
```

which is not very different from the value of 0.078. In the lrtest command ":" refers to the current model and model1 is the model excluding *sex* which was previously stored using estimates store. We could have specified the models in the reverse order as Stata assumes that the model with the lower log likelihood is nested within the other model. Note that it is essential that both models compared in the likelihood ratio test be based on the same sample. If *sex* had missing values, fewer observations would contribute to the model including *sex* than to the nested model excluding *sex*. In this case, we would have to restrict estimation of the nested model to the "estimation sample" of the full model. If the full model has been estimated first, this can be achieved using logistic improve therapy if e(sample).

Retaining the variable *sex* in the model (although it is not significant at the 5% level), the predicted probabilities can be obtained using predict with the pr option

```
predict prob, pr
```

and the four different predicted probabilities may be compared with the observed proportions as follows:

```
table sex, contents(mean prob mean improve freq) ///
by(therapy)
```

therapy and sex	mean(prob)	mean(improve)	Freq.
rec	.4332747	.4296875	128
male	.2843846	.3045478	23
alt	.3163368	.32	125
male	.1938763	.1739113	23

Display 6.5

(see Display 6.5). The agreement is good, so there appears to be no strong interaction between sex and type of therapy. (We could test for an interaction between sex and therapy by using `xi: logistic improve i.therapy*i.sex`.) A more formal assessment of the goodness of fit can be obtained using the command

```
estat gof, table
```

which produces the output shown in Display 6.6. There are four unique covariate patterns given by the combinations of `therapy` and `sex`. The individuals sharing a given covariate pattern are referred to as groups; the last two columns show the corresponding covariate values and the second column gives the predicted probabilities. The observed number of cases in each group is given under `Total` and this is composed of `Obs.1` individuals with a response `improve` equal to 1 and `Obs.2` individuals with a response equal to 0. The expected number of individuals with a 1 response, `Exp.1`, is obtained by multiplying the total number of individuals in each group by the predicted probability for the group. Finally, the expected number of individuals with a 0 response is just `Total - Exp.1`. At the bottom of the output, observed and expected frequencies are compared using Pearson chi-square test. The null hypothesis that the model is correct cannot be rejected here with a *p*-value of 0.73. Note that the test cannot detect if there are important omitted covariates because the table of observed and expected frequencies is aggregated over any covariates not included in the model.

We now fit the proportional odds model using the full ordinal response variable `cutc`:

```
ologit cutc therapy sex, or
```

The results are shown in Display 6.7. The odds ratios represent the estimated effects of `therapy` and `sex` on the odds of being in "com-

Logistic model for improve... goodness-of-fit test

Group	Prob	Obs_1	Exp_1	Obs_0	Exp_0	Total
1	0.1939	4	4.5	19	18.5	23
2	0.2844	7	6.5	16	16.5	23
3	0.3163	40	39.5	85	85.5	126
4	0.4333	55	55.5	73	72.5	128

Group	Prob	therapy	sex
1	0.1939	alt	female
2	0.2844	seq	female
3	0.3163	alt	male
4	0.4333	seq	male

number of observations = 299
 number of covariate patterns = 4
 Pearson chi2(1) = 0.12
 Prob > chi2 = 0.7350

Display 6.6

Ordered logistic regression		Number of obs = 299			
		LR chi2(2)	= 10.91		
		Prob > chi2	= 0.0043		
		Pseudo R2	= 0.0136		
Log likelihood = -394.52832					
outc	Odds Ratio	Std. Err.	z	P> z	[95% Conf. Intervall]
therapy	.559516	.1186973	-2.74	0.006	.3693774 .9479653
sex	.5319366	.1871185	-1.89	0.069	.3314596 1.021694
/cut1	-1.880537	.3820641			-3.609837 -1.109037
/cut2	-2.921603	.3672626			-1.011982 .4276611
/cut3	.758662	.3741486			.0253441 1.49198

Display 6.7

plete remission" (category 4) versus being at best in "partial remission" (categories 1 to 3); or of being at least in "partial remission" (categories 3 and 4) rather than having "progressive disease" or "no change" (categories 1 and 2); or of "no change" or better (categories 2 to 4) versus "progressive disease" (category 1). The second interpretation corresponds to the dichotomization used to fit the binary logistic regression model. Therefore the odds ratio estimates should be similar to those using binary logistic regression shown in Display 6.4 and they are. However, the p-values are lower in the ordinal logistic regression as might be expected because information is lost in dichotomizing the outcome.

We can use the estimates to calculate predicted probabilities. For instance, the predicted probability that a male (*sex*=1) who is receiving sequential therapy (*therapy*=0) will be in complete remission (*outc*=4) is $\bar{\gamma}_3$ (see equation (6.4)):

```
display .5819366*exp(-0.758662)/(1+.5819366*exp(-0.758662))
.21415663
```

However, a much quicker way of computing the predicted probabilities for all four responses and all combinations of explanatory variables is to use the *predict* command:

```
predict p1 p2 p3 p4
```

and to tabulate the results as follows:

```
table sex, contents(mean p1 mean p2 mean p3 mean p4) ///
by(therapy)
```

giving the table in Display 6.8.

therapy and sex	mean(p1)	mean(p2)	mean(p3)	mean(p4)
seq				
male	.2111641	.3506436	.2336566	.2141566
alt				
male	.3150425	.3729235	.178154	.1368799
alt				
female				
seq				
female	.3150425	.3729235	.178154	.1368799
alt				
female				

Display 6.8

6.3.2 Diagnosis of heart attacks

The data in `sck.dat` are read in using

```
infile ck pres abs using sck.dat, clear
```

Each observation represents all subjects with maximum creatine kinase values in the same interval and the variable `ck` contains the lower limits of the intervals. The total number of subjects is `pres+abs`, calculated using

```
generate tot = pres + abs
```

and the number of subjects with the disease is `pres`. The probability of `pres` "successes" in `tot` trials is binomial with "denominator" `tot` and probability π_i , `Binomial(tot, pi_i)`. The programs `logit` and `logistic` are for data where each observation represents a single Bernoulli trial, with binomial "denominator" equal to 1, `Binomial(1, pi_i)`. Another command, `blogit`, can be used to analyze the "grouped" data with "denominators" `tot` as considered here:

```
blogit pres tot ck
```

(see Display 6.9). There is a very significant association between CK

Logistic regression for grouped data						Number of obs	=	380
						LR chi2(1)	=	283.15
						Prob > chi2	=	0.0000
						Pseudo R2	=	0.6013
Log likelihood = -93.886407								
_outcome	Coef.	Std. Err.	Z	P> z	[95% Conf. Interval]			
ck	.0381044	.0040812	8.60	0.000	.0271083	.0431035		
cons	-2.326272	.2993611	-7.77	0.000	-2.913009	-1.730535		

Display 6.9

and the probability of infarct. (Note that the same coefficient and p-value would be obtained using the mid-point of each interval since all intervals are 30 units wide.) We now investigate whether it is reasonable to assume that the log odds depends linearly on CK. Therefore, we plot the observed proportions and predicted probabilities as follows:

```
generate prop = pres/tot
predict pred, pr
label variable prop "observed"
label variable pred "predicted"
```

```
twoway (line pred ck) (scatter prop ck), ///
ytitle("Probability") xtitle(CK)
```

The predict command gives predicted counts by default and therefore the pr option was used to obtain predicted probabilities instead. In the resulting graph in Figure 6.2, the curve fits the data reasonably well, the largest discrepancy being at CK=280. However, the curve for

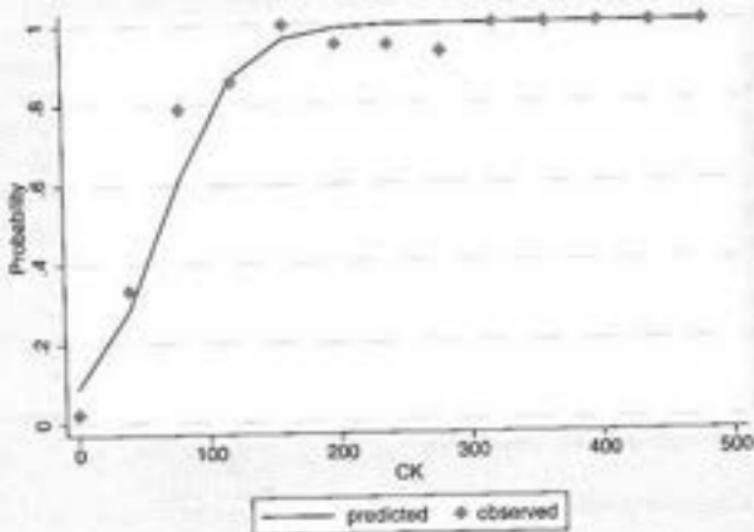


Figure 6.2: Probability of infarct as a function of creatine Kinase levels.

the predicted probabilities is not smooth. Using the plot-type mspline instead of line produces a smooth curve, but a more faithful smooth curve can be obtained using the graph twoway plot-type function command as follows:

```
twoway (function y=1/(1+exp(-_b[_cons]-_b[ck]*x)), ///
range(0 480)) (scatter prop ck), ///
ytitle("Probability") xtitle(CK) ///
legend(order(1 "predicted" 2 "observed"))
```

Here we are using the regression coefficients `_b[_cons]` and `_b[ck]` to calculate the predicted probability as a function of some hypothetical variable `x` varying in the range from 0 to 480. This improved graph is shown in Figure 6.3. Note that we could also use

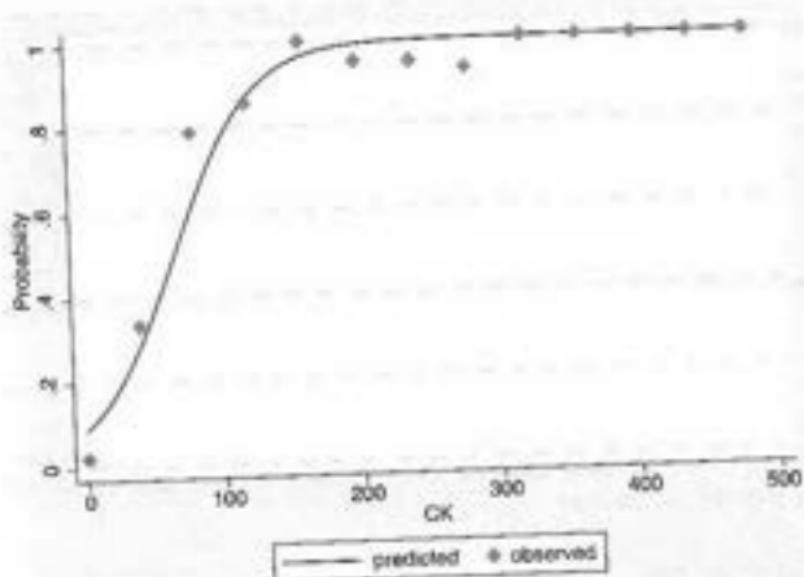


Figure 6.3: Smoother version of Figure 6.2.

the `invlogit()` function to calculate the inverse logit, i.e., function
 $y = \text{invlogit}(\text{b}[\text{cons}] + \text{b}[ck] * x)$.

We will now plot some residuals and then consider the performance of CK as a diagnostic tool, defining the test as positive if CK exceeds a certain threshold. In particular, we will consider the sensitivity (probability of a positive test result if the disease is present) and specificity (probability of a negative test result if the disease is absent) for different thresholds. There are some useful post-estimation commands available for those purposes for use after the `logit` (or `logistic`) command that are not available after `bilogit`. We therefore transform the data into the form required for `logistic`, i.e., one observation per Bernoulli trial with outcome `infct` equal to 0 or 1 so that the number of ones per CK level equals `pres`:

```
expand tot
by ck, sort: generate infct = (.m==pres)
We can reproduce the results of bilogit using logit:
```

```
logit infct ck
```

(see Display 6.10).

To judge if the discrepancy between the observed and expected proportions is acceptable, we can use standardized Pearson residuals for

Logistic regression		Number of obs	=	360	
		LR chi2(1)	=	283.15	
		Prob > chi2	=	0.0000	
Log likelihood = -93.888407		Pseudo R2	=	0.6013	
infct	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
ck	.0351044	.0040813	8.60	0.000	.0271053 .0431035
_cons	-2.326272	.2990611	-7.77	0.000	-2.913009 -1.736635

Display 6.10

each "covariate pattern", i.e., for each combination of values in the covariates (here for each value of CK). These residuals may be obtained and plotted as follows:

```
predict resi, rstandard
twoway scatter resi ck, xlabel(ck)
```

The graph is shown in Figure 6.4. There are several large outliers. The

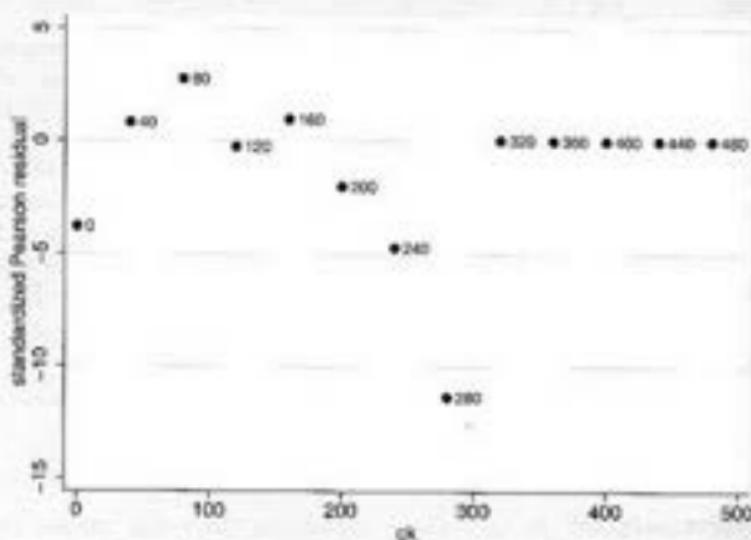


Figure 6.4: Standardized Pearson residuals vs. creatine kinase level.

largest outlier at $CK=280$ is due to one subject out of 14 not having had an infarct although the predicted probability of an infarct is almost 1. (We could also test the goodness-of-fit of the model; see Exercise 6.3.)

We now determine the accuracy of the diagnostic test based on the logistic regression model. A classification table of the predicted diagnosis (using a cut-off of the predicted probability of 0.5) versus the true diagnosis may be obtained using

```
estat classif
```

giving the table shown in Display 6.11. Both the sensitivity and the

Logistic model for infct		True		Total
Classified		D	-D	
+		218	16	234
-		15	114	129
Total		233	130	363

Classified + if predicted $\text{Pr}(D) \geq .5$
 True D defined as infct != 0

Sensitivity	$\text{Pr}(+ D)$	93.48%
Specificity	$\text{Pr}(- D)$	87.69%
Positive predictive value	$\text{Pr}(D +)$	93.07%
Negative predictive value	$\text{Pr}(- D)$	88.37%
False + rate for true -D	$\text{Pr}(+ -D)$	12.31%
False - rate for true D	$\text{Pr}(- D)$	6.82%
False + rate for classified +	$\text{Pr}(-D +)$	6.93%
False - rate for classified -	$\text{Pr}(D -)$	11.43%
Correctly classified		91.39%

Display 6.11

specificity are relatively high. These characteristics are generally assumed to generalize to other populations whereas the positive and negative predictive values (probabilities of the disease being present/absent if the test is positive/negative) depend on the prevalence (or prior probability) of the condition (see for example Sackett *et al.*, 1991).

The use of other probability cut-offs could be investigated using the option `cutoff(#)` in the above command or using the commands `lroc` to plot a ROC-curve (specificity vs. sensitivity for different cut-offs) or `lsens` to plot sensitivity and specificity against cut-off (see Exercise 6.3).

The above classification table may be misleading because we are testing the model on the same sample that was used to derive it. An alternative approach is to compute predicted probabilities for each observation from a model fitted to the remaining observations. This method, called "leave one out" method or jackknifing (see Lachenbruch and Mickey, 1986), can be carried out relatively easily for our data because we only have a small number of covariate and response patterns. Instead of looping through all observations, excluding each observation in the logistic regression command and computing that observation's predicted probability, we can loop through a subset of observations representing all combinations of covariates and responses found in the data.

First, label each unique covariate pattern consecutively in a variable `num` using `predict` with the `number` option:

```
predict num, number
```

Now generate `first`, equal to one for the first observation in each group of unique covariate and response patterns and zero otherwise:

```
by num infct, sort: generate first = (_n==1)
```

(We could also have used `egen first = tag(num infct)`.) Now define `grp`, equal to the cumulative sum of `first`, obtained using the function `sum()`. This variable numbers the groups of unique covariate and response patterns consecutively:

```
generate grp = sum(first)
```

(An alternative way of generating `grp` without having to first create `first` would be to use the command `egen grp = group(num infct)`.) Now determine the number of unique combinations of CK levels and infarct status:

```
summarize grp
```

Variable	Che	Mean	Std. Dev.	Min	Max
grp	360	8.668333	6.625051	1	20

As there are 20 groups, we need to run `logistic` 20 times (for each value of `grp`), excluding one observation from `grp` to derive the model for predicting the probability for all observations in `grp`.

First generate a variable, `nxt`, that consecutively labels the 20 observations to be excluded in turn:

```
generate nxt = first*grp
```

Now build up a variable `ppr` of predicted probabilities as follows:

```

generate prp = 0
forvalues n= 1/20 {
    quietly logistic infct ck if grp!="`n'"
    quietly predict p, pr
    quietly replace prp = p if grp=="`n'"
    drop p
}

```

The purpose of these four commands inside the loop is to

1. derive the model excluding one observation from grp,
2. obtain the predicted probabilities p (predict produces results for the whole sample, not just the estimation sample),
3. set prp to the predicted probability for all observations in grp, and
4. drop p so that it can be defined again in the next iteration.

Here the quietly prefix was used to produce no output.

The classification table for the jackknifed probabilities can be obtained using

```

generate class = (prp>=0.5)
tabulate class infct

```

class	infct		Total
	0	1	
0	114	15	129
1	16	215	231
Total	130	230	360

giving the same result as before, although this will not generally be the case.

6.4 Exercises

6.1 • Treatment of lung cancer

1. Read in the data without using the expand command, and reproduce the result of ordinal logistic regressions by using the appropriate weights.

6.2 • Female psychiatric patients

1. Carry out significance tests for an association between `depress` and `life` for the data described in Chapter 2 using
 - a. ordinal logistic regression with `depress` as dependent variable

- b. logistic regression with `life` as dependent variable.
2. Use `stepwise` together with `logit` to find a model for predicting `life` using the data from Chapter 2 with different sets of candidate variables (see Chapter 3).

6.3 • Diagnosis of heart attacks

1. Test the goodness-of-fit of the logistic regression model with `ck` as the only explanatory variable using a Pearson chi-squared test.
2. To improve the model fit, successively include first a quadratic term of `ck` in the model, then a cubic term, etc., deciding when to stop using Wald tests for the highest order terms.
3. For the chosen model, repeat the goodness-of-fit test.
4. Produce a graph similar to that in Figure 6.2 for the chosen model.
5. Explore the use of `estat classif`, `cutoff(#)`, `lroc`, and `lsens` for the chosen model.

6.4 Psychiatric screening data

Here we consider data from a study of a psychiatric screening questionnaire called the GHQ (General Health Questionnaire). The data are from Der and Everitt (2002). In addition to completing the questionnaire, subjects were diagnosed as clinically depressed or not by a psychiatrist. Here the question of interest is to determine how the probability of being judged depressed (a "case") is related to sex and the GHQ score.

The variables in `screening.dta` are:

- `ghq`: GHQ score
- `sex`: sex (F=female, M=male)
- `cases`: number of cases
- `noncases`: number of non-cases

1. Fit both a linear regression and a logistic regression for the probability of being a case with GHQ score as the single explanatory variable.
2. Plot the predicted probabilities from each model against GHQ score on the same diagram and comment on the two curves.
3. Fit a logistic regression model to the probability of being a case using both GHQ score and sex as explanatory variables. Construct a suitable plot to illustrate the model fitted.
4. Investigate whether the previous model can be improved by including a `sex` \times `GHQ score` interaction.

6.5 Prostate cancer

The data analyzed here arise from a study involving patients with cancer of the prostate (Brown, 1980). The aim was to determine whether a combination of five variables could be used to forecast whether or not the cancer has spread to the lymph nodes, since this forms the basis for the treatment regime that should be adopted. The 53 patients in the study had undergone a laparotomy to determine nodal involvement or not in their case. Here the response variable is binary with zero signifying the absence and unity the presence of nodal involvement.

The variables in `prostate.dta` are:

- `id`: patient identifier
 - `nodal`: nodal involvement (0=no, 1=yes)
 - `age`: age of patient at diagnosis (years)
 - `acid`: level of serum acid phosphatase (in King-Armstrong units)
 - `xray`: result of an X-ray examination (0=negative, 1=positive)
 - `size`: size of the tumour as determined by a rectal examination (0=small, 1=large)
 - `grade`: summary of the pathological grade of the tumour determined from a biopsy (0=less serious, 1=more serious)
1. Carry out a logistic regression with `nodal` as the response and `age` and `acid` as explanatory variables.
 2. Interpret the odds ratios.
 3. Carry out a logistic regression for nodal involvement using all five explanatory variables and investigate which of these five variables are most needed in the model. Use both forward and backward selection procedures.

6.6 Satisfaction with housing conditions

Agresti (1984) discusses data on 1681 residents of twelve areas in Copenhagen that allows investigation of the effect of various factors on satisfaction with the housing conditions. (The data are also described in Madsen, 1976.)

The data are in collapsed form with frequencies given in the variable `freq`. The remaining variables in `housing.dta` are:

- `satisfaction`: level of satisfaction (1=low, 2=medium, 3=high)
- `housing`: type of housing (1=tower blocks, 2=apartments, 3=atrium houses, 4=terraced houses)

- contact: degree of contact with residents (1=low, 2=high)
 - influence: feeling of influence on apartment management (1=low, 2=medium, 3=high)
1. Fit a proportional odds model for satisfaction treating housing as a categorical predictor and influence as a continuous predictor. Make sure to use frequency weights.
 2. Interpret the estimated odds ratios.
 3. Discuss the implicit assumption made in treating influence as continuous.
 4. Use a likelihood ratio test to decide if influence should be treated as categorical instead.
 5. Test for an interaction between housing and contact using both a multivariate Wald test (using `testparm`) and a likelihood ratio test.

Chapter 7

Generalized Linear Models: Australian School Children

7.1 Description of data

This chapter reanalyzes a number of datasets discussed in previous chapters and, in addition, describes the analysis of a new dataset given in Aitkin (1978). These data come from a sociological study of Australian aboriginal and white children. The sample included children from four age groups (final year in primary school and first three years in secondary school) who were classified as slow or average learners. The number of days absent from school during the school year was recorded for each child. The data are given in Table 7.1. The variables are as follows:

- **eth**: ethnic group (A=aboriginal, N=white)
- **sex**: sex (M=male, F=female)
- **age**: class in school (F0, F1, F2, F3)
- **lrn**: average or slow learner (SL=slow learner, AL=average learner)
- **days**: number of days absent from school in one year

One aim of the analysis is to investigate ethnic differences in the mean number of days absent from school while controlling for the other potential predictors **sex**, **age**, and **lrn**.

7.2 Generalized linear models

Previous chapters have described linear (Chapter 3) and logistic regression (Chapter 6). In this chapter, we will describe a more general class of models, called *generalized linear models*, of which linear regression and logistic regression are special cases.

Both linear and logistic regression involve a linear combination of the explanatory variables, called the *linear predictor*, of the form

$$\begin{aligned}\eta_i &= \beta_0 + \beta x_{i1} + \beta x_{i2} + \cdots + \beta x_{ip} \\ &= \mathbf{x}_i' \boldsymbol{\beta}.\end{aligned}\quad (7.1)$$

In both types of regression, the linear predictor determines the expectation μ_i of the response variable. In linear regression, where the response is continuous, μ_i is directly equated with the linear predictor. This is not advisable when the response is dichotomous because in this case the expectation is a probability which must satisfy $0 \leq \mu_i \leq 1$. In logistic regression, the linear predictor is therefore equated with a function of μ_i , the logit, $\eta_i = \log(\mu_i/(1 - \mu_i))$. In generalized linear models, the linear predictor may be equated with any of a number of different functions $g(\mu_i)$ of μ_i , called *link functions*; that is,

$$\eta_i = g(\mu_i).\quad (7.2)$$

In linear regression, the probability distribution of the response variable is assumed to be normal with mean μ_i . In logistic regression a binomial distribution is assumed with probability parameter μ_i . Both the normal and binomial distributions come from the same family of distributions, called the *exponential family*,

$$f(y_i; \theta_i, \phi) = \exp\{(y_i \theta_i - b(\theta_i))/a(\phi) + c(y_i, \phi)\}. \quad (7.3)$$

For example, for the normal distribution,

$$\begin{aligned}f(y_i; \theta_i, \phi) &= \frac{1}{\sqrt{(2\pi\sigma^2)}} \exp\{-(y_i - \mu_i)^2/2\sigma^2\} \\ &= \exp\{(y_i \mu_i - \mu_i^2/2)/\sigma^2 - \frac{1}{2}(y_i^2/\sigma^2 + \log(2\pi\sigma^2))\}\end{aligned}\quad (7.4)$$

so that $\theta_i = \mu_i$, $b(\theta_i) = \theta_i^2/2$, $\phi = \sigma^2$, and $a(\phi) = \phi$.

The parameter θ_i can be written as a function of μ_i and this function is called the *canonical link function*. The canonical link is frequently

Table 7.1 Data in quine-dta presented in four columns to save space. (Taken from Aitkin (1978) with permission of the Royal Statistical Society.)

case	age	sex	days																								
A	54	F	SL	2	A	M	PV	11	A	PV	SL	14	A	C	PV	AL	S	PV	AL	22	A	M	PV				
A	54	PV	AL	3	A	M	PV	13	A	PV	AL	29	A	M	PV	AL	21	A	M	PV	AL	21	A	M	PV		
A	54	PV	SL	6	A	M	PV	6	A	PV	SL	15	A	M	PV	SL	33	A	M	PV	SL	33	A	M	PV		
A	54	PV	AL	14	A	M	PV	50	A	PV	AL	32	A	M	PV	AL	18	A	M	PV	AL	18	A	M	PV		
A	54	PV	SL	67	A	M	PV	14	A	PV	AL	16	A	M	PV	AL	46	A	M	PV	AL	46	A	M	PV		
A	54	PV	AL	17	A	M	PV	52	A	PV	AL	63	A	M	PV	AL	26	A	M	PV	AL	26	A	M	PV		
A	54	PV	AL	8	A	M	PV	23	A	M	PV	23	A	M	PV	SL	2	A	M	PV	SL	2	A	M	PV		
A	54	PV	AL	34	A	M	PV	36	A	M	PV	28	A	M	PV	AL	45	A	M	PV	AL	45	A	M	PV		
A	54	PV	AL	8	A	F	PV	11	A	F	PV	24	A	F	PV	AL	9	A	F	PV	AL	9	A	F	PV		
A	54	PV	SL	3	A	F	PV	51	A	F	PV	6	A	F	PV	SL	32	A	F	PV	SL	32	A	F	PV		
A	54	PV	AL	13	A	F	PV	23	A	F	PV	25	A	F	PV	AL	5	A	F	PV	AL	5	A	F	PV		
A	54	PV	SL	63	A	F	PV	56	A	F	PV	19	A	F	PV	SL	8	A	F	PV	SL	8	A	F	PV		
A	54	PV	AL	11	A	F	PV	17	A	F	PV	19	A	F	PV	SL	17	A	F	PV	SL	17	A	F	PV		
A	54	PV	SL	13	A	F	PV	52	A	F	PV	20	A	F	PV	AL	2	A	F	PV	AL	2	A	F	PV		
A	54	PV	SL	48	A	F	PV	69	A	F	PV	84	A	F	PV	73	AL	5	A	F	PV	73	AL	5	A	F	PV
A	54	PV	AL	8	A	F	PV	2	A	F	PV	3	A	F	PV	AL	36	A	F	PV	AL	36	A	F	PV		
A	54	PV	AL	10	A	F	PV	14	A	F	PV	21	A	F	PV	AL	67	A	F	PV	AL	67	A	F	PV		
A	54	PV	AL	40	N	M	PV	6	N	M	PV	17	N	M	PV	AL	7	N	M	PV	AL	7	N	M	PV		
A	54	PV	AL	0	N	M	PV	70	AL	0	N	PV	2	AL	2	N	M	PV	AL	16	N	M	PV				
A	54	PV	AL	11	N	M	PV	12	N	M	PV	24	N	M	PV	SL	0	N	M	PV	SL	0	N	M	PV		
A	54	PV	AL	11	N	M	PV	31	N	M	PV	35	N	M	PV	SL	11	N	M	PV	SL	11	N	M	PV		
A	54	PV	SL	37	N	M	PV	1	N	M	PV	4	N	M	PV	SL	22	N	M	PV	SL	22	N	M	PV		
A	54	PV	SL	30	N	M	PV	36	N	M	PV	38	N	M	PV	SL	0	N	M	PV	SL	0	N	M	PV		
A	54	PV	AL	1	N	M	PV	5	N	M	PV	20	N	M	PV	AL	16	N	M	PV	AL	16	N	M	PV		
A	54	PV	AL	27	N	M	PV	0	N	M	PV	30	N	M	PV	AL	10	N	M	PV	AL	10	N	M	PV		
A	54	PV	AL	27	N	M	PV	77	N	M	PV	41	N	M	PV	AL	69	N	M	PV	AL	69	N	M	PV		
A	54	PV	AL	14	N	M	PV	19	N	M	PV	11	N	M	PV	AL	26	N	M	PV	AL	26	N	M	PV		
A	54	PV	AL	25	N	M	PV	59	N	M	PV	52	N	M	PV	AL	0	N	M	PV	AL	0	N	M	PV		
A	54	PV	AL	33	N	M	PV	77	N	M	PV	77	N	M	PV	AL	5	N	M	PV	AL	5	N	M	PV		
A	54	PV	SL	5	N	M	PV	5	N	M	PV	5	N	M	PV	SL	12	N	M	PV	AL	1	N	M	PV		
A	54	PV	AL	5	N	M	PV	14	N	M	PV	16	N	M	PV	AL	3	N	M	PV	AL	3	N	M	PV		
A	54	PV	AL	5	N	M	PV	8	N	M	PV	10	N	M	PV	AL	18	N	M	PV	AL	18	N	M	PV		
A	54	PV	AL	1	N	M	PV	7	N	M	PV	9	N	M	PV	AL	15	N	M	PV	AL	15	N	M	PV		
A	54	PV	AL	3	N	M	PV	73	N	M	PV	27	N	M	PV	AL	27	N	M	PV	AL	27	N	M	PV		
A	54	PV	AL	92	N	M	PV	73	N	M	PV	73	N	M	PV	AL	18	N	M	PV	AL	18	N	M	PV		

chosen as the link function (and is the default link in the Stata command for fitting generalized linear models, `glm`), although the canonical link is not necessarily more appropriate than any other link. Table 7.2 lists some of the most common distributions used in generalized linear models and their canonical link functions.

Table 7.2 Probability distributions and their canonical link functions

Distribution	Variance function	Dispersion parameter	Link function	$g(\mu) = \theta(\mu)$
Normal	1	σ^2	identity	μ
Binomial	$\mu(1 - \mu)$	1	logit	$\log(\mu/(1 - \mu))$
Poisson	μ	1	log	$\log(\mu)$
Gamma	μ^2	ν^{-1}	reciprocal	$1/\mu$

The conditional mean and variance of Y_i are given by

$$E(Y_i | \mathbf{x}_i) = b'(\theta_i) = \mu_i \quad (7.5)$$

and

$$\text{var}(Y_i | \mathbf{x}_i) = b''(\theta_i)a(\phi) = V(\mu_i)a(\phi) \quad (7.6)$$

where $b'(\theta_i)$ and $b''(\theta_i)$ denote the first and second derivatives of $b(\cdot)$ evaluated at θ_i , and the variance function $V(\mu_i)$ is obtained by expressing $b''(\theta_i)$ as a function of μ_i . It can be seen from (7.4) that the variance for the normal distribution is simply σ^2 regardless of the value of the mean μ_i , i.e., the variance function is 1.

The data on Australian school children will be analyzed by assuming a Poisson distribution for the number of days absent from school. The Poisson distribution is the appropriate distribution of the number of events observed over a period of time, if these events occur independently in continuous time at a constant instantaneous probability rate (or incidence rate); see for example Clayton and Hills (1993). The Poisson distribution is given by

$$f(y_i; \mu_i) = \mu_i^{y_i} e^{-\mu_i} / y_i! \quad y_i = 0, 1, 2, \dots \quad (7.7)$$

Taking the logarithm and summing over observations, the log likelihood is given by

$$l(\mu; \mathbf{y}) = \sum_i \{(y_i \ln \mu_i - \mu_i) - \ln(y_i!)\} \quad (7.8)$$

so that $\theta_i = \ln \mu_i$, $b(\theta_i) = \exp(\theta_i)$, $\phi = 1$, $a(\phi) = 1$, and $\text{var}(Y_i | \mathbf{x}_i) = \exp(\theta_i) = \mu_i$. Therefore, the variance of the Poisson distribution is not constant, but equal to the mean. Unlike the normal distribution, the Poisson distribution has no separate parameter for the variance and the same is true of the binomial distribution. Table 7.2 shows the variance functions and dispersion parameters for some commonly used probability distributions.

7.2.1 Model selection and measure of fit

Lack of fit may be expressed by the deviance, which is minus twice the difference between the maximized log likelihood of the model and the maximum likelihood achievable, i.e., the maximized likelihood of the full or saturated model. For the normal distribution, the deviance is simply the residual sum of squares. Another measure of lack of fit is the generalized Pearson X^2 .

$$X^2 = \sum_i (y_i - \hat{\mu}_i)^2 / V(\hat{\mu}_i), \quad (7.9)$$

which, for the Poisson distribution, is just the familiar Pearson chi-squared statistic for two-way cross-tabulations (since $V(\hat{\mu}_i) = \hat{\mu}_i$). Both the deviance and Pearson X^2 have asymptotic χ^2 distributions under the null hypothesis. When the dispersion parameter ϕ is fixed (not estimated), an analysis of deviance can be used for comparing nested models. To test the null hypothesis that the restrictions leading to the nested model are true, the difference in deviance between two models is compared with the χ^2 distribution with degrees of freedom equal to the difference in model degrees of freedom.

The Pearson and deviance residuals are defined as the (signed) square roots of the contributions of the individual observations to the Pearson X^2 and deviance respectively. These residuals may be used to assess the appropriateness of the link and variance functions.

A relatively common phenomenon with count data is overdispersion, i.e., the variance is greater than that of the assumed distribution (binomial with denominator greater than 1 or Poisson). This overdispersion may be due to extra variability in the parameter μ_i which has not been completely explained by the covariates. One way of addressing the problem is to allow μ_i to vary randomly according to some distribution and to assume that conditional on μ_i , the response variable follows the binomial (or Poisson) distribution. Such models are called *random effects models*; see also Chapter 9.

A more pragmatic way of accommodating overdispersion in the model is to assume that the variance is proportional to the variance

function, but to estimate the dispersion or scale parameter ϕ rather than assuming the value 1 appropriate for the distributions. For the Poisson distribution, the variance is modeled as

$$\text{var}(Y|\mathbf{x}_i) = \phi\mu_i \quad (7.10)$$

where ϕ is estimated from the deviance or Pearson X^2 . (This is analogous to the estimation of the residual variance in linear regression models from the residual sums of squares.) This parameter is then used to scale the estimated standard errors of the regression coefficients. This approach of assuming a variance function that does not correspond to any probability distribution is an example of the quasi-likelihood approach. If the variance is not proportional to the variance function, robust standard errors can be used as described in the next section. See McCullagh and Nelder (1989) and Hardin and Hilbe (2006) for more details on generalized linear models.

7.2.2 Robust standard errors of parameter estimates

A very useful feature of Stata is that robust standard errors of estimated parameters can be obtained for most estimation commands. In maximum likelihood estimation, the standard errors of the estimated parameters are derived from the Hessian (matrix of second derivatives with respect to the parameters) of the log likelihood. However, these standard errors are correct only if the likelihood is the true likelihood of the data. If this assumption is not correct, for instance due to omission of covariates, misspecification of the link function or probability distribution function, we can still use robust estimates of the standard errors known as the Huber, White, or sandwich variance estimates (for details, see Binder, 1983).

In the description of the robust variance estimator in the *Stata User's Guide* (Section 20.14), it is pointed out that the use of robust standard errors implies a less ambitious interpretation of the parameter estimates and their standard errors than a model-based approach. Instead of assuming that the model is "true" and attempting to estimate "true" parameters, we just consider the properties of the estimator (whatever it may mean) under repeated sampling and define the standard error as its sampling standard deviation.

Another approach to estimating the standard errors without making any distributional assumptions is bootstrapping (Efron and Tibshirani, 1993). If we could obtain repeated samples from the population (from which our data were sampled), we could obtain an empirical sampling distribution of the parameter estimates. In Monte Carlo simulation, the required samples are drawn from the assumed distribution. In

bootstrapping, the sample is resampled "to approximate what would happen if the population were sampled" (Manly, 1997). Bootstrapping works as follows. Take a random sample of n observations (n is the sample size), with replacement, and estimate the regression coefficients. Repeat this a number of times to obtain a sample of estimates. From the resulting sample of parameter estimates, obtain the empirical variance-covariance matrix of the parameter estimates. Confidence intervals may be constructed using the estimated variance or directly from the appropriate centiles of the empirical distribution of parameter estimates. See Manly (1997) and Efron and Tibshirani (1993) for more information on the bootstrap.

7.3 Analysis using Stata

The `glm` command can be used to fit generalized linear models. The syntax is analogous to `logit` and `regress` except that the options `family()` and `link()` are used to specify the probability distribution of the response and the link function, respectively. We first analyze data from the previous chapter to show how linear regression, ANOVA, and logistic regression are performed using `glm` and then move on to the data on Australian school children.

7.3.1 Linear regression

First, we show how linear regression can be carried out using `glm`. In Chapter 3, the U.S. air-pollution data were read in using the instructions

```
infile str10 town so2 temp manuf pop wind precip days ///
    using usair.dat, clear
drop if town=="Chicago"
```

and now we regress `so2` on a number of variables using

```
glm so2 temp pop wind precip, family(gaussian) link(identity)
```

(see Display 7.1). The results are identical to those of the regression analysis. The scale parameter given on the right-hand side above the regression table represents the residual variance given under Residual MS in the analysis of variance table of the regression analysis in Chapter 3. We can estimate robust standard errors using the `vce(robust)` option

```
glm so2 temp pop wind precip, family(gauss) ///
    link(identity) vce(robust)
```

Generalized linear models	No. of obs	=	40				
Optimization : ML	Residual df	=	35				
Deviance = 30150.15199	Scale parameter	=	290.0043				
Pearson = 30150.15199	(1/df) Deviance	=	290.0043				
Variance function: V(u) = 1	(1/df) Pearson	=	290.0043				
Link function : g(u) = u	[Gaussian]						
Log likelihood = -567.4848314	[Identity]						
	AIC	=	8.524242				
	BIC	=	10021.04				
est	BEM						
	Coeff.	Std. Err.	t	P> t		[95% Conf. Interval]	
temp	-1.810123	.4404001	-4.11	0.000	-2.673292	-.9469549	
pop	.0113089	.0074091	1.53	0.127	-.0032126	.0258304	
wind	-3.095284	2.095471	-1.47	0.141	-7.194292	1.023723	
precip	.5660172	.2608601	2.28	0.024	.0743404	1.057694	
cons	131.3388	34.32034	3.83	0.000	64.07126	198.6052	

Display 7.1

(see Display 7.2) giving slightly different standard errors, suggesting that some assumptions may not be entirely satisfied.

7.3.2 ANOVA

We now show how an analysis of variance model can be fitted using `glm`, using the slimming clinic example of Chapter 5. The data are read using

```
infile cond status resp using slim.dat, clear  
and the full, saturated model can be obtained using
```

```
xi: glm resp i.cond*i.status, family(gaussian) link(identity)  
(see Display 7.3). This result is identical to that obtained using the  
command
```

```
xi: regress resp i.cond*i.status
```

(see exercises in Chapter 5).

We can obtain the *F*-statistics for the interaction term by saving the deviance of the above model (residual sum of squares) in a local macro and refitting the model with the interaction removed:

```
local dev1 = e(deviance)
```

```
xi: glm resp i.cond i.status, family(gaussian) link(identity)
```

Generalized linear models
 Optimization : ML
 Deviance = 10150.15189
 Pearson = 10150.15189
 Variance function: V(u) = 1
 Link function : g(u) = u
 Log pseudolikelihood = -167.4848314

No. of obs	= 40
Residual df	= 35
Scale parameter	= 290.0043
(1/df) Deviance	= 290.0043
(1/df) Pearson	= 290.0043
[Gaussian]	
[Identity]	
AIC	= 8.624343
BIC	= 10021.04

so2	Robust					[95% Conf. Interval]
Coef.	Std. Err.	z	Pr> z			
temp	-1.810123	.3280436	-5.62	0.000	-2.453077	-1.16717
pvp	.0113089	.0079634	1.42	0.158	-.0042992	.0269169
wind	-3.085284	1.698542	-1.82	0.069	-6.414366	.2437966
precip	.6660172	.1818484	3.11	0.003	.2096009	.9224335
_cons	131.3386	18.21993	7.21	0.000	95.62816	167.049

Display 7.2

1.0000 _Icond_1<2 (naturally coded; _Icond_1 omitted)
 1. STATUS _Istatus_1<2 (naturally coded; _Istatus_1 omitted)
 1. cond1:STATUS _Icond1status_2<3 (coded as above)

No. of obs	= 34
Residual df	= 30
Scale parameter	= 35.9616
(1/df) Deviance	= 35.9616
(1/df) Pearson	= 35.9616
[Gaussian]	
[Identity]	
AIC	= 6.53046
BIC	= 973.0573

resp	OIR					[95% Conf. Interval]
Coef.	Std. Err.	z	Pr> z			
_Icond_2	.6780002	3.234433	0.21	0.834	-5.661372	7.017373
_Istatus_2	6.126834	3.19204	1.92	0.058	-.1274504	12.38612
_Icond1status_2	-.2688002	4.410437	-0.06	0.952	-6.909799	6.378796
_cons	-7.298	2.68185	-2.72	0.007	-12.55433	-2.04167

Display 7.3

i.cond	_Icond_1-2	(naturally coded; _Icond_1 omitted)				
i.status	_Istatus_1-2	(naturally coded; _Istatus_1 omitted)				
Generalized linear models		No. of obs = 34				
Optimization : ML		Residual df = 31				
		Scale parameter = 34.80576				
Deviance = 1078.97844		(1/df) Deviance = 34.80576				
Pearson = 1078.97844		(1/df) Pearson = 34.80576				
Variance function: V(u) = 1		[Gaussian]				
Link function : g(u) = u		[Identity]				
		AIC = 6.471757				
Log likelihood = -107.0198709		BIC = 969.6613				
<hr/>						
		DEB				
resp	Coeff.	Std. Err.	z	P> z	[95% Conf. Interval]	
_Icond_2	.5352102	2.163277	0.25	0.806	-3.704734	4.775154
_Istatus_2	5.989762	2.167029	2.76	0.006	1.742463	10.23706
_cons	-7.199832	2.094584	-3.44	0.001	-11.30614	-3.094524

Display 7.4

(see Display 7.4). The increase in deviance caused by the removal of the interaction term represents the sum of squares of the interaction term after eliminating the main effects:

```
local dev0 = e(deviance)
local ddev = `dev0' - `dev1'
display `ddev'
.13031638
```

and the F -statistic is simply the mean sum of squares of the interaction term after eliminating the main effects divided by the residual mean square of the full model. The numerator and denominator degrees of freedom are 1 and 30 respectively, so that F and the associated p -value may be obtained as follows:

```
local f = (`ddev'/1)/(`dev1'/30)
display `f'
.00362382

display Ftail(1,30,"f")
.95239704
```

The general method for testing the difference in fit of two nested generalized linear models, using the difference in deviance, is not appropriate here because the scale parameter $\phi = \sigma^2$ was estimated. Note that the z -test in the regression table in Display 7.3, as well as

The chi-squared test performed by `testparm` in `IconX*`, assume sampling distributions that are appropriate if the dispersion parameter is known or for large residual degrees of freedom. Here the *p*-value from the *F*-test is identical to three decimal places to that from the *z*-test.

7.3.3 Logistic regression

We now repeat the logistic regression analysis of Chapter 6 using `glm`. We first read the tumor data as before, without replicating records.

```
infile fr1 fr2 fr3 fr4 using tumor.dat, clear
gen therapy = int((_,n-1)/2)
sort therapy
by therapy: gen sex = _n
reshape long fr, i(therapy sex) j(outc)
gen improve = outc
recode improve 1/2=0 3/4=1
list
```

	therapy	sex	outc	fr	improve
1.	0	1	1	28	0
2.	0	1	2	45	0
3.	0	1	3	29	1
4.	0	1	4	26	1
5.	0	2	1	4	0
6.	0	2	2	12	0
7.	0	2	3	6	1
8.	0	2	4	2	1
9.	1	1	1	41	0
10.	1	1	2	44	0
11.	1	1	3	20	1
12.	1	1	4	20	1
13.	1	2	1	12	0
14.	1	2	2	7	0
15.	1	2	3	3	1
16.	1	2	4	1	1

The `glm` command can be used with the logit link and binomial distribution and with `fr` as frequency weights using

```
glm improve therapy sex [fweight=fr], family(binomial) ///
link(logit)
```

see Display 7.5).

The likelihood ratio test for `sex` can be obtained as follows:

```
local dev1 = e(deviance)
```

Generalized linear models		No. of obs	=	299		
Optimization	: ML	Residual df	=	296		
Deviance	= 381.2634298	Scale parameter	=	1		
Pearson	= 298.7046083	(1/df) Deviance	=	1.288002		
Variance function: V(u) = u*(1-u)		(1/df) Pearson	=	1.009137		
Link function	: g(u) = ln(u/(1-u))					
				[Bernoulli],		
				[Logit]		
Log likelihood	= -190.6317549	AIC	=	1.285195		
		BIC	=	-1306.068		
<hr/>						
Improve	OIM					
	Coef.	Std. Err.	z	P> z		
therapy	-.502914	.3466998	-1.46	0.145	-.9837445	-.0006582
sex	-.6843128	.3714738	-1.85	0.070	-1.382388	.9737629
cosse	.3856098	.4814172	0.80	0.393	-.4969519	1.270571

Display 7.5

```
quietly glm improve therapy [fweight=fr], ///
    family(binomial) link(logit)
local dev0 = e(deviance)
display `dev0'-'`dev1'
3.2459816

display chi2tail(1,`dev0'-'`dev1')
.0673883
```

which gives the same result as in Section 6.3.1 where we used `estimates store` and `lrtest`.

7.3.4 Australian school children

We now move on to analyze the data in Table 7.1 to investigate differences between aboriginal and white children in the mean number of days absent from school after controlling for other covariates. The data are available as a Stata file `quine.dta` and may therefore be read simply by using the command

```
use quine, clear
```

The variables are of type string and can be converted to numeric using the `encode` command as follows:

```
encode eth, gen(ethnic)
drop eth
encode sex, gen(gender)
```

```

drop sex
encode age, gen(class)
drop age
encode lrn, gen(slow)
drop lrn

```

The number of children in each of the combinations of categories of gender, class, and slow can be found using

```

table slow class ethnic, contents(freq) by(gender)

```

-- Display 7.6). This reveals that there were no "slow learners" in

gender and slow	ethnic and class								
	A				B				
	F0	F1	F2	F3	F0	F1	F2	F3	
SL	4	5	1	9	4	6	1	10	
SL	1	10	8		1	11	9		
AL	5	2	7	7	6	2	7	7	
AL	3	3	4		3	7	3		

Display 7.6

-- F3. A table of the means and standard deviations is obtained

```

table slow class ethnic, contents(mean days sd days) ///
by(gender) format(%4.1f)

```

-- Display 7.7), where the format() option causes only a single decimal place to be given. This table suggests that the variance associated with the Poisson distribution is not appropriate here as squaring the standard deviations (to get the variances) results in values that are greater than the means, i.e., there is overdispersion. In this case, the overdispersion probably arises from substantial variability in children's underlying tendency to miss days of school that cannot be fully explained by the variables we have included in the model.

Ignoring the problem of overdispersion for the moment, a generalised linear model with a Poisson family and log link can be fitted using

```

glm days slow class ethnic gender, family(poisson) link(log)

```

gender and slow	ethnic and class								
	A				B				
	F0	F1	F2	F3	F0	F1	F2	F3	
F	AL	21.3	11.4	2.0	14.6	18.5	11.0	1.0	13.5
		17.7	6.5		14.9	10.7	6.9		11.5
	SL	3.0	22.6	36.4		25.0	6.0	6.2	
		18.7	26.5			4.2	5.0		
M	AL	13.0	10.6	27.4	27.1	5.3	3.5	9.1	27.3
		8.0	4.9	14.7	10.4	5.4	0.7	9.5	22.9
	SL	9.0	3.0	37.0		30.0	6.1	29.3	
		6.2	5.2	23.4		32.5	6.1	7.0	

Display 7.7

Iteration 0: log likelihood = -1192.3347
 Iteration 1: log likelihood = -1178.6003
 Iteration 2: log likelihood = -1178.5612
 Iteration 3: log likelihood = -1178.5612
 Generalized linear models
 Optimization : ML
 No. of obs = 148
 Residual df = 145
 Scale parameter = 1
 Deviance = 1768.64529 (1/df) Deviance = 12.54358
 Pearson = 1990.142857 (1/df) Pearson = 14.11448
 Variance function: V(u) = u [Poisson]
 Link function : g(u) = ln(u) [Log]
 AIC = 16.21317
 BIC = 1065.957
 Log likelihood = -1178.56124

days	DIM					
	Coeff.	Std. Err.	<i>t</i>	P> <i>t</i>	[95% Coef.]	Interval
slow	.2661578	.0448715	5.97	0.000	.1787993	.3835164
class	.2094662	.0218345	9.60	0.000	.1666921	.2622434
ethnic	-.5511688	.0418391	-13.17	0.000	-.633372	-.4691058
gender	.2256243	.0418927	5.43	0.000	.1441041	.3071442
_cons	2.336676	.1427925	16.36	0.000	2.066808	2.616545

Display 7.8

→ Display 7.8). The algorithm takes three iterations to converge to the maximum likelihood (or minimum deviance) solution. In the absence of overdispersion, the scale parameters based on the Pearson χ^2 or the deviance should be close to 1. The values of 14.1 and 12.5 (given at the top-right), respectively, therefore indicate that there is overdispersion. Consequently, the confidence intervals are likely to be too narrow. McCullagh and Nelder (1989) use the Pearson χ^2 divided by the degrees of freedom to estimate the scale parameter for the quasi-likelihood method for Poisson models. This may be achieved using the option `scale(x2)`:

```
glm days slow class ethnic gender, family(poisson) ///
link(log) scale(x2)
```

→ Display 7.9). Allowing for overdispersion has no effect on the re-

Generalized linear models		No. of obs	=	146	
Optimisation	: ML	Residual df	=	141	
Deviance	= 1768.64829	Scale parameter	=	1	
Pearson	= 1990.142897	(1/df) Deviance	=	12.54368	
Variance function:	$\text{V}(u) = u$	(1/df) Pearson	=	14.11449	
Link function	: $g(u) = \ln(u)$				
		[Poisson]			
		[Log]			
Log likelihood	= -1178.561184	AIC	=	16.21317	
		BIC	=	1966.957	
<hr/>					
days	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
slow	-0.2661578	.1674519	1.59	0.112	-.0620419 .5943575
class	.2094662	.0819929	2.55	0.011	.0487831 .3701693
ethnic	-.5511688	.1571865	-3.51	0.000	-.8592486 -.243089
gender	.2266343	.1562606	1.44	0.149	-.0805409 .5318896
class	2.336676	.5364606	4.36	0.000	1.285233 3.38812

(Standard errors scaled using square root of Pearson χ^2 -based dispersion)

Display 7.9

gression coefficients, but a large effect on the p-values and confidence intervals so that `gender` and `slow` are now no longer significant at the 5% level. These terms will be removed from the model. The coefficients can be interpreted as the differences in the logs of the predicted mean counts between groups after controlling for the other variables. For example, the log of the predicted mean number of days absent from school for white children is -0.55 lower than that for aborigines after controlling for `slow`, `class`, and `gender`. Exponentiating the coefficients

yields ratios of expected counts (or rate ratios). The `glm` command exponentiates all coefficients and confidence intervals when the `eform` option is used:

```
glm days class ethnic, family(poisson) link(log) ///
    scale(x2) eform
```

(see Display 7.10). Therefore, white children are absent from school

Generalized linear models		No. of obs	=	146
Optimization		Residual df	=	143
Deviance		Scale parameter	=	1
Pearson		(1/df) Deviance	=	12.75162
Variance function: $V(u) = u$		(1/df) Pearson	=	14.62445
Link function: $g(u) = \ln(u)$		[Poisson]		
		[Log]		
Log likelihood		AIC	=	16.56136
		BIC	=	1110.836

days	GLM					[95% Coef. Interval]
	IRR	Std. Err.	z	P> z		
class	1.177095	.0896266	2.15	0.031	1.014872	1.367106
ethnic	.5782831	.0924961	-3.42	0.001	.4226284	.7911836

(Standard errors scaled using square root of Pearson X2-based dispersion).

Display 7.10

about 58% as often as aboriginal children (95% confidence interval from 42% to 79%) after controlling for `class`.

We have treated `class` as a continuous covariate. This implies that the rate ratio for two categories is a constant multiple of the difference in scores assigned to these categories; for example the rate ratio comparing classes F1 and F0 is the same as that comparing F2 and F1. To see whether this appears to be appropriate, we can form the square of `class` and include this in the model:

```
gen class2 = class^2
glm days class class2 ethnic, family(poisson) link(log) ///
    scale(x2) eform
```

(see Display 7.11). This term is not significant at the 5% level so we can return to the simpler model. (Note that the interaction between `class` and `ethnic` is also not significant, see exercises.)

We now look at the residuals for this model. The post-estimation command `predict` that was used for `regress` and `logistic` can be

Generalized linear models		No. of obs	=	148
Optimization	: ML	Residual df	=	142
Deviance	= 1822.560172	Scale parameter	=	1
Pearson	= 2081.259434	(1/df) Deviance	=	12.83493
Variance function: V(u) = u		(1/df) Pearson	=	14.65676
Link function	: g(u) = ln(u)	[Poisson]		
		[Log]		
Log likelihood	= -1205.518628	AIC	=	24.56875
		BIC	=	1114.588

days	OLS		n	P> z	(95% Conf. Interval)
	IRR	Snd. Err.			
class	1.059399	.4643013	0.13	0.893	.4571296 2.455158
class2	1.039513	.0925501	0.25	0.802	.8798806 1.195638
ethnic	.5784844	.092543	-3.42	0.001	.4226825 .7917989

Standard errors scaled using square root of Pearson X2-based dispersion)

Display 7.11

used here as well. To obtain standardized Pearson residuals, use the `pearson` option with `predict` and divide the residuals by the square root of the estimated dispersion parameter stored in `e(dispersp.ps)`:

```
quietly glm days class ethnic, family(poisson) link(log) ///
    scale(x2)
predict resp, pearson
gen stres = resp/sqrt(e(dispersp.ps))
```

The residuals are plotted against the linear predictor using

```
predict xb, xb
twoway scatter stres xb, ytitle("Standardized Residuals")
```

with the result shown in Figure 7.1.

There is one large outlier with a standardized Pearson residual greater than 4. In order to find out which observation this is, we list a number of variables for cases with large standardized Pearson residuals:

```
predict mu, mu
list stres days mu ethnic class if stres>2|stres<-2
```

(see Display 7.12). Case 72, a white primary school child, has a very large residual.

We now also check the assumptions of the model by using robust standard errors:

```
glm days class ethnic, family(poisson) link(log) vce(robust)
```

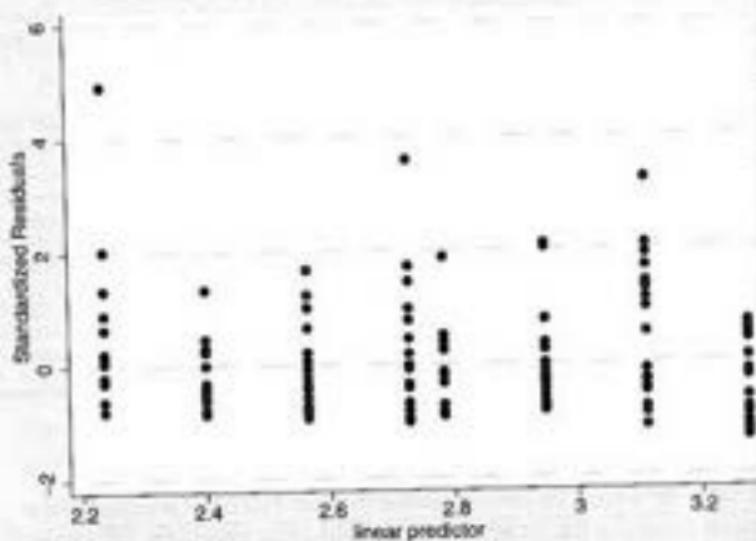


Figure 7.1: Standardized residuals against linear predictor.

	stres	days	age	ethnic	class
45.	2.090936	53	19.07713	A	F1
46.	2.090805	54	19.07713	A	F1
59.	2.070232	60	22.47085	A	F2
59.	3.238862	81	22.47085	A	F2
72.	4.924719	67	9.365361	B	F0
104.	3.588860	69	15.39638	S	F3
109.	2.019514	33	9.365361	S	F0

Display 7.12

(see Display 7.13) giving almost exactly the same p-values as the quasi-

Generalized linear models		No. of obs	=	148
Optimization : ML		Residual df	=	143
Deviance = 1823.481292		Scale parameter	=	1
Pearson = 2091.29704		(1/df) Deviance = 12.75162		
Variance function: V(u) = u		(1/df) Pearson = 14.62445		
Link function : g(u) = ln(u)		[Poisson]		
		[Log]		
Log pseudolikelihood = -1205.979185		AIC	=	16.56136
		BIC	=	1110.828

days	Coef.	Robust Std. Err.	z	P> z	[95% Conf. Interval]
class	.1637288	.0706153	2.14	0.033	.0138658 .313892
ethnic	-.5477436	.1585381	-3.45	0.001	-.8584725 -.2370147
_cons	3.168776	.3065466	10.34	0.000	2.567956 3.769597

Display 7.13

likelihood solution,

```
glm days class ethnic, family(poisson) link(log) scale(x2)
```

(see Display 7.14). We can also use bootstrapping via the `bootstrap` prefix to obtain alternative robust standard errors. Since bootstrapping involves random sampling, we first set the seed of the pseudo random number generator using the `set seed` command so that we can run the sequence of commands again in the future and obtain the same results.

```
set seed 12345678
```

In the `bootstrap` prefix, the statistics are specified for which standard errors are required, here `.b[class]` and `.b[ethnic]`, followed by a comma and any `bootstrap` options, here `reps(500)` to use 500 replicates. Finally the estimation command is specified after a colon.

```
bootstrap ,b[class] ,b[ethnic], reps(500): ///
glm days class ethnic, family(poisson) link(log)
```

(see Display 7.15). The standard errors compare quite well with those using the `vce(robust)` option or using the quasi-likelihood approach.

We could also model overdispersion by assuming a *random effects* model where each child has an unobserved, random proneness to be absent from school. This proneness (called *fruity* in a medical context) multiplies the rate predicted by the covariates so that some children

Generalized linear models		No. of obs	=	146
Optimization : ML		Residual df	=	143
Deviance = 1823.481292		Scale parameter	=	1
Pearson = 2091.29704		(1/df) Deviance	=	12.75163
Variance function: V(u) = u		(1/df) Pearson	=	14.62445
Link function : g(u) = ln(u)		[Poisson]		
		[Log]		
Log likelihood = -1208.979185		AIC	=	18.56120
		BIC	=	1150.826

days	Coeff.	Std. Err.	z	P> z	[95% Conf. Interval]
class	.1637268	.0760048	2.15	0.031	.0147622 .3128954
ethnic	-.5477436	.1599613	-3.42	0.001	-.861262 -.2342252
cons	3.168778	.3170159	10.00	0.000	2.547437 3.790114

(Standard errors scaled using square root of Pearson 32-based dispersion)

Display 7.14

Bootstrap replications (500)						
1	2	3	4	5	50	
.....	100	
.....	150	
.....	200	
.....	250	
.....	300	
.....	350	
.....	400	
.....	450	
.....	500	
Bootstrap results					Number of obs	=
					Replications	=
command: glm days class ethnic, family(poission) link(log)					500	
_bs_1: _b[class]						
_bs_2: _b[ethnic]						
	Observed Coef.	Bootstrap Std. Err.	z	P> z	Normal-based [95% Conf. Interval]	
_bs_1	.1637268	.0708082	2.31	0.021	.0249473	.3029168
_bs_2	-.5477436	.1463108	-3.75	0.000	-.8943115	-.2611757

Display 7.15

have higher or lower rates of absence from school than other children with the same covariates. The observed counts are assumed to have a Poisson distribution conditional on the random effects. If the frailties are assumed to have a gamma distribution, then the (marginal) distribution of the counts has a negative binomial distribution. The negative binomial model can be fitted using `nbinreg` as follows:

```
nbinreg days class ethnic
```

see Display 7.16). Alternatively, the same model can be estimated using `glm` with `family(nbinomial)`.

Negative binomial regression						
						Number of obs = 148
Dispersion = 0.848						LR chi2(2) = 15.77
Log likelihood = -551.24625						Prob > chi2 = 0.0004
						Pseudo R2 = 0.0141
days	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
class	.1506186	.0732832	2.05	0.040	.0088841	.2941489
ethnic	-.5414185	.1578378	-3.43	0.001	-.8507748	-.2320623
cons	3.19392	.3217661	9.93	0.000	2.863266	3.834574
/lnalpha	-.1769664	.1243878			-.4197619	.0678292
alpha	.8386463	.1043178			.6872032	1.079182

Loglikelihood-ratio test of alpha=0: chibar2(01) = 1309.47 Prob>chibar2 = 0.000

Display 7.16

All four methods of analyzing the data lead to the same conclusions. The Poisson model is a special case of the negative binomial model with $\alpha = 0$. The likelihood ratio test for α is therefore a test of the negative binomial against the Poisson distribution. The very small p -value "against Poisson" indicates that there is significant overdispersion. (Note that, as indicated by the expression `chibar2(01)`, the test is based on the correct sampling distribution taking into account that the null hypothesis is on the boundary of the parameter space, see, e.g., Sijders and Bosker, 1999.)

7.4 Exercises

7.1 • Effectiveness of slimming clinics

1. Calculate the F -statistic and difference in deviance for adding `status` to a model already containing `cond` for the data in `slim.dat`.
2. Fit the model using `status` as the only explanatory variable, using robust standard errors. How does this compare with a t -test with unequal variances?

7.2 • Australian school children

1. Carry out a significance test for the interaction between `class` and `ethnic` for the data in `quine.dta`.
2. Excluding the potential outlier (case 72), fit the model with explanatory variables `ethnic` and `class`.
3. Dichotomize days absent from school by classifying 14 days or more as frequently absent. Analyze this new response using the `glm` command with both logit and probit links and the binomial family.
4. Repeat the analyses with the `vce(robust)` option, and compare the robust standard errors with the standard errors obtained using bootstrapping.

See also the Exercises in Chapter 11.

7.3 Wave damage to cargo ships

McCullagh and Nelder (1989) describe data provided by J. Crilley and L. N. Hemingway of Lloyd's Register of shipping concerning the damage caused by waves to the forward section of certain cargo ships. The data are in the form of a table giving the total number of damage incidents by three factors (1) ship type, (2) year of construction and (3) period of operation. (For further discussion of this kind of aggregated data, see the next chapter.) The total number of months in service for each ship type is also given. The purpose of the analysis is to investigate the risk of damage associated with the three factors.

The variables in the dataset `ships.dta` are:

- `damage`: total number of damage incidents
- `type`: ship type (A, B, C, D, or E)
- `construction`: year of construction (1960-64, 1965-69, 1970-74, 1975-79)
- `operation`: period of operation (1960-74 or 1976-79)
- `months`: aggregate number of months in service

Note that `type`, `construction`, and `operation` are string vari-

ables.

- McCullagh and Nelder (1989) consider a log-linear Poisson model with main effects for type, construction and service and with the logarithm of months as an offset (a covariate with regression coefficient set to 1). Fit the model using the `glm` command (see option `offset()`).
- Repeat the analysis above by relaxing the assumption that $\phi = 1$.
- Obtain exponentiated regression coefficients and interpret them.
- Derive scaled Pearson residuals and discuss if there are any potential outliers.
- Consider including an interaction between ship type and year of construction. Note that an *F*-test should be used as demonstrated for a linear model in Section 7.3.2.

7.4 Clotting times of blood

Here we consider data originally published by Hurn et al. (1945) and provided by McCullagh and Nelder (1989). Normal plasma was diluted to nine different concentrations with prothrombin-free plasma and clotting was induced with two different lots of thromboplastin.

The variables in `clotting.dta` are:

- `lot`: lot number (1 or 2)
 - `conc`: concentration of prothrombin-free plasma (in percent)
 - `time`: clotting time
- Following McCullagh and Nelder, use a log transformation of `conc` and specify a gamma distribution and a reciprocal link function $1/\mu_i = \eta_i$ (use the `link(reciprocal)` option). Fit the following sequence of models for the clotting times (1) without covariates, (2) with a main effect of log concentration, (3) with main effects of log concentration and lot and (4) with main effects of log concentration and lot and their interaction. Use *F*-tests as shown in Section 7.3.2 to decide which is the best-fitting model.
 - Calculate predicted mean reaction times and plot them versus concentration, using different line styles for the two lots. Also show the observed clotting times as points on the same graph.

Chapter 8

Summary Measure Analysis of Longitudinal Data: Treatment of Post-Natal Depression

8.1 Description of data

The dataset to be analyzed in this chapter originates from a clinical trial of the use of estrogen patches in the treatment of postnatal depression; full details are given in Gregoire *et al.* (1996). In total, 61 women with major depression, which began within 3 months of childbirth and persisted for up to 18 months postnatally, were allocated randomly to the active treatment or a placebo (a dummy patch); 34 received the former and the remaining 27 received the latter. The women were assessed pretreatment and monthly for six months after treatment on the Edinburgh postnatal depression scale (EPDS), higher values of which indicate increasingly severe depression. The data are shown in Table 8.1; a value of -9 in this table indicates that the observation is missing. The non-integer depression scores result from missing questionnaire items (in this case the average of all available items was multiplied by the total number of items). The variables are

- subj: patient identifier
- group: treatment group (1=estrogen patch, 0=placebo patch)
- pre: pretreatment or baseline EPDS depression score

■ dep1 to dep6: EPDS depression scores for visits 1 to 6

The main question of interest here is whether the estrogen patch is effective at reducing post-natal depression compared with the placebo.

Table 8.1 Data in depress.dat

subj	group	pre	dep1	dep2	dep3	dep4	dep5	dep6
1	0	18	17	18	15	17	14	15
2	0	27	26	23	18	17	12	10
3	0	16	17	14	-9	-9	-9	-9
4	0	17	14	23	17	13	12	12
5	0	15	12	10	8	4	5	5
6	0	20	19	11.54	9	8	6.82	5.05
7	0	16	13	13	9	7	8	7
8	0	28	26	27	-9	-9	-9	-9
9	0	28	26	24	19	13.94	11	9
10	0	25	9	12	15	12	13	20
11	0	24	14	-9	-9	-9	-9	-9
12	0	16	19	13	14	23	15	11
13	0	26	13	22	-9	-9	-9	-9
14	0	21	7	13	-9	-9	-9	-9
15	0	21	18	-9	-9	-9	-9	-9
16	0	22	18	-9	-9	-9	-9	-9
17	0	26	19	13	22	12	18	13
18	0	19	19	7	8	2	5	6
19	0	22	20	15	20	17	15	13.73
20	0	16	7	8	12	10	10	12
21	0	21	19	18	16	13	16	15
22	0	20	16	21	17	21	16	18
23	0	17	15	-9	-6	-9	-9	-9
24	0	22	20	21	17	14	14	10
25	0	19	16	19	-9	-9	-9	-9
26	0	21	7	4	4.19	4.73	3.03	3.45
27	0	18	19	-9	-9	-9	-9	-9
28	1	21	13	12	9	9	13	6
29	1	27	8	17	15	7	5	7
30	1	15	8	12.27	10	10	6	3.96
31	1	24	14	14	13	12	18	15
32	1	15	15	16	11	14	12	8
33	1	17	9	5	3	6	0	2
34	1	20	7	7	7	12	9	6
35	1	18	8	1	1	2	0	1
36	1	28	11	7	3	2	2	2
37	1	21	7	8	6	6.5	4.64	4.97
38	1	18	8	6	4	11	7	6
39	1	27.46	22	27	24	22	24	23
40	1	19	14	12	15	12	9	6
41	1	20	13	10	7	9	11	11
42	1	16	17	26	-9	-9	-9	-9
43	1	21	19	9	9	12	5	7
44	1	23	11	7	5	8	2	3
45	1	23	16	13	-9	-9	-9	-9
46	1	24	16	15	11	11	11	11
47	1	25	20	16	16	9	10	6
48	1	22	15	17.57	12	9	8	6.5
49	1	20	7	2	1	0	0	2
50	1	20	12.13	8	6	3	2	3

Data in depress.dat (continued)

51	1	25	15	24	18	15.19	13	12.52
52	1	18	17	8	2	2	0	1
53	1	26	1	18	10	13	12	10
54	1	29	27	13	9	8	4	5
55	1	17	20	10	8.89	8.49	7.02	6.79
56	1	22	12	-9	-9	-9	-9	-9
57	1	22	15.38	2	4	6	3	3
58	1	23	11	9	10	8	7	4
59	1	17	15	-9	-9	-9	-9	-9
60	1	23	7	12	15	-9	-9	-9
61	1	26	24	-9	-9	-9	-9	-9

8.2 The analysis of longitudinal data

The data in Table 8.1 consist of repeated observations over time on each of the 61 patients; such data are generally referred to as *longitudinal data*, panel data or repeated measurements, and as *cross-sectional time-series* in Stata. There is a large body of methods that can be used to analyze longitudinal data, ranging from the simple to the complex. Some useful references are Diggle *et al.* (2002), Everitt (1995), and Hand and Crowder (1996). In this chapter we concentrate on the following approaches:

- Graphical displays
- Summary measure or response feature analysis

In the next two chapters, more formal modeling techniques will be applied to the data.

8.3 Analysis using Stata

Assuming the data are in an ASCII file, *depress.dat*, as listed in Table 8.1, they may be read into Stata for analysis using the following instructions:

```
infile subj group pre dep1 dep2 dep3 dep4 dep5 dep6 ///
    using depress.dat, clear
mvdecode _all, mv(-9)
```

The second of these instructions converts values of -9 in the data to missing values.

It is useful to begin examination of these data using the *summarize* command to calculate means, variances, etc., within each of the two treatment groups:

```
summarize pre-dep6 if group==0
```

(see Display 8.1).

Variable	Obs	Mean	Std. Dev.	Min	Max
pre	27	20.77778	3.954874	15	28
dep1	27	16.48149	5.279644	7	26
dep2	22	15.88818	6.124177	4	27
dep3	17	14.12882	4.976848	4.19	22
dep4	17	12.27471	6.848791	3	23
dep5	17	11.40294	4.438703	3.03	18
dep6	17	10.89588	4.68157	3.45	20

Display 8.1

summarize pre-dep6 if group==1
 (see Display 8.2). There is a general decline in the depression score

Variable	Obs	Mean	Std. Dev.	Min	Max
pre	34	21.34882	3.574432	15	28
dep1	34	13.36794	5.556373	1	27
dep2	31	11.73677	6.575079	1	27
dep3	29	9.134138	5.475564	1	24
dep4	29	8.827857	4.666663	0	22
dep5	29	7.305295	5.740968	0	24
dep6	29	6.590714	4.730158	1	23

Display 8.2

over time in both groups, with the values in the active treatment group appearing to be consistently lower.

8.3.1 Graphical displays

A useful preliminary step in the analysis of longitudinal data is to graph the observations in some way. The aim is to highlight two particular aspects of the data, namely, how they evolve over time and how the measurements made at different times are related. A number of graphical displays can be used, including:

- separate plots of each subject's responses against time, differen-

- tiating in some way between subjects in different groups
- boxplots of the observations at each time point by treatment group
- a plot of means and standard errors by treatment group for every time point
- a scatterplot matrix of the repeated measurements

To begin, plot the required scatterplot matrix, identifying treatment groups with the labels 0 and 1, using

```
graph matrix pre-dep6, xlabel(group) asymbol(none) ///
    mlabposition(0)
```

The resulting plot is shown in Figure 8.1. The most obvious feature of this diagram is the increasingly strong relationship between the measurements of depression as the time interval between them decreases. This has important implications for the models appropriate for longitudinal data, as we will see in Chapter 10.

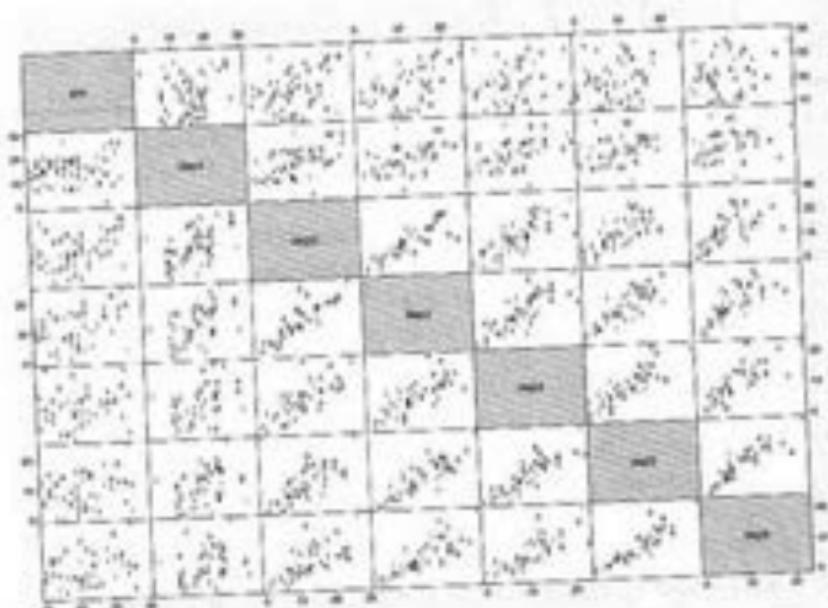


Figure 8.1: Scatter-plot matrix for depression scores at six visits.

To obtain the other graphs mentioned above, the dataset needs to be restructured from its present wide form (one column for each visit)

to the long form (one row for each visit) using the `reshape` command. Before running `reshape`, we will preserve the data using the `preserve` command so that they can later be restored using `restore`:

```
preserve
reshape long dep, i(subj) j(visit)
list in 1/13, clean
```

The first 13 observations of the data in long form are shown in Display 8.3.

	subj	visit	group	pre	dep
1.	1	1	0	18	17
2.	1	2	0	18	18
3.	1	3	0	18	18
4.	1	4	0	18	17
5.	2	5	0	18	14
6.	2	6	0	18	16
7.	2	1	0	27	26
8.	2	2	0	27	23
9.	2	3	0	27	18
10.	2	4	0	27	17
11.	2	5	0	27	12
12.	2	6	0	27	10
13.	3	1	0	16	17

Display 8.3

To inspect the patterns of missing values in this dataset, we first delete observations where `dep` is missing and then use the `xtdes` command:

```
drop if dep=="."
xtdes, i(subj) t(visit)
```

giving the output shown in Display 8.4. We see that 45 subjects have complete data, 8 subjects dropped out after the first visit, 7 after the second, and 1 after the third. This kind of missingness pattern is called "monotonic" because people never return once they have missed a visit.

We will now plot the subjects' individual response profiles over the visits separately for each group using the `by()` option. To obtain the correct group labels with the `by()` option we must first label the values of `group`:

```
label define treat 0 "Placebo" 1 "Estrogen"
label values group treat
```

In each graph we want to connect the points belonging to a given subject, but avoid connecting points of different subjects. A simple way of achieving this is to use the `connect(ascending)` option. Before

xtabs, i(subj) t(visit)							
subj: 1, 2, ..., 61					n =	61	
visit: 1, 2, ..., 6					T =	6	
Delta(visit) = 1; (6-t)+1 = 6							
(subj+visit uniquely identifies each observation)							
Distribution of T_i:	min	5%	25%	50%	75%	95%	max
	1	1	3	6	6	6	6
Freq.	Percent	Cum.	Pattern				
45	73.77	73.77	111111				
8	13.11	86.88	1.....				
7	11.48	98.36	11....				
1	1.64	100.00	111...				
61	100.00		XXXXXX				

Display 8.4

plotting, the data need to be sorted by the grouping variable and by the *x* variable (here visit):

```
sort group subj visit
twoway connected dep visit, connect(ascending) by(group) ///
ytitle(Deprression) xlabel(1/6)
```

The *connect(ascending)* option connects points only so long as visit is ascending. For the first subject (subj=1) this is true; but for the second subject, visit begins at 1 again, so the last point for subject one is not connected with the first point for subject two. The remaining points for this subject are, however, connected and so on. The *xlabel()* option was used here to make the axis range start at 1 instead of 0. The diagram is shown in Figure 8.2. (Some points for different subjects are connected at visit one; this happens when successive subjects have missing data for all subsequent visits so that visit does not decrease when subj increases.) The individual plots reflect the general decline in the depression scores over time indicated by the means obtained using the *summarize* command; there is, however, considerable variability. The phenomenon of "tracking" is apparent whereby some individuals have consistently higher values than other individuals, leading to within-subject correlations. Notice that some profiles are not complete because of missing values.

To obtain the boxplots of the depression scores at each visit for each treatment group, the following instruction can be used:

```
graph box dep, over(visit) over(group, ///
relabel(1 "Placebo group" 2 "Estrogen group"))
```

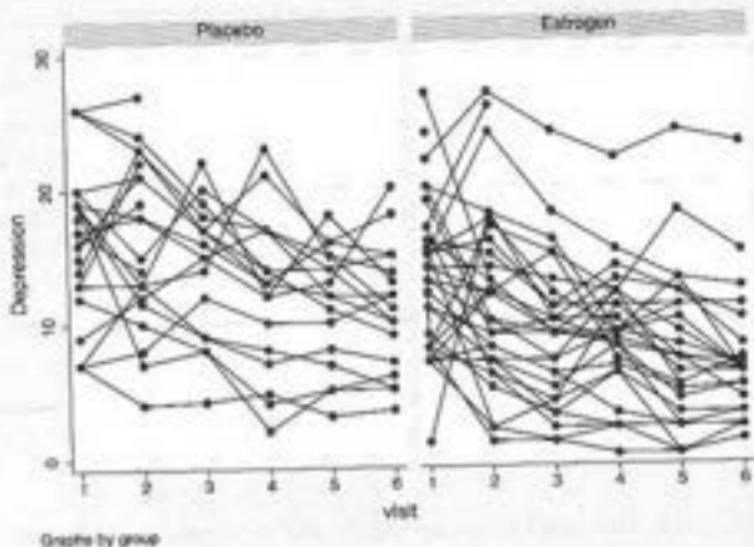


Figure 8.2: Individual response profiles by treatment group.

Here the `over()` options specify two grouping variables, `visit` and `group`, to plot the distributions by visit within groups. The `relabel()` option is used to define labels for the groups. Here "1" refers to the first level of `group` (0 in this case) and "2" to the second. The resulting graph is shown in Figure 8.3. Again, the general decline in depression scores in both treatment groups can be seen and, in the active treatment group, there is some evidence of outliers which may need to be examined. (Figure 8.2 shows that four of the outliers are due to one subject whose response profile lies above the others.)

A plot of the mean profiles of each treatment group, which includes information about the standard errors of each mean, can be obtained using the `collapse` instruction that produces a dataset consisting of selected summary statistics. Here, we need the mean depression score on each visit for each group, the corresponding standard deviations, and a count of the number of observations on which these two statistics are based.

```
collapse (mean) dep (sd) sddep=dep (count) n=dep, ///
by(visit group)
list in 1/10, clean
```

(see Display 8.5). The mean value is now stored in `dep`; but since more than one summary statistic for the depression scores were required, the

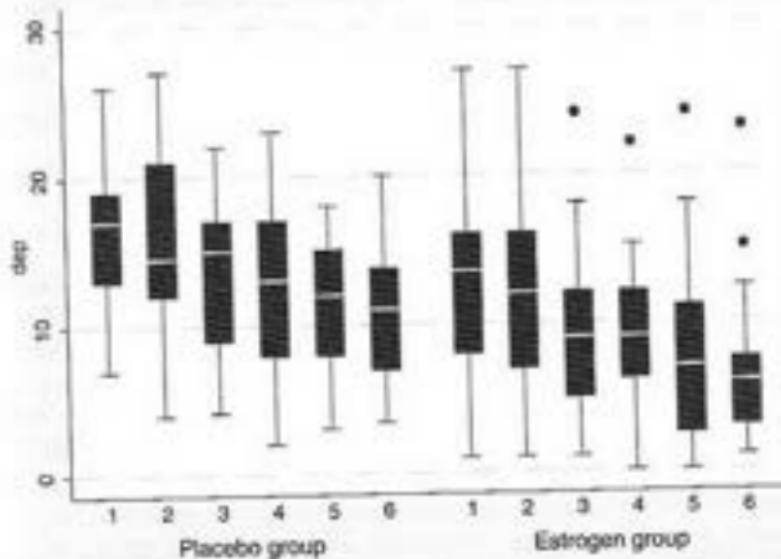


Figure 8.3: Boxplots for six visits by treatment group.

visit	group	dep	sddep	n	
1.	1	Placebo	16.48148	5.279644	27
2.	1	Estrogen	13.35794	5.556373	34
3.	2	Placebo	15.88818	6.124177	22
4.	2	Estrogen	11.73877	6.575079	31
5.	3	Placebo	14.12882	4.974648	17
6.	3	Estrogen	9.134138	5.475564	29
7.	4	Placebo	12.27471	5.848791	17
8.	4	Estrogen	8.827887	4.666653	28
9.	5	Placebo	11.40294	4.438702	17
10.	5	Estrogen	7.309266	5.740968	28

Display 8.5

remaining statistics were given new names in the `collapse` instruction.

The required mean and standard error plots can now be produced as follows:

```
generate high = dep + 2*sddep/sqrt(n)
generate low = dep - 2*sddep/sqrt(n)
twoway (rarea low high visit, bfcolor(gs12) sort) ///
    (connected dep visit, mcolor(black)) ///
    (clcolor(black) sort), by(group) ///
    legend(order(1 "95% CI" 2 "mean depression"))
```

Here `twoway rarea` produces a shaded area between the lines `low` versus `visit` and `high` versus `visit`, the 95% confidence limits for the mean. It is important that the line for the mean is plotted after the shaded area because it would otherwise be hidden underneath it. The `sort` option is used both in the `rarea` and `connected` plots to ensure that the areas and lines are drawn for `visit` in ascending order. The resulting diagram is shown in Figure 8.4.

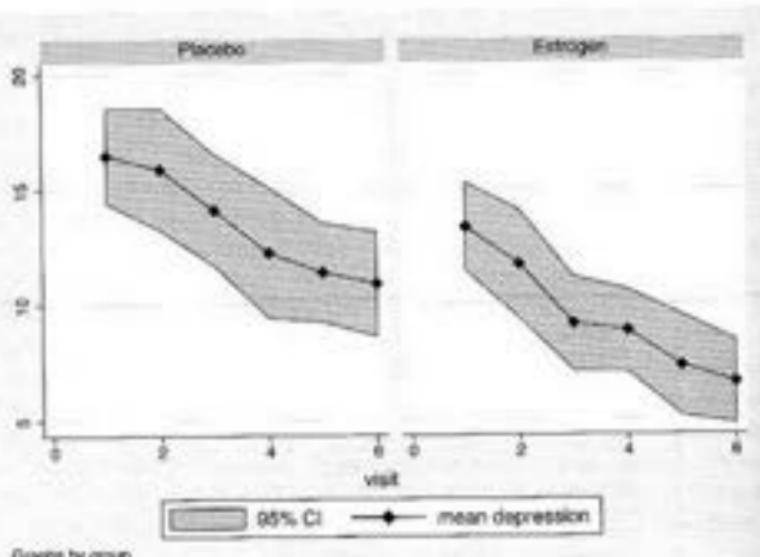


Figure 8.4: Mean and standard error plots; the shaded areas represent ± 2 standard errors.

Table 8.2 Response features suggested in Matthews *et al.* (1990)

Type of data	Property to be compared between groups	Summary measure
Peaked	overall value of response	mean or area under curve
Peaked	value of most extreme response	maximum (minimum)
Peaked	delay in response	time to maximum or minimum
Growth	rate of change of response	linear regression coefficient
Growth	final level of response	final value or (relative) difference between first and last
Growth	delay in response	time to reach a particular value

8.3.2 Response feature analysis

A relatively straightforward approach to the analysis of longitudinal data is that involving the use of *summary measures*, sometimes known as *response feature analysis*. The responses of each subject are used to construct a single number that characterizes some relevant aspect of the subject's response profile. (In some situations more than a single summary measure may be required.) The summary measure needs to be chosen before the analysis of the data. The most commonly used measure is the mean of the responses over time because many investigations, e.g., clinical trials, are most concerned with differences in overall levels rather than more subtle effects. Other possible summary measures are listed in Matthews *et al.* (1990) and are shown here in Table 8.2.

Having identified a suitable summary measure, the analysis of the data generally involves the application of a simple univariate test (usually a *t*-test or its nonparametric equivalent) for group differences on the single measure now available for each subject. For the estrogen patch trial data, the mean over time seems an obvious summary measure. The mean of all non-missing values is obtained (after restoring the data) using

```
restore
egen avg = rowmean(dep1 dep2 dep3 dep4 dep5 dep6)
```

The differences between these means may be tested using a *t*-test assuming equal variances in the populations:

```
ttest avg, by(group)
```

(see Display 8.6). The assumption of equal variances can be relaxed using the *unequal* option:

Two-sample t test with equal variances

Group	Obs	Mean	Std. Err.	Std. Dev.	[95% Conf. Interval]
0	27	14.75606	.8782852	4.563704	12.95071 16.56139
1	34	10.55206	.9187872	5.357404	8.682772 12.42135
combined	61	12.41284	.6923949	5.407777	11.02785 13.79784
diff		4.20399	1.294842		1.653017 6.794964
		diff = mean(0) - mean(1)			t = 3.2467
Ho:	diff = 0				degrees of freedom = 59
Ha:	diff < 0		Ha: diff != 0		Ha: diff > 0
		Pr(T < t) = 0.9990	Pr(T > t) = 0.0019		Pr(T > t) = 0.0010

Display 8.6

ttest avg, by(group) unequal

(see Display 8.7). In each case the conclusion is that the mean depression score is substantially lower in the estrogen group than the placebo group. The difference in mean depression scores is estimated as 4.2 with a 95% confidence interval (assuming equal variances) from 1.6 to 6.8.

Two-sample t test with unequal variances

Group	Obs	Mean	Std. Err.	Std. Dev.	[95% Conf. Interval]
0	27	14.75606	.8782852	4.563704	12.95071 16.56139
1	34	10.55206	.9187872	5.357404	8.682772 12.42135
combined	61	12.41284	.6923949	5.407777	11.02785 13.79784
diff		4.20399	1.271045		1.660343 6.747637
		diff = mean(0) - mean(1)			t = 3.3678
Ho:	diff = 0				Satterthwaite's degrees of freedom = 58.6777
Ha:	diff < 0		Ha: diff != 0		Ha: diff > 0
		Pr(T < t) = 0.9992	Pr(T > t) = 0.0016		Pr(T > t) = 0.0008

Display 8.7

sion score is substantially lower in the estrogen group than the placebo group. The difference in mean depression scores is estimated as 4.2 with a 95% confidence interval (assuming equal variances) from 1.6 to 6.8.

We might also be interested in the rate of change (here decline) of the response. An appropriate summary measure is the regression

coefficient of depression on visit. This can be obtained as a weighted sum of the depression scores at the six visits. However, for subjects who dropped out, the least squares estimator will not be the same as for subjects with complete data. It is therefore considerably easier to ask Stata to estimate a linear regression model for each subject using the **statsby** prefix. First we must reshape the data to long form as before.

```
reshape long dep, i(subj) j(visit)
```

Now we can use the **statsby** command to replace the current dataset by a dataset of summary statistics for each subject:

```
statsby slope=_b[visit] inter=_b[_cons] df=e(df_r), ///
by(group subj) clear: regress dep visit
list in 1/10
```

The second line specifies that **dep** should be regressed on **visit** for each unique combination of **group** and **subject**. The reason for specifying **group** here is so that this variable appears in the summary measure dataset. The first line specifies which results from the regression command should be stored under which variable name. Here **slope** will contain the regression coefficient of **visit**, **inter** the constant, and **df** the residual degrees of freedom. The first ten observations of the new dataset are shown in Display 8.8. To compare the mean slopes for

group	subj	slope	inter	df
1.	1	-.6714286	18	4
2.	2	-3.267143	29.06667	4
3.	3	-3	20	0
4.	4	-1.342857	19.86667	4
5.	5	-1.542857	12.73333	4
6.	6	-2.426	16.39287	4
7.	7	-1.342857	14.2	4
8.	8	1	25	0
9.	9	-3.687429	30.06287	4
10.	10	1.575429	8	4

Display 8.8

subjects who had at least 1 residual degree of freedom, we again use the **ttest** command

```
ttest slope if df>0, by(group)
```

giving the output shown in Display 8.9. There is no evidence for a

Two-sample t test with equal variances

Group	Obs	Mean	Std. Err.	Std. Dev.	[95% Conf. Interval]
0	17	-1.172168	.3237815	1.334985	-1.858554 - .455782
1	29	-1.066108	.2632224	1.417496	-1.605295 - .6269217
combined	46	-1.105304	.2026163	1.373526	-1.513191 - .4974176
diff		-.106667	.4239978		-.9605711 .7684518
		diff = mean(0) - mean(1)			t = -0.2801
Ho: diff = 0					degrees of freedom = 44
Ha: diff < 0				Ha: diff != 0	
Pr(T > t) = 0.4618				Pr(T > t) = 0.8036	
					Ha: diff > 0
					Pr(T > t) = 0.5382

Display 8.9

difference in the mean rate of decline between the two groups.

The summary measure approach to longitudinal data has a number of advantages:

- Appropriate choice of summary measure ensures that the analysis is focused on relevant and interpretable aspects of the data,
- The method is easy to explain and intuitive, and
- To some extent missing and irregularly spaced observations can be accommodated.

However, the method is somewhat *ad hoc*, particularly in its treatment of missing data. For instance, if the summary measure is a mean, but there is actually a decline in the response over time, then the mean of all available data will overestimate the mean for those who dropped out early (a better summary measure in this case is the intercept from a linear regression model). Furthermore, response feature analysis treats all summaries as equally precise even if some are based on fewer observations due to missing data. In the next two chapters we will therefore discuss more formal approaches to longitudinal data, random effects modeling, and generalized estimating equations.

8.4 Exercises

8.1 • Treatment of post-natal depression

1. Produce boxplots corresponding to those shown in Figure 8.3 using the data in wide form.
2. Compare the results of the *t*-tests given in the text with the corresponding *t*-tests calculated only for those subjects having observations on all six post-randomization visits.
3. Repeat the summary measures analysis described in the text using the maximum over time instead of the mean (see `help egen`).
4. Test for differences in the mean over time controlling for the baseline measurement using
 - a change score defined as the difference between the mean over time and the baseline measurement, and
 - b. analysis of covariance of the mean over time using the baseline measurement as a covariate.

See also Exercises in Chapter 9.

8.2 Wage increases

1. For the data described in Exercise 1.2, produce boxplots for the log hourly wage over time by ethnic/racial group.
2. Plot the mean log wage over time by ethnic group, showing the 95% confidence bands as in Figure 8.4.
3. Compare the mean log wages between the three groups using multiple regression with dummy variables.
4. Repeat the analysis above but this time controlling for `educ`, the number of years of schooling.
5. Interpret the findings.

See also Exercise 9.4.

8.3 Jaw growth

In this jaw growth dataset from Pothoff and Roy (1964), eleven boys and sixteen girls had the distance between the center of the pituitary gland and the pterygomaxillary fissure recorded at ages 8, 10, 12, and 14.

The variables in the dataset `growth.dta` are:

- `idnr`: subject identifier
 - `measure`: distance between pituitary and maxillary fissure in millimeters
 - `age`: age in years
 - `sex`: sex (1=boys, 2=girls)
1. Plot the observed growth trajectories, i.e., plot `measure` against

- age, connecting successive observations on the same subject using the connect(ascending) option. Use the by() option to obtain separate graphs by sex.
2. Use the `statsby` prefix to obtain estimated intercepts and slopes for each subject and compare the means of these summary measures between boys and girls using independent samples *t*-tests. To make the intercepts meaningful, subtract 8 from age before running the `statsby` prefix command.

8.4 Treatment of Alzheimer's

The data used here arise from an investigation of the use of lecithin, a precursor of choline, in the treatment of Alzheimer's disease. Traditionally it has been assumed that this condition involves an inevitable and progressive deterioration in all aspects of intellect, self-care, and personality. Recent work suggests that the disease involves pathological changes in the central cholinergic system, which it might be possible to remedy by long-term dietary enrichment with lecithin. In particular, the treatment might slow down or even halt the memory impairment usually associated with the condition. Patients suffering from Alzheimer's disease were randomly allocated to receive either lecithin or placebo for a six-month period. A cognitive test score giving the number of words recalled from a previously studied list was recorded at the start, at one month, at two months, at four months and at six months. (The data are given in Everitt and Pickles, 2004.)

The variables in `alzheimer.dta` are:

- `group`: treatment group (1=placebo, 2=lecithin)
 - `v1` to `v5`: number of words recalled at the start and each subsequent month
1. In these data the clinicians were specifically interested in the maximum value of the response variable over the five measurement occasions. Generate a variable equal to the maximum measurement for each person.
 2. We wish to compare the distribution of the maximum number of words recalled across the five visits between treatment groups. Do the assumptions of an independent samples *t* test appear to be satisfied?
 3. Use an appropriate test for comparing the treatment groups.

Chapter 9

Random Effects Models: Thought Disorder and Schizophrenia

9.1 Description of data

In this chapter we will analyze data from the Madras Longitudinal Schizophrenia Study in which patients were followed up monthly after their first hospitalization for schizophrenia. The study is described in detail in Thara *et al.* (1994). Here we use a subset of the data analyzed by Diggle *et al.* (2002), namely data on thought disorder (1: present, 0: absent) at 0, 2, 6, 8, and 10 months after hospitalization on women only. The thought disorder responses are given as y_{01} to y_{10} in Table 9.1 where a “.” indicates a missing value. The variable **early** is a dummy variable for early onset of disease (1: age-of-onset less than 20 years, 0: age-of-onset 20 years or above). An important question here is whether the course of illness differs between patients with early and late onset. We will also reanalyze the post-natal depression data described in the previous chapter.

9.2 Random effects models

The data listed in Table 9.1 consist of repeated observations on the same subject taken over time and are a further example of a set of longitudinal data. During the last decades, statisticians have considerably

Table 9.1 Data in madras.dta

id	early	y0	y2	y6	y8	y10
1	0	1	1	0	0	0
6	1	0	0	0	0	0
10	0	1	1	0	0	0
13	0	0	0	0	0	0
14	0	1	1	1	1	1
15	0	1	1	0	0	0
16	0	1	0	1	0	0
22	0	1	1	0	0	0
23	0	0	0	1	0	0
25	1	0	0	0	0	0
27	1	1	1	1	0	1
28	0	0	0	-	-	-
31	1	1	1	0	0	0
34	0	1	1	0	0	0
36	1	1	1	0	0	0
43	0	1	1	1	0	0
44	0	0	1	0	0	0
45	0	1	1	0	1	0
46	0	1	1	1	0	0
48	0	0	0	0	0	0
50	0	0	0	1	1	1
51	1	0	1	1	0	0
52	0	1	1	0	1	0
53	0	1	0	0	0	0
56	1	1	0	0	0	0
57	0	0	0	0	0	0
59	0	1	1	0	0	0
61	0	0	0	0	0	0
62	1	1	1	0	0	0
65	1	0	0	0	0	0
66	0	0	0	0	0	0
67	0	0	1	0	0	0
68	0	1	1	1	1	1
71	0	1	1	1	0	0
72	0	1	0	0	0	0
75	1	1	0	0	-	-
76	0	0	1	-	-	-
77	1	0	0	0	0	0
79	0	1	-	-	-	-
80	0	1	1	0	0	0
85	1	1	1	1	0	0
86	0	0	1	-	-	-
87	0	1	0	0	0	-
90	0	1	1	0	0	0

enriched the methodology available for the analysis of such data (see Lindsey, 1999, or Diggle *et al.* 2002) and many of these developments are implemented in Stata.

9.2.1 Normally distributed responses

Longitudinal data require special methods of analysis because the responses at different time points on the same individual may not be independent even after conditioning on the covariates. For a linear regression model this means that the residuals for the same individual are correlated. We can model these residual correlations by partitioning the total residual for subject i at time point j into a subject-specific random intercept or permanent component u_i which is constant over time plus a residual ϵ_{ij} which varies randomly over time. The resulting random intercept model can be written as

$$y_{ij} = \mathbf{x}'_{ij}\beta + u_i + \epsilon_{ij}. \quad (9.1)$$

The random intercept and residual are each assumed to be independently normally distributed with zero means and constant variances τ^2 and σ^2 , respectively. Furthermore, these random terms are assumed to be independent of each other and the covariates \mathbf{x}_{ij} . (It should be noted that moment-based approaches do not require normality assumptions, see, e.g., Wooldridge (2002), Section 10.4.)

The random intercept model implies that the total residual variance is

$$\text{Var}(u_i + \epsilon_{ij}) = \tau^2 + \sigma^2.$$

Due to this decomposition of the total residual variance into a between-subject component τ^2 and a within-subject component σ^2 , the model is sometimes referred to as a variance components model. The covariance between the total residuals at any two time points j and j' on the same subject i is

$$\text{Cov}(u_i + \epsilon_{ij}, u_i + \epsilon_{ij'}) = \tau^2.$$

Note that these covariances are induced by the shared random intercept: for subjects with $u_i > 0$, the total residuals will tend to be larger than the mean (0) and for subjects with $u_i < 0$ they will tend to be smaller than the mean.

It follows from the two relations above that the residual correlations

are given by

$$\text{Cor}(u_i + \epsilon_{ij}, u_i + \epsilon_{ij'}) = \frac{\tau^2}{\tau^2 + \sigma^2}.$$

This intraclass correlation can be interpreted as the proportion of the total residual variance (denominator) that is due to residual variability between subjects (numerator).

The random intercept can be interpreted as the combined effect of all unobserved subject-specific covariates, often referred to as *unobserved heterogeneity*. The random intercepts represent individual differences in the overall mean level of the response after controlling for covariates. Random coefficients of covariates can be used to allow for between-subject heterogeneity in the effects of the covariates. For instance, in longitudinal data, the shape of the response profile may vary between subjects in addition to variability in its vertical position. If the overall shape is linear in time t_{ij} , subjects may differ randomly in their slopes giving a model of the form

$$y_{ij} = \mathbf{x}'_{ij}\beta + u_{0i} + u_{1i}t_{ij} + \epsilon_{ij}. \quad (9.2)$$

Here u_{0i} is a random intercept and u_{1i} a random coefficient or slope of t_{ij} . These random effects are assumed to have a bivariate normal distribution with zero means, variances τ_0^2 and τ_1^2 , and covariance τ_{01} . They are furthermore assumed to be uncorrelated across subjects and uncorrelated with ϵ_{ij} or any of the covariates. If the covariate vector \mathbf{x}_{ij} includes t_{ij} , the corresponding fixed coefficient β_1 represents the mean coefficient of time whereas the random slope u_{1i} represents the deviation from the mean coefficient for subject i . (Not including t_{ij} in the fixed part of the model would imply that the mean slope is zero.) The model can also include nonlinear functions of t_{ij} , typically powers of t_{ij} , whose coefficients may be fixed or random.

The total residual in (9.2) is $u_{0i} + u_{1i}t_{ij} + \epsilon_{ij}$ with variance

$$\text{Var}(u_{0i} + u_{1i}t_{ij} + \epsilon_{ij}) = \tau_0^2 + 2\tau_{01}t_{ij} + \tau_1^2t_{ij}^2 + \sigma^2$$

which is no longer constant over time but *heteroscedastic*. Similarly, the covariance between two total residuals of the same subject,

$$\text{Cov}(u_{0i} + u_{1i}t_{ij} + \epsilon_{ij}, u_{0i} + u_{1i}t_{ij'} + \epsilon_{ij'}) = \tau_0^2 + \tau_{01}(t_{ij} + t_{ij'}) + \tau_1^2t_{ij}t_{ij'}$$

is not constant over time. It should also be noted that both the random intercept variance and the correlation between the random coefficient and random intercept depend on the location of t_{ij} , i.e., re-estimating the model after adding a constant to t_{ij} will lead to different estimates.

General terms for random intercept or random coefficient models are *random effects* models, *mixed effects* or *mixed* models, where "mixed" refers to the presence of both fixed effects β and random effects u_{0i} and u_{1i} . The models are also *hierarchical* or *multilevel* since the elementary observations at the individual time points (level 1) are nested in subjects (level 2). The models discussed in this chapter are appropriate for any kind of clustered or two-level data, not just longitudinal. Other examples of two-level data are people in families, households, neighborhoods, cities, schools, hospitals, firms, etc. In all these types of data, we can generally not assume that responses for subjects in the same cluster are independent after controlling for covariates because there is likely to be unobserved heterogeneity between clusters.

9.2.2 Non-normal responses

For non-normal responses (for example, binary responses) we can extend the generalized linear model discussed in Chapter 7 by introducing a random intercept u_i into the linear predictor,

$$\eta_{ij} = \mathbf{x}'_{ij}\beta + u_i, \quad (9.3)$$

where the u_i are independently normally distributed with mean zero and variance τ^2 . (We have encountered a similar model in Chapter 7, namely the negative binomial model with a log link and Poisson distribution where u_i has a log-gamma distribution, and there is only one observation per subject.) We can further extend the random intercept model to include random coefficients as we did in the previous section.

Unfortunately, such generalized linear mixed models are difficult to estimate. This is because the likelihood involves integrals over the random effects distribution and these integrals generally do not have closed forms. Stata uses numerical integration by adaptive Gauss-Hermite quadrature for random intercept models. A user-written program `gllamm` can be used to estimate random coefficient models by adaptive quadrature. The program can also be used to estimate multilevel models with more than two levels of nesting (Rabe-Hesketh et al., 2005). Note that approximate methods such as penalized quasilikelihood (e.g., Breslow and Clayton, 1993) and its refinements do not tend to work well for data with dichotomous responses and small cluster sizes such as the thought disorder data (see also Rabe-Hesketh et al., 2002).

An important problem with many longitudinal data sets is the occurrence of dropouts, e.g., subjects failing to complete all scheduled visits in the post-natal depression data. A taxonomy of dropouts is given in Diggle et al. (2002). Fortunately, maximum likelihood estimation is consistent as long as the data are missing at random (MAR),

that is, the probability of missingness does not depend on the values that are missing. For example, if the model is correctly specified, we obtain consistent parameter estimates even if the probability of dropping out depends on the responses at earlier time points.

Useful books on random effects modeling include Snijders and Bosker (1999), Verbeke and Molenberghs (2000), Goldstein (2003), and Skrondal and Rabe-Hesketh (2004), as well general books on longitudinal data such as Lindsey (1999). Rabe-Hesketh and Skrondal (2005) is a book on "Multilevel and Longitudinal Modeling Using Stata".

9.3 Analysis using Stata

9.3.1 Post-natal depression data

As an example of continuous responses, we first consider the post-natal depression data analyzed in the previous chapter. The data are read using

```
infile subj group pre dep1 dep2 dep3 dep4 dep5 dep6 ///
        using depress.dat, clear
```

All responses must be stacked in a single variable, including the baseline score `pre`. This is achieved by first renaming `pre` to `dep0` and then using the `reshape` command:

```
rename pre dep0
reshape long dep, i(subj) j(visit)
```

We also define value labels for `group` and change “-9” to missing values

```
label define treat 0 "Placebo" 1 "Estrogen"
label values group treat
mvdecode _all, mv(-9)
```

We now estimate a random intercept model using `xtreg`. (Note that commands for longitudinal data have the prefix `xt` in Stata which stands for cross-sectional time series). First assume that the mean depression score declines linearly from baseline with different slopes in the two groups:

```
generate gr_visit = group*visit
xtreg dep group visit gr_visit, i(subj) mle
```

The syntax is the same as for `regress` except that the `i()` option is used to specify the cluster identifier, here `subj`, and the `mle` option to obtain maximum likelihood estimates.

The estimates of the “fixed” regression coefficients that do not vary over individuals are given in the first part of the table in Display 9.1.

Random-effects ML regression	Number of obs	=	356		
Group variable (i): sub	Number of groups	=	61		
Random effects u_i ~ Gaussian	Obs per group:	min =	2		
		avg =	5.8		
		max =	7		
	LR chi2(3)	=	225.74		
Log likelihood = -1045.7117	Prob > chi2	=	0.0000		
dep	Coeff.	Std. Err.	t	P> t	[95% Conf. Interval]
group	-1.644653	.1163462	-1.41	0.187	-3.924996 -.6356901
visit	-1.531905	.1736977	-8.82	0.000	-1.872346 -1.191464
gr_vis	-.5564469	.2220228	-2.51	0.012	-.9916031 -.1212958
_cons	19.29632	.8717659	22.13	0.000	17.58769 21.00496
/sigma_u	3.560969	.3949951			2.885167 4.425745
/sigma_e	3.948191	.361907			3.643277 4.278624
rho	.4488701	.0688845			.3360721 .5664725

Likelihood-ratio test of sigma_u=0: chibar2(01)= 114.03 Prob>chibar2 = 0.000

Display 9.1

whence the estimates of the standard deviations τ of the random intercept and σ of the residuals are given under /sigma_u and /sigma_e in the second part. The intraclass correlation rho is estimated as 0.45, implying that 45% of the residual variance is between subjects and 55% within subjects. There is a significant interaction between group and visit at the 5% level; the mean decrease in depression score is estimated as 1.53 per visit in the placebo group and 1.53 + 0.56 per visit in the estrogen group. We can obtain the estimated slope of time in the estrogen group with its p-value and confidence interval using lincom:

```
lincom visit + gr_vis
```

(see Display 9.2).

```
( 1) [dep]visit + [dep]gr_vis = 0
```

dep	Coeff.	Std. Err.	t	P> t	[95% Conf. Interval]
(1)	-2.088362	.1304264	-15.09	0.000	-2.359663 -1.817045

Display 9.2

The model assumes that the effect of visit is linear. However, it may well be that the depression score gradually levels off, remaining stable after some period of time. We can investigate this by adding a quadratic term of visit:

```
generate vis2 = visit^2
xtreg dep group visit gr_vis vis2, i(subj) mle
```

The p-value for vis2 in Display 9.3 suggests that the average curve is not linear. To picture the mean curve, we now plot it together with

Random-effects ML regression		Number of obs	=	356
Group variable (i): subj		Number of groups	=	61
Random effects u_i ~ Gaussian		Obs per group: min =	=	2
		avg =	=	5.8
		max =	=	7
		LR chi2(4)	=	268.39
Log likelihood = -1034.3838		Prob > chi2	=	0.0000
dep	Coef.	Std. Err.	t	P> t
group	-1.471391	.139394	-1.29	0.197
visit	-3.787308	.3710868	-10.31	0.000
gr_vis	-.5848499	.2073666	-2.82	0.005
vis2	.3851914	.0569306	6.77	0.000
_cons	29.91077	.8860177	33.60	0.000
/sigma_u	3.5546865	.3869129		2.901173
/sigma_e	3.678709	.1609811		3.394561
rho	.4870545	.0689037		.3737279

Likelihood-ratio test of sigma_u=0: chibar2(01)= 133.51 Prob>chibar2 = 0.000

Display 9.3

the observed individual response profiles:

```
predict pred0, xb
sort subj visit
twoway (line pred0 visit, conn(ascending) lwidth(thick)) ///
        (line dep visit, conn(ascending) lpatt(dash)),    ///
        by(group) ytitle(Depression)                   ///
        legend(order(1 "Fitted mean" 2 "Observed scores"))
```

giving the graph shown in Figure 9.1 which suggests that the response curves tend to level off.

Extending the model to include a random coefficient of visit requires the `xtmixed` command that was introduced in Stata release 9. First we re-estimate the random intercept model using `xtmixed`:

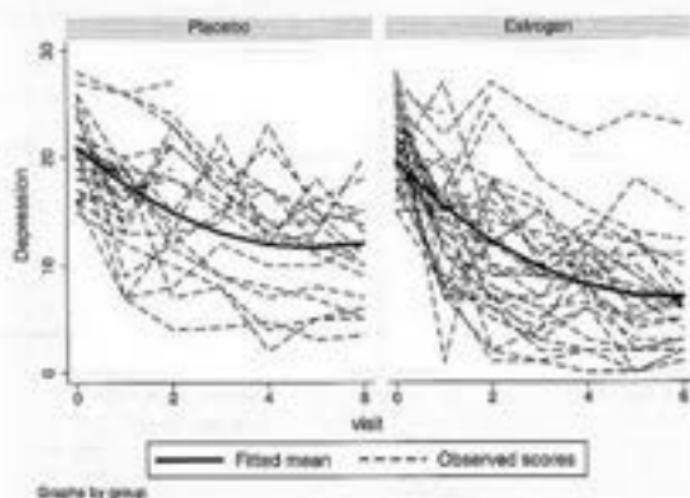


Figure 9.1: Response profiles and fitted mean curves by treatment group.

```
xtmixed dep group visit gr_vis vis2 || subj:, mle
```

Here the fixed part of the model is specified as in all estimation commands, and the random part is specified after the double-bar `||`. First the cluster-identifier is given, followed by a colon. Then all explanatory variables that should have random coefficients varying between clusters are listed. A random intercept is automatically included unless the `nocons` option is used. Here we only require a random intercept, so no variables are listed after `subj:`. After the comma we use the `mle` option to specify maximum likelihood estimation (the default is restricted maximum likelihood estimation).

The output shown in Display 9.4 agrees perfectly with that from `xtreg` with `sd(.cons)` corresponding to `/sigma.u` and `sd(Residual)` to `/sigma.e`.

The most common method of predicting random effects is by their posterior means, their expectations given the observed responses and covariates with the parameter estimates plugged in. These predictions are also known as empirical Bayes predictions, shrinkage estimates or, in linear mixed models, best linear unbiased predictions (BLUP). Adding predictions of the random intercept to the predicted mean response profile gives individual predicted response profiles. These can be obtained after estimation with `xtmixed` using the `predict` command

Mixed-effects ML regression	Number of obs	=	356		
Group variable: subj	Number of groups	=	61		
	Obs per group: min =	2			
	avg =	5.8			
	max =	7			
Log likelihood = -1024.3838	Wald chi2(4)	=	409.15		
	Prob > chi2	=	0.0000		
dep	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
group	-1.471391	.1.138296	-1.29	0.197	-3.704567 .7617843
visit	-3.787308	.3710723	-10.21	0.000	-4.814596 -3.066019
gr_vis	-.5848499	.2973863	-2.82	0.005	-.9933176 -.1783821
vis2	.3851918	.0569337	6.77	0.000	.2736056 .4967776
_cons	20.91077	.8860139	23.60	0.000	19.17422 22.64733
Random-effects Parameters		Estimate	Std. Err.	[95% Conf. Interval]	
subj: Identity	sd(_cons)	3.584666	.3869167	2.901167	4.429391
	sd(Residual)	3.678709	.1509831	3.394557	3.986645
LR test vs. linear regression: chibar2(01) = 133.51 Prob > chibar2 = 0.0000					

Display 9.4

with the **fitted** option:

```
predict pred1, fitted
```

A graph of the individual predicted profiles is obtained using

```
twoway (line pred1 visit, connect(ascending)), by(group) ///
ytitle(Depression) xlabel(0/6)
```

and given in Figure 9.2. It is clear that the mean profiles in Figure 9.1

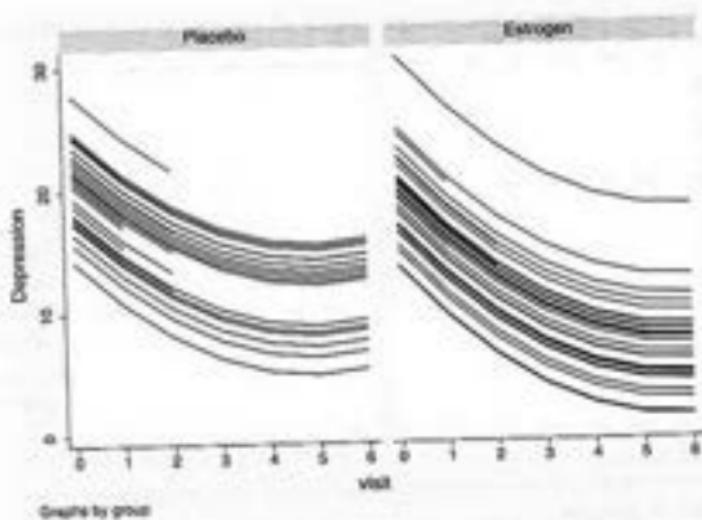


Figure 9.2: Predicted response curves for random intercept model.

have simply been shifted up and down to fit the observed individual profiles more closely.

We can also obtain empirical Bayes predictions of the random intercepts themselves using the **reffects** option

```
predict inter, reffects
```

Unfortunately, **xtmixed** does not produce standard errors of the predictions at the time of writing this book. To obtain these, we will use a user-contributed program **gllamm** (for generalized linear latent and mixed models) described in Rabe-Hesketh *et al.* (2002), (2004b) and Rabe-Hesketh and Skrondal (2005) (see also www.gllamm.org). The program can be obtained from the SSC archive using

```
ssc install gllamm
```

We will first re-estimate the random intercept model using **gllamm**:

```
gllamm dep group visit gr_vis vis2, i(subj) adapt
```

The syntax is ss for `xtreg` except that the `mle` option is not required since `gllamm` always uses maximum likelihood, and we have specified `adapt` to use adaptive quadrature. The estimates are shown in Display 9.5.

```
number of level 1 units = 386
number of level 2 units = 61
```

```
Condition Number = 88.719062
```

```
gllamm model
```

```
log likelihood = -1024.3838
```

dep	Coeff.	Std. Err.	<i>t</i>	P> t	[95% Conf. Interval]
group	-1.471291	.13394	-1.29	0.197	-3.704874 .7617913
visit	-3.787308	.3710907	-10.21	0.000	-4.514632 -3.069983
gr_vis	-.5848499	.2073976	-2.82	0.006	-.9913417 -.178386
vis2	.2855916	.0569339	6.77	0.000	.2738033 .4967799
cons	30.91077	.8860219	33.60	0.000	19.1742 22.64734

Variance at level 1

```
13.632897 (1.1101096)
```

Variances and covariances of random effects

***level 2 (subj)

```
var(1): 12.849637 (2.7739368)
```

Display 9.5

The format of the output for the random part is somewhat different from that of `xtreg` and `xtmixed`. "Variance at level 1" refers to the residual variance σ^2 , whereas `var(1)` under "Variances and covariances of random effects" refers to the random intercept variance with standard errors given in parentheses. (These standard errors are not very useful and neither are the standard errors for the standard deviations reported by `xtreg` or `xtmixed`, since the sampling distributions of the estimates are unlikely to be well approximated by a normal

distribution.)

All estimates from `gllamm` are nearly identical to those using `xtreg` or `xtmixed`. This will not always be the case since `gllamm` uses numerical integration for all models, whereas `xtreg` and `xtmixed` exploit the availability of a closed form likelihood for linear mixed models. (Note that we would not generally recommend using `gllamm` for linear mixed models but do so here to obtain standard errors for the predicted random effects.) In `gllamm` the accuracy of the estimates can be improved by increasing the number of quadrature points for numerical integration from its default of 8 using the `nip()` option.

We can obtain the empirical Bayes predictions of u_i with standard errors using `gllamm`'s prediction command `gllapred` with the `u` option

```
gllapred rand, u  
(means and standard deviations stored in randst1 rands1)
```

which creates new variables `randst1` for the predictions and `rands1` for the standard errors. The standard errors are posterior standard deviations which are equal to the prediction error standard deviations for linear models. In the multilevel literature, these standard errors are also known as "comparative standard errors". First we make sure that the empirical Bayes predictions are close to those previously produced by `xtmixed` and stored in the variable `inter`:

```
assert abs(randst1-inter)<1e-3
```

The predictions are equal to at least three decimal places.

A graph of the predictions with their approximate 95% confidence intervals for the placebo group is then obtained using

```
generate f = visit==0  
sort randst1  
generate rank = sum(f)  
errbar rands1 se rank if visit==0&group==0, scale(2) ///  
xtitle("Rank") ytitle("Random intercept")
```

with the result shown in Figure 9.3. In linear mixed models the predicted random effects should be normally distributed, so we can use graphs to assess the assumption that the "true" random effects are normally distributed. One possibility is a kernel density plot with a normal density superimposed:

```
kdensity randst1 if visit==0, spanschnikov normal ///  
xtitle("Predicted random intercept")
```

In Figure 9.4, the normal density appears to approximate the empirical density well. Note that this diagnostic cannot be used for generalized linear mixed models where the predicted random effects are generally non-normal even under correct specification.

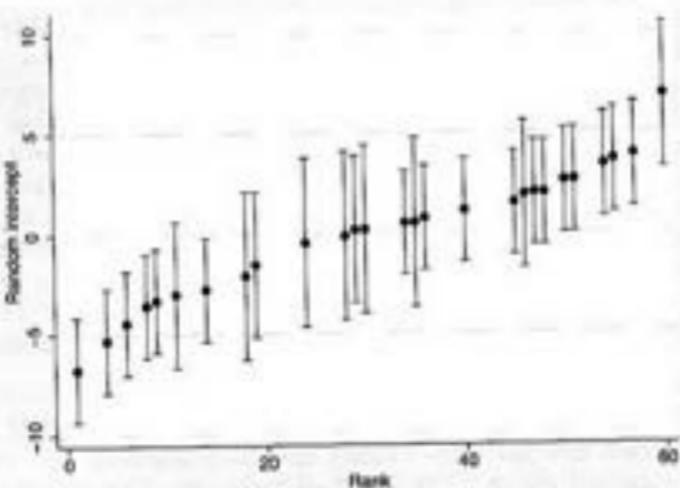


Figure 9.3: Predicted random intercepts and approximate 95% confidence intervals for the placebo group (based on the prediction error standard deviations).

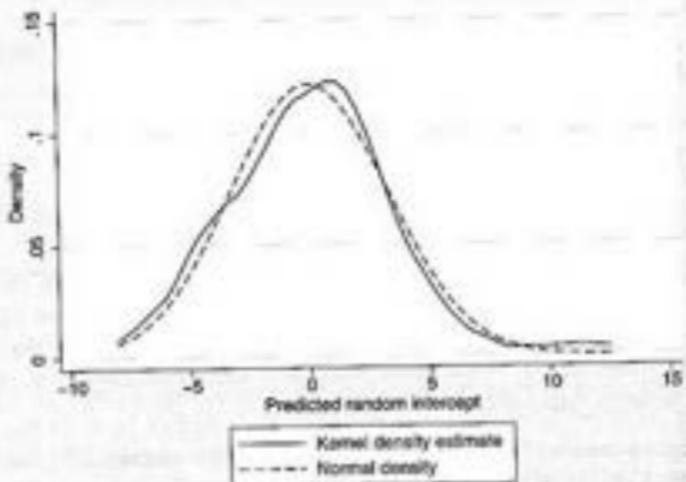


Figure 9.4: Kernel density estimate for empirical Bayes predictions and approximating normal density.

We will now allow the coefficient of `visit` to vary randomly between subjects by including a random slope in the model. This can be done using the `xtmixed` command. (We recommend using `xtmixed` instead of `gllamm` here because `gllamm` is slower and sometimes less accurate than `xtmixed` for linear mixed models.) Now we list a single variable, `visit`, in the random part after `subj:` to request a random coefficient for this variable in addition to a random intercept. By default, `xtmixed` specifies all random effects as mutually independent. We therefore specify the `cov(unstr)` option, abbreviated `cov(unstr)`, to freely estimate the correlation between intercept and slope.

```
xtmixed dep group visit gr_vis vis2 || subj: visit, ///
cov(unstr) mle
```

In Display 9.6 we see that the output under "Random effects parameters" has become more complex. The random intercept standard deviation has been estimated as 3.11, the random slope standard deviation as 0.61, and the correlation between intercepts and slopes as 0.09. The within-subject residual standard deviation (around the subject-specific regression lines) has been estimated as 3.46.

The log likelihood of this model is -1017.27 compared with -1024.38 for the random intercept model. A conventional likelihood ratio test would compare twice the difference in log likelihoods with a chi-squared distribution with two degrees of freedom (for an extra variance and covariance parameter). However, the null hypothesis that the slope has zero variance lies on the boundary of the parameter space (since a variance cannot be negative), and this test is therefore not valid. Snijders and Bosker (1999) suggest dividing the p-value of the conventional likelihood ratio test by two, giving a highly significant result here.

The empirical Bayes predictions for the random coefficient model can be obtained and plotted using

```
predict u*, reffects
twoway scatter u1 u2 if visit==0, xtitle("Intercept") ///
ytitle("Slope")
```

giving the graph in Figure 9.5. We could again assess the normality of the random effects graphically.

The predicted profiles can be computed and plotted using

```
predict pred2, fitted
sort subj visit
twoway (line pred2 visit, connect(ascending)), by(group) ///
ytitle(Deprression) xlabel(0/6)
```

resulting in Figure 9.6 where the curves are now no longer parallel due to the random slopes.

Mixed-effects ML regression	Number of obs	*	356		
Group variable: subj	Number of groups	*	61		
	Obs per group: min	*	2		
	avg	*	5.8		
	max	*	7		
	Wald chi2(4)	*	267.33		
Log likelihood = -1017.2722	Prob > chi2	*	0.0000		
dep	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
group	-1.471561	.021315	-1.44	0.150	-3.473301 -.5301793
visit	-3.779156	.3749763	-10.06	0.000	-4.554693 -.0615582
gr.visit	-.5870936	.2681353	-2.19	0.029	-1.112629 -.4840593
vis2	.3889244	.0636412	7.25	0.000	.2837696
_cons	20.90318	.7971398	26.22	0.000	19.34081 22.46554
Random-effects Parameters	Estimate	Std. Err.	[95% Conf. Interval]		
subj: Unstructured					
sd(visit)	.6003821	.1436785	.3811204	.0647851	
sd(_cons)	3.119386	.4675316	2.339877	4.158581	
corr(visit,,_cons)	.0929159	.2947014	-.4537794	.5887945	
sd(Residual)	3.468239	.1559451	3.16571	3.777799	
LR test vs. linear regression:	chi2(3) =	147.73	Prob > chi2 =	0.0000	
Note: LR test is conservative and provided only for reference					

Display 9.6

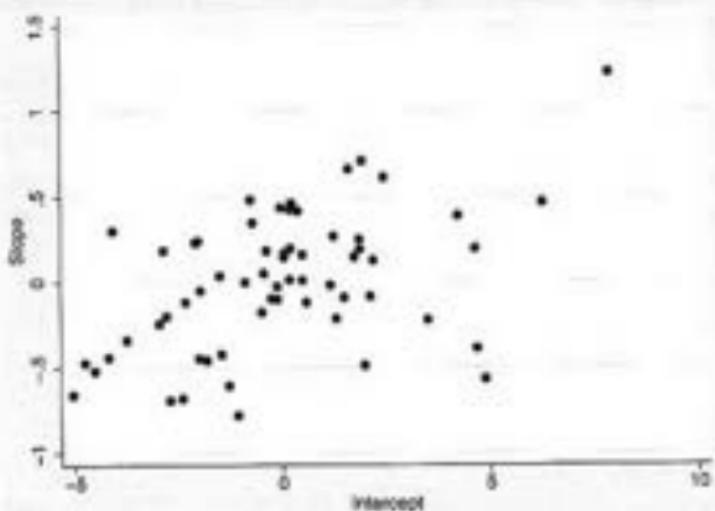


Figure 9.5: Scatterplot of predicted intercepts and slopes.

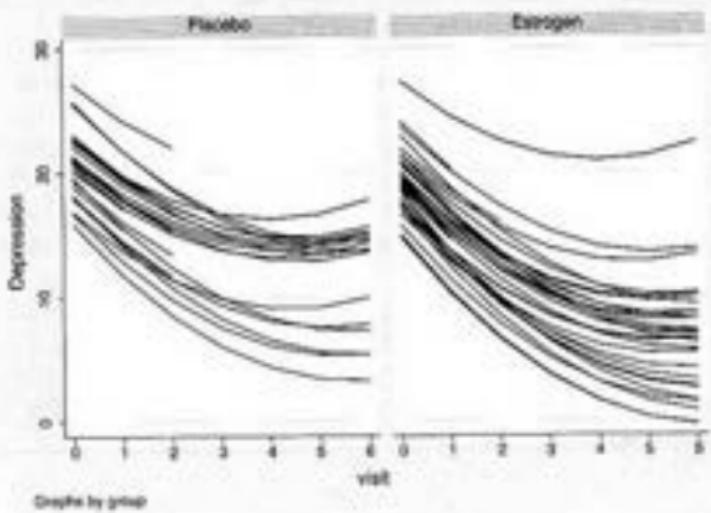


Figure 9.6: Predicted response profiles for random coefficient model.

Finally, we can assess the fit of the model by plotting both observed and predicted profiles in a trellis graph containing a separate scatterplot for each subject. For the placebo group the command is

```
twoway (line pred2 visit) (connect dep visit, lpat(dash)) ///
    if group==0, by(ssbj, style(compact)) ///
    ytitle(Deprression) legend(order(1 "Fitted" 2 "Observed"))
```

and similarly for the treatment group. The resulting graphs are shown in Figure 9.7. The model appears to represent the data reasonably well.

9.4 Thought disorder data

The thought disorder data are read in using

```
use madras, clear
```

Next we stack the dichotomous responses *y*0 to *y*10 into a single variable *y*, and create a new variable *month* taking on values 0 to 10 using

```
reshape long y, i(id) j(month)
```

We wish to investigate how the risk of having thought disorder evolves over time and whether there are differences between early and late onset patients. An obvious first model to estimate is a logistic random intercept model with fixed effects of *month*, *early* and their interaction. This can be done using Stata's *xtlogit* command:

```
generate month_early = month*early
xtlogit y month early month_early, i(id) or
```

The output is shown in Display 9.7. The *or* option was used to obtain odds ratios in the first part of the table. These suggest that there is a decrease in the odds of having thought disorder over time. However, patients with early onset schizophrenia do not differ significantly from late onset patients in their odds of thought disorder at the time of hospitalization ($OR=1.05$) nor do their odds change at a significantly different rate over time ($OR=0.94$). The log of the random intercept standard deviation is estimated as 1.02 and the standard deviation itself as 1.63. Here *rho* is the estimated intraclass correlation for the latent responses,

$$\hat{\rho} = \frac{\tau^2}{\tau^2 + \pi^2/3};$$

see the latent response formulation of the ordinal logit model in Chapter 6.

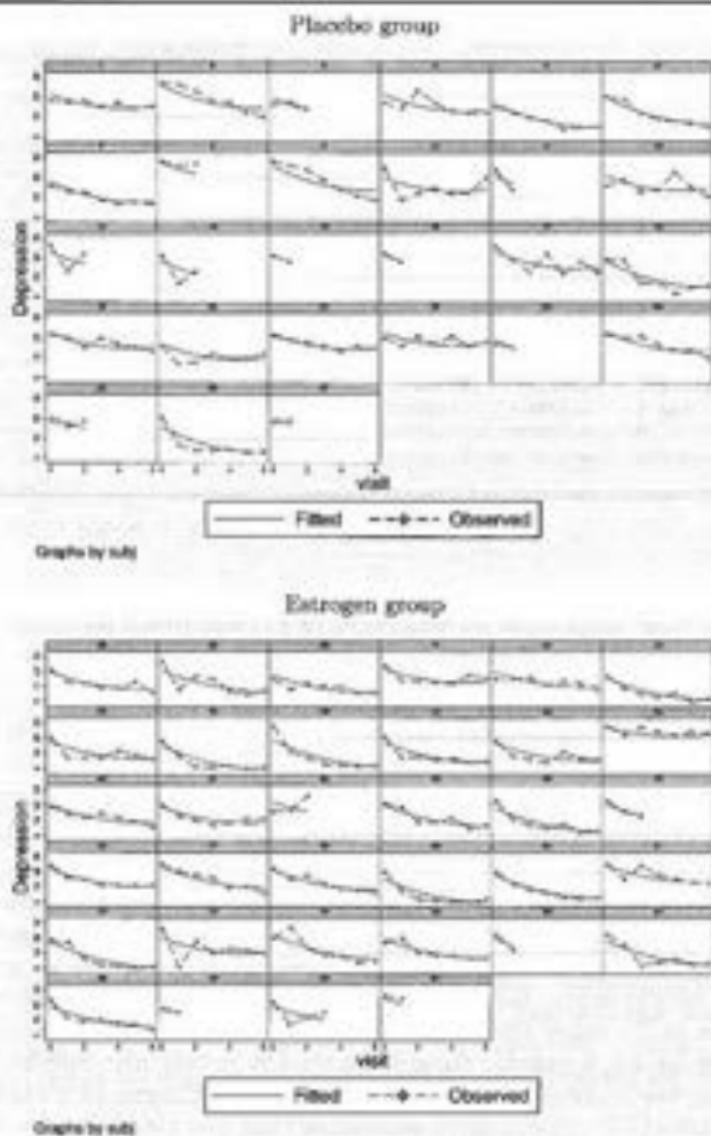


Figure 9.7: Observed and predicted response profiles for random coefficient model.

Random-effects logistic regression	Number of obs	=	248		
Group variable (i): id	Number of groups	=	44		
Random effects u_i ~ Gaussian	Obs per group: min	=	1		
	avg	=	5.5		
	max	=	6		
	Valid chisq(3)	=	37.72		
Log likelihood = -124.77883	Prob > chi2	=	0.0000		
y	OR	Std. Err.	<i>t</i>	P> <i>t</i>	[95% Conf. Interval]
month	.6886815	.0607932	-5.29	0.000	.577066 .7768862
early	1.047054	.9092773	0.05	0.960	.19068854 5.74335
month_early	.9358536	.1302074	-0.48	0.634	.7124893 1.239042
/lnsig2u	.9798043	.4214692			-1537477 1.005881
sigma_u	1.632557	.3439488			1.079906 2.466822
rho	.4474342	.1042017			.26171 .6490836
Likelihood-ratio test of rho=0: chibar2(01) =	26.04	Prob > chibar2 = 0.000			

Display 9.7

The same model can be estimated in **gllamm** which allows posterior means and other predictions to be computed using **gllapred**:

```
gllamm y month early month_early, i(id) link(logit) ///
family(binom) adapt eform
estimates store mod1
```

Here we used syntax very similar to that of **glm** with the **link()**, **family()**, and **eform** options and stored the estimates for later using **estimates store**. In Display 9.8 we can see that the estimates are quite close to those using **xtlogit**. However, there are some small discrepancies because the two programs use different algorithms and have different defaults for the number of quadrature points (12 in **xtlogit** and 8 in **gllamm**). Increasing the number of quadrature points for **gllamm** to 12 using the **nip(12)** option gives virtually the same estimates as for 8 points, suggesting that 8 points are sufficient. Increasing the number of quadrature points for **xtlogit** to 20 using the **intpoints(20)** option gives estimates that are closer to the **gllamm** estimates above.

We now include random slopes of **month** in the model,

$$\eta_{ij} = \mathbf{x}'_{ij} + u_{0i} + u_{1i}t_{ij}.$$

When there are several random effects (here intercept and slope), we have to define an equation for each of them to specify the variable multi-

number of level 1 units = 264

number of level 2 units = 44

Condition Number = 19.833627

gllamm model

log likelihood = -124.74702

<i>y</i>	exp(b)	Std. Err.	<i>t</i>	P> t	[95% Conf. Interval]
month	.6677456	.0616849	-5.22	0.000	.5737548 .7771337
early	1.047086	.919712	0.05	0.958	.1872069 5.856504
month_early	.935053	.1309963	-0.48	0.632	.7105372 1.295606

Variances and covariances of random effects

***level 2 (id)

var(1): 2.755484 (1.242613)

Display 9.8

plying the random effect. The random intercept u_{0i} in equation (9.2) is not multiplied by anything, or equivalently it's multiplied by 1, whereas the random coefficient u_{1i} is multiplied by t_{ij} , the variable month. We therefore define the equations as follows:

```
generate cons = 1
eq inter: cons
eq slope: month
```

The syntax for the equations is simply `eq label: varlist`, where *label* is an arbitrary equation name. We can now run gllamm with two extra options, `nrf(2)` to specify that there are two random effects and `eqs(inter slope)` to define the variables multiplying the random effects:

```
gllamm y month early month_early, i(id) nrf(2) ///
eqs(inter slope) link(logit) family(binom) ///
adapt eform
estimates store mod2
```

giving the output in Display 9.9. The log likelihood has decreased by about 3.5 suggesting that the random slope is needed. The fixed effects estimates are very similar to those for the random intercept model. The

estimated variances of the intercept and slope are denoted $\text{var}(1)$ and $\text{var}(2)$ respectively, and their covariance $\text{cov}(2,1)$. (The first random effect is the random intercept since `inter` was the first equation in the `eqs()` option.) We see that the estimated correlation between the random intercepts and slopes at the time of hospitalization is -0.71 . Note that both the random intercept variance and the covariance and correlation refer to the situation when `month` is zero and change if we translate `month` by adding or subtracting a constant.

number of level 1 units = 244
number of level 2 units = 44

Condition Number = 26.803719

gllamm model

log likelihood = -121.19976

γ	exp(b)	Std. Err.	β	P> z	[95% Conf. Interval]
month	.4631122	.0769516	-3.96	0.000	.4696698 .7744683
early	1.039168	1.262063	0.03	0.975	.0679712 11.02232
month_early	.9377842	.1937734	-0.31	0.756	.6286614 1.405975

Variances and covariances of random effects

***level 2 (id)

```
var(1): 7.1611437 (.0335342)  
cov(2,1): -.70697229 (.53384332) cov(2,2): -.70296393  
var(2): .14123952 (.004197?)
```

Display 9.9

We now produce some graphs of the model predictions, considering first the random intercept model. For the post-natal depression data the mean profile in Figure 9.1 was simply equal to the fixed part of the random intercept model $\mathbf{x}'_{ij}\bar{\beta}$ since the mean of the random intercept is zero. In the logistic model, things are more complicated because the probability of thought disorder given the random intercept (the *subject-*

specific probability) is a nonlinear function of the random intercept:

$$\Pr(y_{ij} = 1|u_i) = \frac{\exp(\mathbf{x}'_{ij}\beta + u_i)}{1 + \exp(\mathbf{x}'_{ij}\beta + u_i)}. \quad (9.4)$$

The population averaged probability is therefore not equal to the above with $u_i = 0$:

$$\int \Pr(y_{ij} = 1|u_i) g(u_i) du_i \neq \frac{\exp(\mathbf{x}'_{ij}\beta)}{1 + \exp(\mathbf{x}'_{ij}\beta)}, \quad (9.5)$$

where $g(u_i)$ is the normal probability density function of u_i . For this reason the coefficients β , representing the conditional or subject-specific effects of covariates, for a given value of the random effect, cannot be interpreted as population averaged or marginal effects. The marginal effects tend to be closer to zero or "attenuated". (Note that here the term "marginal effects" means population averaged effects, a very different notion than "marginal effects" in econometrics as computed by the Stata command `mfx`.)

In `gllapred` we can use the `mu` and `marg` options to obtain the marginal probabilities on the left-hand side of (9.5) by numerical integration and the `mu` and `us()` options to obtain the conditional probabilities in (9.4) for given values of u_i . To obtain smooth curves, we first create a new dataset where `month` increases gradually from 0 to 10 and `early` equals 1:

```
replace month = 10*(_n-1)/(_N-1)
replace early = 1
replace month_early = month
```

Now we can obtain marginal probabilities for the random intercept model by first restoring the estimates and then using `gllapred`:

```
estimates restore mod1
gllapred probm1, mu marg
```

To calculate conditional probabilities, we must first define variables equal to the values at which we wish to evaluate u_i (0 and ± 1.7 , approximately one standard deviation). The variable names must end on "-1" since the random intercept is the first (and here the only) random effect:

```
generate m1 = 0
generate l1 = -1.7
generate u1 = 1.7
gllapred probcl_m, mu us(m)
gllapred probcl_l, mu us(l)
```

```
gllapred probc1_u, mu us(u)
drop m1 l1 u1
```

Here the `us(m)` option specifies that, if there were more than one random effect, the values would be in `m1`, `m2`, etc. Before producing the graph, we will make predictions for the random coefficient model:

```
estimates restore mod2
gllapred probm2, mu marg
generate m1 = 0
generate m2 = 0
generate l1 = -2.7
generate l2 = -0.4
generate u1 = 2.7
generate u2 = 0.4
generate u11 = 2.7
generate u12 = -0.4
generate lui = -2.7
generate lui2 = 0.4
gllapred probc2_m, mu us(m)
gllapred probc2_l, mu us(l)
gllapred probc2_u, mu us(u)
gllapred probc2_u1, mu us(u1)
gllapred probc2_u2, mu us(u2)
```

We have produced five conditional predictions: one with both random effects equal to 0 and four for all combinations of high and low values of the random intercept and slope. The graphs are obtained using

```
label variable month ///
"Number of months since hospitalization"
twoway (line probm1 month) ///
(line probc1_m month, lpatt(shortdash)) ///
(line probc1_l month, lpatt(dash)) ///
(line probc1_u month, lpatt(dash)), ///
legend(order(1 "Marginal" 2 "Fixed part" 3 "Conditional")) ytitle("Predicted probability")
```

and similarly for the random coefficient model; see Figures 9.8 and 9.9. In Figure 9.8 it is clear that the marginal or population averaged curve is flatter than the conditional or subject-specific curves. The dotted curve for $u_i = 0$ represents the curve of an average individual since 0 is the mean of the random effects distribution. Note that this curve of an average or typical individual differs from the population averaged curve! In Figure 9.9, we can see how different the trajectories for different patients can be in a random coefficient model. Again, the curve of the average individual differs from the averaged curve. Although the conditional predictions for the two models are quite different, the marginal predictions are nearly the same as can be seen by plotting the two marginal curves on the same graph:

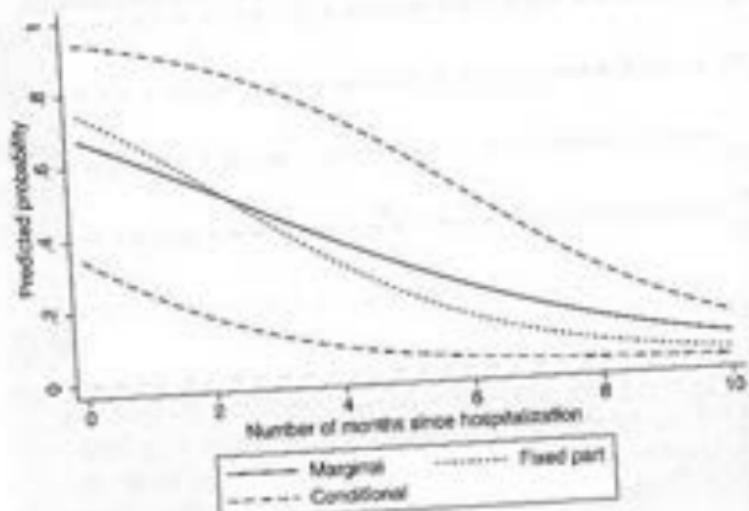


Figure 9.8: Marginal and conditional predicted probabilities for random intercept model. The dotted curve is the conditional probability when $u_{0i} = 0$.

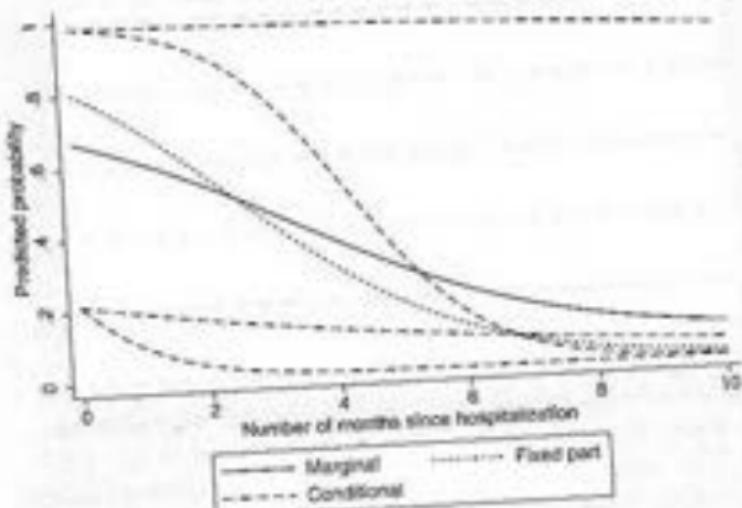


Figure 9.9: Conditional and marginal predicted probabilities for random coefficient model. The dotted curve is the conditional probability when $u_{0i} = 0$ and $u_{1i} = 0$.

```

twoway (line probm1 month)
    (line probm2 month, lpatt(dash)),
    legend(order(1 "Random intercept"
    2 "Random int. & slope")) ylabel(0(.2)1) ///
    ytitle("Marginal probability of thought disorder")

```

(see Figure 9.10). Here the `ylabel()` option was used to extend the *y*-axis range to 1 and produce appropriate axis labels to make this graph comparable with Figures 9.8 and 9.9.

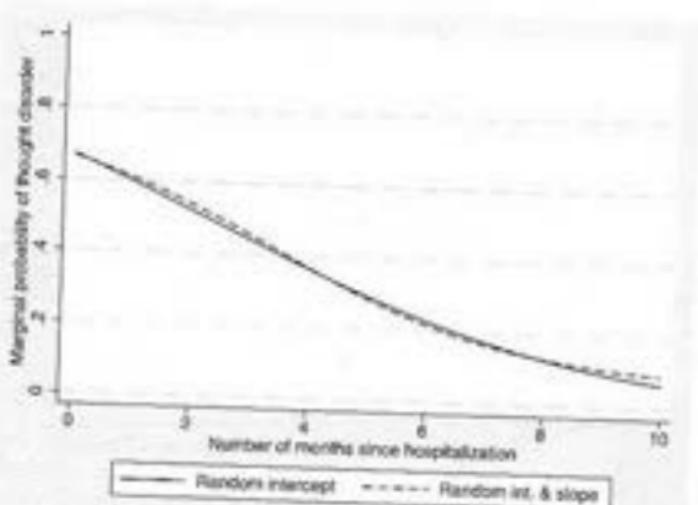


Figure 9.10: Marginal predicted probabilities for random intercept and random coefficient models.

Note that `gllamm` can be used for a wide range of models with random effects and other latent variables such as factors, including (multilevel) structural equation models and latent class models with many different response types as well as mixed responses (see Skrondal and Rabe-Hesketh, 2003; 2004; Rabe-Hesketh *et al.*, 2003; 2004a; 2004b). The webpage <http://www.gllamm.org> gives more references and up-to-date information.

9.5 Exercises

9.1 • Thought disorder and schizophrenia

- For the thought disorder data, produce graphs of the predicted probabilities for the individual patients separately for early and late onset (similar to Figures 9.2 and 9.6 for the post-natal depression data). Hint: use `gllapred` with the `mu` option (not `marg` or `us()`) to obtain posterior mean probabilities.

9.2 • Australian school children

- Analyze the Australian school children data described in Chapter 7 using a Poisson model with a random intercept for each child and compare the estimates with those of the negative binomial model estimated in Section 7.3.4, where the exponentiated random intercept (frailty) has a gamma distribution instead of a log-normal distribution.

9.3 Jaw growth

- For the jaw growth data described in Exercise 8.3, estimate the following random intercept model

$$y_{ij} = \beta_0 + \beta_1 x_i + \beta_2 t_i + u_i + \epsilon_{ij}$$

where x_i is a dummy variable for being male and t_i is age – 8.

- Interpret all the parameter estimates.
- Extend the model to allow boys and girls to differ in their mean rate of growth and interpret the regression coefficients.
- Extend the model further by including a random coefficient of t_i and use a likelihood ratio test to choose between this model and the previous model.
- For the chosen model, plot the predicted growth trajectories of the children (based on parameter estimates and empirical Bayes predictions of the random effects) by gender.

9.4 Wage increases

Here we consider the data used in Exercises 1.2 and 8.2, and will make use of the following additional variables:

- `nr`: person identifier
- `educ`: years of schooling
- `exper`: labor market experience ($\text{Age} - 6 - \text{educ}$)

- `expersq`: labor market experience squared
 - `married`: dummy variable for being married
 - `union`: dummy variable for being a member of a union (i.e., wage being set in collective bargaining agreement)
1. Estimate a model for `lwage` with `black`, `hispanic`, `educ`, `exper`, `expersq`, `married`, and `union` as explanatory variables and with a random intercept for subjects.
 2. Interpret the estimates.
 3. An alternative approach to panel data that is popular in econometrics is to specify fixed intercepts for subjects instead of random ones; this can be accomplished by using `xtreg` with the `fe` option (which is equivalent to, but much more efficient than using dummy variables for subjects). Fit the fixed effects version of the model above.
 4. Explain why some variables are dropped by Stata and compare the regression coefficients of the remaining variables with those estimated for the random intercept model.

9.5 Epileptic seizures and chemotherapy

1. Analyze the epileptic seizure data introduced in the next chapter using a Poisson model with the same fixed part as specified in Section 10.3.2 and with a random intercept for subjects. Use `gllamm` with the `adapt` option. Make sure you are using enough quadrature points by comparing estimates with different numbers of quadrature points (`nlp(0)` option).
2. Add a random slope for `post` and use a likelihood ratio test to decide whether or not to retain this model.
3. For the chosen model, plot the model-implied marginal relationship between the expected epilepsy rate and visit for 25-year olds in the two treatment groups. (Hint: use `gllapred` with the options `mu`, `marg`, and `nooffset` and plot the predictions for subjects 3 and 46.)

Chapter 10

Generalized Estimating Equations: Epileptic Seizures and Chemotherapy

10.1 Description of data

In a clinical trial reported by Thall and Vail (1990), 59 patients with epilepsy were randomized to groups receiving either the anti-epileptic drug progabide or a placebo in addition to standard chemotherapy. The number of seizures was counted over four two-week periods. In addition, a baseline seizure rate was recorded for each patient, based on the eight-week prerandomization seizure count. The age of each patient was also recorded. The main question of interest is whether the treatment progabide reduces the frequency of epileptic seizures compared with placebo. The data are shown in Table 10.1. (These data also appear in Hand *et al.*, 1994.)

Table 10.1 Data in *epil.dta*

subj	id	y1	y2	y3	y4	treat	base	age
1	104	5	3	3	3	0	11	31
2	106	3	5	3	3	0	11	30
3	107	2	4	0	5	0	6	25
4	114	4	4	1	4	0	8	36
5	116	7	18	9	21	0	66	22
6	118	5	2	8	7	0	27	29

Table 10.1 Data in epil.dta (continued)

7	123	6	4	0	2	0	12	31
8	126	40	20	23	12	0	52	42
9	130	5	6	6	5	0	23	37
10	135	14	13	6	0	0	10	28
11	141	26	12	5	22	0	52	36
12	145	12	6	8	4	0	33	24
13	201	4	4	6	2	0	18	23
14	202	7	9	12	14	0	42	36
15	205	16	24	10	9	0	87	26
16	206	11	0	0	5	0	50	26
17	210	0	0	3	3	0	18	28
18	213	37	29	28	29	0	111	31
19	215	3	5	2	5	0	18	32
20	217	3	0	6	7	0	20	21
21	219	3	4	3	4	0	12	29
22	220	3	4	3	4	0	9	21
23	222	2	3	3	5	0	17	32
24	226	8	12	2	8	0	28	25
25	227	18	24	76	25	0	55	30
26	230	2	1	2	1	0	9	40
27	234	3	1	4	2	0	10	19
28	238	13	15	13	12	0	47	22
29	101	11	14	9	8	1	76	18
30	102	8	7	9	4	1	38	32
31	103	0	4	3	0	1	19	20
32	108	3	6	1	3	1	10	30
33	110	2	6	7	4	1	19	18
34	111	4	3	1	3	1	24	24
35	112	22	17	19	16	1	31	30
36	113	5	4	7	4	1	14	35
37	117	2	4	0	4	1	11	27
38	121	3	7	7	7	1	67	20
39	122	4	18	2	5	1	41	22
40	124	2	1	1	0	1	7	28
41	128	0	2	4	0	1	22	25
42	129	5	4	0	3	1	13	40
43	137	11	14	25	15	1	46	33
44	139	10	5	3	8	1	36	21
45	143	19	7	6	7	1	38	35
46	147	1	1	2	3	1	7	25
47	203	6	10	8	8	1	36	26
48	204	2	1	0	0	1	11	25
49	207	102	65	72	63	1	151	22
50	208	4	3	2	4	1	22	32
51	209	8	6	5	7	1	41	25
52	211	1	3	1	5	1	32	35
53	214	18	11	28	13	1	56	21
54	218	6	3	4	0	1	24	41

Table 10.1 Data in epil.dta (continued)

55	221	3	5	4	3	1	16	32
56	225	1	23	19	8	1	22	26
57	228	2	3	0	1	1	25	21
58	232	0	0	0	0	1	13	36
59	236	1	4	3	2	1	12	37

10.2 Generalized estimating equations

In this chapter we consider an approach to the analysis of longitudinal data that is very different from random effects modeling described in the previous chapter. Instead of attempting to model the dependence between responses on the same individuals as arising from between-subject heterogeneity represented by random intercepts and possibly random slopes, we will concentrate on estimating the marginal mean structure, treating the dependence as a nuisance.

10.2.1 Normally distributed responses

If we suppose that a normally distributed response is observed on each individual at T time points, then the basic regression model for longitudinal data becomes (cf. equation (3.3))

$$\mathbf{y}_i = \mathbf{X}_i\boldsymbol{\beta} + \boldsymbol{\epsilon}_i \quad (10.1)$$

where $\mathbf{y}'_i = (y_{i1}, y_{i2}, \dots, y_{iT})$, $\boldsymbol{\epsilon}'_i = (\epsilon_{i1}, \epsilon_{i2}, \dots, \epsilon_{iT})$, \mathbf{X}_i is a $T \times (p+1)$ design matrix, and $\boldsymbol{\beta}' = (\beta_0, \dots, \beta_p)$ is a vector of regression parameters. The residual terms are assumed to have a multivariate normal distribution with a covariance matrix of some particular form that is a function of (hopefully) a small number of parameters. Maximum likelihood estimation can be used to estimate both the parameters in (10.1) and the parameters structuring the covariance matrix (details are given in Jennrich and Schluchter, 1986). The latter are often not of primary interest (they are often referred to as nuisance parameters), but using a covariance matrix that fails to match that of the repeated measurements can lead to inefficient estimates and invalid standard errors for the parameters that are of concern, namely the $\boldsymbol{\beta}$ in (10.1).

If each non-replicated element of the covariance matrix is treated as a separate parameter, giving an unstructured covariance matrix, and if there are no missing data, then this approach is essentially equivalent to multivariate analysis of variance for longitudinal data (see Everitt, 2001). However, it is often more efficient to impose some meaningful structure on the covariance matrix. The simplest (and most unrealis-

tic) structure is *independence* with all off-diagonal elements (the covariances) equal to zero, and typically all diagonal elements (the variances) equal to each other. Another commonly used simple structure, known as *compound symmetry* (for example, see Wiener, 1971), requires that all covariances are equal and all variances are equal. This is just the correlation structure of a linear random intercept model described in the previous chapter except that the random intercept model also requires that the correlation be positive.

Other correlation structures include autoregressive structures where the correlations decrease with the distance between time points. Whatever the assumed correlation structure, all models may be estimated by maximum likelihood.

10.2.2 Non-normal responses

Unfortunately, it is generally not straightforward to specify a multivariate model for non-normal responses. One solution, discussed in the previous chapter, is to induce residual dependence among the responses using random effects. An alternative approach is to give up the idea of a model altogether by using *generalized estimating equations* (GEE) as introduced by Liang and Zeger (1986). Generalized estimating equations are essentially a multivariate extension of the quasi-likelihood approach discussed in Chapter 7 (see also Wedderburn, 1974). In GEE the parameters are estimated using "estimating equations" resembling the score equations for maximum likelihood estimation of the linear model described in the previous section. These estimating equations only require specification of a link and variance function and a correlation structure for the observed responses conditional on the covariates. As in the quasi-likelihood approach, the parameters can be estimated even if the specification does not correspond to any statistical model.

The regression coefficients represent marginal effects, i.e., they determine the population averaged relationships. Liang and Zeger (1986) show that the estimates of these coefficients are valid even when the correlation structure is incorrectly specified. Correct inferences can be obtained using robust estimates of the standard errors based on the sandwich estimator for clustered data (e.g., Binder, 1983; Williams, 2000). The parameters of the correlation matrix, referred to as the *working correlation matrix*, are treated as *nuisance parameters*. However, Lindsey and Lambert (1998) and Crouchley and Davies (1999) point out that estimates are no longer consistent if "endogenous" covariates such as baseline responses are included in the model. Fortunately, inclusion of the baseline response as a covariate does yield consistent estimates of treatment effects in clinical trial data such as

the epilepsy data considered here (see Crouchley and Davies, 1999) as long as the model does not contain a baseline by treatment interaction.

There are some important differences between GEE and random effects modeling. First, while random effects modeling is based on a statistical model and typically maximum likelihood estimation, GEE is an estimation method that is not based on a statistical model. Second, there is an important difference in the interpretation of the regression coefficients. In random effects models, the regression coefficients represent conditional or subject-specific effects for given values of the random effects. For GEE, on the other hand, the regression coefficients represent marginal or population averaged effects. As we saw in the thought disorder data in the previous chapter, conditional and marginal relationships can be very different. Either may be of interest; for instance patients are likely to want to know the subject-specific effect of treatments, whereas health economists may be interested in population averaged effects. Whereas random effects models allow the marginal relationship to be derived, GEE does not allow derivation of the conditional relationship. Note that conditional and marginal relationships are the same if an identity link is used and, in the case of random intercept models (no random coefficients), if a log link is specified (see Diggle et al., 2002). Third, GEE is often preferred because, in contrast to the random effects approach, the parameter estimates are consistent even if the correlation structure is misspecified (although this is true only if the mean structure is correctly specified). Fourth, while maximum likelihood estimation of a correctly specified model is consistent if data are missing at random (MAR), this is not the case for GEE which requires that responses are missing completely at random (MCAR), or that missingness depends only on the covariates included in the model. See Hardin and Hilbe (2002) for a thorough introduction to GEE.

10.3 Analysis using Stata

The generalized estimating equations approach, as described in Liang and Zeger (1986), is implemented in Stata's `xtgee` command. The main components which have to be specified are:

- the assumed distribution of the response variable (given the covariates), specified in the `family()` option – this determines the variance function;
- the link between the response variable and its linear predictor, specified in the `link()` option, and

- the structure of the working correlation matrix, specified in the `correlation()` option.

In general, it is not necessary to specify the `link()` option since, as for the `glm` command, the default link is the canonical link for the specified family.

Since the `xtgee` command will often be used with the `family(gauss)` option, together with the identity link function, we will illustrate this option on the post-natal depression data used in the previous two chapters before moving on to deal with the epilepsy data in Table 10.1.

10.3.1 Post-natal depression data

The data are obtained using

```
infile subj group dep0 dep1 dep2 dep3 dep4 dep5 dep6 ///
        using depress.dat, clear
reshape long dep, i(subj) j(visit)
mvedecode _all, mv(-9)
```

To begin, we fit a model that regresses `dep` on `group`, `visit`, their interaction and `visit` squared as in the previous chapter but under the unrealistic assumption of independence. The necessary command written out in its fullest form is

```
generate gr_vis = group*visit
generate vis2 = visit^2
xtgee dep group visit gr_vis vis2, i(subj) t(visit) ///
corr(indep) link(iden) family(gauss)
```

(see Display 10.1). Here, the fitted model is simply a multiple regression model for 365 observations which are assumed to be independent of one another; the estimated scale parameter is just the residual mean square, and the deviance is equal to the residual sum of squares. The estimated regression coefficients and their associated standard errors indicate that the `group` by `visit` interaction is significant at the 5% level. However, treating the observations as independent is unrealistic and will almost certainly lead to poor estimates of the standard errors. Standard errors for between-subject factors (here `group`) are likely to be underestimated because we are treating observations from the same subject as independent, thus increasing the apparent sample size; standard errors for within-subject factors (here `visit`, `gr_vis`, and `vis2`) are likely to be overestimated since we are not controlling for residual between-subject variability.

We therefore now abandon the assumption of independence and estimate a correlation matrix having compound symmetry (i.e., constraining the correlations between the observations at any pair of time

GE population-averaged model			Number of obs	=	356
Group variable:	subj		Number of groups	=	61
LINK:	identity		Obs per group: min =		2
Family:	Gaussian		avg =		5.8
Correlation:	independent		max =		7
			Wald chi2(4)	=	269.61
Scale parameter:	26.89935		Prob > chi2	=	0.0000
Pearson chi2(306):	9676.17		Deviance	=	9676.17
Dispersion (Pearson):	26.89935		Dispersion	=	26.89935
dep	Coef.	Std. Err.	N	P>tstat	[95% Conf. Interval]
group	-1.506834	.9033647	-1.61	0.108	-3.345995 .3323974
visit	-3.842465	.5091836	-7.56	0.000	-4.887447 -2.853483
gr_vis	-.6090744	.277417	-2.20	0.028	-1.152802 -.0653471
vis2	.3904383	.079783	4.89	0.000	.2340665 .5466102
_cons	20.96633	.7826299	26.79	0.000	19.4314 23.49928

Display 10.1

points to be equal). Such a correlation structure is specified using `corr(exchangeable)`, or the abbreviated form `corr(exc)`. The model can be fitted as follows:

```
xtgee dep group visit gr_vis vis2, i(subj) t(visit) ///
corr(exc) link(iden) fam(gauss)
```

Instead of specifying the subject and time identifiers using the options `i()` and `t()`, we can also declare the data as being of the form `xt` (for cross-sectional time series) as follows:

```
iis subj
tis visit
```

and omit the `i()` and `t()` options from now on. Since both the link and the family correspond to the default options, the same analysis may be carried out using the shorter command

```
xtgee dep group visit gr_vis vis2, corr(exc)
```

(see Display 10.2). After estimation, `estat wcorrelation` reports the estimated working "within" correlation matrix

```
estat wcorrelation, format(%6.4g)
```

which is shown in Display 10.3. Here the `format()` option was used to reduce the number of decimal places and therefore avoid rows of the matrix wrapping over two lines.

Note that the standard error for `group` has increased whereas those for `visit`, `gr_vis`, and `vis2` have decreased as expected. The estimated within-subject correlation matrix is compound symmetric. This

GEE population-averaged model		Number of obs	=	356
Group variable:	subj	Number of groups	=	61
Link:	identity	Obs per group: min	=	2
Family:	Gaussian	avg	=	5.8
Correlation:	exchangeable	max	=	7
Scale parameter:	36.92728	Vald chisq(4)	=	421.11
		Prob > chisq	=	0.0000
dep	Coef.	Std. Err.	z	P> z
group	-1.470155	1.162063	-1.27	0.206
visit	-3.785601	.3648346	-10.38	0.000
gr.visit	-.5837938	.2040388	-2.86	0.004
visit2	.3850221	.0659384	6.88	0.000
_cons	20.90907	.901062	23.20	0.000
				[95% Conf. Interval]

Display 10.2

Estimated within-subj correlation matrix R:

	c1	c2	c3	c4	c5	c6	c7
r1	1						
r2	.515	1					
r3	.515	.515	1				
r4	.515	.515	.515	1			
r5	.515	.515	.515	.515	1		
r6	.515	.515	.515	.515	.515	1	
r7	.515	.515	.515	.515	.515	.515	1

Display 10.3

structure is frequently not acceptable since correlations between pairs of observations widely separated in time will often be lower than for observations closer together. This pattern was apparent from the scatterplot matrix given in Chapter 8.

To allow for such a pattern of correlations among the repeated observations, we can move to an autoregressive structure. For example, in a first-order autoregressive specification the correlation between time points r and s is assumed to be $\rho^{|r-s|}$. The necessary instruction for fitting the previously considered model but with this first-order autoregressive structure for the correlations is

```
xtgee dep group visit gr_vis vis2, corr(ar1)
```

GEF population-averaged model			Number of obs	=	356
Group and time vars:	subj visit		Number of groups	=	61
Link:	identity		Obs per group:	min =	2
Family:	Gaussian		avg =	8.6	
Correlation:	AR(1)		max =	7	
Scale parameter:	27.50248		Wald chi2(4)	=	213.85
			Prob > chi2	=	0.0000
dep	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
group	-.539061	1.277002	-0.42	0.673	-3.041938 1.963616
visit	-6.061961	.4741261	-12.57	0.000	-4.991227 -3.132695
gr_vis	-.7835801	.3332716	-2.35	0.019	-1.43478 -.1283796
vis2	.4207375	.0693395	6.07	0.000	.2848346 .5566404
_cons	21.10085	.9732406	21.68	0.000	19.19334 23.06837

Display 10.4

The estimates of the regression coefficients and their standard errors in Display 10.4 have changed but not substantially. The estimated within-subject correlation matrix may again be obtained using

```
estat wcorrelation, format(%6.4g)
```

(see Display 10.5) which has the expected pattern in which correlations decrease substantially as the separation between the observations increases.

Other correlation structures are available for `xtgee`, including the option `correlation(unstructured)` in which no constraints are placed on the correlations. (This is essentially equivalent to multivariate analysis of variance for longitudinal data, except that the variance is assumed to be constant over time.) It might appear that using this option

Estimated within-subj correlation matrix R:

	c1	c2	c3	c4	c5	c6	c7
r1	1						
r2	.6475	1					
r3	-.4192	-.6475	1				
r4	.2714	.4192	.6475	1			
r5	.1757	.2714	.4192	.6475	1		
r6	.1138	.1757	.2714	.4192	.6475	1	
r7	.0737	.1138	.1757	.2714	.4192	.6475	1

Display 10.5

would be the most sensible one to choose for all data sets. This is not, however, the case since it necessitates the estimation of many nuisance parameters. This can, in some circumstances, cause problems in the estimation of those parameters of most interest, particularly when the sample size is small and the number of time points is large.

10.3.2 Epilepsy data

We now analyze the epilepsy data using a similar model as for the depression data, but using the Poisson distribution and log link. The data are available in a Stata file `epil.dta` and can be read using

```
use epil, clear
```

We will treat the baseline measure as one of the responses:

```
generate y0 = baseline
```

Some useful summary statistics can be obtained using

```
summarize y0 y1 y2 y3 y4 if treat==0
```

```
summarize y0 y1 y2 y3 y4 if treat==1
```

(see Displays 10.6 and 10.7).

We see that the number of observations is constant over time so there appears to be no dropout. The means and standard deviations of `y0` are larger than for the other responses because seizures were counted over an 8-week period at baseline and over 2-week periods at the subsequent visits. The largest value of `y1` in the pro gabide group seems out of step with the other maximum values and may indicate an outlier. Some graphics of the data may be useful for investigating this possibility further, but first it is convenient to reshape the data from its present "wide" form to the "long" form. We now reshape the data as follows:

Variable	Obs	Mean	Std. Dev.	Min	Max
y0	28	30.78571	26.10429	6	131
y1	28	9.387143	10.13689	0	40
y2	28	8.385714	8.164318	0	29
y3	28	8.785714	14.67262	0	76
y4	28	7.964286	7.627835	0	29

Display 10.6

Variable	Obs	Mean	Std. Dev.	Min	Max
y0	31	31.6129	37.98175	7	151
y1	31	8.580645	18.24667	0	102
y2	31	8.419355	11.88966	0	86
y3	31	8.129032	13.89422	0	72
y4	31	6.709677	11.26404	0	63

Display 10.7

```
reshape long y, i(subj) j(visit)
sort subj treat visit
list in 1/12, clean
```

(see Display 10.8).

	subj	visit	id	y	treat	baseline	age
1.	1	0	104	11	0	11	31
2.	1	1	104	5	0	11	31
3.	1	2	104	3	0	11	31
4.	1	3	104	3	0	11	31
5.	1	4	104	3	0	11	31
6.	2	0	106	11	0	11	30
7.	2	1	106	3	0	11	30
8.	2	2	106	5	0	11	30
9.	2	3	106	3	0	11	30
10.	2	4	106	3	0	11	30
11.	3	0	107	6	0	6	25
12.	3	1	107	2	0	6	25

Display 10.8

Perhaps the most useful graphical display for investigating the data is a set of graphs of individual response profiles. Since we are planning

to fit a Poisson model with the log link to the data, we take the log transformation before plotting the response profiles. (We need to add a positive number, say 1, because some seizure counts are zero.)

```
generate ly = log(y+1)
```

However, the baseline measure represents seizure counts over an 8-week period, compared with 2-week periods for each of the other time points. We therefore divide the baseline count by 4:

```
replace ly = log(y/4+1) if visit==0
```

and then plot the log-counts:

```
twoway connect ly visit if treat==0, by(subj, ///
    style(compact)) ytitle("Log count")
twoway connect ly visit if treat==1, by(subj, ///
    style(compact)) ytitle("Log count")
```

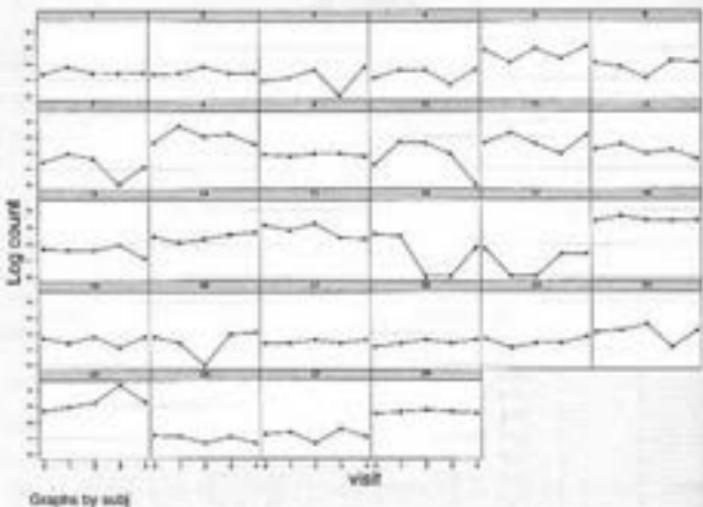


Figure 10.1: Response profiles in placebo group.

The resulting graphs are shown in Figures 10.1 and 10.2. There is no obvious improvement in the progabide group. Subject 49 had more epileptic fits overall than any other subject and might perhaps be considered an outlier (see Exercise 10.2).

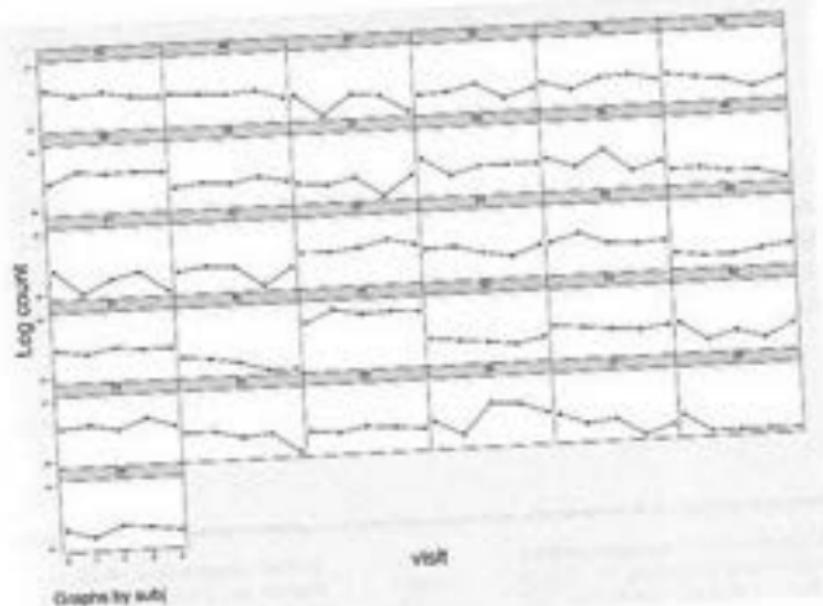


Figure 10.2: Response profiles in the treated group.

As discussed in Chapter 7, the most plausible distribution for count data is often the Poisson distribution. The Poisson distribution is specified in `xtgee` models using the option `family(poisson)`. The log link is implied (since it is the canonical link). The baseline counts were obtained over an 8-week period whereas all subsequent counts are over 2-week periods. To model the seizure rate in counts per week, we must therefore use the log observation period $\log(p_i)$ as an offset (a covariate with regression coefficient set to 1). The model for the mean count μ_{ij} then becomes

$$\log(\mu_{ij}) = \mathbf{x}'_{ij}\boldsymbol{\beta} + \log(p_i),$$

so that the rate is modeled as

$$\mu_{ij}/p_i = \exp(\mathbf{x}'_{ij}\boldsymbol{\beta}).$$

We can compute the required offset using

```
generate lnobs = cond(visit==0,ln(8),ln(2))
```

Following Diggle et al. (2002), we will allow the log rate to change by a treatment group-specific constant after the baseline assessment. The necessary covariates, an indicator for the post-baseline visits and an

interaction between that indicator and treatment group, are created using

```
generate post = visit>0
generate tr_post = treat*post
```

We will also control for the age of the patients. The summary tables for the selcure data given on page 210 provide strong empirical evidence that there is overdispersion (the variances are greater than the means), and this can be incorporated using the `scale(x2)` option to allow the dispersion parameter ϕ to be estimated (see also Chapter 7).

```
iim subj
xtgee y age treat post tr_post, corr(exc) family(pois) ///
offset(inobs) scale(x2)
```

GEE population-averaged model				Number of obs		=	295
Group variable:		subj		Number of groups		=	59
Link:		log		Obs per group: min		=	5
Family:		Poisson		avg		=	5.0
Correlation:		exchangeable		max		=	5
Scale parameter:		18.48008		Wald chi2(4)		=	6.43
				Prob > chi2		=	0.2458
y	Coef.	Std. Err.	*	P> z	[95% Conf. Interval]		
age	-0.0322513	.0148644	-2.17	0.030	-.061385	-	-.0031176
treat	-.0577873	.201945	-0.09	0.930	-.4136922		.3780176
post	.1107981	.1800635	0.74	0.460	-.183321		.4049173
tr_post	-.1036687	.213317	-0.49	0.627	-.5217742		.3144129
_cons	2.265255	.4400816	5.15	0.000	1.402711		8.1278
lnobs (offset)							

(Standard errors scaled using square root of Pearson X2-based dispersion)

Display 10.9

The output assuming an exchangeable correlation structure is given in Display 10.9, and the estimated correlation matrix is obtained using `xtcorr`.

```
estat wcorrelation
```

(see Display 10.10).

In Display 10.9, the parameter ϕ is estimated as 18.5, indicating severe overdispersion in these data. We briefly illustrate how important it was to allow for overdispersion by omitting the `scale(x2)` option:

Estimated within-subj correlation matrix R:

	c1	c2	c3	c4	c5
r1	1				
r2	.7685773	1			
r3	.7685773	.7685773	1		
r4	.7685773	.7685773	.7685773	1	
r5	.7685773	.7685773	.7685773	.7685773	1

Display 10.10

```
xtgee y age treat post tr_post, corr(exc) family(pois) ///
offset(lnobs)
```

GEE population-averaged model				Number of obs	=	295
Group variable:	subj	Number of groups	=	59		
Link:	log	Obs per group: min	=	5		
Family:	Poisson	avg	=	5.0		
Correlation:	exchangeable	max	=	5		
Scale parameter:		Wald chi2(4)	=	100.38		
		Prob > chi2	=	0.0000		
y	Coeff.	Std. Err.	z	P> z	[95% Conf. Interval]	
age	-.0322613	.0034578	-9.33	0.000	-.0390284	-.0254742
treat	-.0177873	.0469765	-0.38	0.705	-.1008596	.074285
post	.1107981	.0349279	3.17	0.002	.04238	.1792163
tr_post	-.1036807	.0496219	-2.09	0.037	-.2009378	-.0068235
_cons	2.265355	.102372	22.13	0.000	2.05461	2.466901
lnobs	(offset)					

Display 10.11

The results given in Display 10.11 show that the standard errors are now much smaller than before. Even if overdispersion had not been suspected, this error could have been detected by using the vce(robust) option (see Chapter 7):

```
xtgee y age treat post tr_post, corr(exc) family(pois) ///
offset(lnobs) vce(robust)
```

The results of the robust regression in Display 10.12 are remarkably similar to those of the overdispersed Poisson model, suggesting that the latter is a reasonable "model" for the data.

GEE population-averaged model		Number of obs	=	295
Group variable:	subj	Number of groups	=	59
Link:	log	Obs per group: min =		5
Family:	Poisson	avg =		6.0
Correlation:	exchangeable	max =		5
Scale parameter:	1	Wald chi2(4)	=	6.85
		Prob > chi2	=	0.1442
(Std. Err. adjusted for clustering on subj)				

y	Semi-robust				[95% Conf. Interval]
	Coeff.	Std. Err.	t	P> t	
age	-.0322513	.0148746	-2.17	0.030	-.0614049 - .0030976
treat	-.0177873	.216051	-0.08	0.934	-.4412395 - .4066649
post	.1107981	.1170963	0.95	0.344	-.1187064 - .3403027
tr_post	-.1036807	.2154436	-0.48	0.630	.8259426 - .3188811
case	2.265255	.4371124	5.18	0.000	1.408531 3.12199
lnobs	(offset)				

Display 10.12

The estimated coefficient of *tr_post* represents the estimated difference in the change in log seizure rate from baseline to post randomization between the placebo and pro gabide groups. In the placebo group there is an increase in the log seizure rate of 0.1108, and in the pro gabide group there is an increase of only 0.007 ($= 0.1108 - .1037$). However, the difference is not significant ($p=0.63$). The exponential of the interaction coefficient gives an estimated incidence rate ratio, here the ratio of the relative increase in seizure rate for the treated patients compared with the control patients. The exponentiated coefficient and the corresponding confidence interval can be obtained directly using the *eform* option in *xtgee*:

```
xtgee y age treat post tr_post, corr(exc) ///
family(pois) offset(lnobs) scale(x2) eform
```

The results in Display 10.13 indicate that the relative increase in seizure rate is 10% lower in the treated group compared with the control group, with a 95% confidence interval from 41% lower to 37% greater.

However, before interpreting these estimates, we should perform some diagnostics. Standardized Pearson residuals can be useful for identifying potential outliers (see equation (7.9)). These can be found by first using the *predict* command to obtain predicted counts, subtracting the observed counts, and dividing by the estimated standard deviation $\sqrt{\hat{\phi}\mu_{ij}}$, where $\hat{\phi}$ is the estimated dispersion parameter:

```
quietly xtgee y treat baseline age visit, corr(exc) ///
```

GEE population-averaged model			Number of obs	=	295
Group variable:		subj	Number of groups	=	59
Link:		log	Dif per group: min	=	5
Family:		Poisson	avg	=	8.0
Correlation:		exchangeable	max	=	5
Scale parameter:		18.48008	Wald chi2(4)	=	5.43
			Prob > chi2	=	0.2468
β	IRR	Std. Err.	z	P> z	[95% Conf. Interval]
age	.9602632	.6143927	-2.17	0.030	.9404611 .9968373
treat	.98237	.1983847	-0.09	0.930	.6612706 1.469389
post	1.117169	.1676464	0.74	0.460	.8325009 1.499179
TF_post	.9016131	.1923008	-0.49	0.627	.5934667 1.369455
lambda (offset)					

(Standard errors scaled using square root of Pearson X2-based dispersion)

Display 10.13

```
family(pois) scale(x2)
predict pred, mu
generate pres = (y-pred)/sqrt(e(chi2_dis)*pred)
```

Boxplots of those residuals at each visit are obtained using

```
sort visit
graph box stpres, medtype(line) over(visit,      ///
    relabel(1 "visit 1" 2 "visit 2" 3 "visit 3" ///
    4 "visit 4"))
```

The resulting graph is shown in Figure 10.3. Pearson residuals greater than 4 are certainly a cause for concern, so we can check which subjects they belong to using

```
list subj id if stpres>4
```

subj	id
41.	49 207
90.	49 207
176.	49 207
178.	49 207
185.	28 227
292.	49 207

Subject 49 appears to be an outlier due to extremely large counts as we saw in Figure 10.2. Subject 25 also has an unusually large count at visit 3. It would be a good idea to repeat the analysis without subject 49 to see how much the results are affected by this unusual subject (see Exercise 10.2). This can be viewed as a sensitivity analysis.

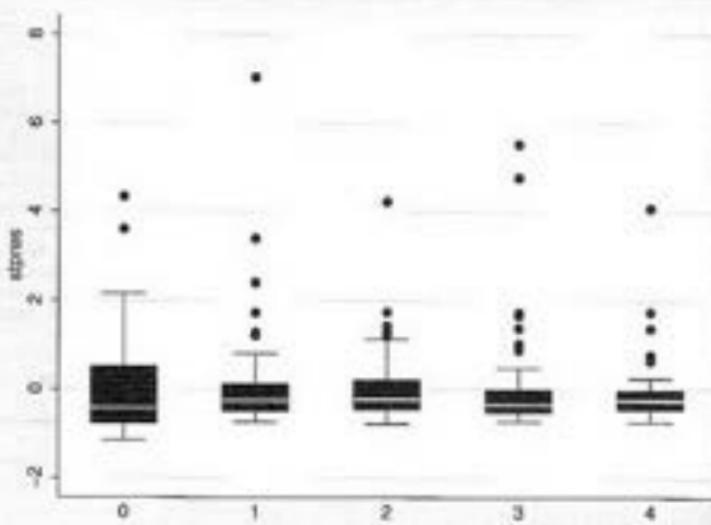


Figure 10.3: Standardized Pearson residuals.

10.4 Exercises

10.1 • Treatment of post-natal depression

1. For the depression data, compare the results of GEE with a compound symmetric structure with ordinary linear regression where standard errors are corrected for the within-subject correlation using:
 - a. the options, `vce(robust)` `cluster(subj)`, to obtain the sandwich estimator for clustered data (see help for `regress` and
 - b. bootstrapping, by sampling subjects with replacement. This may be achieved using the `bootstrap` prefix, together with the option `cluster(subj)`.

10.2 Epileptic seizures and chemotherapy

1. Explore other possible correlation structures for the seizure data in the context of a Poisson model. Examine the robust standard errors in each case.
2. Repeat the above analyses, but excluding subject 49 (who appears to be an outlier). Compare the results.

10.3 Thought disorder and schizophrenia

- For the thought disorder data discussed in the previous chapter, estimate the effect of `early`, `month` and their interaction on the logit of thought disorder using GEE with an exchangeable correlation structure. Use robust standard errors.
- Interpret the estimates.
- Plot the predicted probability over time for early onset women (using `graph twoway function`, see Section 6.3.2), and compare the curve with the curves in Figure 9.10.

10.4 Driver education

In a randomized experiment to investigate if driver education reduces the number of collisions and traffic violations of teenagers (Stock *et al.*, 1983), eligible high school students were randomized to three groups: safe performance curriculum (SPC), pre-driver license curriculum (PDL), and control. Whereas the SPC was a 70-hour state of-the-art program, the PDL was a 30-hour course containing only the minimum training required to pass the driving test. The control group received no training through the school system and was taught by the parents and/or private training schools only. During three years of follow-up, the occurrence of collisions and moving violations were obtained using records from the state Department of Motor Vehicles. (The data are from Davis, 2002.)

The variables in `drivers.dta` are:

- `program`: group (string variable with values SPC, PDL, and Control)
 - `gender`: gender (string variable with values Male and Female)
 - `col1` to `col13`: indicator for at least one collision or moving violation during years 1 to 3
 - `num`: number of times the response-covariate pattern occurred
- Prepare the data for analysis using GEE. (Hint: make sure to expand the data first using `expand num`, then reshape to long.)
 - Investigate the effect of time, program, gender, and the program by gender interaction on the odds of at least one collision or moving violation using generalized estimating equations with a logit link and unstructured correlations. Use robust standard errors throughout this exercise.
 - Perform a Wald test for the interaction terms and remove them if the test is not significant at the 5% level.

4. By inspecting the estimated correlation matrix, choose the correlation structure that appears to be most appropriate and estimate the model with that correlation structure.
5. Interpret the odds ratio estimates for the final model.

Chapter 11

Some Epidemiology

11.1 Description of data

This chapter illustrates analysis of different epidemiological designs, namely cohort studies and matched as well as unmatched case-control studies. Four datasets will be used which are presented in the form of cross-tabulations in Tables 11.1 to 11.4. (Tables 11.1 and 11.4 are taken from Clayton and Hills (1993) with permission of their publisher, Oxford University Press.)

The data in Table 11.1 result from a *cohort study* which investigated the relationship between diet and ischemic heart disease (IHD). Here we consider low energy intake as a risk factor since it is highly associated with lack of physical exercise. The table gives frequencies of IHD by ten-year age-band and exposure to a high or low calorie diet. The total person-years of observation are also given for each cell.

The dataset in Table 11.2 is the result of a *case-control study* investigating whether keeping a pet bird is a risk factor for lung cancer. This dataset is given in Hand *et al.* (1994).

The datasets in Tables 11.3 and 11.4 are from matched case-control studies, the first with a single matched control and the second with three matched controls. Table 11.3 arises from a matched case-control study of endometrial cancer where cases were matched on age, race, date of admission, and hospital of admission to a suitable control not suffering from cancer. Past exposure to conjugated estrogens was determined. The dataset is described in Everitt (1994). Finally, the data in Table 11.4, described in Clayton and Hills (1993), arise from a case-control study of breast cancer screening. Women who had died of breast cancer were matched with three control women. The screening

history of each control was assessed over the period up to the time of diagnosis of the matched case.

11.2 Introduction to epidemiology

Epidemiology can be described as the study of diseases in populations, in particular the search for causes of disease. For ethical reasons, subjects cannot be randomized to possible risk factors in order to establish whether these are associated with an increase in the incidence of disease, and therefore epidemiology is based on observational studies. The most important types of studies in epidemiology are cohort studies and case-control studies. We will give a very brief description of the design and analysis of these two types of studies, following closely the explanations and notation given in the excellent book, *Statistical Models in Epidemiology*, by Clayton and Hills (1993).

11.2.1 Cohort studies

In a cohort study, a group of subjects free of the disease is followed up, and the presence of risk factors as well as the occurrence of the disease of interest are recorded. This design is illustrated in Figure 11.1. An

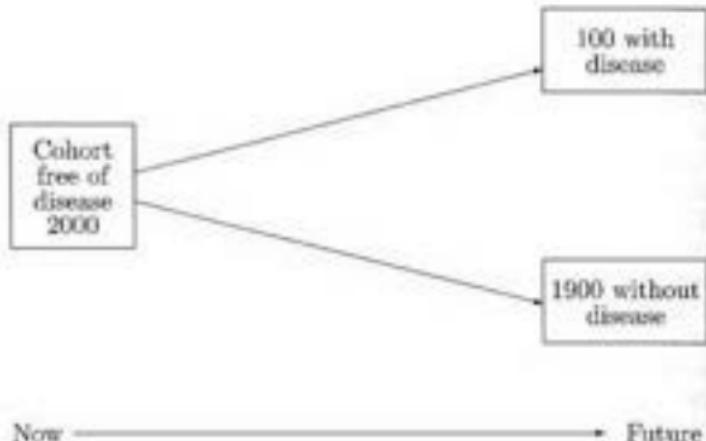


Figure 11.1: Cohort study.

Table 11.1 Number of IHD cases and person-years of observation by age and exposure to low energy diet

Age	Exposed		Unexposed	
	< 2750 kcal	≥ 2750 kcal	Cases	Pers-ys
40-49	2	311.9	4	607.9
50-59	12	878.1	5	1272.1
60-69	14	657.5	8	888.9

Table 11.2 Number of lung cancer cases and controls who keep a pet bird

Kept pet birds	Cases	Controls
Yes	98	101
No	141	328
Total	239	429

Table 11.3 Frequency of exposure to oral conjugated estrogens among cases of endometrial cancer and their matched controls

Cases	+	Controls		Total
		+	-	
+	12	43	55	
-	7	121	128	
Total	19	164	183	

Table 11.4 Screening history in subjects who died of breast cancer and 3 matched controls

Status of the case	Number of controls screened			
	0	1	2	3
Screened	1	4	3	1
Unscreened	11	10	12	4

example of a cohort study is the study described in the previous section where subjects were followed up to monitor the occurrence of ischemic heart disease in two risk groups, those with high and low energy intake, giving the results in Table 11.1.

The incidence rate of the disease λ may be estimated by the number of new cases of the disease D during a time interval divided by the person-time of observation Y , the sum of all subjects' periods of observation during the time interval:

$$\hat{\lambda} = \frac{D}{Y}.$$

This is the maximum likelihood estimator of λ assuming that D follows a Poisson distribution (independent events occurring at a constant probability rate in continuous time) with mean λY , where Y is treated as fixed.

The most important quantity of interest in a cohort study is the incidence rate ratio (or relative risk), the ratio λ_1/λ_0 of incidence rates for those exposed to a risk factor and those not exposed to the risk factor (subscripts 1 and 0 denote exposed and unexposed, respectively). The incidence rate ratio may be estimated by

$$\hat{\theta} = \frac{D_1/Y_1}{D_0/Y_0}.$$

This estimator can be derived by maximizing the conditional (binomial) likelihood that there were D_1 cases in the exposed group conditional on there being a total of $D = D_0 + D_1$ cases.

However, a potential problem in estimating this rate ratio is confounding arising from systematic differences in prognostic factors between the exposure groups. This problem can be dealt with by dividing the cohort into groups or strata according to prognostic factors and assuming that the rate ratio for exposed and unexposed subjects is the same across strata. If there are D^* cases and Y^* person-years of observation in stratum s , then the common rate ratio may be estimated using the method of Mantel and Haenszel by

$$\hat{\theta}_{MH} = \frac{\sum_s D_s^* Y_s^*/Y^*}{\sum_s D_0^* Y_0^*/Y^*}.$$

Note that the strata might not correspond to groups of subjects. For example, if the confounder is age, subjects who cross from one age-band into the next during the study contribute parts of their periods of observation to different strata. This is how Table 11.1 was constructed.

A more general way of controlling for confounding variables is to use Poisson regression to model the number of occurrences of disease or "failures". Such an approach allows inclusion of several covariates. The complexity of the model can be decided by model selection criteria, often leading to smoothing through the omission of higher order interactions. If a log link is used, the expected number of failures can be made proportional to the person-years of observation by adding the log of the person-years of observation to the linear predictor as an offset (an explanatory variable with regression coefficient set to 1), giving

$$\log[\mathbb{E}(D)] = \log(Y) + \mathbf{x}'\boldsymbol{\beta}.$$

Exponentiating the equation and dividing by Y gives

$$\frac{\mathbb{E}(D)}{Y} = \exp(\mathbf{x}'\boldsymbol{\beta})$$

as required.

11.2.2 Case-control studies

If the incidence rate of a disease is small, a cohort study requires a large number of person-years of observation making it very expensive. A more feasible type of study in this situation is a case-control study in which cases of the disease of interest are compared with non-cases, often called controls, with respect to exposure to possible risk factors in the past. The basic idea of case-control studies is shown in Figure 11.2. The assumption here is that the probability of selection into the study is independent of the exposures of interest. The data in Table 11.2 derive from a case-control study in which cases with lung cancer and healthy controls were interviewed to ascertain whether they had been "exposed" to a pet bird.

Let D and H be the number of cases and controls, respectively, and let the subscripts 0 and 1 denote "unexposed" and "exposed". Since the proportion of cases was determined by the design, it is not possible to estimate the relative risk of disease comparing exposed and unexposed subjects. However, the odds of exposure in the cases or controls can be estimated, and the ratio of these odds is equal to the odds ratio of being a case in the exposed group compared with the unexposed group

$$\frac{D_1/D_0}{H_1/H_0} = \frac{D_1/H_1}{D_0/H_0}.$$

We model the (log) odds of being a case using logistic regression with

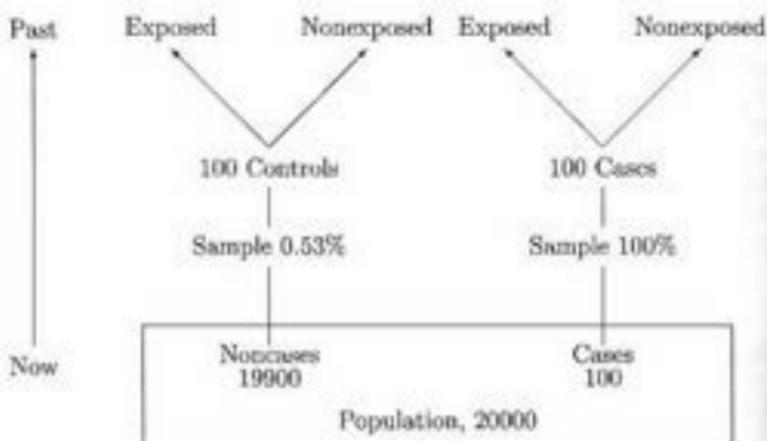


Figure 11.2: Case-control study.

the exposure as an explanatory variable. Then the coefficient of the exposure variable is an estimate of the desired log odds ratio even though the estimate of the odds (which depends on the constant) is determined by the proportion of cases in the study. Logistic regression is the most popular method for estimating adjusted odds ratios for risk factors of interest after controlling for confounding variables.

11.2.3 Matched case-control studies

A major difficulty with case-control studies is to find suitable controls who are similar enough to the cases (so that differences in exposure can reasonably be assumed to be due to their association with the disease) without being overmatched, which can result in very similar exposure patterns. The problem of finding controls who are sufficiently similar is often addressed by matching controls individually to cases according to important variables such as age and sex. Examples of such matched case-control studies are given in Tables 11.3 and 11.4. In the screening study, matching had the following additional advantage noted in Clayton and Hills (1993). The screening history of controls could be determined by considering only the period up to the diagnosis of the case, ensuring that cases did not have a decreased opportunity for screening because they would not have been screened after their diagnosis.

The statistical analysis has to take account of the matching. Two methods of analysis are McNemar's test in the simple case of 2×2 tables and conditional logistic regression in the case of several controls per case and/or several explanatory variables. Since the case-control sets have been matched on variables that are believed to be associated with disease status, the sets can be thought of as strata with subjects in one stratum having higher or lower odds of being a case than those in another stratum after controlling for the exposures. A logistic regression model would have to accommodate these differences by including a parameter α_c for each case-control set c , so that the log odds of being a case for subject i in case-control set c would be

$$\log(\Omega_{ci}) = \log(\alpha_c) + \mathbf{x}'_i \boldsymbol{\beta}. \quad (11.1)$$

However, this would result in too many parameters to be estimated (the incidental parameter problem). Furthermore, the parameters α_c are of no interest to us.

In conditional logistic regression, the nuisance parameters α_c are eliminated as follows. In a 1:1 matched case-control study, ignoring the fact that each set has one case, the probability that subject 1 in the set is a case and subject 2 is a noncase is

$$\Pr(1) = \frac{\Omega_{c1}}{1 + \Omega_{c1}} \times \frac{1}{1 + \Omega_{c2}},$$

and the probability that subject 1 is a noncase and subject 2 is a case is

$$\Pr(2) = \frac{1}{1 + \Omega_{c1}} \times \frac{\Omega_{c2}}{1 + \Omega_{c2}},$$

However, conditional on there being one case in a set, the probability of subject 1 being the case is simply

$$\frac{\Pr(1)}{\Pr(1) + \Pr(2)} = \Omega_{c1}/(\Omega_{c1} + \Omega_{c2}) = \frac{\exp(\mathbf{x}'_1 \boldsymbol{\beta})}{\exp(\mathbf{x}'_1 \boldsymbol{\beta}) + \exp(\mathbf{x}'_2 \boldsymbol{\beta})}, \quad (11.2)$$

since α_c cancels out; see equation (11.1). The expression on the right-hand side of equation (11.2) is the contribution of a single case-control set to the conditional likelihood of the sample. Similarly, it can be shown that if there are k controls per case and the subjects within each case-control set are labeled 1 for the case and 2 to $k+1$ for the

controls then the log likelihood becomes

$$\sum_i \log \left(\frac{\exp(\mathbf{x}_i'\beta)}{\sum_{j=1}^{k+1} \exp(\mathbf{x}_j'\beta)} \right).$$

11.3 Analysis using Stata

11.3.1 Cohort study

There is a collection of instructions in Stata, the `epitab` commands, that may be used to analyze small tables in epidemiology. These commands either refer to variables in an existing dataset or can take cell counts as arguments (i.e., they are immediate commands; see Chapter 1).

The first cohort dataset in Table 11.1 is given in a file as tabulated and may be read using

```
infile str5 age num1 py1 num0 py0 using ibd.dat, clear
```

Here the number of cases and person years have been named `num1` and `py1` in the exposed group and `num0` and `py0` in the unexposed group. We can stack the responses for both groups into variables `num` and `py` using the `reshape` command after producing an identifier `agegr` for the rows in the data which correspond to age groups.

```
generate agegr = _n  
reshape long num py, i(agegr) j(exposed)
```

Ignoring `agegr`, the incidence rate ratio may be estimated using

```
ir num exposed py
```

giving the table in Display 11.1. The incidence rate ratio of ischemic heart disease, comparing low energy with high energy intake, is estimated as 2.46 with a 95% confidence interval from 1.29 to 4.78. (Note that we could report the reciprocals of these figures if we wished to consider high energy intake as the risk factor.) The terms `(exact)` imply that the confidence intervals are exact (no approximation was used).

Controlling for age using the `epitab` command

```
ir num exposed py, by(age)
```

(see Display 11.2) gives very similar estimates as shown in the row labeled M-H combined (the Mantel-Haenszel estimate).

Another way of controlling for age is to carry out Poisson regression with the log of `py` as an offset. The exponentiated offset `py` may be

	exposed Exposed	Unexposed	Total
num 77	28 1857.5	17 2768.9	45 4626.4
Incidence Rate	.015074	.0061396	.0097268
	Point estimate	[95% Conf. Interval]	
Inc. rate diff.	.0088346	.0026342	.0152346
Inc. rate ratio	2.455204	1.297757	4.781095 (exact)
Attr. frac. ex.	.5927019	.22944	.7908429 (exact)
Attr. frac. pop	.3687923		
(midp) Pr(x>=28) =		0.0016 (exact)	
(midp) 2*Pr(x>=28) =		0.0031 (exact)	

Display 11.1

age	IRR	[95% Conf. Interval]	N-H Weight
40-49	.9748111	.0881524	0.790694
50-59	3.476871	1.14018	12.59783
60-69	2.33045	.9123878	8.411597
Crude	2.455204	1.297757	4.781095 (exact)
N-H combined	2.403914	1.306881	4.421829
Test of homogeneity (N-H) chi2(2) =	1.57	Pr>chi2 = 0.4666	

Display 11.2

specified using the `exposure(py)` option. To obtain exponentiated coefficients, we use the `irr` ("incidence rate ratio") option:

`xi: poisson num exposed i.age, exposure(py) irr`
 (see Display 11.3) showing that there is an estimated age-adjusted in-

<code>i.age</code>		<code>.Iage_1-3</code>		<code>(.Iage_1 for age==40-49 omitted)</code>			
Poisson regression					Number of obs = 6		
					LR chi2(3) = 12.91		
					Prob > chi2 = 0.0048		
Log likelihood = -11.898238					Pseudo R2 = 0.3616		
var		IRR	Std. Err.	<i>z</i>	P> z	[95% Conf. Interval]	
<code>exposed</code>		2.386096	.7350226	2.82	0.005	1.304609	4.364108
<code>.Iage_2</code>		1.137701	.5408325	0.27	0.786	.4481154	2.868461
<code>.Iage_3</code>		1.907803	.9218379	1.50	0.134	.8066976	4.935362
<code>py</code> (exposure)							

Display 11.3

cidence rate ratio of 2.39 with a 95% confidence interval from 1.30 to 4.36. The coefficients of `.Iage_2` and `.Iage_3` show that the incidence increases with age (although the rate ratios for age groups are not significant at the 5% level) as would be expected. Another way of achieving the same result is using `glm` with the Poisson distribution and log link with the log of person-years of follow-up specified as an offset using the `offset()` option

```
generate lpy = ln(py)
xi: glm num exposed i.age, family(poisson) link(log) ///
  offset(lpy) eform
```

Here the `eform` option is used to obtain exponentiated coefficients (incidence rate ratios instead of their logarithms). An advantage of this modeling approach is that we can investigate the possibility of an interaction between `exposed` and `age`. If there were more age categories, we could attempt to model the effect of age as a smooth function.

11.3.2 Case-control study

We will analyze the case-control study using the "immediate" command `cci`. The following notation is used for `cci`:

	Exposed	Unexposed
Cases	a	b
Noncases	c	d

where the quantities a , b , etc. in the table are specified in alphabetical order, i.e.,

cci a b c d

(See help epitab for the arguments required for other immediate epitab commands.) The bird data may therefore be analyzed as follows:

cci 98 141 101 328

giving the output in Display 11.4. The odds ratio of lung cancer, com-

	Exposed	Unexposed	Total	Exposed
Cases	98	141	239	0.4100
	101	328	429	0.2354
Total	199	469	668	0.2979
	Point estimate		[95% Conf. Interval]	
Ddds ratio			1.580938	3.218756 (exact)
Aver. frac. ex.	2.267145		.367463	.689321 (exact)
Aver. frac. pop	.5566624			
	.2283779			
	<i>chi2(1) = 22.37 Pr>chi2 = 0.0000</i>			

Display 11.4

paring those with pet birds with those without pet birds, is estimated as 2.26 with an exact 95% confidence interval from 1.58 to 3.22. The *p*-value for the null hypothesis of no association between pet birds and lung cancer is < 0.001. This *p*-value is based on a chi-squared test; an exact *p*-value could be obtained using the *exact* option.

11.3.3 Matched case-control studies

The matched case-control study with one control per case may be analyzed using the immediate command mcci which requires four numbers a to d defined as

		Controls	
Cases	Exposed	Exposed	Unexposed
	Unexposed	a	b

which corresponds to the layout of Table 11.3. The required command therefore is

```
mcci 12 43 7 121
```

The results in Display 11.5 suggest that there is an increased odds of endometrial cancer in subjects exposed to oral conjugated estrogens (odds ratio = 2.89, 95% confidence interval from 1.89 to 4.44).

Cases	Controls		Total
	Exposed	Unexposed	
Exposed	12	43	55
	7	121	128
Total	19	164	183

McNemar's chi2(1) = 28.92 Prob > chi2 = 0.0000
Exact McNemar significance probability = 0.0000

Proportion with factor

Cases	.2005464	
Controls	.1038251	[95% Conf. Interval]
difference	.1967213	.1210924 .2723502
ratio	2.894737	1.885462 4.444269
rel. diff.	.2196122	.1448549 .2941685
odds ratio	6.142657	2.730772 16.18458 (exact)

Display 11.5

The matched case-control study with three controls per case cannot be analyzed using `epitab`. Instead, we will use conditional logistic regression. We need to convert the data in Table 11.4 into the form required for conditional logistic regression; that is, one observation per subject (including cases and controls); an indicator variable, `cancer`, for cases; another indicator variable, `screen`, for screening and a third variable, `caseid`, an identifier for each case-control set of four women.

First, read the data which are in the form shown in Table 11.4.

```
infile vi-v4 using screen.dat, clear
```

Then transpose the data so that the first column contains frequencies for unscreened cases (variable `ncases0`) and the second for screened

cases (variable `ncases1`). This can be achieved by first defining the string variable `_varname` to contain the required variable names and then using the `xpose` command:

```
generate _varname = cond(_n==1,"ncases1","ncases0")
xpose, clear
list
```

(see Display 11.6).

	<code>ncases1</code>	<code>ncases0</code>
1.	1	11
2.	4	10
3.	3	12
4.	1	4

Display 11.6

The four rows in this transposed dataset correspond to 0, 1, 2, and 3 matched controls who have been screened. We will define a variable `nconstr` taking on these four values. We can then stack the two columns into a single variable `ncases` and create an indicator `casescr` for whether or not the case was screened using the `reshape` command:

```
generate nconscr = _n-1
reshape long ncases, i(nconscr) j(casescr)
list
```

(see Display 11.7). The next step is to replicate each of the records `ncases` times so that we have one record per case-control set. Then define the variable `caseid`, and expand the dataset four times in order to have one record per subject. The four subjects within each case-control set are arbitrarily labeled 0 to 3 in the variable `control` where 0 stands for "the case" and 1, 2, and 3 for the controls.

```
expand ncases
sort casescr nconscr
generate caseid = _n
expand 4
quietly by caseid, sort: generate control = _n-1
list in 1/8
```

(see Display 11.8). Now `screen`, the indicator whether a subject was

	score1	change	cases
1.	0	0	11
2.	0	1	1
3.	1	0	10
4.	1	1	4
5.	2	0	12
6.	2	1	3
7.	3	0	4
8.	3	1	1

Display 11.7

	score1	change	score2	caseid	control
1.	0	0	11	1	0
2.	0	0	11	1	1
3.	0	0	11	1	2
4.	0	0	11	1	3
5.	0	0	11	2	0
6.	0	0	11	2	1
7.	0	0	11	2	2
8.	0	0	11	2	3

Display 11.8

screened, is defined to be 0 except for the cases who were screened and for as many controls as were screened according to ncasescr. The variable cancer is 1 for cases and 0 otherwise.

```
generate screen = 0
replace screen = 1 if control==0&casescr==1 /* the case */
replace screen = 1 if control==1&ncasescr>0
replace screen = 1 if control==2&ncasescr>1
replace screen = 1 if control==3&ncasescr>2
generate cancer = control==0
```

We can reproduce Table 11.4 by temporarily collapsing the data (using `preserve` and `restore` to revert to the original data) as follows:

```
preserve
collapse (sum) screen (mean) casescr, by(caseid)
generate ncasescr = screen - casescr
tabulate casescr ncasescr
restore
```

(see Display 11.9).

(mean) casescr	0	1	2	3	Total
0	11	10	12	4	37
1	1	4	3	1	9
Total	12	14	15	5	46

Display 11.9

We are now ready to carry out conditional logistic regression:

```
clogit cancer screen, group(caseid) or
```

(see Display 11.10). Screening therefore seems to be protective of death from breast cancer, reducing the odds to about a third (95% confidence interval from 0.13 to 0.69).

Conditional (fixed-effects) logistic regression	Number of obs	=	184
	Likelihood chi2(1)	=	9.18
	Prob > chi2	=	0.0026
Log likelihood = -89.181618	Pseudo R2	=	0.0719
cancer	Odds Ratio	Std. Err.	z
0.00000	.2995582	.1278366	-2.82
			0.006
			.1297867
			.6914047

Display 11.10

11.4 Exercises

11.1 • Estrogens and endometrial cancer

- Carry out conditional logistic regression to estimate the odds ratio for the data in Table 11.3. The data are given in the same form as in the table in a file called `estrogen.dat`.

11.2 • Low energy diet and heart disease

- For the data `ihd.dat`, use the `iri` command to estimate the incidence rate ratio of IHD comparing subjects with low and high energy diets without controlling for age.
- Use Poisson regression to test whether the effect of exposure to a low energy diet on incidence of IHD differs between age groups.

11.3 Oral contraceptive use and myocardial infarction

Mann *et al.* (1968) analyzed the data shown in Table 11.5 which are also given in Rothman (1986). The data come from a case-control study to investigate the effect of oral contraceptive use on myocardial infarction. Cases and controls are also classified by age group (< 40 and ≥ 40).

The variable in `oral.dta` are:

- `case`: dummy variable for case (myocardial infarction) versus noncase
 - `oral`: dummy variable for oral contraceptive use
 - `age`: age group (0=Age<40, 1=Age ≥ 40)
 - `num`: number of women with given values of `case`, `oral`, and `age`
- Use the `cc` command to estimate the odds ratio for myoc...

Table 11.5 Number of cases of myocardial infarction and controls who did and did not use oral contraceptives by age group

	Younger Age < 40		Older Age ≥ 40	
	User	Nonuser	User	Nonuser
Case	21	26	18	88
Control	17	59	7	95

dial infarction comparing those who have and have not used oral contraceptives. Also estimate the age-adjusted odds ratio using the same command.

2. Discuss why the adjusted and unadjusted odds ratios are not the same here.
3. Now use logistic regression to estimate the age-adjusted odds ratio.

11.4 Induced abortion and ectopic pregnancy

In a matched case-control study, 18 women with ectopic pregnancies were individually matched according to age, number of pregnancies, and husband's education with four controls. All women had had at least one previous pregnancy, and the exposure of interest is having had at least one induced abortion. The data given in Table 11.6 come from Trichopoulos *et al.* (see Miettinen, 1969) and were previously analyzed by Rothman (1986).

Table 11.6 Pattern of exposure (to induced abortion) of 4 controls individually matched to exposed and unexposed cases of ectopic pregnancy

Status of the case	Number of exposed controls				
	0	1	2	3	4
Exposed	3	5	3	0	1
Unexposed	5	1	0	0	0

The variables in `ectopic.dta` are:

- `exposed`: dummy variable for case being exposed
 - `numcon`: number of matched controls who were exposed
 - `numgroups`: number of matched case-control groups with a given status of the case and a given number of matched controls who were exposed
1. Use conditional logistic regression to investigate if there is an association between induced abortion and ectopic pregnancy.
 2. Interpret the estimated odds ratio and confidence interval.

Chapter 12

Survival Analysis: Retention of Heroin Addicts in Methadone Maintenance Treatment

12.1 Description of data

The data to be analyzed in this chapter are on 131 heroin addicts in two different clinics receiving methadone maintenance treatment to help them overcome their addiction. Early dropout is an important problem with this treatment. We will therefore analyze the time from admission to termination of treatment (in days), given as `time` in Table 12.1. For patients still in treatment when these data were collected, `time` is the time from admission to the time of data collection. The variable `status` is an indicator for whether `time` refers to dropout (1) or end of study (0). Possible explanatory variables for retention in treatment are maximum methadone dose and a prison record as well as which of two clinics the addict was treated in. These variables are called `dose`, `prison`, and `clinic`, respectively. The data were first analyzed by Caplehorn and Bell (1991) and also appear in Hand *et al.* (1994).

Table 12.1 Data in `heroin.dat`

<code>id</code>	<code>clinic</code>	<code>status</code>	<code>time</code>	<code>prison</code>	<code>dose</code>	<code>id</code>	<code>clinic</code>	<code>status</code>	<code>time</code>	<code>prison</code>	<code>dose</code>
1	1	1	428	0	50	132	2	0	633	0	70
2	1	1	275	1	55	133	2	1	661	0	40
3	1	1	262	0	55	134	2	1	392	1	20
4	1	1	187	0	30	135	2	1	15	1	60

Table 12.1 Data in heroin.dat (continued)

id	clinic	status	time	prison	class	id	clinic	status	time	prison	class
5	1	1	259	1	66	137	2	0	583	0	70
6	1	1	714	0	55	138	2	0	869	0	80
7	1	1	638	1	65	143	2	0	1052	0	80
8	1	0	796	1	60	144	2	0	944	0	80
9	1	1	662	0	50	145	2	0	881	0	80
10	1	1	363	1	65	146	2	1	190	0	50
11	1	0	161	1	80	148	2	1	79	0	40
12	1	1	826	1	50	149	2	0	884	1	50
13	1	1	523	0	55	150	2	1	170	0	40
14	1	1	612	0	70	153	2	1	286	0	45
15	1	1	212	1	60	156	2	0	358	0	60
16	1	1	398	1	60	158	2	0	326	1	60
17	1	1	771	1	75	159	2	0	769	1	40
18	1	1	511	1	80	160	2	1	162	0	40
19	1	1	532	0	80	161	2	0	864	1	80
21	1	1	624	1	80	162	2	1	268	1	70
22	1	1	269	1	60	163	2	0	611	1	60
23	1	1	341	1	60	164	2	1	322	0	55
24	1	1	299	0	55	165	2	0	1076	1	80
25	1	0	826	0	80	166	2	0	2	1	40
26	1	1	262	1	65	168	2	0	788	0	70
27	1	0	566	1	45	169	2	0	575	0	80
28	1	1	368	1	55	170	2	1	109	1	70
30	1	1	302	1	50	171	2	0	730	1	60
31	1	0	602	0	60	172	2	0	790	0	50
32	1	1	652	0	80	173	2	0	456	1	70
33	1	1	255	0	65	175	2	1	231	1	60
34	1	0	564	0	60	176	2	1	143	1	70
36	1	1	384	1	55	177	2	0	86	1	40
37	1	1	755	1	65	178	2	0	1021	0	80
38	1	1	393	0	55	179	2	0	684	1	80
39	1	0	787	0	80	180	2	1	878	1	60
40	1	1	729	0	60	181	2	1	216	0	60
41	1	1	559	1	60	182	2	0	858	0	60
42	1	1	837	0	60	183	2	1	268	1	40
43	1	1	412	0	65	184	2	0	223	0	40
44	1	0	581	0	70	186	2	0	583	0	100
45	1	1	523	0	60	187	2	0	496	0	40
46	1	1	594	1	60	188	2	1	383	0	55
48	1	1	785	1	80	189	2	1	126	1	75
49	1	1	774	1	65	190	2	1	17	1	60
50	1	1	560	0	65	192	2	1	350	0	60
51	1	1	160	0	35	193	2	0	531	1	65
52	1	1	482	0	30	194	2	0	317	1	50
53	1	1	518	0	65	195	2	1	461	1	75
54	1	1	683	0	50	196	2	1	37	0	60
55	1	1	147	0	65	197	2	1	167	1	55
57	1	1	563	1	70	198	2	1	358	0	45
58	1	1	616	1	60	199	2	1	49	0	60
59	1	1	869	0	60	200	2	1	487	1	40
60	1	1	857	0	60	201	2	1	127	0	20
61	1	1	180	1	70	202	2	1	7	1	40
62	1	1	452	0	60	203	2	1	29	1	60
63	1	1	766	0	60	204	2	1	62	0	40
64	1	1	496	0	65	205	2	0	150	1	60
65	1	1	258	1	60	206	2	1	223	1	60
66	1	1	181	1	60	207	2	0	129	1	40
67	1	1	388	0	60	208	2	0	204	1	65
68	1	0	439	0	80	209	2	1	129	1	50
69	1	0	563	0	75	210	2	1	381	0	65
70	1	1	337	0	60	211	2	1	176	0	55
71	1	0	613	1	60	212	2	1	30	0	60
72	1	1	193	1	80	213	2	1	41	0	60
73	1	0	405	0	80	214	2	0	543	0	40
74	1	1	667	0	80	215	2	0	210	1	50

Table 12.1 Data in heroin.dat (continued)

id	clinic	status	time	prison	dose	id	clinic	status	time	prison	dose
75	1	0	905	0	50	216	1	1	193	1	70
76	1	1	247	0	70	217	1	1	454	0	55
77	1	1	821	0	80	218	1	1	367	0	45
78	1	1	821	1	75	219	1	1	348	1	60
79	1	0	517	0	45	220	1	0	28	0	50
80	1	0	346	1	60	221	1	0	325	0	40
81	1	1	294	0	65	222	1	0	175	1	60
82	1	1	244	1	60	223	2	1	149	1	50
83	1	1	50	1	60	224	1	1	546	1	50
84	1	1	256	1	55	225	1	1	84	0	45
85	1	1	212	0	40	226	1	0	283	1	60
86	1	1	96	0	70	227	1	1	533	0	55
87	1	1	522	0	80	228	1	1	207	1	50
88	1	1	522	1	70	229	1	1	216	0	50
89	1	1	879	0	35	230	1	0	28	0	50
90	1	0	408	0	50	231	1	1	67	1	50
91	2	0	840	0	80	232	1	0	62	1	60
92	2	0	148	1	65	233	1	0	111	0	55
93	1	1	168	0	65	234	1	1	257	1	60
94	1	1	589	0	80	235	1	1	138	1	55
95	1	0	511	0	50	236	1	0	342	0	60
96	1	1	205	0	50	237	1	1	43	0	45
97	1	0	478	1	75	238	1	0	531	1	45
98	1	1	237	0	45	239	1	0	95	0	40
99	1	1	517	0	70	240	1	1	145	1	55
100	1	1	749	0	70	241	1	1	59	0	50
101	1	1	150	1	80	242	1	0	53	0	50
102	1	1	465	0	65	243	1	0	193	1	50
103	2	1	708	1	60	244	1	0	2	1	60
104	2	0	713	0	50	245	1	1	157	1	60
105	2	0	145	0	50	246	1	1	75	1	55
106	2	1	450	0	55	247	1	1	19	1	40
109	2	0	555	0	80	248	1	1	35	0	50
110	2	1	460	0	50	249	2	0	394	1	50
111	2	0	53	1	60	250	1	1	117	0	40
113	2	1	122	1	60	251	1	1	175	1	60
114	2	1	35	1	40	252	1	1	180	1	60
118	2	0	532	0	70	253	1	1	314	0	70
119	2	0	684	0	65	254	1	0	480	0	30
120	2	0	769	1	70	255	1	0	323	1	60
121	2	0	501	0	70	256	2	1	280	0	50
122	2	0	769	1	40	257	1	1	264	0	50
123	2	0	605	1	100	258	2	1	366	0	55
124	2	0	802	1	80	259	2	0	531	1	50
125	2	0	932	1	80	260	1	1	59	1	45
126	2	0	587	0	110	261	1	1	33	1	60
127	2	1	28	0	40	262	2	1	540	0	50
128	2	0	72	1	60	263	2	0	351	0	65
129	2	0	641	0	70	264	1	1	90	0	40
131	2	0	387	0	70	265	1	1	47	0	45

The data can be described as *survival* data, although the "endpoint" is not death in this case, but dropout from treatment. From engineering applications, another commonly used term for the endpoint is "failure". Duration or survival data can generally not be analyzed by conventional methods such as linear regression. The main reason for this is that some durations are usually right-censored; that is, the endpoint of interest has not occurred during the period of observation and all that is known about the duration is that it exceeds the observation period. In the present dataset, this applies to all observations

where `status` is 0. Another reason why conventional linear regression would not be appropriate is that survival times tend to have positively skewed distributions. A third reason is that time-varying covariates, such as the time of year, could not be handled. In the next section, we therefore describe methods specifically developed for survival data.

12.2 Survival analysis

12.2.1 Introduction

The survival time T may be regarded as a random variable with a probability distribution $F(t)$ and probability density function $f(t)$. An obvious quantity of interest is the probability of surviving to time t or beyond, the *survivor function* or survival curve $S(t)$, which is given by

$$S(t) = P(T \geq t) = 1 - F(t). \quad (12.1)$$

A further function which is of interest for survival data is the *hazard function*. This represents the instantaneous failure rate, that is, the probability that an individual experiences the event of interest at a time point given that the event has not yet occurred. It can be shown that the hazard function is given by

$$h(t) = \frac{f(t)}{S(t)}. \quad (12.2)$$

the instantaneous probability of failure at time t divided by the probability of surviving up to time t . Note that the hazard function is just the incidence rate discussed in Chapter 11. It follows from equations (12.1) and (12.2) that

$$-\frac{d \log(S(t))}{dt} = h(t),$$

so that

$$S(t) = \exp(-H(t)), \quad (12.3)$$

where $H(t)$ is the integrated hazard function, also known as the *cumulative hazard function*.

12.2.2 Kaplan-Meier estimator

The Kaplan-Meier estimator is a nonparametric estimator of the survivor function $S(t)$. If all the failure times, or times at which the event occurs in the sample, are ordered and labeled $t_{(j)}$ such that $t_{(1)} \leq t_{(2)} \cdots \leq t_{(n)}$, the estimator is given by

$$\hat{S}(t) = \prod_{j: t_{(j)} \leq t} \left(1 - \frac{d_j}{n_j}\right),$$

where d_j is the number of individuals who experience the event at time $t_{(j)}$, and n_j is the number of individuals who have not yet experienced the event at that time and are therefore still "at risk" of experiencing it (including those censored at $t_{(j)}$). The product is over all failure times less than or equal to t .

12.2.3 Cox regression

We can compare survival in different subgroups by plotting the Kaplan-Meier estimators of the group-specific survivor functions and applying simple significance tests (such as the log-rank test). However, when there are several explanatory variables, and in particular when some of these are continuous, it is much more useful to use a regression method such as Cox regression. Here the hazard function for individual i is modeled as

$$h_i(t) = h_0(t) \exp(\mathbf{x}_i' \boldsymbol{\beta}), \quad (12.4)$$

where $h_0(t)$ is the baseline hazard function, $\boldsymbol{\beta}$ are regression coefficients, and \mathbf{x}_i covariates. The baseline hazard is the hazard when all covariates are zero, and this quantity is left unspecified. This nonparametric treatment of the baseline hazard combined with a parametric representation of the effects of covariates gives rise to the term *semiparametric model*. The main assumption of the model is that the hazard of any individual i is a time-constant multiple of the hazard function of any other individual j , the factor being $\exp((\mathbf{x}_i - \mathbf{x}_j)' \boldsymbol{\beta})$, the hazard ratio or incidence rate ratio. This property is called the *proportional hazards assumption*. The exponentiated regression coefficients can therefore be interpreted as hazard ratios when the corresponding explanatory variables increase by one unit if all other covariates remain constant.

The parameters $\boldsymbol{\beta}$ are estimated by maximizing the partial log likelihood

likelihood given by

$$\sum_f \log \left(\frac{\exp(\mathbf{x}_f' \boldsymbol{\beta})}{\sum_{i \in r(f)} \exp(\mathbf{x}_i' \boldsymbol{\beta})} \right) \quad (12.5)$$

where the first summation is over all failures f , and the second summation is over all subjects $r(f)$ who are still at risk at the time of failure, the "risk set". It can be shown that this log likelihood is a log profile likelihood (i.e., the log of the likelihood in which the baseline hazard parameters have been replaced by functions of $\boldsymbol{\beta}$ which maximize the likelihood for fixed $\boldsymbol{\beta}$). Note also that the likelihood in equation (12.5) is equivalent to the likelihood for matched case-control studies described in Chapter 11 if the subjects at risk at the time of a failure (the risk set) are regarded as controls matched to the case failing at that point in time (see Clayton and Hills, 1993).

The baseline hazard function may be estimated by maximizing the full log likelihood with the regression parameters evaluated at their estimated values, giving nonzero values only when a failure occurs. Integrating the hazard function gives the cumulative hazard function

$$H_i(t) = H_0(t) \exp(\mathbf{x}_i' \boldsymbol{\beta}), \quad (12.6)$$

where $H_0(t)$ is the integral of $h_0(t)$. The survival curve may be obtained from $H(t)$ using equation (12.3). This leads to the Kaplan-Meier estimator when there are no covariates.

It follows from equation (12.3) that the survival curve for a Cox model is given by

$$S_i(t) = S_0(t)^{\exp(\mathbf{x}_i' \boldsymbol{\beta})}. \quad (12.7)$$

The log of the cumulative hazard function predicted by the Cox model is given by

$$\log(H_i(t)) = \log H_0(t) + \mathbf{x}_i' \boldsymbol{\beta}, \quad (12.8)$$

so that the log cumulative hazard functions of any two subjects i and j are parallel with constant difference given by $(\mathbf{x}_i - \mathbf{x}_j)' \boldsymbol{\beta}$.

Stratified Cox regression can be used to relax the assumption of proportional hazards for a categorical predictor. The partial likelihood of a stratified Cox model has the same form as equation (12.5) except that the risk set $r(f)$ for each failure is now confined to subjects in the same stratum as the subject contributing to the numerator.

Survival analysis is described in Allison (1984), Clayton and Hills (1993), Collett (2003), and Klein and Moeschberger (2003). Cleves *et al.* (2004) discuss survival analysis using Stata.

12.3 Analysis using Stata

The data are available as an ASCII file called `heroin.dat` on the disk accompanying Hand *et al.* (1994). Since the data are stored in a two-column format with the set of variables repeated twice in each row, as shown in Table 12.1, we have to use `reshape` to bring the data into the usual form:

```
infile id1 clinic1 status1 time1 prison1 dose1 ///
       id2 clinic2 status2 time2 prison2 dose2 ///
       using heroin.dat, clear
generate row=_n
reshape long id clinic status time prison dose, ///
  i(row) j(col)
drop row col
```

Before fitting any survival models, we declare the data as being of the form `st` (for survival time) using the `stset` command

```
stset time, failure(status)

failure event: status != 0 & status < .
obs. time interval: (0, time]
exit on or before: failure

238 total obs.
0 exclusions

238 obs. remaining, representing
150 failures in single record/single failure data
25812 total analysis time at risk, at risk from t =
earliest observed entry t =
last observed exit t = 1075
```

and look at a summary of the data using

```
stsum, by(clinic)

failure _d: status
analysis time _t: time

      incidence          no. of
clinic   time at risk    rate#    subjects | Survival time
                                                25%    50%    75%
1        59558    .0020484     163     182     425     682
2        36254    .0007723      75     280      -      -
total    95812    .0015635     238     312     504     821
```

There are 238 subjects in total with a median survival time of 504 days. If the incidence rate (i.e., the hazard function) could be assumed to be constant, it would be estimated as 0.0016 per day (which corresponds to 0.57 per year). Overall, 25% of subjects remain in the clinic at least 212 days, but this differs considerably by clinic (192 subjects in clinic 1 and 280 in clinic 2). In fact, in clinic 2, less than 50% of people had dropped out by the end of follow up so that the median survival time and the time until 75% of subjects have dropped out are not given.

The Kaplan-Meier estimator of the survivor functions for the two clinics are obtained and plotted using

```
sts graph, by(clinic)
```

giving the graph in Figure 12.1. Dropout seems to occur at a faster rate

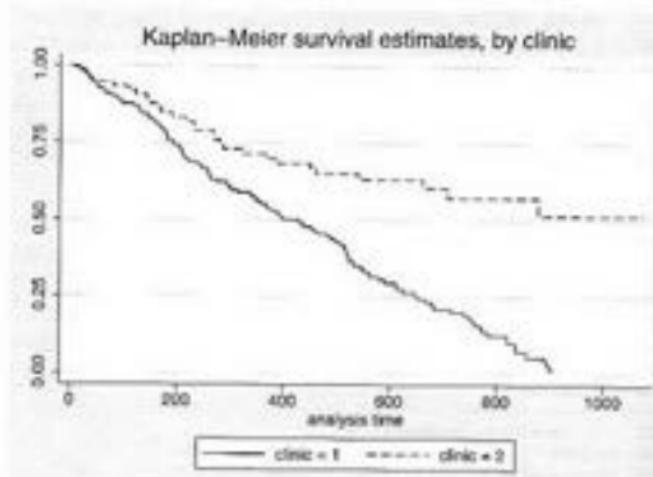


Figure 12.1: Kaplan-Meier survival curves.

in clinic 1. According to Caplehorn and Bell (1991), the more rapid decline in the proportion remaining in clinic 1 compared with clinic 2 may be due to the policy of clinic 1 to attempt to limit the duration of maintenance treatment to two years.

To investigate the effects of dose and prison on survival, we will use Cox regression. We will allow the hazard functions for the two clinics to be non-proportional. A Cox regression model with clinics as strata can be estimated using the `stcox` command with the `strata` option:

option:

```
stcox dose prison, strata(clinic)
```

giving the output shown in Display 12.1. Therefore, subjects with a

failure _d: status		analysis time _t: time		Stratified Cox regr. -- Efron method for ties				
No. of subjects	=	238	Number of obs	=	238			
No. of failures	=	150						
Time at risk	=	95812			LR chi2(2) = 33.94			
Log likelihood	=	-597.714			Prob > chi2 = 0.0000			
_t	Haz. Ratio	Std. Err.	Z	Pri>	[95% Conf. Interval]			
dose	.9654655	.0062438	-5.44	0.000	.953309	.977777		
prison	1.475192	.2491827	2.30	0.021	1.059418	2.054138		

Stratified by clinic

Display 12.1

prison history are 47.5% more likely to drop out at any given time given that they remained in treatment until that time) than those without a prison history. For every increase in methadone dose by one unit (1 mg), the hazard is multiplied by 0.965. This coefficient is very close to one, but this may be because one unit of methadone dose is not a large quantity. Even if we know little about methadone maintenance treatment, we can assess how much one unit of methadone dose is by finding the sample standard deviation:

summarize dose		Obs	Mean	Std. Dev.	Min	Max
Variable						
dose		238	80.39916	14.45013	20	130

indicating that a unit is not much at all; subjects often differ from each other by 10 to 15 units. To find the hazard ratio of two subjects differing by one standard deviation, we need to raise the hazard ratio to the power of one standard deviation, giving $0.9654655^{14.45013} = 0.60179167$. We can obtain the same result (with greater precision) using the stored macros `b[dose]` for the log hazard ratio and `r[Var]` for the variance,

```
display exp(.b[dose]*sqrt(r(Var)))
.60179167
```

In the above calculation, we simply rescaled the regression coefficient before taking the exponential. To obtain this hazard ratio in the Cox regression, we need to standardize `dose` to have unit standard deviation. In the command below we also standardize to mean zero, although this will make no difference to the estimated coefficients (only the baseline hazards are affected):

```
egen zdose = std(dose)
```

We repeat the Cox regression with `zdose` instead of `dose`:

```
stcox zdose prison, strata(clinic)
```

(see Display 12.2). The coefficient of `zdose` is identical to that calcu-

		Failure _d: status			
		analysis time _t: time			
		Stratified Cox regr. --- Breslow method for ties			
No. of subjects =	238			Number of obs =	238
No. of failures =	160				
Time at risk =	86812				
				LR chi2(2) =	33.94
Log likelihood =	-897.714			Prob > chi2 =	0.0000
_t	Haz. Ratio	Std. Err.	Z	P> z	[95% Conf. Interval]
zdose	.6017887	.0662195	-5.44	0.000	.5010998 .7227087
prison	1.475192	.2491827	2.30	0.021	1.059418 2.064138
Stratified by clinic					

Display 12.2

lated previously and may now be interpreted as indicating a decrease of the hazard by 40% when the methadone dose increases by one standard deviation.

Assuming the variables `prison` and `zdose` satisfy the proportional hazards assumption (see Section 12.3.1), we now present the model graphically. To do this, we will plot the predicted survival curves separately for the two clinics and for those with and without a prison record where `zdose` is evaluated at its clinic-specific mean. Such a graph may be produced by using `stcox` with the `bases()` option to generate a variable containing the predicted baseline survival function and then applying equation (12.7) to obtain the predicted survival functions for particular covariate values:

```
stcox zdose prison, strata(clinic) bases(s)
```

```
egen mdose = mean(zdose), by(clinic)
generate surv = exp(_b[zdose]*mdose + _b[prison]*prison)
```

Note that these survival functions represent predicted values for subjects having the clinic-specific mean dose. We now transform time to time in years, and plot the survival curves separately for each clinic:

```
generate tt = time/365.26
label variable tt "time in years"
label define clin 1 "Clinic 1" 2 "Clinic 2"
label values clinic clin

sort clinic time
twoway (line surv tt if prison==0, connect(stairstep)) ///
    (line surv tt if prison==1, connect(stairstep)) ///
    lpatt(dash), by(clinic) ylabel(0(.2)1) ///
    legend(order(1 "Prison record" 2 "No prison record"))
```

Here the `connect(stairstep)` option was used to produce the step-shaped survival curves shown in Figure 12.2.

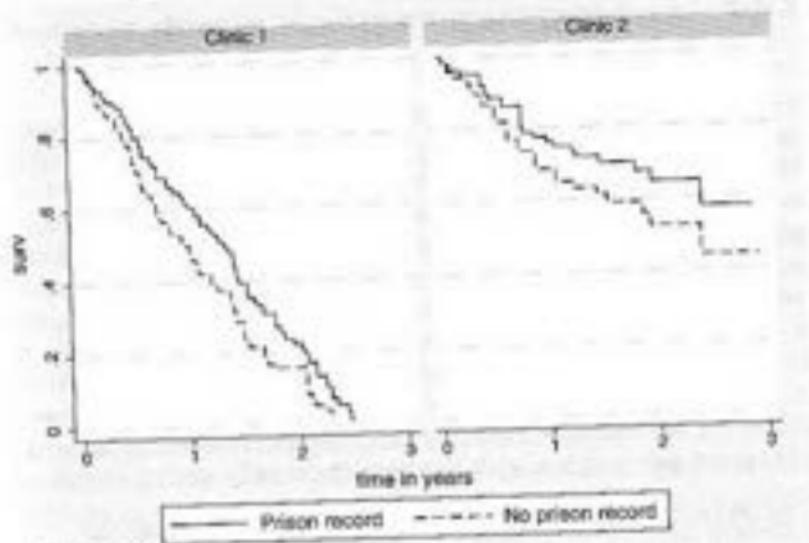


Figure 12.2: Survival curves.

The partial likelihood is appropriate for continuous survival times, which should theoretically never take on the same value for any two

individuals or be *tied*. However, in practice survival times are measured in discrete units, days in the present example, and ties will frequently occur. For example, subjects 61 and 252 in clinic 1 both dropped out after 180 days. By default `stcox` uses *Breslow's method* for ties, where the risk sets in (12.5) contain all subjects who failed at or after the failure time of the subject contributing to the numerator. For a group of subjects with tied survival times, the contributions to the partial likelihood therefore each have the same denominator. However, risk sets usually decrease by one after each failure. In *Efron's method*, contributions to the risk set from the subjects with tied failure times are therefore downweighted in successive risk sets. In the *exact* method (referred to as "exact marginal log likelihood" in [ST] `stcox`, the Stata reference manual for survival analysis), the contribution to the partial likelihood from a group of tied survival times is the sum, over all possible orderings (or permutations) of the tied survival times, of the contributions to the partial likelihood corresponding to these orderings. Efron's method can be obtained using the `efrom` option and the *exact* method using the `exactm` option (see Exercise 12.1).

12.3.1 Assessing the proportional hazards assumption

12.3.1.1 Graphical methods

We now discuss methods for assessing the proportional hazards assumption. A graphical approach is available for categorical predictors if there are sufficient observations for each value of the predictor. In this case the model is first estimated by stratifying on the categorical predictor of interest, thus not making any assumption regarding the relationship between the baseline hazards for different values of the predictor or strata. The log cumulative baseline hazards for the strata are then derived from the estimated model and plotted against time. According to equation (12.8), the resulting curves should be parallel if the proportional hazards assumption holds. Here we demonstrate this method for the variable `clinic`. The cumulative baseline hazard can be obtained using `stcox` with the `basech()` option as follows:

```
quietly stcox zdose prison, strata(clinic) basech(ch)
```

We now compute and then plot the logarithm of the cumulative baseline hazard function using

```
generate lh = log(ch)
sort time
twoway (line lh time if clinic==1, connect(stairstep)) ///
    (line lh time if clinic==2, connect(stairstep) ///
    lpatt(dash)), xtitle("Time in days") ///
```

```
ytitle("Log cumulative hazard") /////
legend(order(1 "Clinic 1" 2 "Clinic 2"))
```

giving the graph shown in Figure 12.3. Clearly, the curves are not

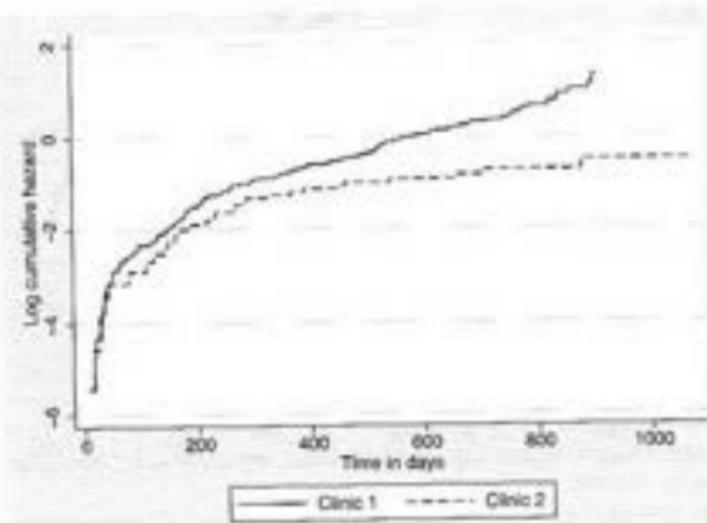


Figure 12.3: Log of minus the log of the survival functions for the two clinics estimated by stratified Cox regression.

parallel, and we will therefore continue treating the clinics as strata. Note that a quicker way of producing a similar graph would be to use the stphplot command as follows:

```
stphplot, strata(clinic) adjust(zdose prison) zero ///
 xlabel(1/7)
```

Here the `adjust()` option specifies the covariates to be used in the Cox regression, and the `zero` option specifies that these covariates are to be evaluated at zero. As a result, minus the logs of the cumulative baseline hazard functions (stratified by clinic) are plotted against the log of the survival time, see Figure 12.4.

To determine whether the hazard functions for those with and without a prison history are proportional, we could split the data into four strata by `clinic` and `prison`. However, as the strata get smaller, the estimated survival functions become less precise (because the risk sets in equation (12.5) become smaller). Also, a similar method could not be used to check the proportional hazard assumption for the continuous

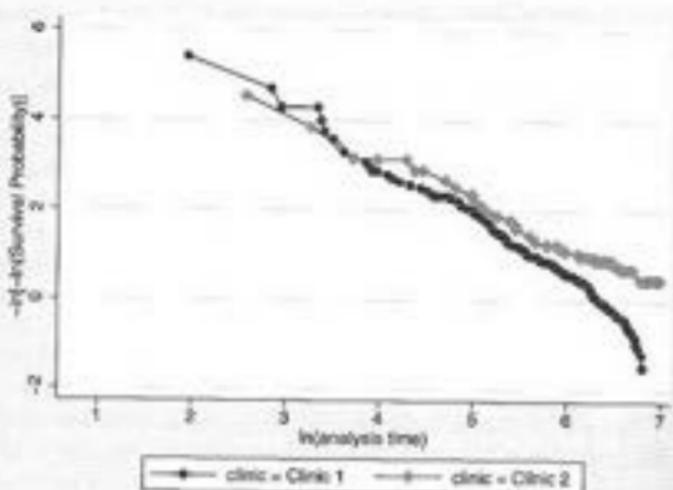


Figure 12.4: Minus the log of minus the log of the survival functions for the two clinics versus log survival time, estimated by stratified Cox regression.

variable `zdose` without splitting it into arbitrary categories.

12.3.1.2 Time-varying covariates

Another way of testing the proportional hazards assumption of `zdose`, say, is to introduce a time-varying covariate equal to the interaction between time (since admission) and `zdose`, thus allowing the effect of `zdose` to change over time. To estimate this model, the terms in equation (12.5) need to be evaluated for values of the time-varying covariates at the times of the failure in the numerator. These values are not available for the denominator in the present dataset since each subject is represented only once, at the time of their own failure (and not at all previous failure times). One possibility is to create the required dataset (see below). A simpler option is to simply use the `stcox` command with the `tvc()` and `txp()` options:

```
stcox zdose prison, strata(clinic) tvc(zdose) ///
txp({_t-504}/365.25)
```

The `tvc()` option specifies the variable that should interact with (a function of) time and the `txp()` option specifies the function of time to be used. Here we have simply subtracted the median survival time so that the coefficient of `zdose` can be interpreted as the effect of

`zdose` at the median survival time. We have divided by 365.25 to see by how much the effect of `zdose` changes between intervals of one year. The output is shown in Display 12.3 where the estimated increase

```

failure _d: status
analysis time _t: time
Stratified Cox regr. -- Efron's method for ties
No. of subjects =      338                      Number of obs =      338
No. of failures =      150
Time at risk =        98812
LR chi2(3) =          34.78
Prob > chi2 =        0.0000
Log likelihood = -597.29131

```

	<i>t</i>	Haz. Ratio	Std. Err.	<i>s</i>	Pival	[95% Conf. Interval]
z _b		.6442974	.0767535	-3.69	0.000	.5101348 .8137442
zdose		1.481193	.2429978	2.33	0.020	1.064036 2.061099
z _c		1.147853	.1720104	0.92	0.387	.6557175 1.539722

Stratified by clinic

Note: Second equation contains variables that continuously vary with respect to time; variables are interacted with current values of $(_t-504)/365.25$.

Display 12.3

in the hazard ratio for `zdose` is 15% per year. This small effect is not significant at the 5% level which is confirmed by carrying out the likelihood ratio test as follows:

```

estimates store model1
quietly stcox zdose prison, strata(clinic)
lrtest model1
likelihood-ratio test
(Assumption: model1 nested in model2)
LR chi2(1) =      0.85
Prob > chi2 =    0.3579

```

giving a very similar p-value as before and confirming that there is no evidence that the effect of dose on the hazard varies with time. A similar test can be carried out for `prison` (see Exercise 12.1).

Although Stata makes it very easy to include an interaction between a variable and a function of time, inclusion of other time-varying covariates, or of more than a single time-varying covariate, requires an expanded version of the current dataset. In the expanded dataset each subject's record should appear (at least) as many times as that subject contributes to a risk set in equation (12.5), with the time variable equal

to the corresponding failure times. This can be achieved very easily using the `stsplit` command, but only after defining an `id` variable using `stset`:

```
stset time, failure(status) id(id)
stsplit, at(failures) strata(clinic)
```

The `stsplit` command generates new time and censoring variables `_t` and `_d`, respectively. For subject 103, these are listed using

```
sort id _t
list _t _d if id==103, clean noobs
```

giving the values shown in Display 12.4. The last value of `_t` (708)

	<code>_t</code>	<code>_d</code>
13	0	
26	0	
35	0	
41	0	
79	0	
109	0	
122	0	
143	0	
149	0	
161	0	
170	0	
190	0	
216	0	
231	0	
202	0	
268	0	
289	0	
286	0	
323	0	
366	0	
389	0	
450	0	
460	0	
540	0	
661	0	
708	1	

Display 12.4

is just the value of the original variable `time`, the subject's survival or censoring time, whereas the previous values are all unique survival times (at which failures occurred) in the same stratum (clinic 2) which are less than the subject's own survival time. These "invented" survival times are times beyond which the subject survives, so the censoring variable `_d` is set to zero for all invented times and equal to `status` for

the original survival time. This new survival dataset is equivalent to the original one, and we obtain the same results as before if we run

```
stcox zdose prison, strata(clinic)
```

(output not shown). However, we can now create time-varying covariates making use of the new time variable *.t*. To assess the proportional hazards assumption for *zdose*, we generate an interaction between *zdose* and the linear transformation of *.t* we used previously:

```
generate tdoze = zdose*(t-504)/365.25
```

We now fit the Cox regression, allowing the effect of *zdose* to vary with time:

```
stcox zdose tdoze prison, strata(clinic)
```

giving the same result as before in Display 12.5.

	failure _d: status	analysis time _t: time	id: id				
Stratified Cox regr. -- Breslow method for ties							
No. of subjects =	238			Number of obs =	11228		
No. of failures =	150						
Time at risk =	95812			LR chi2(2) =	33.94		
Log likelihood = -597.714				Prob > chi2 =	0.0000		
_t	Haz. Ratio	Std. Err.	N	P> z	[95% Conf. Interval]		
zdose	.6017887	.0863190	-8.44	0.000	.5010996	.7227067	
prison	1.476192	.2481627	2.30	0.021	1.069418	9.064138	
							Stratified by clinic

Display 12.5

We can restore the original data using the *stjoin* command after deleting any time-varying covariates (apart from *.t* and *.d*):

```
drop tdoze
stjoin
```

A test of proportional hazards based on rescaled Schoenfeld or efficient score residuals (see below), suggested by Grambsch and Therneau (1994), is also available using the *estat phtest* command (see for example Cleves et al., 2004).

12.3.2 Residuals

It is a good idea to produce some residual plots, for example a graph of the deviance residuals against the linear predictor. In order for predict to be able to compute the deviance residuals, we must first store the martingale residuals (see for example Collett, 2003) using stcox with the mgale() option:

```
stcox mdoce prison, strata(clinic) mgale(mart)
predict devr, deviance
```

A scatterplot is produced using

```
predict xb, xb
twoway scatter devr xb, mlabel(id) msabpos(0) ///
msymbol(none)
```

with the result shown in Figure 12.5. There appear to be no serious outliers.

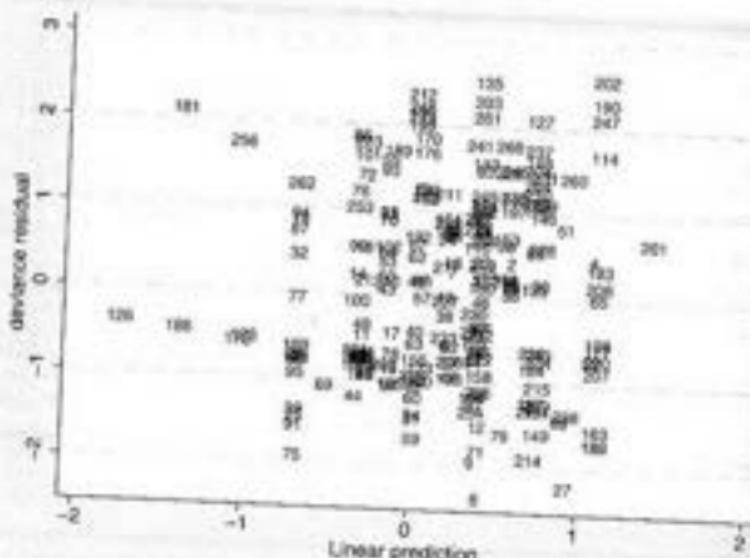


Figure 12.5: Deviance residuals for survival analysis.

Another type of residual is the Schoenfeld or efficient score residual, defined as the first derivative of the partial log likelihood function with respect to an explanatory variable. The score residual is large in absolute value if a case's explanatory variable differs substantially from

the explanatory variables of subjects whose estimated risk of failure is large at the case's time of failure or censoring. Since our model has two explanatory variables, we can compute the efficient score residuals for *zdose* and *prison* and store them in *score1* and *score2* using *stcox* with the *esr* option:

```
stcox zdose prison, strata(clinic) esr(score*)
```

These residuals can be plotted against survival time using

```
twoway scatter score1 tt, xlabel(id) xlabelpos(0) msymbol(none)
```

and similarly for *score2*. The resulting graphs are shown in Figures 12.6 and 12.7. Subject 89 has a low value of *zdose* (-1.75) compared with other subjects at risk of failure at such a late time. Subjects 8, 27, 12, and 71 drop out relatively late considering that they have a police record, whereas others remaining beyond their time of dropout (or censoring) tend to have no police record.

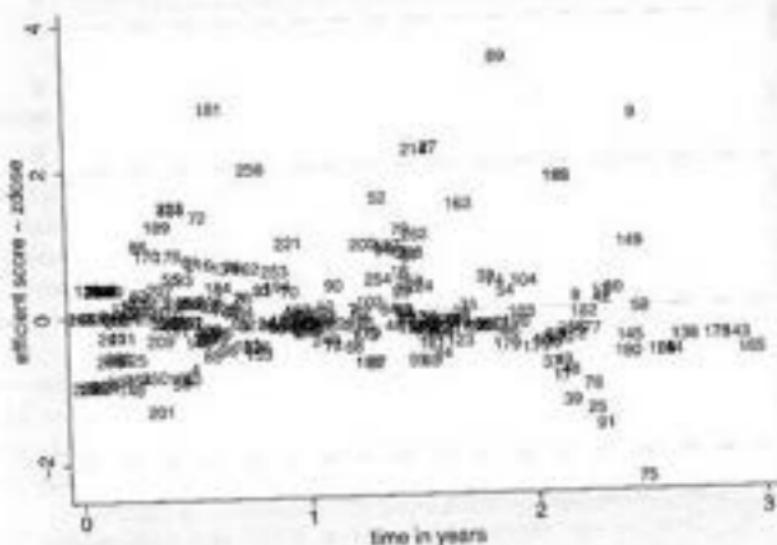


Figure 12.6: Score residuals for *zdose*.

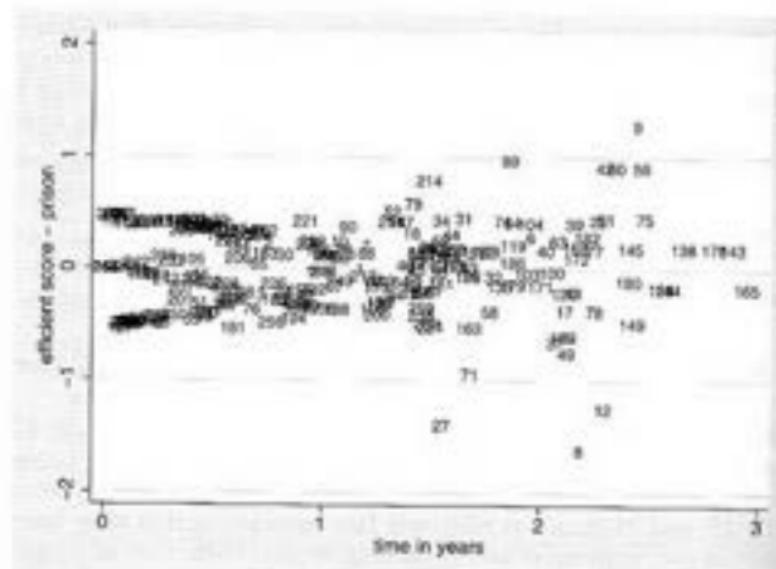


Figure 12.7: Score residuals for prison.

12.4 Exercises

12.1 • Retention of heroin addicts in methadone maintenance treatment

1. In the original analysis of this data, Caplehorn and Bell (1991) judged that the hazards were approximately proportional for the first 450 days (see Figure 12.3). They therefore analyzed the data for this time period using `clinic` as a covariate instead of stratifying by clinic. Repeat this analysis, using `prison` and `dose` as further covariates.
2. Following Caplehorn and Bell (1991), repeat the above analysis treating `dose` as a categorical variable with three levels (< 60 , $60 - 79$, ≥ 80), and plot the predicted survival curves for the three dose categories when `prison` and `clinic` take on one of their values.
3. Test for an interaction between `clinic` and methadone dose, treating `dose` as both continuous and categorical.
4. For the model treating `dose` as categorical and containing no interaction, compare the estimates using three different methods of handling ties: the Breslow, Efron, and exact methods.
5. Check the proportional hazards assumption for `prison` using

a graphical method.

12.2 Survival of patients with primary biliary cirrhosis

The data to be analyzed here come from a Mayo Clinic trial conducted between 1974 and 1984 and accompany the book by Therneau and Grambsch (2000). They were previously analyzed by Dickson et al. (1989) and Fleming and Harrington (1991). Patients with primary biliary cirrhosis were randomized to receive either D-penicillamine or placebo. The treatment was found not to be effective in prolonging survival, and the data, supplemented with an additional 106 patients not participating in the trial, have been used to develop a model for survival in a "natural history setting".

The variables in `pbc.dta` that will be used here are:

- `id`: subject identifier
 - `futime`: number of days between registration and the earlier of liver transplant, death, or end of follow-up
 - `status`: status (0=alive, 1=liver transplant, 2=dead)
 - `age`: age in days
 - `edema`: presence of edema (0=no edema and no therapy for edema, 0.5=edema, either not treated or resolved by treatment, 1=edema despite treatment)
 - `bilir`: serum bilirubin concentration (mg/dl)
 - `prothr`: prothrombin time (seconds)
 - `album`: albumin concentration (mg/dl)
1. Form the natural logarithms of `bilir`, `prothr`, and `album` and convert `age` to age in years.
 2. Fit a Cox regression model with the above transformed variables and the categorical variable `edema` as covariates. Treat liver transplantation as censoring.
 3. Interpret the estimated hazard ratios.
 4. Relax the proportional hazards assumption for `age` using an interaction with analysis time and use a likelihood ratio test to assess the assumption.
 5. Perform an analogous test for `edema`. Also use a graphical method for assessing the proportional hazards assumption for `edema`.
 6. Produce and plot efficient score residuals for all covariates in the model.

12.3 Duration of UN peacekeeping missions

Box-Steffensmeier and Jones (2004) provide a dataset on the duration of UN peacekeeping missions between 1948 and 2001. The data were originally analyzed by Green *et al.* (1998).

The variables in `un.dta` used for this exercise are:

- `duration`: duration of peacekeeping mission until the earlier of the completion date or 2001
 - `complete`: dummy variable for peacekeeping mission being completed
 - `contype`: the type of conflict that led to the peacekeeping mission (1=civil war, 2=interstate conflict, 3=internationalized civil war)
1. Produce Kaplan-Meier survival curves by type of contest.
 2. Fit a Cox proportional hazards model to investigate the effect of type of conflict on the duration on UN peacekeeping missions. Use the exact method for handling ties. Interpret the estimates.
 3. Plot the model-implied survival curves for the three types of conflict and compare them with the Kaplan-Meier curves.
 4. Test the null hypothesis that type of contest does not affect the duration of peacekeeping missions using a Wald test based on the estimates from the Cox regression and using a log rank test (see `help sts test`) and compare the results.

12.4 Treatment of prostate cancer

Here we consider data from a clinical trial for the treatment of prostate cancer that were previously analyzed by Collett (2003) and Everitt and Pickles (2004). (This is a subset of data analyzed by Andrews and Herzberg, 1985.) Patients were randomized to 1mg per day of diethylstilbestrol (DFS) or placebo and their survival was recorded in months.

The variables in `survprost.dta` are:

- `time`: time from the start of the trial to death or censoring (in months)
- `status`: dummy variable for death (versus censoring)
- `treatment`: dummy variable for DFS treatment (versus placebo)
- `age`: age at the start of the trial (in years)
- `haem`: serum haemoglobin level in gml/100ml
- `gleason`: a combined index of tumor stage and grade (a larger index indicates a more advanced tumor)

1. Fit a Cox regression model to estimate the unadjusted hazard ratio for DFS treatment versus placebo. Interpret the estimated treatment effect.
2. Use a forward selection procedure to select additional covariates, in addition to treatment, among the other variables listed above. Note that you should force treatment to be in the model (see help `stepwise`).
3. Interpret the estimates.
4. For the selected model, plot the model-implied survival curves by treatment group, evaluating the covariates at their mean.

Chapter 13

Maximum Likelihood Estimation: Age of Onset of Schizophrenia

13.1 Description of data

Table 13.1 gives the ages of onset of schizophrenia (determined as age on first admission) for 99 women. These data will be used to investigate whether there is any evidence for the subtype model of schizophrenia (see Lewine, 1981), according to which there are two types of schizophrenia characterized by early and late onset.

13.2 Finite mixture distributions

The most common type of finite mixture distribution for continuous responses is a mixture of univariate normal distributions of the form

$$f(y_i; \mathbf{p}, \boldsymbol{\mu}, \boldsymbol{\sigma}) = p_1 g(y_i; \mu_1, \sigma_1) + p_2 g(y_i; \mu_2, \sigma_2) + \cdots + p_k g(y_i; \mu_k, \sigma_k),$$

where $g(y; \mu, \sigma)$ is the normal or Gaussian density with mean μ and standard deviation σ .

$$g(y; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left\{-\frac{1}{2}\left(\frac{y-\mu}{\sigma}\right)^2\right\}, \quad (13.1)$$

Table 13.1 Age of onset of schizophrenia data in onset.dat

Ages of onset									
20	30	21	23	30	25	13	19		
16	25	20	25	27	43	6	21		
15	26	23	21	23	23	34	14		
17	18	21	16	35	32	48	53		
51	48	29	25	44	23	36	58		
28	51	40	43	21	48	17	23		
28	44	28	21	31	22	56	60		
15	21	39	26	28	23	21	20		
43	39	40	26	50	17	17	23		
44	30	35	20	11	18	39	27		
28	30	34	33	30	29	46	36		
58	28	30	28	37	31	29	32		
48	49	30							

and p_1, \dots, p_k are mixing probabilities with $\sum_{j=1}^k p_j = 1$.

The parameters $\mathbf{p}' = (p_1, \dots, p_k)$, $\boldsymbol{\mu}' = (\mu_1, \dots, \mu_k)$, and $\boldsymbol{\sigma}' = (\sigma_1, \dots, \sigma_k)$ are usually estimated by maximum likelihood. Standard errors can be obtained in the usual way from the observed information matrix (i.e., from the inverse of the Hessian matrix, the matrix of second derivatives of the log likelihood). Determining the number of components k in the mixture is more problematic since the conventional likelihood ratio test cannot be used to compare models with different k , a point we will return to later.

For a short introduction to finite mixture modeling, see Everitt (1990); a more comprehensive account is given in McLachlan and Peel (2000). Maximum likelihood estimation using Stata is described in detail by Gould *et al.* (2006).

13.3 Analysis using Stata

Stata has a command called `m1`, which can be used to maximize a user-specified log likelihood using the Newton-Raphson algorithm. The algorithm is iterative. Starting with initial parameter values, the program evaluates the first and second derivatives of the log likelihood at the parameter values to find a new set of parameter values where the likelihood is likely to be greater. The derivatives are then evaluated for the new parameters to update the parameters again, etc., until the maximum has been found (where the first derivatives are zero and the second derivatives negative). The EM algorithm, an alternative to Newton-Raphson, is often believed to be superior for finite mixture

models. However, in our experience the implementation of Newton-Raphson in Stata works very well for these models.

To use `ml`, the user must write a program that evaluates the log likelihood and possibly its derivatives. The `ml` command provides four methods: `d0`, `d1`, `d2`, and `lf`. The `d0` method does not require the user's program to evaluate any derivatives of the log likelihood, `d1` requires first derivatives only, and `d2` requires both first and second derivatives. When `d0` is chosen, first and second derivatives are found numerically; this makes this alternative slower and less accurate than `d1` or `d2`.

The simplest approach to use is `lf` which also does not require any derivatives to be programmed. Instead, the structure of most likelihood problems is used to increase both the speed and the accuracy of the numerical differentiation. Whereas `d0` to `d2` can be used for any maximum likelihood problem, method `lf` can only be used if the likelihood satisfies the following two criteria:

1. The units of observations in the dataset are independent, i.e., the log likelihood is the sum of the log likelihood contributions of the units.
2. The log likelihood contributions have a *linear form*, i.e., they are (not necessarily linear) functions of linear predictors of the form $\eta_i = x_{1i}\beta_1 + \cdots + x_{ki}\beta_k$.

The first restriction is usually met, an exception being longitudinal or multilevel data. The second restriction is not as severe as it appears because there may be several linear predictors, as we shall see later. These restrictions allow `lf` to evaluate derivatives efficiently and accurately using chain rules. For example, first derivatives are obtained as

$$\frac{\partial \ell_i}{\partial \beta_1} = \frac{\partial \ell_i}{\partial \eta_i} x_{1i},$$

where ℓ_i is the log likelihood contribution from the i th observation. All that is required are the derivatives with respect to the linear predictor(s) from which the derivatives with respect to the individual parameters follow automatically.

In this chapter, we will give only a brief introduction to maximum likelihood estimation using Stata, restricting our examples to the `lf` method. We recommend the book on *Maximum Likelihood Estimation with Stata* by Gould *et al.* (2006) for a thorough treatment of the topic.

We will eventually fit a mixture of normals to the age of onset data, but will introduce the `ml` procedure by a series of simpler models.

13.3.1 Single normal density

To begin with, we will fit a normal distribution with standard deviation fixed at 1 to a set of simulated data. A (pseudo)-random sample from the normal distribution with mean 5 and standard deviation 1 can be obtained using the instructions

```
clear
set obs 100
set seed 12345678
generate y = invnormal(uniform()) + 5
```

where the purpose of the `set seed` command is simply to ensure that the same data will be generated each time we repeat this sequence of commands. We use `summarize` to confirm that the sample has a mean close to 5 and a standard deviation close to 1.

Variable	Obs	Mean		Std. Dev.	Min	Max
y	100	5.002311		1.053095	2.112869	7.351898

First, we will define a program, `mixing0`, to evaluate the log likelihood contributions when called from `ml`. The program must have two arguments; the first is the variable name where `ml` will look for the computed log likelihood contributions; the second is the variable name containing the "current" value of the linear predictor during the iterative maximization procedure:

```
capture program drop mixing0
program mixing0
    version 9.2
    args lj xb

    tempname s

    scalar `s' = 1
    quietly replace `lj' = ln(normalden($ML_y1,"xb","s"))
end
```

After giving names `lj` and `xb` to the arguments passed to `mixing0` by `ml`, `mixing0` defines a temporary name stored in the local macro `s` and subsequently defines a scalar having that name and taking the value 1. This scalar represents the standard deviation used in the next command. Using temporary names avoids any confusion with variables that may exist in the dataset. The final command returns the log of the normal density in '`lj`' as required by the calling program. Here we used the `normalden(y, μ, σ)` function to calculate the normal density in equation (13.1) where `y` is the dependent variable whose

name is stored in the global macro `ML_y1` by `ml`. The mean μ is just the linear predictor '`xb`', and the standard deviation σ is the scalar '`s`'.

Note that all variables defined within a program to be used with `ml` should be of storage type `double` to enable `ml` to estimate accurate numerical derivatives. Scalars should be used instead of local macros to hold constants as scalars have higher precision.

The commands above can be typed into a do-file. After running the commands, define the model using the command

```
ml model lf mixing0 (xb: y=)
```

which specifies the method as `lf` and the program to evaluate the log likelihood contributions as `mixing0`. The response and explanatory variable are given by the "equation" in parentheses. Here the name before the colon, `xb`, is the name of the equation, the variable after the colon, `y`, is the response variable, and the variables after the "`=`" are the explanatory variables contributing to the linear predictor. No explanatory variables are given here, so a constant only model will be fitted.

As a result of this model definition, the global `ML_y1` will be equal to `y` and in `mixing0 "xb"` will be equal to the intercept parameter (the mean) that is going to be estimated by `ml`.

Now we maximize the log likelihood using the command

```
ml maximize, noheader
```

giving the results shown in Display 13.1. The program converged in

	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
<code>_cons</code>	5.002311	.1	50.02	0.000	4.806314 5.196307

Display 13.1

three iterations and the maximum likelihood estimate of the mean is equal to the sample mean of 5.002311. If we were interested in observing

the value of the mean parameter in each iteration, we could use the `trace` option in the `ml maximize` command. We used the `noheader` option to suppress output relating to the likelihood ratio test against a "null" model since we have not specified such a model.

We will now extend the program step by step until it can be used to estimate a mixture of two Gaussians. The first step is to allow the standard deviation to be estimated. Since this parameter does not contribute linearly to the linear predictor used to estimate the mean, we must define another linear predictor by specifying another equation with no dependent variable in the `ml model` command. Assuming that the program to evaluate the log likelihood contributions is called `mixing1`, the `ml model` command becomes:

```
ml model lf mixing1 (xb: y=) (lsd:)
```

The new equation has the name `lsd`, has no dependent variable (since `y` is the only dependent variable), and the linear predictor is simply a constant. A short-form of the above command is

```
ml model lf mixing1 (xb: y=) /lsd
```

We intend to use `lsd` as the log standard deviation. Estimating the log of the standard deviation will ensure that the standard deviation itself is positive. We will now modify the function `mixing0` so that it has an additional argument for the log standard deviation:

```
capture program drop mixing1
program mixing1
    version 9.2
    args l1j xb ls

    tempvar s

    quietly generate double `s' = exp(`ls')
    quietly replace `l1j' = ln(normalden($ML_y1, `xb', `s'))
end
```

We now define a temporary variable `s` instead of a scalar because the linear predictor "`ls`" is a variable. This is because the linear predictor is defined in the `ml model` command and could in principle contain covariates (see Exercise 13.2) and hence differ between units. The temporary variable name will not clash with any existing variable names, and the variable will automatically be deleted when the program has finished running. Running

```
ml maximize, noheader
```

gives the output shown in Display 13.2.

The standard deviation estimate is obtained by exponentiating the estimated coefficient of `.cons` in equation `lsd`. Instead of typing `display`

```

initial: log likelihood = -1397.9454
alternative: log likelihood = -534.98948
rescale: log likelihood = -301.15405
rescale eq: log likelihood = -180.56802
Iteration 0: log likelihood = -180.56802
Iteration 1: log likelihood = -147.55179
Iteration 2: log likelihood = -146.56996
Iteration 3: log likelihood = -146.5647
Iteration 4: log likelihood = -146.56469

```

	<i>t</i>	Coeff.	Std. Err.	<i>t</i>	P> t	[95% Conf. Interval]
sb						
_cons	5.002311	.1047886	47.74	0.000	4.796943	.5207679
lsd						
_cons	.0467082	.0707107	0.66	0.509	-.0918821	.1952946

Display 13.2

$\exp(0.0467082)$, we can use the following syntax for accessing coefficients and their standard errors:

```

display [lsd]_b[_cons]
.04670823

```

```

display [lsd]_se[_cons]
.07071068

```

We can also omit ".b" from the first expression, and compute the required standard deviation using

```

display exp([lsd]_[_cons])
.0478163

```

This is smaller than the sample standard deviation from `summarize` because the maximum likelihood estimate of the standard deviation is given by

$$\hat{\sigma} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2}, \quad (13.2)$$

where n is the sample size, whereas the factor $\frac{1}{n-1}$ is used in `summarize`. Since n is 100 in this case, the maximum likelihood estimate must be "blown up" by a factor of $\sqrt{100/99}$ to obtain the sample standard deviation:

```

display exp([lsd]_[_cons])*sqrt(100/99)

```

1.083095

13.3.2 Two-component mixture model

The program can now be extended to estimate a mixture of two Gaussians. To allow us to test the program on data from a known distribution, we will simulate a sample from a mixture of two Gaussians with standard deviations equal to 1 and means equal to 0 and 5 and with mixing probabilities $p_1 = p_2 = 0.5$. This can be done in two stages; first randomly allocate observations to groups (variable z) with probabilities p_1 and p_2 , and then sample from the different component densities according to group membership.

```
clear
set obs 300
set seed 12345678
generate z = cond(uniform()<0.5,1,2)
generate y = invnorm(uniform())
replace y = y + 5 if z==2
```

We now need five linear predictors, one for each parameter to be estimated: μ_1 , μ_2 , σ_1 , σ_2 , and p_1 (since $p_2 = 1 - p_1$). As before, we can ensure that the estimated standard deviations are positive by taking the exponential inside the program. The mixing proportion p_1 must lie in the range $0 \leq p_1 \leq 1$. One way of ensuring this is to interpret the linear predictor as representing the log odds (see Chapter 6) so that p_1 is obtained from the linear predictor of the log odds, $lo1$ using the transformation $1/(1+\exp(-lo1))$. The program now becomes

```
capture program drop mixing2
program mixing2
    version 9.2
    args l1 xb1 xb2 lo1 ls1 ls2

    tempvar f1 f2 p s1 s2

    quietly {
        generate double `s1' = exp(`ls1')
        generate double `s2' = exp(`ls2')
        generate double `p' = 1/(1+exp(-`lo1'))

        generate double `f1' = normalden($ML_y1,'xb1','s1')
        generate double `f2' = normalden($ML_y1,'xb2','s2')
        replace `l1' = ln(`p'*`f1' + (1-`p')*`f2')
    }
end
```

Here we have applied quietly to a whole block of commands by enclosing them in braces.

Stata simply uses zeroes as starting or initial values for all parameters. However, it is not advisable here to start with the same initial value for both component means. Therefore the starting values should be set using the `ml init` commands as follows:

```
ml model lf mixing2 (xb1: y=) /xb2 /lo1 /lsd1 /lsd2
ml init 1 6 0 0.2 -0.2, copy
ml maximize, noheader
```

In the `ml init` command, the first two values are initial values for the means, the third for the log odds, and the fourth and fifth for the logs of the standard deviations.

The results are shown in Display 13.3 where the standard deviations

	<i>β</i>	Coeff.	Std. Err.	<i>t</i>	P> <i>t</i>	[95% Conf. Intervall]
xb1 _cons	-0.0487452	.0806427	-0.57	0.571	.503893	.1123117
xb2 _cons	4.983717	.0663482	73.72	0.000	4.814476	6.182268
lo1 _cons	.1122615	.1172032	0.96	0.338	-.1174728	.3419588
lsd1 _cons	-.027468	.0627909	-0.44	0.662	-.1805164	.0955804
lsd2 _cons	-.0173444	.0663055	-0.26	0.794	-.1473008	.1126121

Display 13.3

are estimated as

```
display exp(flsd1[_cons])
.57290682
```

and

```
display exp([lsd2] [_cons])
.9928062
```

and the probability of membership in group 1 as

```
display 1/(1 + exp(-[lo1] [_cons]))
.52803094
```

The maximum likelihood estimates agree quite closely with the true parameter values.

Alternative estimates of the mixing probability and means and standard deviations, treating group membership as known (usually not possible!), are obtained using

table z, contents(freq mean y sd y)

z	Freq.	mean(y)	sd(y)
1	180	-0.0006726	1.003264
2	140	5.0122068	.9642397

and these are also similar to the maximum likelihood estimates. The maximum likelihood estimate of the proportion (0.528) is closer to the realized proportion $160/300 = 0.533$ than the "true" proportion 0.5.

The standard errors of the estimated means are given in the regression table in Display 13.3 (0.081 and 0.086). We can estimate the standard errors of the standard deviations and of the probability from the standard errors of the log standard deviations and log odds using the delta method (see for example Agresti, 2002, pages 577-581). According to the delta method, if $y = f(x)$, then approximately, $se(y) = |f'(x)|se(x)$ where $f'(x)$ is the first derivative of $f(x)$ with respect to x evaluated at the estimated value of x . For the standard deviation, $sd = \exp(lsd)$, so that, by the delta method,

$$se(sd) \approx sd \times se(lsd). \quad (13.3)$$

For the probability, $p = 1/(1 + \exp(-lo))$, so that

$$se(p) = p(1 - p) \times se(lo). \quad (13.4)$$

However, an even easier way of obtaining and displaying a function of coefficients with the correct standard error in Stata is using the **nlcom** command:

```
nlcom (sd1: exp([lsd1] [_cons])) (sd2: exp([lsd2] [_cons])) ///
(p: invlogit([lo1] [_cons]))
```

```

sd1: exp([lsd1]_cons)
sd2: exp([lsd2]_cons)
p: invlogit([lci]_cons)

```

T	Coeff.	Std. Err.	t	P> t	(95% Conf. Interval)
sd1	.9729058	.0610799	15.93	0.000	.8631913 1.09382
sd2	.9928062	.0651654	15.08	0.000	.8550833 1.110527
p	.8280309	.0292087	18.08	0.000	.4707829 .588279

Display 13.4

In each set of parentheses, we specify a label, followed by a colon and then an expression defining the function of stored estimates we are interested in. Here we used the `invlogit()` function to obtain the probability from the log odds. Stata then uses numerical derivatives to work out the correct standard error using the delta-method giving the results shown in Display 13.4. The z-statistic, p-value, and confidence interval should be ignored unless it is reasonable to assume a normal sampling distribution for the derived parameter.

We now apply the same program to the age of onset data. The data can be read in using

```

infile y using onset.dat, clear
label variable y "age of onset of schizophrenia"

```

A useful graphical display of the data is a histogram produced using

```
histogram y, bin(12)
```

which is shown in Figure 13.1.

It seems reasonable to use initial values of 20 and 45 for the two means. In addition, we will use a mixing proportion of 0.5 and log standard deviations of 2 as initial values.

```

ml model lf mixing2 (y: y=) /xb2 /lci /lsd1 /lsd2
ml init 20 45 0 2 2, copy
ml maximize, noheader

```

The output is given in Display 13.5. The means are estimated as 24.8 and 46.4, the standard deviations as 6.5 and 7.1, and the mixing proportions as 0.74 and 0.26, for groups 1 and 2, respectively. The approximate standard errors may be obtained as before; see Exercise 13.1.

We will now plot the estimated mixture density together with a kernel density estimate. Instead of using the command `kdensity`, we will use `twoway kdensity`, allowing us to add the mixture density onto the same graph:

```

initial: log likelihood = -391.61148
rescale: log likelihood = -391.61148
rescale eq: log likelihood = -391.61148
Iteration 0: log likelihood = -391.61148 (not concave)
Iteration 1: log likelihood = -374.66715 (not concave)
Iteration 2: log likelihood = -374.36498
Iteration 3: log likelihood = -373.94187
Iteration 4: log likelihood = -373.67092
Iteration 5: log likelihood = -373.66896
Iteration 6: log likelihood = -373.66896

```

	<i>B</i>	Coef.	Std. Err.	<i>t</i>	P> <i>t</i>	(95% Conf. Interval)
zbl	_cons	24.79772	1.133545	21.88	0.000	22.57801 27.01942
zb2	_cons	46.44685	2.740866	16.95	0.000	41.07485 51.81888
lei	_cons	1.034415	.3697502	2.80	0.005	.3097179 1.759112
indi	_cons	1.877886	.1261092	14.89	0.000	1.630629 2.124866
led2	_cons	1.955009	.256892	7.61	0.000	1.451455 2.458543

Display 13.5

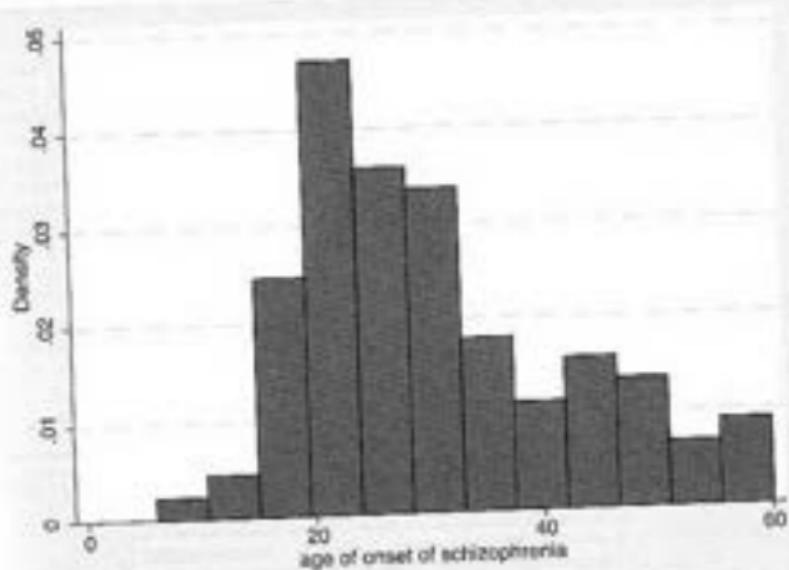


Figure 13.1: Histogram of age of onset of schizophrenia in women.

```

graph twoway (kdensity y, width(3))           ///
(function invlogit([lo1],_cons)               ///
* normalden(x,[xb1],_cons,exp([lad1],_cons))  ///
+ (1-invlogit([lo1],_cons))                   ///
* normalden(x,[xb2],_cons,exp([lad2],_cons)),   ///
range(y) lpatt(dash),                         ///
xtitle("Age") ytitle("Density")               ///
legend(order(1 "Kernel density" 2 "Mixture model"))

```

In Figure 13.2 the two estimates of the density are surprisingly similar. (Admittedly, we did specify `width(3)` for the half-width of the kernel because this gave a good fit!)

The histogram and kernel density estimate suggest that there are two subpopulations. To test this more formally, we could also fit a single normal distribution and compare the likelihoods. However, as mentioned earlier, the conventional likelihood ratio test is not valid here. Wolfe (1971) suggests, on the basis of a limited simulation study, that the difference in minus twice the log likelihood for a model with k components compared with a model with $k + 1$ components has approximately a χ^2 distribution with $2\nu - 2$ degrees of freedom: here ν is the number of extra parameters in the $k + 1$ component mixture. The log likelihood of the current model may be accessed using `e(l1)`. We

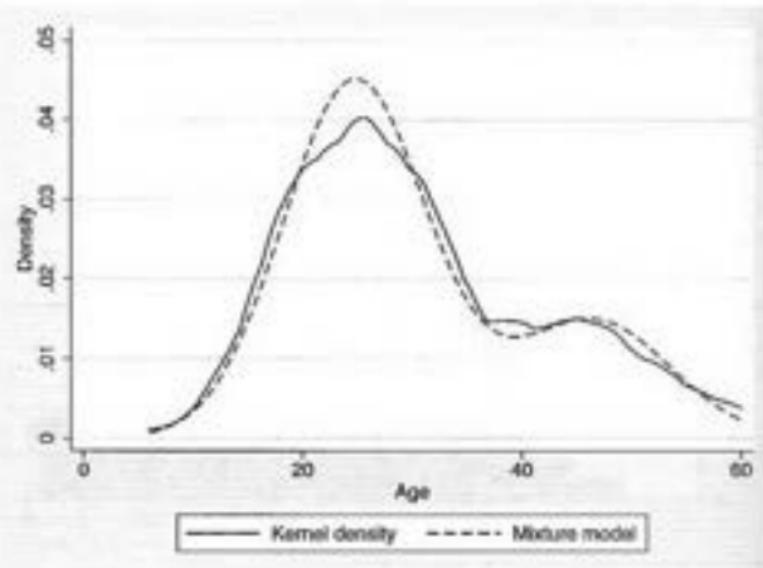


Figure 13.2: Kernel and mixture model densities for the age of onset data.

store this in a local macro

```
local ll = e(ll)
```

and fit the single normal model using the program `mixing1` as follows:

```
m1 model lf mixing1 (xb: y) /lsd
m1 init 30 1.9, copy
m1 maximize, nobheader
```

with the result shown in Display 13.6. Comparing the log likelihoods using the method proposed by Wolfe,

```
local chi2 = 2*(`ll'-e(ll))
display chi2tail(4,`chi2')
.00060994
```

confirms that there appear to be two subpopulations.

```

initial: log likelihood = -429.14924
rescale: log likelihood = -429.14924
rescale sq: log likelihood = -429.14924
Iteration 0: log likelihood = -429.14924
Iteration 1: log likelihood = -383.94844
Iteration 2: log likelihood = -383.4
Iteration 3: log likelihood = -383.39689
Iteration 4: log likelihood = -383.39689

```

y	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
rb_cons	30.47478	.1680485	186.07	0.000	28.18346 32.76603
led_cons	2.463747	.0710669	34.53	0.000	2.314458 2.593035

Display 13.6

13.4 Exercises

13.1 • Two-component mixture of normals

1. Create a do-file with the commands necessary to fit the two-component mixture of normals discussed in this chapter.
2. Add commands to the end of the do-file to calculate the standard deviations and mixing probability and the standard errors of these parameters. What are the standard errors of the estimates for the age of onset data?
3. Simulate values from two normals trying out different values for the various parameters. Do the estimated values tend to lie within two estimated standard errors from the "true" values?

13.2 • Heteroscedastic linear regression

1. Use the program `mixing1` to fit a linear regression model to the slimming data from Chapter 5 with `status` as the only explanatory variable.
2. Compare the estimated residual standard deviation with the root mean squared error obtained using the `regress` command.
3. Use the same program again to fit a linear regression model where the variance is allowed to differ between the groups defined by `status`. (Hint: Modify the `(led:)` equation in the `ml model` command; see page 268). Is there any evidence

for heteroscedasticity? How do the results compare with those of `sdtest`?

13.3 • Three-component mixture of normals

1. Extend the program `mixing2` to fit a mixture of three normals and test this on simulated data (hint: use transformations $p1 = 1/d$, $p2 = \exp(\text{lo1})/d$, and $p3 = \exp(\text{lo2})/d$ where $d = 1 + \exp(\text{lo1}) + \exp(\text{lo2})$).

13.4 Two-component mixture of Poisson distributions

Hasselblad (1969) fitted a two-component mixture of Poisson distributions to the number of death notices of women aged 80 and over published in the *New York Times* between 1910 and 1912. This classic dataset is tabulated below.

Number of notices	0	1	2	3	4	5	6	7	8	9
Frequency	162	267	271	185	111	61	27	8	3	1

The model can be written as

$$\begin{aligned} \Pr(Y=y) &= p_1 \exp(y \ln(\mu_1) - \mu_1)/y! \\ &\quad + (1-p_1) \exp(y \ln(\mu_2) - \mu_2)/y!, \end{aligned}$$

where μ_1 and μ_2 are the means for the two components and p_1 is the probability of belonging to the first component. Hasselblad obtained the estimates $\hat{p}_1 = 0.3599$, $\hat{\mu}_1 = 1.2561$, and $\hat{\mu}_2 = 2.6634$. This solution was interpreted as indicating a different pattern of deaths in the winter (component 1) and summer (component 2).

1. Write a program to evaluate the log-likelihood contribution for the `m1` command with method `lf`. Note that the model is defined only for $\mu_1 > 0$ and $\mu_2 > 0$. You can use the function `lnfactorial()` to evaluate $\ln(y!)$.
2. Enter the data. (Hint: Use the `expand` command.)
3. Fit the model. Any discrepancies between your results and those given above are likely due to programming errors. You can rule out that it is due to poor starting values by using values close to the required results. Revise your program if necessary.

4. Calculate the expected frequencies for each number of notices and compare them graphically with the observed frequencies. (Hist: You can use the program that evaluates the log-likelihood contributions to calculate the log probabilities for different values of y by defining the global ML.y1 appropriately and passing the name of an existing variable and the estimated parameters to the program as arguments.)

13.5 Latent class model

Van der Heijden *et al.* (1992) analyze data collected by the Netherlands Ministry of Justice to investigate differences in involvement in crime among young people from four ethnic groups: Moroccans, Turks, Surinamese, and Dutch. The Dutch sample consisted of people who lived in the same streets as the people from the other ethnic groups. Three crime measures were gathered from police records: Property crime, aggression against persons, and vandalism.

The variables in `crime.dta` are:

- `vandalism`: indicator for being arrested for vandalism
- `aggress`: indicator for aggression against a person
- `property`: indicator for being arrested for property crime
- `ethnicity`: ethnic group (1=Moroccan, 2=Turkish, 3=Surinamese, 4=Dutch)
- `age`: age group (1=12-13, 2=14-15, 3=16-17)

Assuming that there are subpopulations differing in their pattern of responses on the three delinquency items, we will consider a finite mixture model with two components for simplicity. What is different between this model and the other finite mixture models discussed in this chapter is that there are three responses per subject. Assuming that responses are independent within the latent classes (an assumption known as conditional or local independence), the model can be written as

$$\Pr(y_i) = p_1 \prod_{c=1}^3 \frac{\exp(\alpha_{ci}y_i)}{1 + \exp(\alpha_{ci})} + (1 - p_1) \prod_{c=1}^3 \frac{\exp(\alpha_{ci}y_i)}{1 + \exp(\alpha_{ci})}$$

where α_{ci} is the log odds that the i th response y_i equals one for a person in latent class c . Here y_i multiplies α_{ci} in the numerator to produce a numerator equal to $\exp(\alpha_{ci})$ if $y_i = 1$ and equal to 1 if $y_i = 0$ as required.

1. Write a program to evaluate the log-likelihood contributions

for `ml`. Note that three response variables must be specified in the `ml model`. You can simply use equations of the form `(name1: vandalis=)`, `(name1: aggressi=)`, and `(name1: property=)` to specify the three response variables and one linear predictor for each, and then use equations without response variables for all the remaining linear predictors. The names of the response variables will be stored in the globals `ML_y1`, `ML_y2`, and `ML_y3`.

To test the program, create a variable `junk` equal to 1 and specify `junk` as the variable name for the log-likelihood contributions, passing values of zero for all linear predictors. The result should be that `junk` equals $\ln(.5^3) = -2.0794415$ for all observations.

2. Estimate the model. If non-convergence occurs, change the starting values and use the option `search(norescale)` in the `ml maximize` command. Assuming that one population will be more delinquent than the other, reasonable starting values are negative values for α_1 , and positive values for α_2 (or vice versa).
3. Interpret the estimates.
4. Modify the `ml model` command to allow the latent class membership probability p_i to depend on the covariates `ethnicity` and `age` via a logistic regression model.
5. Interpret the estimates.

Chapter 14

Principal Components Analysis: Hearing Measurement Using an Audiometer

14.1 Description of data

The data in Table 14.1 are adapted from those given in Jackson (1991), and relate to hearing measurements with an instrument called an audiometer. An individual is exposed to a signal of a given frequency with an increasing intensity until the signal is perceived. The lowest intensity at which the signal is perceived is a measure of hearing loss, calibrated in units referred to as decibel loss in comparison with a reference standard for that particular instrument. Observations are obtained one ear at a time for a number of frequencies. In this example, the frequencies used were 500 Hz, 1000 Hz, 2000 Hz, and 4000 Hz. The limits of the instrument are -10 to 99 decibels. (A negative value does not imply better than average hearing; the audiometer had a calibration "zero", and these observations are in relation to that.)

Table 14.1 Data in hear.dat (taken from Jackson (1991) with permission of his publisher, John Wiley & Sons)

id	500	1000	1500	2000	500	1000	2000	4000
1	-9	5	10	15	0	5	5	15
2	-5	0	-10	0	0	5	5	15
3	-5	0	15	15	0	0	5	15
4	-5	0	-10	-10	-10	-5	-10	10

Table 14.1 Data in hear.dat (continued)

5	-5	-5	-10	10	0	-10	-10	50
6	5	5	5	-10	0	5	0	20
7	0	0	0	20	5	5	5	10
8	-10	-10	-10	-5	-10	-5	0	5
9	0	0	0	40	0	0	-10	10
10	-5	-5	-10	20	-10	-5	-10	15
11	-10	-5	-5	5	5	0	-10	5
12	5	5	20	25	-5	-5	5	15
13	0	0	-10	15	-10	-10	-10	20
14	5	15	5	60	5	5	0	50
15	5	0	5	15	5	-5	0	25
16	-5	-5	5	30	5	5	5	25
17	0	-10	0	20	0	-10	-10	25
18	5	0	0	50	10	10	5	65
19	-10	0	0	15	-10	-5	5	15
20	-10	-10	-5	0	-10	-5	-5	5
21	-5	-5	-5	25	-5	-5	-10	20
22	5	15	5	20	5	5	5	25
23	-10	-10	-10	25	-5	-10	-10	25
24	-10	0	5	15	-10	-5	5	20
25	0	0	0	20	-5	-5	10	30
26	-10	-5	0	15	0	0	0	10
27	0	0	5	50	5	0	5	40
28	-5	-5	-5	25	-5	5	5	20
29	0	15	0	20	10	-5	0	20
30	-10	-5	0	15	-5	0	10	20
31	-10	-10	5	10	0	0	20	20
32	-5	5	20	25	-5	0	5	20
33	0	5	0	10	-10	0	0	0
34	-10	-10	-10	45	-10	-10	5	45
35	-5	10	30	45	-5	10	35	60
36	-5	-5	-5	30	-5	0	10	60
37	-10	-5	-5	45	-10	-5	-5	50
38	5	5	5	25	-5	-5	5	40
39	-10	-10	-10	6	-10	-10	-10	5
40	10	20	15	19	25	20	10	20
41	-10	-10	-10	20	-10	-10	0	5
42	5	1	-5	40	5	10	0	45
43	-10	0	10	20	-10	0	15	10
44	-10	-10	10	10	-10	-10	5	0
45	-5	-5	-10	35	-5	0	-10	10
46	5	5	10	25	10	5	5	20
47	5	0	10	70	-5	5	15	40
48	5	10	0	15	5	10	0	20
49	-5	-5	5	-10	-10	-5	0	20
50	-5	0	20	50	-10	0	5	20
51	-10	-10	-10	5	-10	-10	-5	5
52	5	10	20	25	0	5	15	0
53	-10	-10	50	25	-10	-10	-10	40
54	5	10	0	-10	0	5	-5	15
55	15	20	10	60	20	20	0	25
56	-10	-10	-10	5	-10	-10	-5	10
57	-5	-5	-10	30	0	-5	-10	15
58	-5	-5	0	5	-5	-5	0	10
59	-5	5	5	40	0	0	0	10
60	5	10	30	20	5	5	30	60
61	5	5	0	10	-5	5	0	10
62	0	5	10	35	0	0	5	20
63	-10	-10	-10	0	-5	0	-5	0
64	-10	-5	-5	20	-10	-10	-5	5
65	5	10	0	25	5	5	0	15
66	-10	0	5	60	-10	-5	0	65
67	5	10	40	55	0	5	20	40
68	-5	-10	-10	20	-5	-10	-10	15
69	-5	-5	-5	20	-5	0	0	0
70	-5	-5	-5	5	-5	0	0	5

Table 14.1 Data in bear.dat (continued)

71	0	10	40	60	-5	0	-20	50
72	-5	-5	-5	-5	-5	-5	-5	5
73	0	5	45	50	0	10	15	50
74	-5	-5	10	25	-10	-5	25	60
75	0	-10	0	60	15	0	5	50
76	-5	0	10	35	-10	0	0	15
77	5	0	0	15	0	5	5	25
78	15	15	5	35	10	15	-5	0
79	-10	-10	-10	5	-5	-5	-5	5
80	-10	-10	-5	15	-10	-10	-5	5
81	0	-5	5	35	-5	-5	5	15
82	-5	-5	-5	10	-5	-5	-5	5
83	-5	-5	-10	-10	0	-5	-10	0
84	5	10	10	20	-5	0	0	10
85	-10	-10	-10	5	-10	-5	-10	20
86	5	5	10	0	0	5	5	5
87	-10	0	-5	-10	-10	0	0	-10
88	-10	-10	10	15	0	0	5	15
89	-5	0	10	25	-5	0	5	10
90	5	0	-10	-35	10	0	0	0
91	0	0	5	15	5	0	0	5
92	-5	0	-5	0	-5	-5	-10	0
93	-5	5	-10	45	-5	0	-5	25
94	-10	-5	0	20	-10	5	-10	10
95	-10	-5	0	5	-10	-5	-5	5
96	5	0	5	0	5	0	5	15
97	-10	-20	5	40	-10	-5	-10	5
98	10	20	15	55	0	0	5	75
99	-5	5	5	20	-5	5	5	40
100	-5	-5	-10	-10	-5	0	15	10

14.2 Principal component analysis

Principal component analysis is one of the oldest but still most widely used techniques of multivariate analysis. Originally introduced by Pearson (1901) and independently by Hotelling (1933), the basic idea of the method is to try to describe the variation of the variables in a set of multivariate data as parsimoniously as possible using a set of derived uncorrelated variables, each of which is a particular linear combination of those in the original data. In other words, principal component analysis is a transformation from the observed variables, y_{11}, \dots, y_{pi} to new variables z_{1i}, \dots, z_{pi} where

$$\begin{aligned} z_{1i} &= a_{11}y_{1i} + a_{12}y_{2i} + \dots + a_{1p}y_{pi} \\ z_{2i} &= a_{21}y_{1i} + a_{22}y_{2i} + \dots + a_{2p}y_{pi} \\ &\vdots = \vdots + \vdots + \vdots + \vdots + \vdots \\ z_{pi} &= a_{p1}y_{1i} + a_{p2}y_{2i} + \dots + a_{pp}y_{pi}. \end{aligned} \tag{14.1}$$

The new variables are derived in decreasing order of importance. The coefficients a_{11} to a_{1p} for the first principal component are derived so that the sample variance of g_{1i} is as large as possible. Since this variance could be increased indefinitely by simply increasing the co-

efficients, a restriction must be placed on them, generally that their sum of squares is one. The coefficients defining the second principal component y_{2i} are determined to maximize its sample variance subject to the constraint that the sum of squared coefficients equals 1 and that the sample correlation between y_{1i} and y_{2i} is 0. The other principal components are defined similarly by requiring that they are uncorrelated with all previous principal components. It can be shown that the required coefficients are given by the eigenvectors of the sample covariance matrix of y_{11}, \dots, y_{pn} , and their variances are given by the corresponding eigenvalues. In practice components are often derived from the correlation matrix instead of the covariance matrix, particularly if the variables have very different scales. The analysis is then equivalent to calculation of the components from the original variables after these have been standardized to unit variance.

The usual objective of this type of analysis is to assess whether the first few components account for a substantial proportion of the variation in the data. If they do, they can be used to summarize the data with little loss of information. This may be useful for obtaining graphical displays of the multivariate data or for simplifying subsequent analysis. The principal components can be interpreted by inspecting the eigenvectors defining them. Here it is often useful to multiply the elements by the square root of the corresponding eigenvalue in which case the coefficients represent correlations between an observed variable and a component. A detailed account of principal component analysis is given in Everitt and Dunn (2001).

14.3 Analysis using Stata

The data can be read in from an ASCII file *bear.dat* as follows:

```
infile id 1500 11000 12000 14000 r500 r1000 r2000 r4000 ///
      using bear.dat, clear
summarize
```

(see Display 14.1).

Before undertaking a principal component analysis, some graphical exploration of the data may be useful. A scatterplot matrix, for example, with points labeled with a subject's identification number can be obtained using

```
graph matrix 1500-r4000, mlabel(id) msymbol(none) ///
      mlabposition(0) half
```

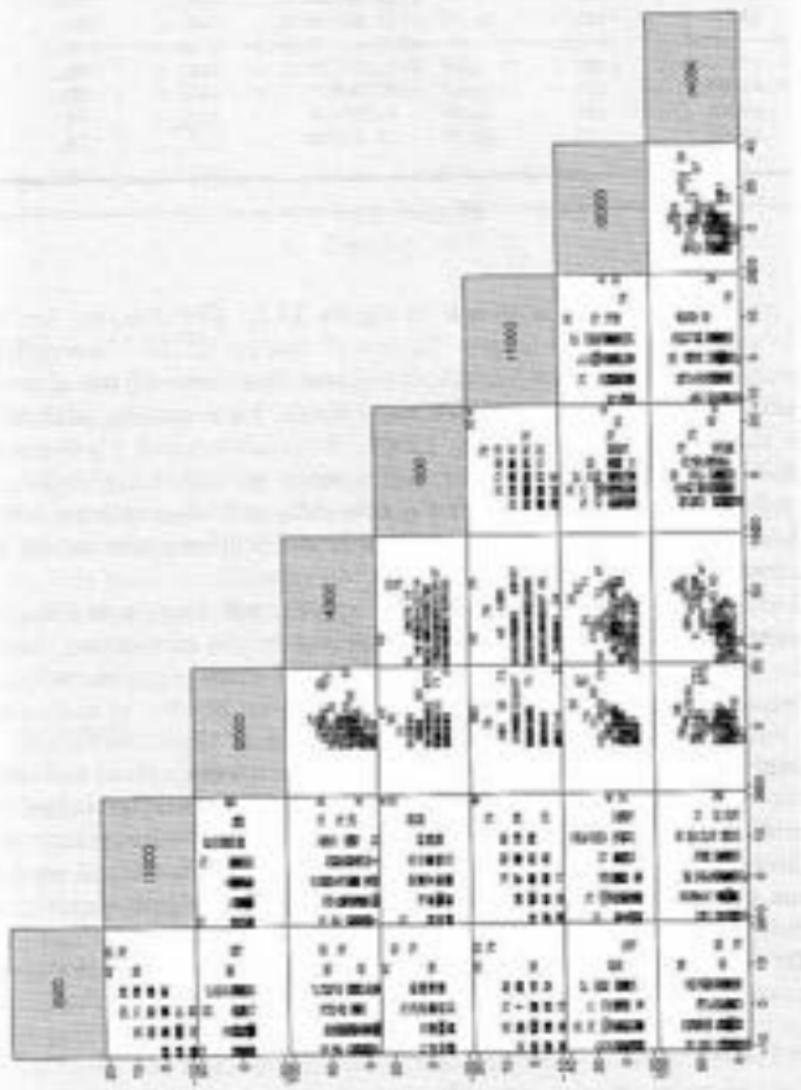


Figure 14.1: Scatterplot matrix of hearing loss at different frequencies for the left and right ear.

Variable	Obs	Mean	Std. Dev.	Min	Max
ld	100	50.5	29.01148	-1	100
l500	100	-2.6	6.406843	-10	13
l1000	100	-.5	7.871231	-10	29
l2000	100	2.45	11.04463	-10	50
l4000	100	21.35	19.61569	-10	70
r500	100	-2.6	7.123796	-10	25
r1000	100	-.7	6.396811	-10	20
r2000	100	1.85	9.257675	-10	35
r4000	100	30.95	19.43264	-10	75

Display 14.1

The resulting diagram is shown in Figure 14.1. The diagram looks a little "odd" due to the largely "discrete" nature of the observations. Some of the individual scatterplots suggest that some of the observations might perhaps be regarded as outliers; for example, individual 53 in the plot involving 12000, r2000. This subject has a score of 50 at this frequency in the left ear, but a score of -10 in the right ear. It might be appropriate to remove this subject's observations before further analysis, but we shall not do this and will continue to use the data from all 100 individuals.

As mentioned in the previous section, principal components may be extracted from either the covariance matrix or the correlation matrix of the original variables. A choice needs to be made since there is not necessarily any simple relationship between the results in each case. The summary table shows that the variances of the observations at the highest frequencies are approximately nine times those at the lower frequencies; consequently, a principal component analysis using the covariance matrix would be dominated by the 4000 Hz frequency. But this frequency is not more clinically important than the others, and so, in this case, it seems more reasonable to use the correlation matrix as the basis of the principal component analysis.

To find the correlation matrix of the data requires the following instruction:

```
correlate l500-r4000
```

and the result is given in Display 14.2. Note that the highest correlations occur between adjacent frequencies on the same ear and between corresponding frequencies on different ears.

The *pca* command can be used to obtain the principal components of this correlation matrix:

	1500	11000	12000	14000	r500	r1000	r2000
1500	1.0000						
11000	0.7775	1.0000					
12000	0.3247	0.4437	1.0000				
14000	0.2684	0.2749	0.3984	1.0000			
r500	0.6963	0.5815	0.1795	0.1790	1.0000		
r1000	0.6416	0.7070	0.3632	0.2632	0.6634	1.0000	
r2000	0.2399	0.3606	0.5910	0.3193	0.1575	0.4181	1.0000
r4000	0.2264	0.2109	0.3596	0.6783	0.1421	0.2249	0.4044
		r4000					
r4000		1.0000					

Display 14.2

pca 1500-r4000

which gives the results shown in Display 14.3. (The principal components of the covariance matrix can be obtained using the covariance option.)

An informal rule for choosing the number of components to represent a set of correlations is to use only those components with eigenvalues greater than one, i.e., those with variances greater than the average. Here, this leads to retaining only the first two components. Another informal indicator of the appropriate number of components is the *scree plot*, a plot of the eigenvalues against their rank. A scree plot may be obtained using

screeplot

with the result shown in Figure 14.2. The number of eigenvalues above a distinct "elbow" in the scree plot is usually taken as the number of principal components to select. From Figure 14.2, this would again appear to be two. The first two components account for 68% of the variance in the data.

Examining the eigenvectors defining the first two principal components, we see that the first accounting for 48% of the variance has coefficients that are all positive and all approximately the same size. This principal component essentially represents the overall hearing loss of a subject and implies that individuals suffering hearing loss at certain frequencies will be more likely to suffer this loss at other frequencies as well. The second component, accounting for 20% of the variance, contrasts high frequencies (2000 Hz and 4000 Hz) and low frequencies (500 Hz and 1000 Hz). It is well known in the case of normal hearing that hearing loss as a function of age is first noticeable in the higher frequencies.

Principal components/correlations		Number of obs.	=	100		
		Number of comp.	=	8		
		Trace	=	8		
		Rho	=	1.0000		
Rotation: (unrotated = principal)						
Component	Eigenvalue	Difference	Proportion	Cumulative		
Comp1	3.82375	2.18915	0.4780	0.4780		
Comp2	1.63459	.725852	0.2043	0.6823		
Comp3	.999042	.409528	0.1138	0.7959		
Comp4	.499514	.122081	0.0624	0.8584		
Comp5	.377433	.0383341	0.0472	0.9055		
Comp6	.339098	.07930671	0.0424	0.9479		
Comp7	.261011	.105451	0.0326	0.9805		
Comp8	.155561	,	0.0194	1.0000		
Principal components (eigenvectors)						
Variable	Comp1	Comp2	Comp3	Comp4	Comp5	Comp6
1800	0.4091	-0.3126	0.1359	-0.2722	-0.1665	0.4153
11000	0.4242	-0.2301	-0.0933	-0.3628	-0.4990	-0.0847
12000	0.3271	0.3007	-0.4777	-0.4872	0.5033	0.0904
14000	0.2850	0.4468	0.4711	-0.1798	0.0990	-0.5129
r500	0.3611	-0.3874	0.2394	0.3045	0.6283	0.1775
r1000	0.4160	-0.2367	-0.0588	0.3645	-0.6861	-0.5486
r2000	0.3090	0.3228	-0.5384	0.5169	-0.5623	0.1255
r4000	0.2696	0.4972	0.4150	0.1976	-0.1787	0.4289
Variable	Comp7	Comp8	Unexplained			
1800	0.2628	-0.6006	0			
11000	-0.0292	0.6133	0			
12000	-0.2793	-0.0640	0			
14000	0.4354	-0.0298	0			
r500	0.1275	0.3660	0			
r1000	-0.4618	-0.3428	0			
r2000	0.4476	0.0293	0			
r4000	-0.6709	0.0747	0			

Display 14.3

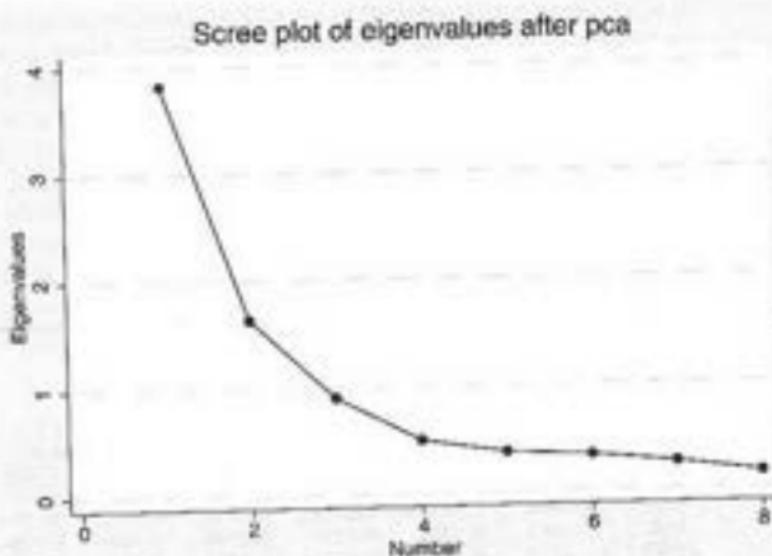


Figure 14.2: Scree plot.

Scores for each individual on the first two principal components might be used as a convenient way of summarizing the original eight-dimensional data. Such scores are obtained by applying the elements of the corresponding eigenvector to the standardized values of the original observations for an individual. The necessary calculations can be carried out using the `predict` command with the `score` option:

```
predict pc1 pc2, score
```

(see Display 14.4).

The new variables `pc1` and `pc2` contain the scores for the first two principal components, and the output lists the coefficients used to form these scores. For principal component analysis, these coefficients are just the elements of the eigenvectors in Display 14.3. The principal component scores can be used to produce a useful graphical display of the data in a single scatterplot, which may then be used to search for structure or patterns in the data, particularly the presence of clusters of observations (see Everitt et al., 2001). Such a principal component plot is obtained using

```
twoway scatter pc2 pc1, xlabel(id)
```

(Note the `scoreplot` command can be used to produce the same graph without first storing scores as new variables in the dataset). The result-

(6 components skipped)

Scoring coefficients

sum of squares(column-loading) = 1

Variable	Comp1	Comp2	Comp3	Comp4	Comp5	Comp6
z500	0.4091	-0.3126	0.1359	-0.2722	-0.1665	0.4168
z1000	0.4242	-0.2301	-0.0933	-0.3528	-0.4998	-0.0847
z2000	0.3271	0.3007	-0.4777	-0.4872	0.5033	0.0404
z4000	0.2880	0.4688	0.4711	-0.1796	0.0990	-0.5129
r500	0.3611	-0.3874	0.2394	0.3045	0.6283	0.1776
r1000	0.4160	-0.2367	-0.0868	0.3645	-0.0861	-0.5446
r2000	0.3090	0.3238	-0.5384	0.5169	-0.1623	0.1255
r4000	0.2696	0.4972	0.4150	0.1976	-0.1757	0.4589

Variable	Comp7	Comp8
z500	0.2828	-0.6005
z1000	-0.0292	0.6133
z2000	-0.2793	-0.0640
z4000	0.4354	-0.0298
r500	0.1275	0.3660
r1000	-0.6813	-0.3428
r2000	0.4476	0.0293
r4000	-0.4709	0.0747

Display 14.4

ing diagram is shown in Figure 14.3. Here, the variability in differential hearing loss for high versus low frequencies (pc2) is greater among subjects with higher overall hearing loss, as would be expected.

Note that the distances between observations in this graph approximate the Euclidean distances between the (standardized) variables, i.e., the graph is a multidimensional scaling solution. In fact, the graph is the classical scaling (or principal coordinate) scaling solution to the Euclidean distances (see Everitt and Dunn, 2001, or Everitt and Rabe-Hesketh, 1997). If other variables such as age were available, it would be interesting to investigate their relationship with the principal components (see Exercise 14.2).

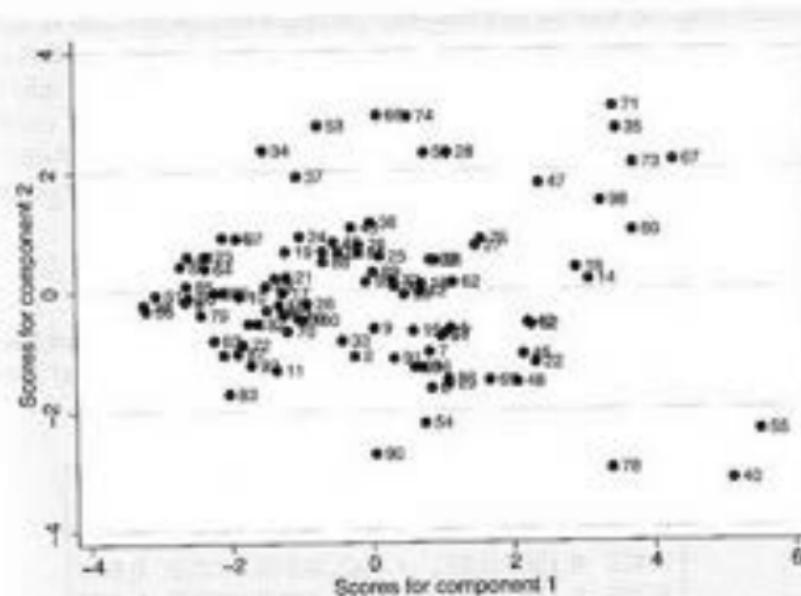


Figure 14.3: Principal component plot.

14.4 Exercises

14.1 • Hearing measurement using an audiometer

1. Rerun the principal component analysis described in this chapter using the covariance matrix of the observations. Compare

- the results with those based on the correlation matrix.
- Interpret components 3 through 8 in the principal components analysis based on the correlation matrix.
 - Create a scatterplot matrix of the first five principal component scores.

14.2 • Determinants of pollution in U.S. cities

- Apply principal component analysis to the air pollution data analyzed in Chapter 3, excluding the variable `so2`, and plot the first two principal components (i.e., the two-dimensional classical scaling solution for Euclidean distances between standardized variables).
- Regress `so2` on the first two principal components and add a line corresponding to this regression (the direction of steepest increase in `so2` predicted by the regression plane) into the multidimensional scaling solution. If the principal components are denoted p_1 and p_2 and the corresponding estimated regression coefficients $\hat{\beta}_1$ and $\hat{\beta}_2$, add the line $p_1\hat{\beta}_2/\hat{\beta}_1$ versus p_1 to the scatterplot of p_2 versus p_1 .

14.3 Characteristics of criminals

The correlation matrix below was calculated from measurements of seven physical characteristics in each of 3000 convicted criminals (MacDonnell, 1902).

1.000	0.402	0.396	0.301	0.305	0.339	0.340
0.402	1.000	0.618	0.150	0.135	0.206	0.183
0.396	0.618	1.000	0.321	0.289	0.363	0.345
0.301	0.150	0.321	1.000	0.846	0.759	0.661
0.305	0.135	0.289	0.846	1.000	0.797	0.800
0.339	0.206	0.363	0.759	0.797	1.000	0.736
0.340	0.183	0.345	0.661	0.800	0.736	1.000

The characteristics were (in the same order as for the correlation matrix) (1) Head length; (2) Head breadth; (3) Face breadth; (4) Left finger length; (5) Left forearm length; (6) Left foot length; (7) Height. The correlation matrix is contained in the ASCII file `criminals.txt`.

- Find the principal components (Hint: convert the data to a matrix using `mkmat` and then use the `pcamat` command with the `names()` option).
- Plot a scree plot and discuss how many principal components

appear to be required.

3. Interpret the principal components.

14.4 Nutritional contents of food

Hartigan (1975) provides data on the nutritional content of different foodstuffs (the quantity involved is always three ounces).

The variables in `nutrition.dta` are:

- `food`: type of food (string variable)
 - `energy`: energy content (calories) in 3 ounces of the foodstuff
 - `protein`: protein (grams) in 3 ounces of the foodstuff
 - `fat`: fat (grams) in 3 ounces of the foodstuff
 - `calcium`: calcium (milligrams) in 3 ounces of the foodstuff
 - `iron`: iron (milligrams) in 3 ounces of the foodstuff
1. Create a scatterplot matrix of the data labeling the foodstuffs appropriately in each panel. Use only the first two characters of the strings in `food` as labels.
 2. On the basis of this diagram undertake what you think is an appropriate principal components analysis.
 3. Produce a principal component plot with two-character labels for the foodstuffs.
 4. Try to interpret the first two principal components.

Chapter 15

Cluster Analysis: Tibetan Skulls and Determinants of Pollution in U.S. Cities

15.1 Description of data

The first set of data to be used in this chapter is shown in Table 15.1. These data, collected by Colonel L.A. Waddell, were first reported in Morant (1923) and are also given in Hand *et al.* (1994). The data consist of five measurements on each of 32 skulls found in the southwestern and eastern districts of Tibet. The five measurements (all in millimeters) are as follows:

- y_1 : greatest length of skull
- y_2 : greatest horizontal breadth of skull
- y_3 : height of skull
- y_4 : upper face length
- y_5 : face breadth, between outermost points of cheek bones

The main question of interest about these data is whether there is any evidence of different types or classes of skull.

The second set of data that we shall analyze in this chapter is the air pollution data introduced previously in Chapter 3 (see Table 3.1). Here we shall investigate whether the clusters of cities found are predictive of air pollution.

Table 15.1 Tibetan skull data

<i>y</i> ₁	<i>y</i> ₂	<i>y</i> ₃	<i>y</i> ₄	<i>y</i> ₅
190.5	152.5	145.0	73.5	136.5
172.5	132.0	125.5	63.0	121.0
167.0	130.0	125.5	69.5	119.5
169.5	150.5	133.5	64.5	128.0
175.0	138.5	126.0	77.5	135.5
177.5	142.5	142.5	71.5	131.0
179.5	142.5	127.5	70.5	134.5
179.5	138.0	133.5	73.5	132.5
173.5	135.5	130.5	70.0	133.5
162.5	139.0	131.0	62.0	126.0
178.5	135.0	136.0	71.0	124.0
171.5	148.5	132.5	65.0	146.5
180.5	139.0	132.0	74.5	134.5
183.0	149.0	121.5	76.5	142.0
169.5	130.0	131.0	68.0	119.0
172.0	140.0	136.0	70.5	133.5
170.0	126.5	134.5	66.0	118.5
182.5	136.0	138.5	76.0	134.0
179.5	135.0	128.5	74.0	132.0
191.0	140.5	140.5	72.5	131.5
184.5	141.5	134.5	76.5	141.5
181.0	142.0	132.5	79.0	136.5
173.5	136.5	126.0	71.5	136.5
188.5	130.0	143.0	79.5	136.0
175.0	153.0	130.0	76.5	142.0
196.0	142.5	123.5	76.0	131.0
200.0	139.5	143.5	82.5	146.0
185.0	134.5	140.0	81.5	137.0
174.5	143.5	132.5	74.0	136.5
195.5	144.0	138.5	78.5	144.0
197.0	131.5	135.0	80.5	139.0
182.5	131.0	135.0	68.5	136.0

15.2 Cluster analysis

Cluster analysis is a generic term for a set of (largely) exploratory data analysis techniques that seek to uncover groups or clusters in data. The term exploratory is important since it explains the largely absent "p-value", ubiquitous in many other areas of statistics. Clustering methods are primarily intended for generating rather than testing hypotheses. A detailed account of what is now a very large area is given in Everitt et al. (2001).

The most commonly used class of clustering methods contains those methods that lead to a series of nested or hierarchical classifications of the observations, beginning at the stage where each observation is regarded as forming a single-member "cluster" and ending at the stage where all the observations are in a single cluster. The complete hierarchy of solutions can be displayed as a tree diagram known as a dendrogram. In practice, most users will be interested not in the whole dendrogram, but in selecting a particular number of clusters that is optimal in some sense for the data. This entails "cutting" the dendrogram at some particular level.

Most hierarchical methods operate not on the raw data, but on an inter-individual distance matrix calculated from the raw data. The most commonly used distance measure is Euclidean and is defined as:

$$d_{ij} = \sqrt{(y_{1i} - y_{1j})^2 + (y_{2i} - y_{2j})^2 + \cdots + (y_{pi} - y_{pj})^2}, \quad (15.1)$$

where y_{1i} to y_{pi} are the variables for individual i .

A variety of hierarchical clustering techniques arise because of the different ways in which the distance between a cluster containing several observations and a single observation, or between two clusters, can be defined. The inter-cluster distances used by three commonly applied hierarchical clustering techniques are:

- Single linkage clustering: distance between the closest pair of observations, where one member of the pair is in the first cluster and the other in the second cluster, and
- Complete linkage clustering: distance between the most remote pair of observations where one member of the pair is in the first cluster and the other in the second cluster.
- Average linkage: average of distances between all pairs of observations where one member of the pair is in the first cluster and the other in the second cluster.

An alternative approach to clustering to that provided by the hierarchical methods described above is k -means clustering. Here the data are partitioned into a specified number of groups set by the user

by an iterative process in which, starting from an initial set of cluster means, each observation is placed into the group to whose mean vector it is closest (generally in the Euclidean sense). After each iteration, new group means are calculated and the procedure repeated until no observations change groups. The initial group means can be chosen in a variety of ways. In general, the method is applied to the data for different numbers of groups and then an attempt is made to select the number of groups that provides the best fit for the data.

Important issues that need to be considered when using clustering in practice include how to scale the variables before calculating the chosen distance matrix, which particular method of cluster analysis to use, and how to decide on the appropriate number of groups in the data. These and many other practical problems of clustering are discussed in Everitt *et al.* (2001).

15.3 Analysis using Stata

15.3.1 Tibetan skulls

Assuming the data in Table 15.1 are contained in a file `tibetan.dat`, they can be read into Stata using the instruction

```
infile y1 y2 y3 y4 y5 using tibetan.dat, clear
generate id = _n
```

Here we have also generated an identifier variable `id` for the skulls. To begin it is good practice to examine some graphical displays of the data. With multivariate data such as the measurements on skulls in Table 15.1 a scatterplot matrix is often helpful and can be generated as follows:

```
graph matrix y1-y5
```

The resulting plot is shown in Figure 15.1. A few of the individual scatterplots in Figure 15.1 are perhaps suggestive of a division of the observations into distinct groups, for example that for `y4` (upper face height) versus `y5` (face breadth).

We shall now apply each of single linkage, complete linkage, and average linkage clustering to the data using Euclidean distance as the basis of each analysis. Here the five measurements are all on the same scale, so that standardization before calculating the distance matrix is probably not needed (but see the analysis of the air pollution data described later). The necessary Stata commands are

```
cluster singlelinkage y1-y5, name(s1)
cluster dendrogram
```

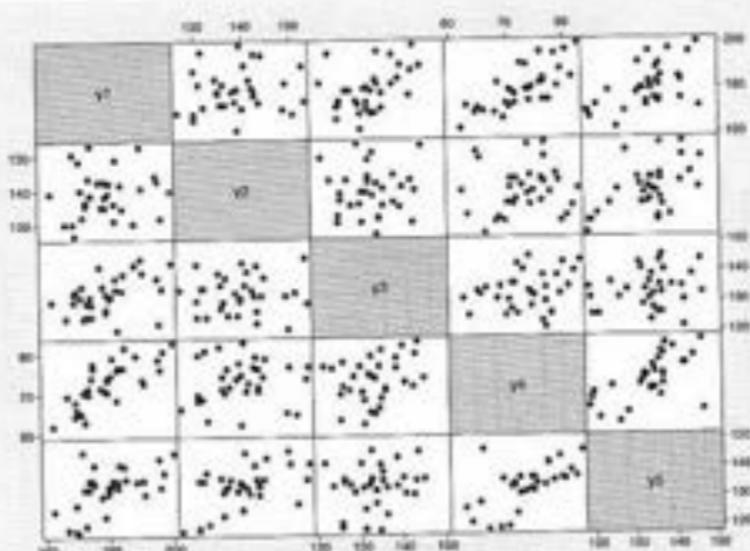


Figure 15.1: Scatterplot matrix of Tibetan skull data

```

cluster completemalinkage y1-y5, name(cl)
cluster dendrogram
cluster averagelinkage y1-y5, name(al)
cluster deadrogram

```

Here the `name()` option is used to attach a name to the results from each cluster analysis. The resulting three dendograms are shown in Figures 15.2, 15.3, and 15.4.

The single linkage dendrogram illustrates one of the common problems with this technique, namely its tendency to incorporate observations into existing clusters rather than begin new ones, a property generally referred to as chaining (see Everitt *et al.*, 2001, for full details). The complete linkage and average linkage dendograms show more evidence of cluster structure in the data, although this structure appears to be different for each method, a point we shall investigate later.

In most applications of cluster analysis the researcher will try to determine the solution with the optimal number of groups, i.e., the number of groups that "best" fits the data. Estimating the number of groups in a cluster analysis is a difficult problem without a completely satisfactory solution – see Everitt *et al.* (2001). Two stopping rules are provided in Stata, the Calinski and Harabasz pseudo F -statistic

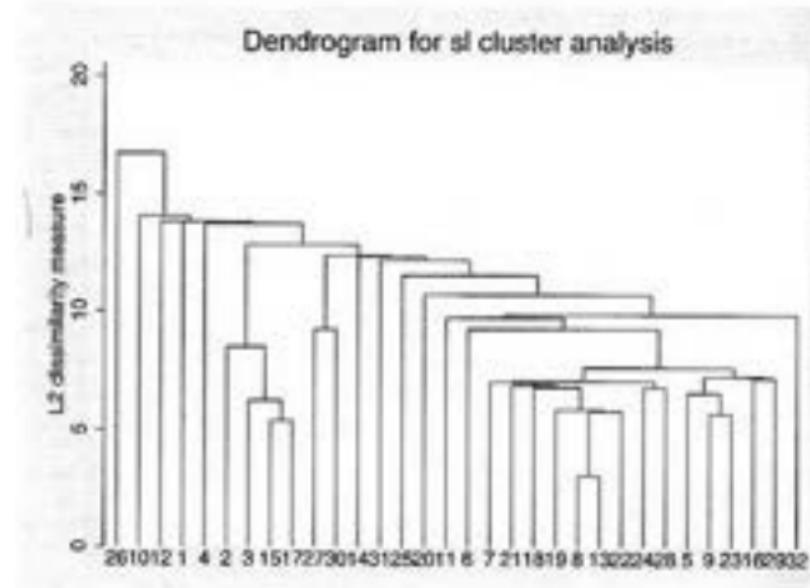


Figure 15.2: Dendrogram using single linkage.

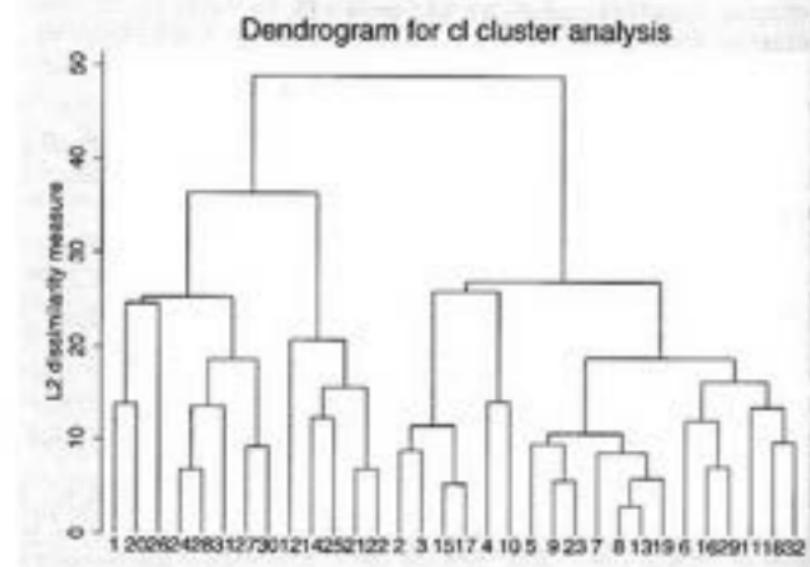


Figure 15.3: Dendrogram using complete linkage.

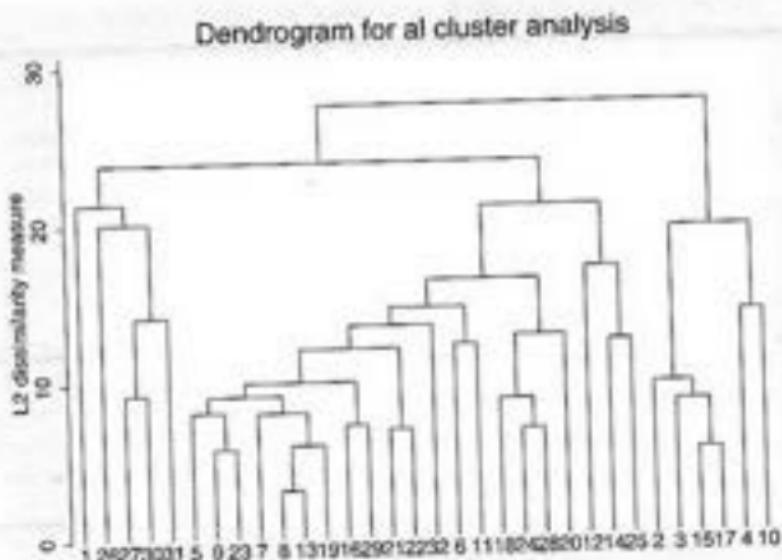


Figure 15.4: Dendrogram using average linkage.

(Calinski and Harabasz, 1974) and the Duda and Hart index (Duda and Hart, 1973); see [MV] cluster stop for details. For both these rules, larger values indicate more distinct clustering.

Here we shall illustrate the use of the Duda and Hart index in association with the three clustering techniques applied above. The Stata commands are

```
cluster stop sl, rule(duda) groups(1/5)
```

(see Display 15.1),

```
cluster stop cl, rule(duda) groups(1/5)
```

(see Display 15.2), and

```
cluster stop al, rule(duda) groups(1/5)
```

(see Display 15.3).

Distinct clustering is generally considered to be indicated by large values of the Duda and Hart index and small values of the Duda and Hart pseudo χ^2 -squared. Adopting this approach, the results from single linkage clustering do not suggest any distinct cluster structure largely because of the chaining phenomenon. The results associated with complete linkage clustering suggest a five-group solution and those from the average linkage method suggest perhaps three or four groups.

Number of clusters	Duda/Hart	
	$J_e(2)/J_e(1)$	pseudo T-squared
1	0.9612	1.54
2	0.9357	1.99
3	0.9430	1.69
4	0.9327	1.95
5	0.9380	1.72

Display 15.1

Number of clusters	Duda/Hart	
	$J_e(2)/J_e(1)$	pseudo T-squared
1	0.5685	14.98
2	0.6073	7.11
3	0.5803	13.34
4	0.3356	7.92
5	0.7006	2.86

Display 15.2

Number of clusters	Duda/Hart	
	$J_e(2)/J_e(1)$	pseudo T-squared
1	0.6732	14.63
2	0.7182	9.37
3	0.5959	2.03
4	0.7200	7.38
5	0.3731	3.36

Display 15.3

To see which skulls are placed in which groups we can use the `cluster generate` command. For example, to examine the five group solution given by complete linkage we use

```
cluster generate g5cl = groups(5), name(c1)
sort g5cl id
forvalues i=1/5 {
    display " "
    display "cluster `i`"
    list id if g5cl=="`i`", noobs noheader separator(0)
}
```

Here we use a `forvalues` loop to list `id` for each cluster. The `noobs` option suppresses line-numbers; the `noheader` option suppresses variable names; and the `separator(0)` option suppresses separator lines. The resulting output is shown in Display 15.4. The numbers of observations in each group can be tabulated using

```
tabulate g5cl
```

giving the table in Display 15.5.

It is often helpful to compare the mean vectors of each of the clusters. The necessary code to find these is:

```
table g5cl, contents(mean y1 mean y2 mean y3 mean y4 ///
mean y5) format(%4.1f)
```

The skulls in cluster 1 are characterized by being relatively long and narrow. Those in cluster 2 are, on average, shorter and broader. Cluster 3 skulls appear to be particularly narrow, and those in cluster 4 have short upper face length. Skulls in cluster 5 might perhaps be considered "average".

The scatterplot matrix of the data used earlier to allow a "look" at the raw data is also useful for examining the results of clustering the data. For example, we can produce a scatterplot matrix with observations identified by their cluster number from the three group solution from average linkage:

```
cluster generate g3al = groups(3), name(c1)
graph matrix y1-y5, xlabel(g3al) xlabelpos(0) msymbol(i)
```

(see Figure 15.5). The separation between the three groups is most distinct in the panel for greatest length of skull `y1` versus face breadth `y5`.

The data as originally collected by Colonel Waddell were thought to consist of two types of skulls; the first type, skulls 1-17, came from graves in Sikkim and the neighboring areas of Tibet. The remaining 15 skulls were picked up on a battlefield in the Lhasa district and were believed to be those of native soldiers from the eastern province of

cluster 1

1
20
24
26
27
28
30
33

cluster 2

12
14
21
22
28

cluster 3

2
3
15
17

cluster 4

4
10

cluster 5

5
6
7
8
9
11
13
16
18
19
23
29
32

Display 15.4

g5c1	Freq.	Percent	Cum.
1	8	25.00	25.00
2	5	15.63	40.63
3	4	12.50	53.13
4	2	6.25	59.38
5	13	40.63	100.00
Total	32	100.00	

Display 15.5

g5c1	mean(y1)	mean(y2)	mean(y3)	mean(y4)	mean(y5)
1	192.9	139.4	138.6	78.1	135.0
2	179.0	146.8	130.2	74.7	141.7
3	169.8	129.6	129.1	68.6	119.5
4	166.0	144.8	132.9	63.3	127.0
5	177.6	137.9	133.7	72.6	133.4

Display 15.6

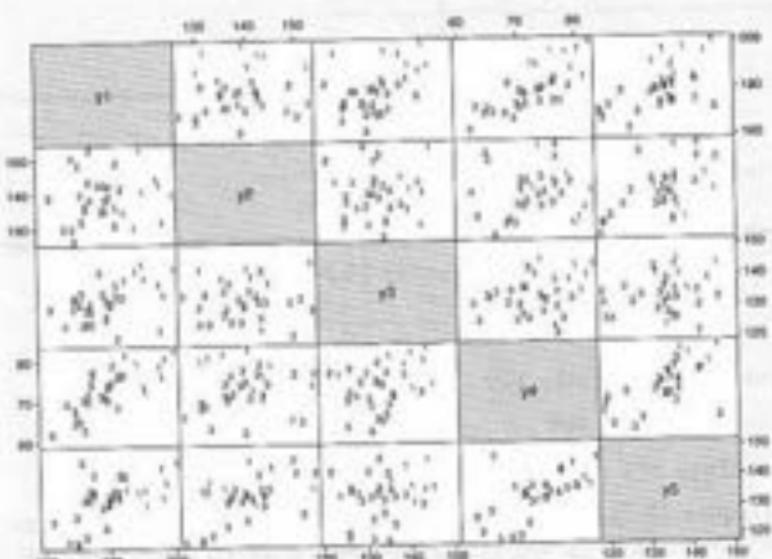


Figure 15.5: Scatterplot matrix with observations identified by their cluster number.

Khaans. These skulls were of particular interest because it was thought at the time that Tibetans from Khaans might be survivors of a particular fundamental human type, unrelated to the Mongolian and Indian types which surrounded them. We can compare this classification with the two group solutions given by each of the three clustering methods by cross-tabulating the corresponding categorical variables containing group membership. The Stata code for this is:

```
generate cl2 = cond(id<=17,1,2)
cluster generate g2al = groups(2), name(al)
cluster generate g2cl = groups(2), name(cl)
cluster generate g2al = groups(2), name(al)
tabulate cl2 g2al, row
```

(see Display 15.7).

Key				Total
		g2al		
		1	2	
cl2				
1		0	17	17
		0.00	100.00	100.00
2		1	14	15
		6.67	93.33	100.00
Total		1	31	32
		3.13	96.88	100.00

Display 15.7

tabulate cl2 g2cl, row

(see Display 15.8), and

tabulate cl2 g2al, row

(see Display 15.9). The two group solution from single linkage consists of 31 observations in one group and only a single observation in the second group, again illustrating the chaining problem associated with this method. The complete linkage solution provides the closest match to the division originally proposed for the skulls (with group labels interchanged).

Key	
frequency	
row percentage	

c12	g2ai		Total
	1	2	
1	3	14	17
	17.65	82.35	100.00
2	10	5	15
	66.67	33.33	100.00
Total	13	19	32
	40.63	59.36	100.00

Display 15.8

Key	
frequency	
row percentage	

c12	g2ai		Total
	1	2	
1	11	6	17
	64.71	35.29	100.00
2	15	0	15
	100.00	0.00	100.00
Total	26	6	32
	81.25	18.75	100.00

Display 15.9

15.3.2 Determinants of pollution in U.S. cities

In this section we shall apply k-means clustering to the air pollution data from Chapter 3. We will use the variables temp, manuf, pop, wind, precip, and days for the cluster analysis. Since these variables have very different metrics we shall begin by standardizing them.

```
infile str10 town so2 temp manuf pop wind precip days ///
    using usair.dat, clear
foreach var of varlist temp manuf pop wind precip days {
    egen s'var' = std('var')
}
```

We will now use the k-means algorithm to divide the data into 2, 3, 4, and 5 groups using the default for choosing initial cluster centers, namely the random selection of k unique observations from among those to be clustered. To be able to reproduce the results here, we set the random number seed using the option `start(krandom(234))`. The commands are

```
cluster kmeans stemp smanuf spop swind sprecip sdays. ///
    k(2) start(krandom(234)) name(cluster2)
cluster kmeans stemp smanuf spop swind sprecip sdays. ///
    k(3) start(krandom(234)) name(cluster3)
cluster kmeans stemp smanuf spop swind sprecip sdays. ///
    k(4) start(krandom(234)) name(cluster4)
cluster kmeans stemp smanuf spop swind sprecip sdays. ///
    k(5) start(krandom(234)) name(cluster5)
```

We can now use the Calinski and Harabasz approach to selecting the optimal number of groups:

```
cluster stop cluster2
```

Number of clusters	Calinski/ Harabasz pseudo- F
2	11.63

```
cluster stop cluster3
```

Number of clusters	Calinski/ Harabasz pseudo- F
3	11.68

cluster stop cluster4

Number of clusters	Calinski/ Harabasz pseudo-F
4	14.64

cluster stop cluster5

Number of clusters	Calinski/ Harabasz pseudo-F
5	14.01

The largest value of the Calinski and Harabasz index corresponds to the four group solution. Details of this solution can be found from

```
sort cluster4 town
forvalues i=1/4 {
    display " "
    display "cluster `i'"
    list town if cluster4=="`i'", noobs noheader separator(0)
}

```

The output is shown in Display 15.10. Note that Chicago is in a cluster of its own indicating again (as in Chapter 3) that this appears to be an outlier. It is worthwhile repeating the cluster analysis with Chicago removed to see how much the composition of the remaining clusters changes (see Exercise 15.2), but here we will continue interpreting the current solution. We can use the tabstat command to tabulate the cluster means (the table command can only tabulate up to five statistics):

```
tabstat temp manuf pop wind precip days, by(cluster4) ///
nototal format(%4.1f)
```

(see Display 15.11).

We will now compare pollution levels (the annual mean concentration of sulphur dioxide so2) between these five clusters of towns. The means and standard deviations can be tabulated using

```
table cluster4, contents(mean so2 sd so2) format(%4.1f)
```

(see Display 15.12). Clusters 2, Chicago, has extremely high pollution levels and cluster 1 has much higher pollution levels than clusters 3 and 4. A more formal analysis of differences in pollution levels among

cluster 1

Albany
Baltimore
Buffalo
Charleston
Cincinnati
Cleveland
Columbus
Des Moines
Detroit
Hartford
Indianapolis
Kansas City
Louisville
Milwaukee
Minneapolis
Omaha
Philadelphia
Pittsburgh
Portland
Seattle
St. Louis
Washington
Wilmington

cluster 2

Chicago

cluster 3

Albuquerque
Dallas
Denver
Phoenix
Salt Lake City
St. Paul
Wichita

cluster 4

Atlanta
Boston
Jacksonville
Lubbock
Memphis
Miami
Nashville
Newark
Norfolk
Richmond

 Summary statistics: mean

by categories of: cluster4

cluster4	temp	meanad	pop	wind	precip	days
1	51.5	476.7	885.4	9.7	36.7	127.0
2	50.6	3344.0	3369.0	10.4	34.4	122.0
3	58.5	295.6	479.1	9.3	38.6	70.9
4	64.2	363.2	476.6	9.0	49.8	113.0

Display 15.11

cluster4	mean(so2)	sd(so2)
1	37.8	21.1
2	110.0	
3	13.6	7.0
4	16.5	7.9

Display 15.12

the clusters can be undertaken using a one-way analysis of variance. (Note that the variable so2 did not contribute to the cluster analysis. If it had, it would be circular and invalid to carry out the analysis of variance.) We will first log-transform so2 since the standard deviations appear to increase with the mean.

```
generate lso2 = ln(so2)
anova lso2 cluster4
```

(see Display 15.13). The analysis shows that the clusters differ significantly in their average pollution levels, $F_{3,37} = 13.21$, $p < 0.001$.

15.4 Exercises

15.1 • Tibetan skulls

1. Repeat the analyses of the Tibetan skull data described in this chapter using the Manhattan distance measure rather than the Euclidean. Compare the two sets of results.

	Number of obs =	41	R-squared =	0.5172
	Root MSE =	16.8678	Adj R-squared =	0.4781
Source	Partial SS	df	MS	F
Model	11397.949	3	3799.31634	13.21
cluster4	11397.949	3	3799.31634	13.21
Residual	19639.9534	37	527.566309	
Total	22037.9024	40	550.947561	

Display 15.13

15.2 • Determinants of pollution in U.S. cities

1. Repeat the k -means cluster analysis of the air pollution data with Chicago removed. Does this lead to a very different solution?
2. Investigate the use of other options for determining an initial partition when applying k -means to the air pollution data (still with Chicago removed).
3. Compare the results from k -medians cluster analysis applied to the air pollution data with those from k -means (still with Chicago removed).

15.3 Romano-British pottery

Tubb et al. (1990) provide the chemical composition of 48 specimens of Romano-British pottery, determined by atomic absorption spectrophotometry. In addition to the chemical composition of the pots, the kiln site at which the pottery was found is also noted.

The variables in *pottery.dta* are:

- no: identification number of pottery
 - kiln: identification number of kiln (site) where pottery was found
 - al2o3, fe2o3, mgo, na2o, k2o, tio2, mno, bao amount of Al₂O₃, Fe₂O₃, MgO, CaO, Na₂O, K₂O, TiO₂, MnO, and BaO respectively
1. Apply k -means clustering to the chemical data remembering that the variables are on very different scales.
 2. Use the Calinski and Harabasz approach for finding the best number of clusters (try between two and six clusters).

- Once you have chosen a particular solution, assess whether there is any association between the kiln site and the distinct compositional groups found by cluster analysis.

15.4 Life expectancies

Keyfitz and Flicger (1971) provide data on life expectancy (in the 1960s) in years by country, age, and sex. Here life expectancy refers to the mean number of extra years of life for people of a given sex who have reached a given age.

The variables in `life.dta` are:

- `country`: country (string variable)
 - `short`: abbreviation of country
 - `m0` to `m75`: men, ages 0, 25, 50, and 75
 - `w0` to `w75`: women, ages 0, 25, 50, and 75
- Apply complete linkage, average linkage, and single linkage cluster analysis and generate variables of group membership for the four-cluster solutions.
 - Perform principal components analysis based on the covariance matrix of the life expectancy variables. Interpret the first three components.
 - Plot the four-group solution using complete linkage in the space of the first two principal components using different marker symbols for the four groups and the strings in the variable `short` as labels. Produce analogous graphs for average linkage and single linkage.

Appendix: Answers to Selected Exercises

Chapter 1

1.1 Some data manipulation

2. Assuming that the data are stored in the directory `c:\user`,
`cd c:\user`
`insheet using test.dat, clear`
4. `label define s 1 male 2 female`
`label values sex s`
5. `generate id = _n`
6. `rename v1 time1`
`rename v2 time2`
`rename v3 time3`
or
`forvalues i = 1/3 {`
 `rename var'i' time'i'`
`}`
7. `reshape long time, i(id) j(occ)`
8. `egen d = mean(time), by(id)`
`replace d = (time-d)^2`
9. `drop if occ==3&id==2`

Chapter 2

2.1 Female psychiatric patients

1. insheet using fem.dat, clear
table depress, contents(mean weight)
2. foreach var in iq age weight {
 table life, contents(mean `var' sd `var')
}
3. graph bar (count) id,
 over(sex, relabel(1 "no" 2 "yes")) ///
 over(life, relabel(1 "non-suicidal" 2 "suicidal")) ///
 ytitle(Percentages by group (suicidal versus not)) ///
 asyvars percent showyvars legend(off)
4. search mann
help ranksum
5. ranksum weight, by(life)
6. twoway (scatter iq age if life==1, msymbol(circle) ///
 mcolor(black) jitter(2)) ///
 (scatter iq age if life==2, msymbol(x) ///
 mcolor(black) jitter(2)),
 legend(order(1 "no" 2 "yes")) ///
7. spearman age iq
8. Save the commands in the Review window and edit the file using the Do-file Editor or any text editor, e.g., Notepad. Add the commands given in the do-file template in Section 1.11, and save the file with the extension .do. Run the file by typing the command
do filename.

Chapter 4

4.1 Treating hypertension

1. infile bp11 bp12 bp13 bp01 bp02 bp03 using bp.raw, clear
Now follow the commands on pages 88 to 89. There is no need to redefine labels, but if you wish to do so, first issue the command
label drop _all.
2. graph box bp, over(drug)
graph box bp, over(diet)
graph box bp, over(biofeed)
4. sort id
save bp
infile id age using age.dat, clear
sort id

```
merge id using bp
anova bp drug diet biofeed age, continuous(age)
```

Chapter 5

5.1 Effectiveness of slimming clinics

1. infile manual exper resp using slim.dat, clear
anova resp manual*exper exper manual, sequential
2. generate dmanual = manual - 1
generate dexper = exper - 1
generate dinter = dmanual*dexper
regress resp dmanual dexper dinter
3. xi: regress resp i.manual*i.exper
4. char manual[omit] 2
char exper[omit] 2
xi: regress resp i.manual*i.exper

Chapter 6

6.1 Treatment of lung cancer

1. infile fr1 fr2 fr3 fr4 using tumor.dat, clear
generate therapy = int(_n-1)/2
sort therapy
by therapy: generate sex = _n
label define t 0 seq 1 alt, modify
label values therapy t
label define s 1 male 2 female, modify
label values sex s
reshape long fr, i(therapy sex) j(outc)
ologit outc therapy sex [fweight=fr], table

6.2 Female psychiatric patients

1. a. insheet using fem.dat, clear
replace sleep=. if sleep==3
recode sleep 1=2 2=1
ologit depress life
b. replace life = life - 1
logistic life depress
2. Even if we use very lenient inclusion and exclusion criteria,

```
stepwise, pr(0.3) pe(0.2) forward: ///
    logit life depress anxie iq sex sleep
only depress is selected. If we exclude depress from the list of
candidate variables, anxiety and sleep are selected.
```

6.3 Diagnosis of heart attacks

1. infile ck pres abs using sck.dat, clear
generate tot = pres + abs
expand tot
by ck, sort: generate infct = (_n<=pres)
logit infct ck
estat gof, table
 2. generate ck2 = ck^2
logit infct ck ck2
generate ck3 = ck^3
logit infct ck ck2 ck3
generate ck4 = ck^4
logit infct ck ck2 ck3 ck4
logit infct ck ck2 ck3
 3. estat gof, table
 4. infile ck pres abs using sck.dat, clear
generate tot = pres + abs
generate prop = pres/tot
- twoway (function y = invlogit(_b[_cons]+_b[ck]*x ///
+ _b[ck2]*x^2+_b[ck3]*x^3), range(0 480)) ///
(scatter prop ck), xtitle(CK) ///
legend(order(1 "predicted" 2 "observed"))

Chapter 7

7.1 Effectiveness of slimming clinics

1. infile cond status resp using slim.dat, clear
xi: glm resp i.cond i.status, fam(gauss) link(id)
local dev1 = e(deviance)
xi: glm resp i.cond, fam(gauss) link(id)
local dev0 = e(deviance)
local ddev = `dev0' - `dev1'
/* F-test equivalent to anova cond status, sequential */
local f = (`ddev'/1)/(`dev1'/31)
display `f'

```

display fprob(1,31,'f')
/* difference in deviance */
display `ddev'
display chiprob(1, `ddev')
2. regress resp status, vce(robust)
ttest resp, by(status) unequal

```

7.2 Australian school children

- use quine, clear
 encode eth, gen(ethnic)
 drop eth
 encode sex, gen(gender)
 drop sex
 encode age, gen(class)
 drop age
 encode lrm, gen(slow)
 drop lrm
 generate cleth = class*ethnic
 glm days class ethnic cleth, family(poisson) link(log)
- glm days class ethnic if stres<4, family(poisson) link(log)
 or, assuming the sort order of the data has not changed,
 glm days class ethnic if _n!=72, family(poisson) link(log)
- generate abs = cond(days>=14,1,0)
 glm abs class ethnic, family(binomial) link(logit)
 glm abs class ethnic, family(binomial) link(probit)
- glm abs class ethnic, family(binomial) link(logit) ///
 vce(robust)
 glm abs class ethnic, family(binomial) link(probit) ///
 vce(robust)
 bootstrap _b[class] _b[ethnic], reps(500): ///
 glm abs class ethnic, family(binomial) link(logit)
 bootstrap _b[class] _b[ethnic], reps(500): ///
 glm abs class ethnic, family(binomial) link(probit)

Chapter 8

8.1 Treatment of post-natal depression

- infile subj group pre dep1 dep2 dep3 dep4 dep5 dep6 ///
 using depress.dat, clear
 mvdecode _all, mv(-9)
 graph box dep1-dep6, by(group)

2. We can obtain the mean over visits for subjects with complete data using the simple command (data in "wide" form)

```
generate av2 = (dep1+dep2+dep3+dep4+dep5+dep6)/6
```

For subjects with missing data, av2 will be missing whereas the egen function rowmean() would return the mean of all available data. The t-tests are obtained using

- ```
ttest av2, by(group)
ttest av2, by(group) unequal
```
3. egen max = rowmax(dep1-dep6)  
 ttest max, by(group)
4. a. generate diff = avg - pre  
     ttest diff, by(group)  
 b. anova avg group pre, continuous(pre)

## Chapter 9

### 9.1 Thought disorder and schizophrenia

```
use madras, clear
reshape long y, i(id) j(month)
label variable month ///
 "Number of months since hospitalization"
generate month_early = month*early
label define e 0 "Late onset" 1 "Early onset"
label values early e

gllamm y month early month_early, i(id) ///
 link(logit) family(binom) eform adapt
gllapred prob1, mu
sort id month
twoway (line prob1 month, connect(ascending)), ///
 by(early) ytitle(Predicted probability)

generate cons = 1
eq slope: month
eq inter: cons
gllamm y month early month_early, i(id) nrf(2) ///
 eqs(inter slope) link(logit) family(binom) ///
 eform adapt
gllapred prob2, mu
sort id month
```

```
twoway (line prob2 month, connect(ascending)), ///
by(early) ytitle(Predicted probability)
```

## 9.2 Australian school children

```
1. use ..\data\quine, clear
encode eth, gen(ethnic)
drop eth
encode sex, gen(gender)
drop sex
encode age, gen(class)
drop age
encode lrm, gen(slow)
drop lrm

generate id=_n
gllamm days class ethnic, i(id) adapt ///
family(poisson) link(log)
```

## Chapter 10

### 10.1 Treatment of post-natal depression

```
a. infile subj group dep0 dep1 dep2 dep3 ///
dep4 dep5 dep6 using depress.dat, clear
mvdecode _all, mv(-9)
reshape long dep, i(subj) j(visit)
generate gr_vis = group*visit
xtgee dep group visit gr_vis, i(subj) cor(exch)
regress dep group visit gr_vis, vce(robust) ///
cluster(subj)
b. bootstrap _b[group] _b[visit] _b[gr_vis], ///
cluster(subj) reps(500): ///
regress dep group visit gr_vis
```

## Chapter 11

### 11.1 Estrogens and endometrial cancer

```
1. infile v1-v2 using estrogen.dat, clear
generate _varname = cond(_n==1,"ncases1","ncases0")
xpose, clear
generate casestr = 2-_n
reshape long ncases, i(casestr) j(casestr)
```

```

expand mcases
sort casestr conestr
generate caseid = _n
expand 2
by caseid, sort: generate control = _n-1
* (dummy for control: 1=cont., 0=caso)
generate estr = 0
replace estr = 1 if control==0&casestr==1
replace estr = 1 if control==1&conestr>0
generate cancer = cond(control==0,1,0)

preserve
collapse (sum) estr (mean) casestr , by(caseid)
generate conestr = estr - casestr
tabulate casestr conestr
restore

clogit cancer estr, group(caseid) or

```

### 11.2 Low energy diet and heart disease

1. infile str5 age num1 py1 num0 py0 using ihd.dat,clear  
generate agegr = \_n  
reshape long num py, i(agegr) j(exposed)

```

table exposed, contents(sum num sum py)
iri 28 17 1857.5 2768.9

```

2. Keeping the data from the previous exercise:

```

xi: poisson num i.age*exposed, exposure(py) irr
testparm _IageX*

```

The interaction is not statistically significant at the 5% level.

## Chapter 12

### 12.1 Retention of heroin addicts in methadone maintenance treatment

1. We consider anyone still at risk after 450 days as being censored at 450 days and therefore need to make the appropriate changes to status and time before running Cox regression.

```

use heroin, clear
replace status = 0 if time>450

```

```

replace time = 450 if time>450
egen zdose = std(dose)
stset time status
stcox zdose prison clinic

```

2. The model is fitted using

```

generate dosecat = 0 if dose<.
replace dosecat = 1 if dose>=60 & dose<.
replace dosecat = 2 if dose>=80 & dose<.
xi: stcox i.dosecat i.prison i.clinic, bases(s)

```

The survival curves for no prison record, clinic 1 are obtained and plotted using

```

generate s0 = s if dosecat==0
generate s1 = s^(exp(_b[_Idosecat_1])) if dosecat==1
generate s2 = s^(exp(_b[_Idosecat_2])) if dosecat==2
sort time
graph twoway (line s0 time, connect(stairstep)) ///
 (line s1 time, connect(stairstep) lpat(dash)) ///
 (line s2 time, connect(stairstep) lpat(dot)), ///
 legend(order(1 "<60" 2 "60-79" 3 ">=80"))

```

Note that the baseline survival curve is the survival curve for someone whose covariates are all zero. If we had used `clinic` instead of `i.clinic` above, this would have been meaningless; we would have had to exponentiate `s0`, `s1`, and `s2` by `.b[clinic]` to calculate the survival curves for clinic 1.

3. Treating dose as continuous:

```

generate clindose = clinic*zdose
stcox zdose clinic clindose prison

```

Treating dose as categorical:

```

xi: stcox i.dosecat*i.clinic i.prison
testparm _IdosX*

```

4. xi: stcox i.dosecat i.prison i.clinic

```

xi: stcox i.dosecat i.prison i.clinic, efron
xi: stcox i.dosecat i.prison i.clinic, exactm

```

It makes almost no difference which method is used.

5. stcox zdose prison, strata(clinic) tvc(prison) ///

```

 texp((t-504)/365.26)
estimates store model1

```

```

quietly stcox zdose prison, strata(clinic)

```

```

lrtest model1 .

```

## Chapter 13

### 13.1 Two-component mixture of normals

```
2. nlcom (sd1: exp([lsd1] [_cons])) ///
 (sd2: exp([lsd2] [_cons])) ///
 (p: invlogit([lo1] [_cons]))
```

giving estimates (standard errors) 6.54 (0.82) and 7.06 (1.81) for the standard deviations and 0.74 (0.07) for the probability.

### 13.2 Heteroscedastic linear regression

1. The only thing that is different from fitting a normal distribution with constant mean is that the mean is now a linear function of status so that the `ml model` command changes as shown below:

```
infile cond status resp using slim.dat, clear
ml model lf mixing1 (xb: resp = status) /lsd
ml maximize, noheader
```

2. In linear regression, the mean squared error is equal to the sum of squares divided by the degrees of freedom,  $n - 2$ . The maximum likelihood estimate is equal to the sum of squares divided by  $n$ . We can therefore get the root mean square error for linear regression using

```
disp exp([lsd] [_cons])*sqrt(e(N)/(e(N)-2))
```

Note that the standard errors of the regression coefficients need to be corrected by the same factor, i.e.,

```
disp [xb] _se[status]*sqrt(e(N)/(e(N)-2))
```

Compare this with the result of

```
regress resp status
```

3. Repeat the procedure above but replace the `ml model` command by

```
ml model lf mixing1 (resp = status) (lsd: status)
```

The effect of status on the standard deviation is significant ( $p = 0.003$ ) which is not too different from the result of

```
sttest resp, by(status)
```

### 13.3 Three-component mixture of normals

1. capture program drop mixing3
 program mixing3
 version 9.2
 args l1 xb1 xb2 xb3 lo1 lo2 ls1 ls2 ls3

```

tempvar f1 f2 f3 p1 p2 p3 s1 s2 s3 d

quietly {
 generate double `s1' = exp(`ls1')
 generate double `s2' = exp(`ls2')
 generate double `s3' = exp(`ls3')
 generate double `d' = 1 + exp(`lo1') + exp(`lo2')
 generate double `p1' = 1/`d'
 generate double `p2' = exp(`lo1')/`d'
 generate double `p3' = exp(`lo2')/`d'

 generate double `f1' = normalden($ML_y1,"xb1","s1")
 generate double `f2' = normalden($ML_y1,"xb2","s2")
 generate double `f3' = normalden($ML_y1,"xb3","s3")

 replace `lj' = ln(`p1'*`f1' ///
 + `p2'*`f2' + `p3'*`f3')
}
end

clear
set obs 300
set seed 12345678
generate z = uniform()
generate y = invnormal(uniform())
replace y = y + 5 if z<1/3
replace y = y + 10 if z>2/3&z>=1/3
ml model lf mixing3 (xb1: y) ///
 /xb2 /xb3 /lo1 /lo2 /lsd1 /lsd2 /lsd3
ml init 0 5 10 0 0 0 0, copy
ml maximize, noheader trace

```

## Chapter 14

### 14.1 Hearing measurement using an audiometer

1. infile id 1500 11000 12000 14000 r500 r1000 ///
 r2000 r4000 using hear.dat, clear
 pca 1500-r4000, cov
 predict npc1 npe2, score
 twoway scatter npe2 npe1, mlabel(id)

```
3. capture drop pc*
pca 1500-r4000
predict pc1-pc5, score
graph matrix pc1-pc5
```

#### 14.2 Determinants of pollution in U.S. cities

```
1. infile str10 town so2 temp manuf pop wind precip ///
 days using usair.dat, clear
pca temp manuf pop wind precip days
predict pc1 pc2, score
scatter pc2 pc1, mlabel(town) mlabpos(0) msymbol(i)
2. regress so2 pc1 pc2
generate regline = pc1*_b[pc2]/_b[pc1]
twoway (line regline pc1) ///
(scatter pc2 pc1, mlabel(town) mlabpos(0) msymbol(i))
```

## Chapter 15

### 15.1 Tibetan skulls

```
1. infile y1 y2 y3 y4 y5 using tibetan.dat, clear
cluster singlalinkage y1-y5, name(s1) manhattan
cluster completelinkage y1-y5, name(c1) manhattan
cluster averagelinkage y1-y5, name(a1) manhattan
Then plot dendrograms, etc.
```

### 15.2 Determinants of pollution in U. S. cities

```
1. infile str10 town so2 temp manuf pop wind precip ///
 days using usair.dat, clear
foreach v of varlist temp manuf pop wind precip days {
 egen s`v' = std(`v')
}
drop if town=="Chicago"
```

Repeating the analysis from Page 308 to 309 (with the same random number seed) leads to a 5-cluster solution with the former clusters 1 and 4 being split up and cluster 3 nearly remaining intact.

```
2. cluster kmeans stemp smanuf spop swind sprecip ///
 sdays, k(5) name(cluster5) start(segments)
The start(segments) option splits the sample into k (here 5) equal groups and uses the means as starting values.
3. cluster kmedians stemp smanuf spop swind sprecip ///
 sdays, k(5) name(cluster5)
```

---

## References

---

- Acock, A. C. 2006. *A Gentle Introduction to Stata*. College Station, TX: Stata Press.
- Agresti, A. 1984. *Analysis of Ordinal Categorical Data*. New York: Wiley.
- Agresti, A. 1996. *Introduction to Categorical Data Analysis*. New York: Wiley.
- Agresti, A. 2002. *Categorical Data Analysis (Second Edition)*. Hoboken, NJ: Wiley.
- Aitkin, M. 1978. The analysis of unbalanced cross-classifications. *Journal of the Royal Statistical Society, Series A*, 41, 195–223.
- Allison, P. D. 1984. *Event History Analysis. Regression for Longitudinal Event Data*. Sage University Paper Series on Quantitative Applications in the Social Sciences. Newbury Park, CA: Sage.
- Altman, D. G. 1990. *Practical Statistics for Medical Research*. London: Chapman & Hall.
- Andrews, D. F., & Herzberg, A. M. 1985. *Data: A Collection of Problems from Many Fields for the Student and Research Worker*. New York: Springer-Verlag.
- Binder, D. A. 1983. On the variances of asymptotically normal estimators from complex surveys. *International Statistical Review*, 51, 279–292.
- Boniface, D. R. 1995. *Experimental Design and Statistical Methods*. London: Chapman & Hall.
- Box-Steffensmeier, J. M., & Jones, B. S. 2004. *Event History Modeling: A Guide to Social Statistics*. Cambridge: Cambridge University Press.

- Breslow, N. E., & Clayton, D. G. 1993. Approximate inference in generalized linear mixed models. *Journal of the American Statistical Association*, **88**, 9–25.
- Brown, B. W. 1980. Prediction analysis for binary data. In: Miller, R. J., Efron, B., Brown, B. W., & Moses, L. E. (eds), *Biostatistics Casebook*. New York: Wiley.
- Calinski, T., & Harabasz, J. 1974. A dendrite method for cluster analysis. *Communications in Statistics – Theory and Methods A*, **3**, 1–27.
- Caspary, J., & Bell, J. 1991. Methadone dosage and the retention of patients in maintenance treatment. *The Medical Journal of Australia*, **154**, 195–199.
- Chatterjee, S., & Hadi, A. S. 2006. *Regression Analysis by Example (Fourth Edition)*. New York: Wiley.
- Clayton, D. G., & Hills, M. 1993. *Statistical Models in Epidemiology*. Oxford: Oxford University Press.
- Cleves, M., Gould, W. W., & Gutierrez, R. 2004. *An Introduction to Survival Analysis Using Stata (Revised Edition)*. College Station, TX: Stata Press.
- Collett, D. 2002. *Modelling Binary Data (Second Edition)*. London: Chapman & Hall/CRC.
- Collett, D. 2003. *Modelling Survival Data in Medical Research (Second Edition)*. Boca Raton, FL: Chapman & Hall/CRC.
- Cook, R. D. 1977. Detection of influential observations in linear regression. *Technometrics*, **19**, 15–18.
- Cook, R. D., & Weisberg, S. 1982. *Residuals and Influence in Regression*. London: Chapman & Hall.
- Cook, T. D., & Campbell, D. T. 1979. *Quasi-Experimentation*. Boston: Houghton-Mifflin.
- Cox, D. R., & Solomon, P. J. 2003. *Components of Variance*. Boca Raton, FL: Chapman & Hall /CRC.
- Cox, N. J. 2002a. Speaking Stata: How to face lists with fortitude. *The Stata Journal*, **2**, 202–222.
- Cox, N. J. 2002b. Speaking Stata: How to move step by step. *The Stata Journal*, **2**, 86–102.
- Crouchley, R., & Davies, R. B. 1999. A comparison of population average and random effects models for the analysis of longitudinal count data with base-line information. *Journal of the Royal Statistical Society, Series A*, **162**, 331–347.

- Davis, C. S. 2002. *Statistical Methods for the Analysis of Repeated Measurements*. New York: Springer.
- Der, G., & Everitt, B. S. 2002. *A Handbook of Statistical Analyses using SAS (Second Edition)*. London: Chapman & Hall/CRC.
- Dickson, E. R., Grambsch, P. M., Fleming, T. R., Fisher, L. D., & Langworthy, A. 1989. Prognosis in primary biliary cirrhosis: Model for decision making. *Hepatology*, **10**, 1-7.
- Diggle, P. J., Heagerty, P. J., Liang, K.-Y., & Zeger, S. L. 2002. *Analysis of Longitudinal Data*. Oxford: Oxford University Press.
- Duda, R. O., & Hart, P. E. 1973. *Pattern Classification and Scene Analysis*. Chichester: Wiley.
- Efron, B., & Tibshirani, R. J. 1993. *An Introduction to the Bootstrap*. London: Chapman & Hall.
- Everitt, B. S. 1992. *The Analysis of Contingency Tables (Second Edition)*. London: Chapman & Hall.
- Everitt, B. S. 1994. *Statistical Methods for Medical Investigations*. London: Edward Arnold.
- Everitt, B. S. 1996. The analysis of repeated measures: A practical review with examples. *The Statistician*, **44**, 113-135.
- Everitt, B. S. 1996. An introduction to finite mixture distributions. *Statistical Methods in Medical Research*, **5**, 107-127.
- Everitt, B. S. 2001. *Statistics for Psychologists: An Intermediate Course*. Mahwah, NJ: Lawrence Erlbaum.
- Everitt, B. S., & Dunn, G. 2001. *Applied Multivariate Data Analysis (Second Edition)*. London: Edward Arnold.
- Everitt, B. S., & Pickles, A. 2004. *Statistical Aspects of the Design and Analysis of Clinical Trials*. London: Imperial College Press.
- Everitt, B. S., & Rabe-Hesketh, S. 1997. *The Analysis of Proximity Data*. London: Edward Arnold.
- Everitt, B. S., Landau, S., & Leese, M. 2001. *Cluster Analysis (Fourth Edition)*. London: Edward Arnold.
- Fleming, T. R., & Harrington, D. P. 1991. *Counting Process and Survival Analysis*. New York: Wiley.
- Goldstein, H. 2003. *Multilevel Statistical Models (Third Edition)*. London: Arnold.
- Gould, W., Pitblado, J., & Sribney, W. 2006. *Maximum Likelihood Estimation with Stata (Third Edition)*. College Station, TX: Stata Press.

- Grambsch, P., & Therneau, T. 1994. Proportional hazards tests and diagnostics based on weighted residuals. *Biometrika*, **81**, 515–526.
- Green, D. M., Kahl, C., & Diehl, P. F. 1998. The price of peace: A predictive model of UN peacekeeping fiscal costs. *Policy Studies Journal*, **26**, 620–635.
- Gregoire, A. J. P., Kumar, R., Everitt, B. S., Henderson, A. F., & Studd, J. W. W. 1996. Transdermal oestrogen for the treatment of severe post-natal depression. *The Lancet*, **347**, 930–934.
- Hand, D. J., & Crowder, M. J. 1996. *Practical Longitudinal Data Analysis*. London: Chapman & Hall.
- Hand, D. J., Daly, F., Lunn, A. D., McConway, K. J., & Ostrowski, E. 1994. *A Handbook of Small Data Sets*. London: Chapman & Hall.
- Hardin, J., & Hilbe, J. 2002. *Generalized Estimating Equations*. Boca Raton, FL: Chapman & Hall/CRC.
- Hardin, J., & Hilbe, J. 2006. *Generalized Linear Models and Extensions (Second Edition)*. College Station, TX: Stata Press.
- Harrell, F. E. 2001. *Regression Modeling Strategies. With Application to Linear Models, Logistic Regression and Survival Analysis*. New York: Springer.
- Hartigan, J. 1975. *Clustering Algorithms*. New York: Wiley.
- Holtbrugge, W., & Schumacher, M. 1991. A comparison of regression models for the analysis of ordered categorical data. *Applied Statistics*, **40**, 249–259.
- Hotelling, H. 1933. Analysis of a complex of statistical variables into principal components. *Journal of Educational Psychology*, **24**, 498–520.
- Hout, M., Duncan, O. D., & Sobel, M. E. 1987. Association and Heterogeneity: Structural Models of Similarities and Differences. Pages 146–184 of: Clogg, C. C. (ed), *Sociological Methodology 1987*. Washington, DC: American Sociological Association.
- Hurn, M. W., Barker, N. W., & Magath, T. D. 1945. The determination of prothrombin time following the administration of dicumarol with specific reference to thromboplastin. *Journal of Laboratory & Clinical Medicine*, **30**, 432–447.
- Jackson, J. E. 1991. *A User's Guide to Principal Components*. New York: Wiley.
- Jennrich, R. I., & Schluchter, M. D. 1986. Unbalanced repeated measures models with unstructured covariance matrices. *Biometrics*, **42**, 805–820.

- Keyfitz, N., & Flieger, W. 1971. *Population: Facts and Methods of Demography*. San Francisco, CA: Freeman.
- Klein, J. P., & Moeschberger, M. L. 2003. *Survival Analysis: Techniques for Censored and Truncated Data (Second Edition)*. New York: Springer.
- Klemchuk, H. P., Bond, L. A., & Howell, D. C. 1990. Coherence and correlates of level 1 perspective taking in children. *Merrill-Palmer Quarterly*, 36, 369-387.
- Kohler, U., & Kreuter, F. 2005. *Data Analysis Using Stata*. College Station, TX: Stata Press.
- Lachenbruch, P., & Mickey, R. M. 1986. Estimation of error rates in discriminant analysis. *Technometrics*, 10, 1-11.
- Lewis, A. Y., & Shakun, M. F. 1976. *Policy Sciences: Methodology and Cases*. Oxford: Pergamon Press.
- Lewine, R. R. J. 1981. Sex differences in schizophrenia: timing or subtypes? *Psychological Bulletin*, 90, 432-444.
- Liang, K.-Y., & Zeger, S. L. 1986. Longitudinal data analysis using generalized linear models. *Biometrika*, 73, 13-22.
- Lindsey, J. K. 1999. *Models for Repeated Measurements (Second Edition)*. Oxford, UK: Oxford University Press.
- Lindsey, J. K., & Lambert, P. 1998. On the appropriateness of marginal models for repeated measurements in clinical trials. *Statistics in Medicine*, 17, 447-469.
- Long, J. S., & Freese, J. 2006. *Regression Models for Categorical Dependent Variables Using Stata (Second Edition)*. College Station, TX: Stata Press.
- Madsen, M. 1901. On criminal anthropometry and the identification of criminals. *Biometrika*, 1, 177-227.
- Mallows, C. L. 1973. Some comments on  $C_p$ . *Technometrics*, 15, 661-667.
- Manly, B. F. J. 1997. *Randomisation, Bootstrap and Monte Carlo Methods in Biology*. London: Chapman & Hall.
- Mann, J. I., Inman, W. H. W., & Thorogood, M. 1986. Oral contraceptive use in older women and fatal myocardial infarction. *British Medical Journal*, 2, 193-199.
- Mann, L. 1981. The baiting crowd in episodes of threatened suicide. *Journal of Personality and Social Psychology*, 41, 703-709.
- Matthews, J. N. S., Altman, D. G., Campbell, M. J., & Royston, P. 1990. Analysis of serial measurements in medical research. *British Medical Journal*, 300, 230-235.

- Maxwell, S. E., & Delaney, H. D. 1990. *Designing Experiments and Analysing Data*. Belmont, CA: Wadsworth.
- McCullagh, P., & Nelder, J. A. 1989. *Generalized Linear Models (Second Edition)*. London: Chapman & Hall.
- McLachlan, G., & Peel, D. A. 2000. *Finite Mixture Models*. New York, NY: Wiley.
- Miettinen, O. S. 1969. Individual matching with multiple controls in the case of all-or-none response. *Biometrics*, **25**, 339–355.
- Miles, J., & Shevlin, M. 2001. *Applying Regression and Correlation*. London: Sage Publications.
- Mitchell, M. 2004. *A Visual Guide to Stata Graphics*. College Station, TX: Stata Press.
- Morant, G. M. 1923. A first study of the Tibetan skull. *Biometrika*, **14**, 193–260.
- Nelder, J. A. 1977. A reformulation of linear models. *Journal of the Royal Statistical Society (Series A)*, **140**, 48–63.
- Newton, H., & Cox, N. J. (eds). 2006. *Thirty-three Stata Tips*. College Station, TX: Stata Press.
- Pearson, K. 1901. On lines and planes of closest fit to points in space. *Philosophical Magazine, Series 6*, **2**, 559–572.
- Pothoff, R. F., & Roy, S. N. 1964. A generalized multivariate analysis of variance model useful especially for growth curve problems. *Biometrika*, **51**, 313–326.
- Rabe-Hesketh, S., & Skrondal, A. 2005. *Multilevel and Longitudinal Modeling using Stata*. College Station, TX: Stata Press.
- Rabe-Hesketh, S., Skrondal, A., & Pickles, A. 2002. Reliable estimation of generalized linear mixed models using adaptive quadrature. *The Stata Journal*, **2**, 1–21.
- Rabe-Hesketh, S., Pickles, A., & Skrondal, A. 2003. Correcting for covariate measurement error in logistic regression using nonparametric maximum likelihood estimation. *Statistical Modelling*, **3**, 215–232.
- Rabe-Hesketh, S., Skrondal, A., & Pickles, A. 2004a. Generalized multilevel structural equation modeling. *Psychometrika*, **69**, 167–190.
- Rabe-Hesketh, S., Skrondal, A., & Pickles, A. 2004b. *GLLAMM Manual*. Tech. rept. 160. U.C. Berkeley Division of Biostatistics Working Paper Series. Downloadable from <http://www.bepress.com/ucbbiostat/paper160/>.

- Rabe-Hesketh, S., Skrondal, A., & Pickles, A. 2005. Maximum likelihood estimation of limited and discrete dependent variable models with nested random effects. *Journal of Econometrics*, **128**, 301–323.
- Rawlings, J. O., Pantula, S. G., & Dickey, D. A. 1998. *Applied Regression Analysis: A Research Tool (Second Edition)*. New York: Springer.
- Rothman, K. J. 1986. *Modern Epidemiology*. Boston: Little, Brown & Company.
- Sackett, D. L., Haynes, R. B., Guyatt, G. H., & Tugwell, P. 1991. *Clinical Epidemiology*. Massachusetts: Little Brown & Company.
- Skrondal, A., & Rabe-Hesketh, S. 2003. Multilevel logistic regression for polytomous data and rankings. *Psychometrika*, **68**, 267–287.
- Skrondal, A., & Rabe-Hesketh, S. 2004. *Generalized Latent Variable Modeling: Multilevel, Longitudinal, and Structural Equation Models*. Boca Raton, FL: Chapman & Hall/CRC.
- Snijders, T. A. B., & Bosker, R. J. 1999. *Multilevel Analysis*. London: Sage.
- Sokal, R. R., & Rohlf, F. J. 1981. *Biometry*. San Francisco: W. H. Freeman.
- Sprent, P., & Smeeton, N. C. 2001. *Applied Nonparametric Statistical Methods (Third Edition)*. Boca Raton, FL: Chapman & Hall / CRC.
- StataCorp. 2005a. *Getting Started with Stata 9 for Windows Manual*. College Station, TX: Stata Press.
- StataCorp. 2005b. *Stata 9 Base Reference Manual*. College Station, TX: Stata Press.
- StataCorp. 2005c. *Stata 9 Data Management Reference Manual*. College Station, TX: Stata Press.
- StataCorp. 2005d. *Stata 9 Graphics Reference Manual*. College Station, TX: Stata Press.
- StataCorp. 2005e. *Stata 9 Longitudinal/Panel Data Reference Manual*. College Station, TX: Stata Press.
- StataCorp. 2005f. *Stata 9 Programming Reference Manual*. College Station, TX: Stata Press.
- StataCorp. 2005g. *Stata 9 Quick Reference and Index*. College Station, TX: Stata Press.
- StataCorp. 2005h. *Stata 9 User's Guide*. College Station, TX: Stata Press.

- Stock, J. R., Weaver, J. K., Ray, H. W., Brink, J. R., & Sadoff, M. G. 1983. *Evaluation of Safe Performance Secondary School Driver Education Curriculum Demonstration Project*. Washington, DC: U.S. Department of Transportation, National Highway Traffic Safety Administration.
- Thall, P. F., & Vail, S. C. 1990. Some covariance models for longitudinal count data with overdispersion. *Biometrics*, **46**, 657–671.
- Thara, R., Henrietta, M., Joseph, A., Rajkumar, S., & Eaton, W. 1994. Ten year course of schizophrenia - the Madras Longitudinal study. *Acta Psychiatrica Scandinavica*, **90**, 329–336.
- Therneau, T. M., & Grambsch, P. M. 2000. *Modeling Survival Data*. New York: Springer.
- van Belle, G., Fisher, L. D., Heagerty, P. J., & Lumley, T. S. 2004. *Biostatistics: A Methodology for the Health Sciences (Second Edition)*. New York: Wiley.
- van der Heijden, P. G. M., Mooijaart, A., & de Leeuw, J. 1992. Constrained latent budget analysis. Pages 279–320 of: Clogg, C. C. (ed), *Sociological Methodology*, vol. 22. Oxford: Blackwell.
- Vella, F., & Verbeek, M. 1998. Whose wages do unions raise? A dynamic model of unionism and wage rate determination for young men. *Journal of Applied Econometrics*, **13**, 163–183.
- Verbeke, G., & Molenberghs, G. 2000. *Linear Mixed Models for Longitudinal Data*. New York, NY: Springer.
- Wedderburn, R. W. M. 1974. Quasi-likelihood functions, generalized linear models, and the Gauss-Newton method. *Biometrika*, **61**, 439–447.
- Williams, R. L. 2000. A note on robust variance estimation for cluster-correlated data. *Biometrics*, **56**, 645–646.
- Winer, B. J. 1971. *Statistical Principles in Experimental Design*. New York: McGraw-Hill.
- Wolfe, J. H. 1971. *A Monte Carlo study of the sampling distribution of the likelihood ratio for mixtures of multinormal distributions*. Technical Bulletin, vol. STB72-2. San Diego: Naval Personnel and Training Research Laboratory.
- Wooldridge, J. M. 2002. *Econometric Analysis of Cross Section and Panel Data*. Cambridge, MA: The MIT Press.

---

# Index

---

- accessing results, 24, 36  
  `b[username], 24  
  e(), 36  
  e(b), 24  
  e(chi2.dif), 216  
  e(deviance), 140  
  e(dispersp.ps), 149  
  e(l1), 275  
  {eqname}.b[username], 269  
  {eqname}.se[username], 269  
  r(), 36  
  r(mean), 36  
  r(Var), 247  
adaptive quadrature, 177  
adjusted  $R^2$ , 67  
ado-file, 38  
algebraic expression, 16  
analysis of variance, 85–90, 101  
ANCOVA, *see* analysis of covariance  
ANOVA, *see* analysis of variance  
autoregressive structure, 209  
average linkage, 297  
awights, 23  
  
backward elimination, 71  
bar chart, 52  
baseline hazard function, 243  
binomial distribution, 137  
bootstrapping, 138  
boxplot, 48, 163, 217  
  
canonical link, 134  
case-control, 222  
chi-squared test, 46, 51  
classification table, 127  
closing Stata, 8  
cluster analysis, 295–313  
cohort studies, 222  
collinearity, 69  
command  
  adopath, 38  
  anova, 91, 104, 311  
    regress option, 107  
    sequential option, 106  
  assert, 47  
  bologit, 123  
  bootstrap, *see* prefix command, bootstrap  
  by, *see* prefix command, by  
  capture, *see* prefix command, capture  
  cc, 236  
  cci, 230  
  clear, 10  
  cllogit, 235  
  cluster averagelinkage, 298  
  cluster completemlinkage, 298  
  cluster generate, 303  
  cluster kmeans, 308  
    start() option, 308  
  cluster singlelinkage, 298  
  codebook, 47  
  collapse, 22, 164, 166, 235  
  correlate, 54, 67, 286  
  cprplot, 79  
  decode, 13  
  destring, 13  
  display, 13, 14, 30, 122

do, 35, 37  
 drop, 20, 22, 66  
 egen, 19, 20, 30, 167, 248, 308  
 encode, 13, 144  
 epitab, 228  
 estat classif, 127  
     cutoff() option, 127  
 estat gof, 120  
 estat phtest, 255  
 estat rif, 69  
 estat wcorrelation, 207  
     format() option, 207  
 estimates store, 118, 253  
 exit, 36  
     clear option, 9  
 expand, 22, 125, 233, 278  
 foreach, 48, 74, 79, 308  
 format, 12  
 forvalues, 303  
 generate, 19  
 glamm, 60, 183, 185, 192  
     adapt option, 184  
     eform option, 192  
     eqs() option, 193  
     nlp() option, 192  
     nr() option, 193  
 glasper, 185, 192  
     marg option, 195  
     mu option, 195  
     u option, 185  
     us() option, 195  
 glm, 136, 139, 230  
     eform option, 148, 230  
     family() option, 139  
     link() option, 139  
     link(reciprocal) option, 155  
     offset() option, 155, 230  
     scale(x2) option, 147  
     vce(robust) option, 139  
 glm, 139–153  
 global, 36  
 graph, 24–30  
     legend() option, 26  
 graph bar  
     asbyvars option, 53  
     legend(off) option, 53  
     over() option, 28, 53  
     showyvars option, 53  
 graph box, 49, 163, 217  
     over() option, 28, 163, 164  
     relabel() option, 164  
     yline() option, 49  
 graph matrix, 55, 64, 161, 284, 298,  
     303  
     jitter() option, 55  
     mlabel() option, 65  
     mlabposition() option, 65  
 graph twoway, 26, 56, 123, 180, 212,  
     249  
     by() option, 28  
     legend() option, 27  
     xtitle() option, 26, 27  
     ylabel() option, 198  
     ytitle() option, 26, 27  
 graph twoway connected  
     connect(ascending) option, 163  
 graph twoway function, 124  
 graph twoway kdensity, 273  
 graph twoway line  
     connect(ascending) option, 162  
     connect(stairstep) option, 249  
 graph twoway mspline, 124  
 graph twoway rareg, 166  
 graph twoway scatter, 149, 289  
     lpatt() option, 27  
     mlabel() option, 77  
     msymbol() option, 27  
 help, 8  
 help whatnew, 39  
 histogram, 273  
 iin, 207  
 infile, 11, 64, 88, 104, 116, 139,  
     228, 284, 298  
 insheet, 11, 47, 56  
 ir, 228  
 iri, 236  
 kdensity, 185, 273  
 keep, 20, 22  
 ktau, 56  
 label define, 12  
 label values, 12, 89, 178  
 label variable, 11  
 lf, 267  
 lincom, 179  
 list, 12, 18, 162  
     clean option, 162  
     noheader option, 303  
     noobs option, 303

- separator() option, 303  
 local, 36, 140  
 log, 7  
 logistic, 118  
 legit, 117  
     or option, 118  
 lookfor, 15  
 lroc, 127  
 lrtest, 119, 144, 253  
 lsens, 127  
 matrix, 31  
 mcci, 231, 232  
 memory, 11  
 merge, 22  
 nksat, 292  
 nl, 39, 264  
 nl init, 271  
 nl maximize, 267  
     noheader option, 268  
     trace option, 268  
 nl model, 267  
 more, 75  
 svdecode, 12, 178, 206  
 svencode, 12  
 nbreg, 153  
 set cd, 40  
 set from, 40  
 set install, 40  
 nlcom, 272  
 odbc, 11  
 ologit, 120  
 outfile, 11  
 outsheets, 11  
 pca, 286  
     covariance option, 287  
 pcamat, 292  
     names() option, 292  
 pmorn, 77  
 poisson, 230  
     exposure() option, 230  
     irr option, 230  
 predict, 23, 77, 119, 122, 128, 149,  
     180, 181, 216, 256, 289  
 cooked option, 79  
 fitted option, 183  
 number option, 128  
 pearson option, 149  
 pr option, 119, 124  
 reffects option, 183  
 rstandard option, 77  
 score option, 269  
 preserve, 22, 162, 235  
 program, 38  
 program define, 37, 265  
 program drop, 38, 266  
 pwcorr, 54  
     obs option, 54  
     sig option, 54  
 qnorm, 50  
 quietly, *sw* prefix command, quietly  
 recode, 12, 20, 48, 106  
     generate() option, 117  
 regress, 66, 96, 106  
 rename, 11  
 replace, 10, 19  
 reshape, 20–21, 88, 116, 162, 178,  
     190, 206, 210, 228, 233, 245  
     i() option, 89  
 restore, 22, 162, 235  
 robvar, 50  
 rvfplot, 74  
 rvpplot, 74  
 sampei, 30  
 save, 9  
 scalar, 266  
 scoreplot, 289  
 screwpot, 287  
 search, 8  
 serrbar, 185  
 set memory, 10  
 set more off, 36  
 set obs, 206  
 set scheme, 29  
 set seed, 151, 266  
 sort, 17, 163, 233, 250  
 sec install, 40, 183  
 stataby, *see* prefix command, stataby  
 stcox, 246  
     basehazard() option, 250  
     bases() option, 248  
     wsr option, 257  
     mgale() option, 256  
     strata() option, 246  
     texp() option, 252  
     tvc() option, 252  
 stepwise, *see* prefix command, stepwise  
 stjoin, 255  
 stphplot, 251

- stset, 245, 254  
 stsplit, 254  
 stsum, 245  
 summarize, 36, 159, 210, 286, 284  
 syntax, 38  
 sysuse, 10  
 table, 89, 108, 116, 119, 145, 272,  
     303  
     contents() option, 89  
     format() option, 145  
 tabstat, 48, 309  
     statistics() option, 48  
 tabulate, 51, 104, 235  
     exact option, 52  
     expected option, 52  
     nofreq option, 52  
     row option, 52  
 tempfile, 266  
 tempvar, 268  
 testparm, 95, 143  
 tis, 207  
 ttest, 51, 167  
     by() option, 167  
     unequal option, 51, 167  
 ttesti, 31  
 twoway, see command, graph twoway  
 update all, 39  
 use, 9, 10, 144, 210  
     clear option, 10  
 version 9.2, 35  
 xi, see prefix command, xi  
 xpose, 22, 233  
 xtdes, 162  
 xtgee, 205, 206, 209, 213, 216  
     corr(exchangeable) option, 207  
     correlation() option, 206  
     correlation(unstructured) op-  
         tion, 209  
     efrom option, 216  
     family() option, 205  
     i() option, 207  
     link() option, 206  
     scale(x2) option, 214  
     t() option, 207  
     vce(robust) option, 215  
 xtlogit, 190  
     intpoints() option, 192  
 xtmixed, 180  
 covariance(unstructured) op-  
         tion, 187  
     mls option, 181  
     nocons option, 181  
 xtreg, 178  
     fe option, 200  
 commenting out lines, 35  
 complete linkage, 297  
 compound symmetry, 204  
 conditional likelihood, 224  
 conditional logistic regression, 227, 237,  
     235  
 contingency table, see two-way table  
 Cook's distance, 79  
 correlation, 67  
 correlation matrix, 286  
 Cox regression, 243  
 cross-sectional time series, 178  
 cumulative hazard function, 242, 244  
  
 data  
     age of onset of schizophrenia, 263,  
         277  
     Australian school children, 133, 154,  
         199  
     auto pollution filter noise, 96  
     clotting times of blood, 155  
     crowd reactions to threatened sui-  
         cide, 60  
     determinants of pollution in U.S.  
         cities, 61, 82, 292, 308, 312  
     diagnosis of heart attacks, 111, 130  
     driver education, 219  
     duration of UN peacekeeping mis-  
         sions, 260  
     effectiveness of slimming clinics, 101,  
         108, 153, 277  
     efficiency of cycling, 98  
     epileptic seizures and chemotherapy,  
         209, 201, 218  
     estrogens and endometrial cancer,  
         236  
     extroversion and car care, 82  
     female psychiatric patients, 43, 57,  
         129  
     hearing measurement using an audiometer, 281, 291  
     induced abortion and ectopic preg-  
         nancy, 237

- invasion of acacia trees by ants, 59  
 jaw growth, 171, 199  
 life expectancies, 313  
 low energy diet and heart disease, 236  
 maternal behavior in rats, 99  
 mortality from skin cancer, 58, 83  
*New York Times* death notices, 278  
 oral contraceptive use and myocardial infarction, 236  
 prostate cancer, 131  
 psychiatric screening data, 130  
 retention of heroin addicts in methadone maintenance treatment, 239, 258  
 role-taking in young children, 109  
 Roman-British pottery, 312  
 satisfaction with housing conditions, 131  
 sexual satisfaction, 50  
 survival of patients with primary biliary cirrhosis, 259  
 systolic blood pressure, 109  
 thought disorder and schizophrenia, 173, 199, 219  
 Tibetan skulls, 295, 311  
 treating hypertension, 85  
 treatment of Alzheimer's, 172  
 treatment of lung cancer, 111, 129  
 treatment of post-natal depression, 157, 170, 218  
 treatment of prostate cancer, 260  
 wage increases, 41, 171, 199  
 water hardness, 83  
 wave damage to cargo ships, 154
- data browser, 4  
 data editor, 4  
 data management, 19-22  
 date format, 12  
 delta method, 118, 272  
 dendrogram, 297  
 deviance, 137  
 deviance residuals, 256  
 dichotomous, 46  
 dictionary file, 87  
 do-file, 34, 35  
 do-file editor, 6, 35  
 double, see storage type, double  
 dummy variables, 108
- egen function  
 group(), 128  
 rowmean(), 20, 167  
 seq(), 116  
 std(), 248, 308  
 tag(), 128  
 total(), 20  
 eigenvalues, 284  
 empirical Bayes, 181  
 epidemiology, 221  
 equation, 267, 268  
 estimation command, 22-23  
 anova, see command, anova  
 clogit, see command, clogit  
 gllamm, see command, gllamm  
 gls, see command, gls  
 logistic, see command, logistic  
 logit, see command, logit  
 ologit, see command, ologit  
 poisson, see command, poisson  
 regress, see command, regress  
 stcox, see command, stcox  
 xtmixed, see command, xtmixed  
 xtgee, see command, xtgee  
 xtlogit, see command, xtlogit  
 xtreg, see command, xtreg  
 Euclidean distance, 297
- F-test, 105  
 FAQ, 2  
 finite mixture distribution, 263  
 Fisher's exact test, 52  
 forward selection, 78  
 frailty, 151  
 function  
 chi2tail(), 16, 144, 276  
 coadd(), 20, 270  
 date(), 13  
 exp(), 16, 269  
 Ptail(), 142  
 invlogit(), 125, 273  
 invnormal(), 16, 266  
 lm(), 273  
 lnfactorial(), 278  
 log(), 16, 212  
 normal(), 16  
 normalden(), 266  
 scalar(), 266  
 sqrt(), 16, 269

- `substr()`, 16, 65
- `sum()`, 128
- `uniform()`, 16, 206
- `fweights`, 23, 143
- GEE, see generalized estimating equations
- generalized estimating equations, 205
- generalized linear mixed model, 177
- generalized linear model, 133–153
- global macro, 267
- graphics, see command, graph
- hazard function, 243
- hazard ratio, 243
- help, 1
- help file, 8
- hierarchical sums of squares, see sequential sums of squares
- histogram, 273
- immediate command, 30
  - `cci`, see command, `cci`
  - `iri`, see command, `iri`
  - `acci`, see command, `acci`
  - `sampai`, see command, `sampai`
  - `ttesti`, see command, `ttesti`
- incidence rate, 224
- incidence rate ratio, 224
- indexed variable, 16
- information matrix, 204
- initial values, 273
- interaction, 86
- interaction diagram, 91
- interval scale, 46
- intraclass correlation, 176
- jackknifing, 128
- k-means clustering, 297
- Kendall's tau-b, 47
- least squares, 63
- leave one out method, 128
- likelihood ratio, 117
- linear predictor, 134, 225, 268
- linear regression, see multiple regression
- link functions, 134
- local macro, 18, 31, 140, 206
- log file, 5
- log likelihood, 264
- log odds, 270
- log transformation, 93, 212
- logical expression, 15
- logistic regression, 111–129, 227
  - `logit`, 113
- longitudinal data, 157–172, 201
- looping, 17
- lowess, 79
- main effects, 86
- Mann-Whitney *U*-test, 46
- Mantel-Haenszel estimate, 228
- matched case-control studies, 236, 231
- matching, 226
- matrix, 31
- maximum likelihood estimation, 204, 224, 263–276
- McNemar's test, 227
- mean profiles, 164
- missing values, 159
- multilevel, 177
- multiple correlation coefficient, 64, 67
- multiple regression, 61–82
- multivariate data, 283
- negative binomial, 153
- NetCourses, 2
- Newton-Raphson algorithm, 264
- normal probability plot, 77
- odds, 113
- odds ratio, 118, 225
- offset, 213, 225
- ordinal, 46
- ordinal logistic regression, 114
- overdispersion, 137, 147, 153, 214, 215
- pairwise correlation, 54
- partial log likelihood, 244
- partial residual plot, 79
- PCA, see principal components analysis
- Pearson  $X^2$ , 137, 147
- Pearson correlation, 47
- Pearson residual, 125, 216
- person-time of observation, 234
- plot, see command, graph
- plug-in, 39
- Poisson distribution, 136, 137, 210, 213, 224

- Poisson regression, 225, 228  
 post-estimation command, 23–24  
     predict, see command, estat  
     estimates, see command, estimates  
     estat, see command, estat  
     gllapred, see command, gllapred  
     lincom, see command, lincom  
     lrtest, see command, lrtest  
     nlcom, see command, nlcom  
     predict, see command, predict  
     test, see command, test  
     testparm, see command, testparm  
 prefix command, 14  
     bootstrap, 151  
         reps() option, 151  
     by, 14, 17, 128, 233  
         sort option, 17, 128  
     capture, 35, 38, 266  
     quietly, 35, 266, 271  
     statsby, 169  
     stepwise, 72, 73  
         pe() option, 72  
         pr() option, 73  
         xi, 95  
 principal components analysis, 281–291  
 profile likelihood, 244  
 programming, 34–39, 263–280  
 proportional hazards, 243  
 proportional odds model, see ordinal logistic regression  
 pwights, 23  
 Q-Q plot, 50  
 quasi-likelihood, 138, 147, 204  
 random effects, 151  
 random effects model, 137, 173–199  
 random intercept model, 175  
 reading data, 9  
     infile, see command, infile  
     insheet, see command, insheet  
     use, see command, use  
 regression  
     conditional logistic, see conditional logistic regression  
     lineal, see multiple regression  
     logistic, see logistic regression  
     negative binomial, see negative binomial regression  
     ordinal logistic, see ordinal logistic regression  
     Poisson, see Poisson regression  
     regression coefficient, 63, 67  
     regression diagnostics, 77–82  
     relative risk, 224  
     REML, see restricted maximum likelihood  
         residuals, 73, 75, 125, 148, 256–257  
         response feature analysis, 167  
         response profiles, 164, 211  
         right-censored, 241  
         robust standard errors, 138, 149  
         ROC-curve, 127  
     saving data, 9  
         outfile, see command, outfile  
         outsheet, see command, outsheets  
         save, see command, save  
     scatterplot matrix, 58, 64, 161, 284  
     screener plot, 287  
     search, 7  
     sensitivity, 127  
     sequential sums of squares, 103, 105  
     single linkage, 297  
     SJ, see Stata Journal  
     Spearman rank correlation, 47  
     specificity, 127  
     SPI, see plug-in  
     SSC, see Statistical Software Components  
     standard error, 207  
     standardized residual, 77, 126  
     Stata Journal, 39  
     Stata Technical Bulletin, 39  
     Stata web page, 2  
     Statistical Software Components, 39  
     STB, see Stata Technical Bulletin  
     stepwise regression, 71  
     stopping Stata, 8  
     storage type, 11  
         double, 267  
     stratified Cox model, 244  
     survival analysis, 239  
     survival data, 241  
     survivor function, 242  
     temporary name, 266  
     t-test, 46, 50, 167

- two-way table, 51
- Type I sums of squares, *see* sequential sums of squares
- Type III sums of squares, *see* unique sums of squares
- unique sums of squares, 104
- updating Stata, 39–40
- variance inflation factors, 67
- weights, 14, 23, 143

*A Handbook of*  
**Statistical  
 Analyses**  
*Using Stata*

*Fourth Edition*

Sophia Rabe-Hesketh and Brian S. Everitt

With each new release of Stata, a comprehensive resource is needed to highlight the improvements as well as discuss the fundamentals of the software. Fulfilling this need, *A Handbook of Statistical Analyses Using Stata, Fourth Edition* has been fully updated to provide an introduction to Stata Version 9. This edition covers many new features of Stata, including a new command for mixed models and a new matrix language.

Each chapter describes the analysis appropriate for a particular application, focusing on the medical, social, and behavioral fields. The authors begin each chapter with descriptions of the data and the statistical techniques to be used. The methods covered include descriptive statistics, simple tests, analysis of variance, multiple linear regression, logistic regression, generalized linear models, survival analysis, random effects models, and cluster analysis. The core of the book centers on how to use Stata to perform analyses and how to interpret the results. The chapters conclude with several exercises based on datasets from different disciplines.

**Features**

- Demonstrates how a wide variety of statistical analyses, including descriptive statistics, graphics, model estimation and diagnostics, can be performed using Stata
- Features the new mixed-models estimation command `xtmixed` and the new matrix language Mata
- Incorporates numerous examples that use real-life data to explain the applications of the methods described
- Includes diverse datasets in many exercises as well as selected solutions in the appendix to build a better understanding of the methodology

**Sophia Rabe-Hesketh** is Professor of Educational Statistics at the Graduate School of Education at the University of California, Berkeley, USA, and Chair of Social Statistics at the Institute of Education, University of London, UK.

**Brian S. Everitt** is Emeritus Professor of Statistics at the Institute of Psychiatry, King's College, University of London, UK.



**Chapman & Hall/CRC**

Taylor & Francis Group  
 an Informa business

[www.taylorandfrancisgroup.com](http://www.taylorandfrancisgroup.com)

6000 Broken Sound Parkway, NW  
 Suite 300, Boca Raton, FL 33487

270 Madison Avenue  
 New York, NY 10016

2 Park Square, Milton Park,  
 Abingdon, Oxon OX14 4BN, UK

C7567

ISBN 1-58488-756-7



9781584887560

[www.crcpress.com](http://www.crcpress.com)